

# Java Parallel Stream Internals

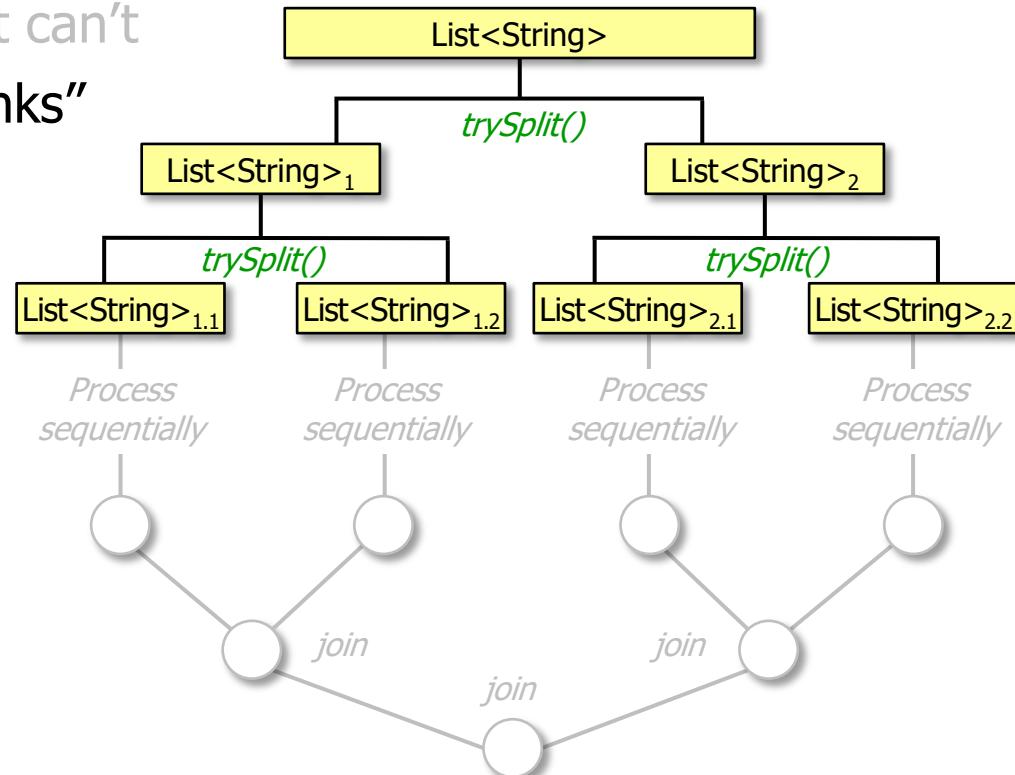
---

## Partitioning

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

- Understand parallel stream internals, e.g.
  - Know what can change and what can't
  - Partition a data source into "chunks"



---

# Partitioning a Parallel Stream

# Partitioning a Parallel Stream

- A “splittable iterator” (Spliterator) partitions a Java parallel stream into chunks.

## Interface Spliterator<T>

### Type Parameters:

T - the type of elements returned by this Spliterator

### All Known Subinterfaces:

`Spliterator.OfDouble`, `Spliterator.OfInt`, `Spliterator.OfLong`,  
`Spliterator.OfPrimitive<T,T_CONS,T_SPLITR>`

### All Known Implementing Classes:

`Spliterators.AbstractDoubleSpliterator`,  
`Spliterators.AbstractIntSpliterator`,  
`Spliterators.AbstractLongSpliterator`,  
`Spliterators.AbstractSpliterator`

## public interface Spliterator<T>

An object for traversing and partitioning elements of a source. The source of elements covered by a Spliterator could be, for example, an array, a `Collection`, an IO channel, or a generator function.

A Spliterator may traverse elements individually (`tryAdvance()`) or sequentially in bulk (`forEachRemaining()`).

# Partitioning a Parallel Stream

- We've shown how a Spliterator can *traverse* elements in a source.

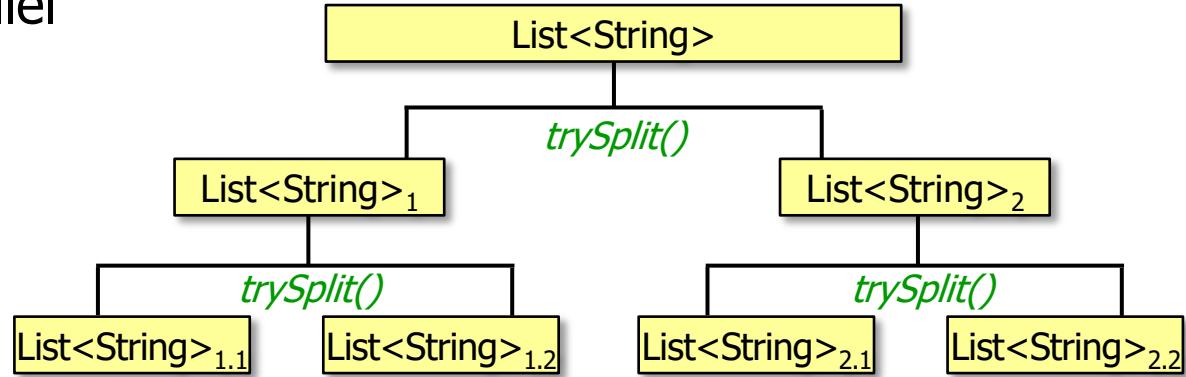
```
List<String> quote = Arrays.asList
    ("This ", "above ", "all- ",
     "to ", "thine ", "own ",
     "self ", "be ", "true", ",\n",
     ...);

for(Spliterator<String> s =
    quote.spliterator();
    s.tryAdvance(System.out::print)
    != false;
)
    continue;
```

# Partitioning a Parallel Stream

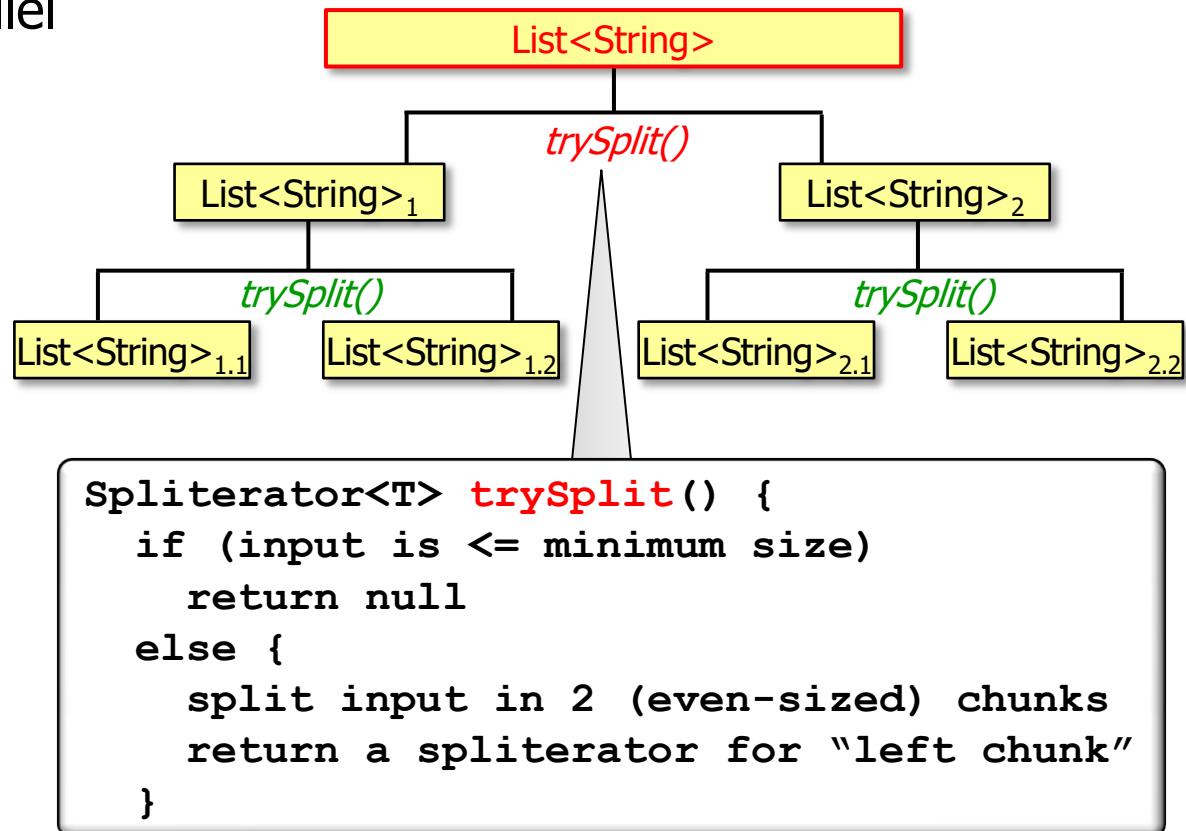
---

- We now outline how a parallel Spliterator can *partition* all elements in a source.



# Partitioning a Parallel Stream

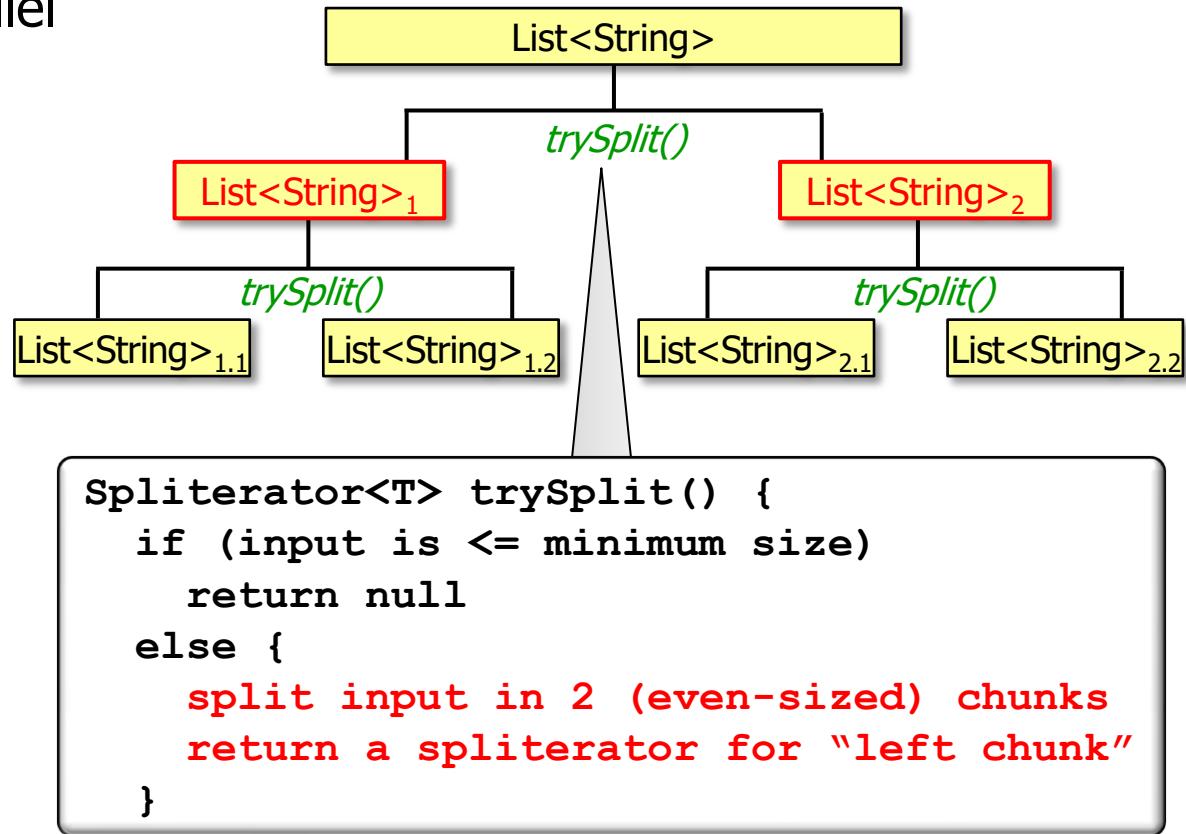
- We now outline how a parallel Spliterator can *partition* all elements in a source.



The streams framework calls a Spliterator's `trySplit()` method, not a user's app.

# Partitioning a Parallel Stream

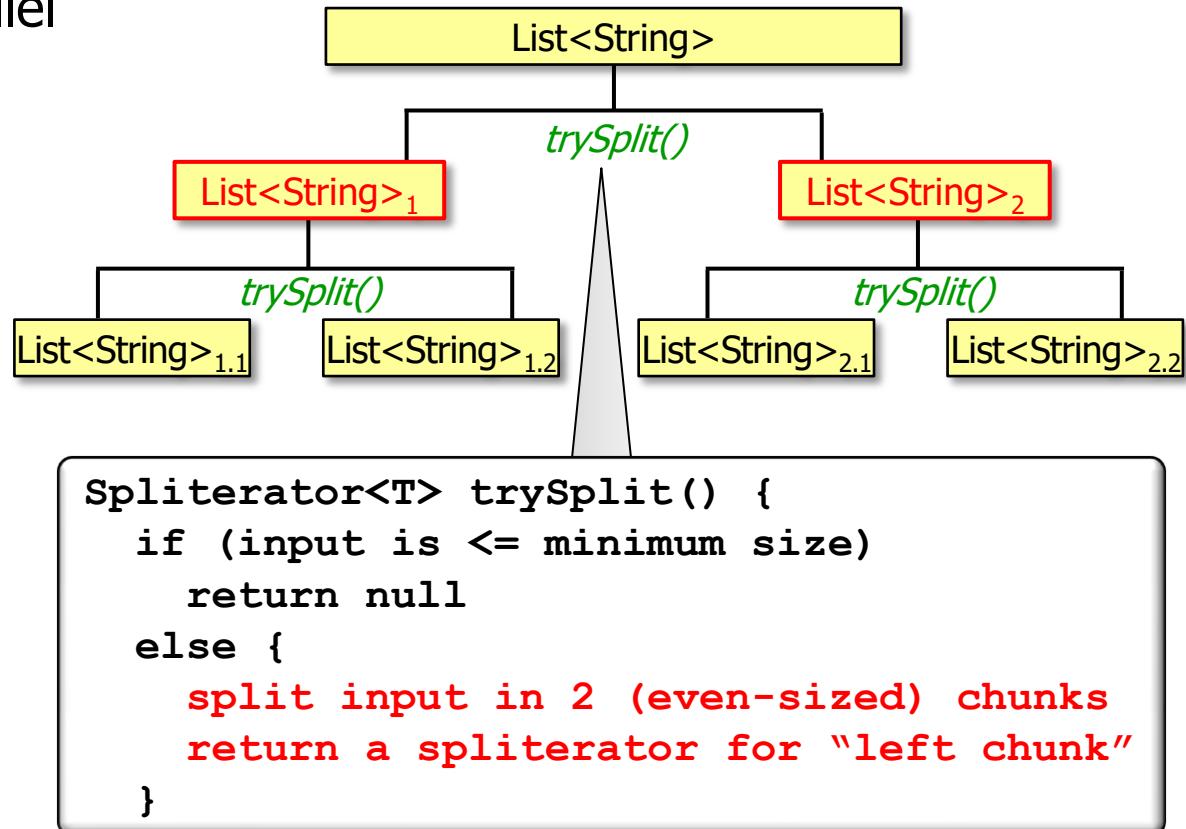
- We now outline how a parallel Spliterator can *partition* all elements in a source.



trySplit() attempts to split the input evenly (if it's not  $\leq$  the minimum size)

# Partitioning a Parallel Stream

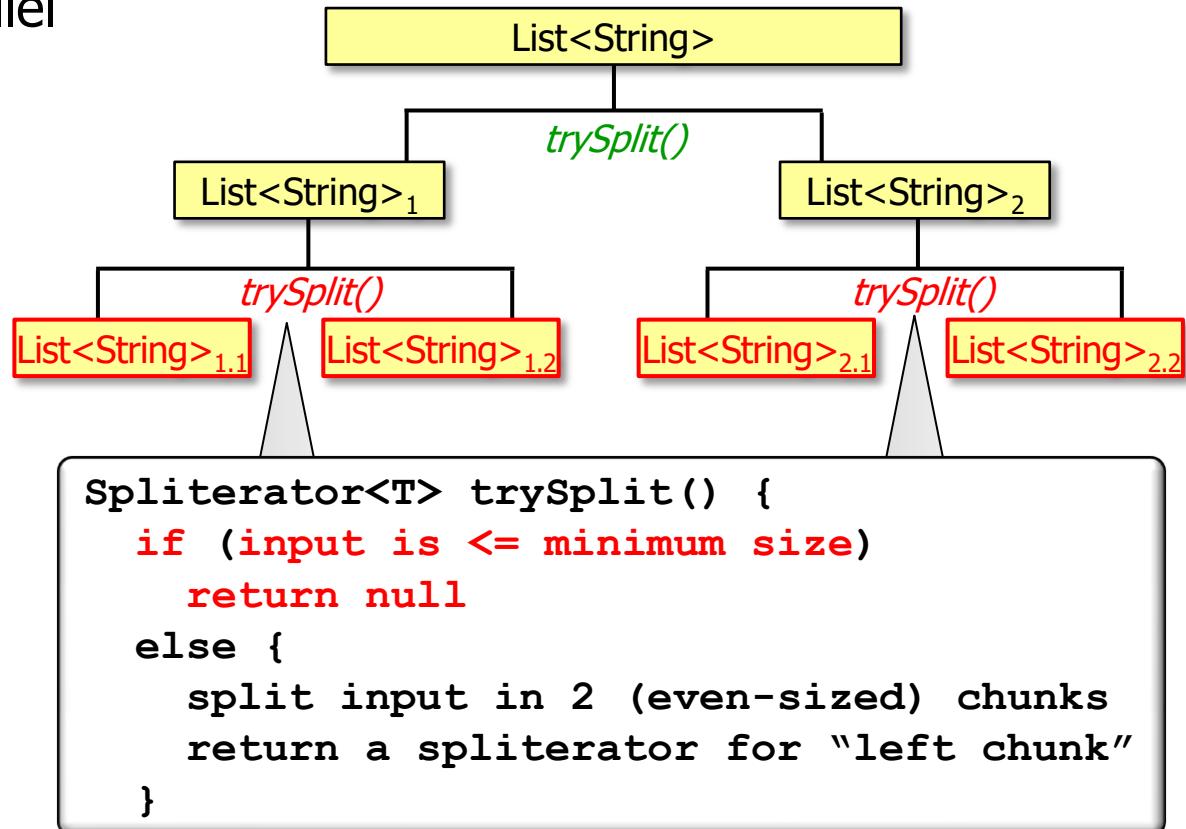
- We now outline how a parallel Spliterator can *partition* all elements in a source.



A Spliterator usually needs no synchronization nor does it need a “join” phase!

# Partitioning a Parallel Stream

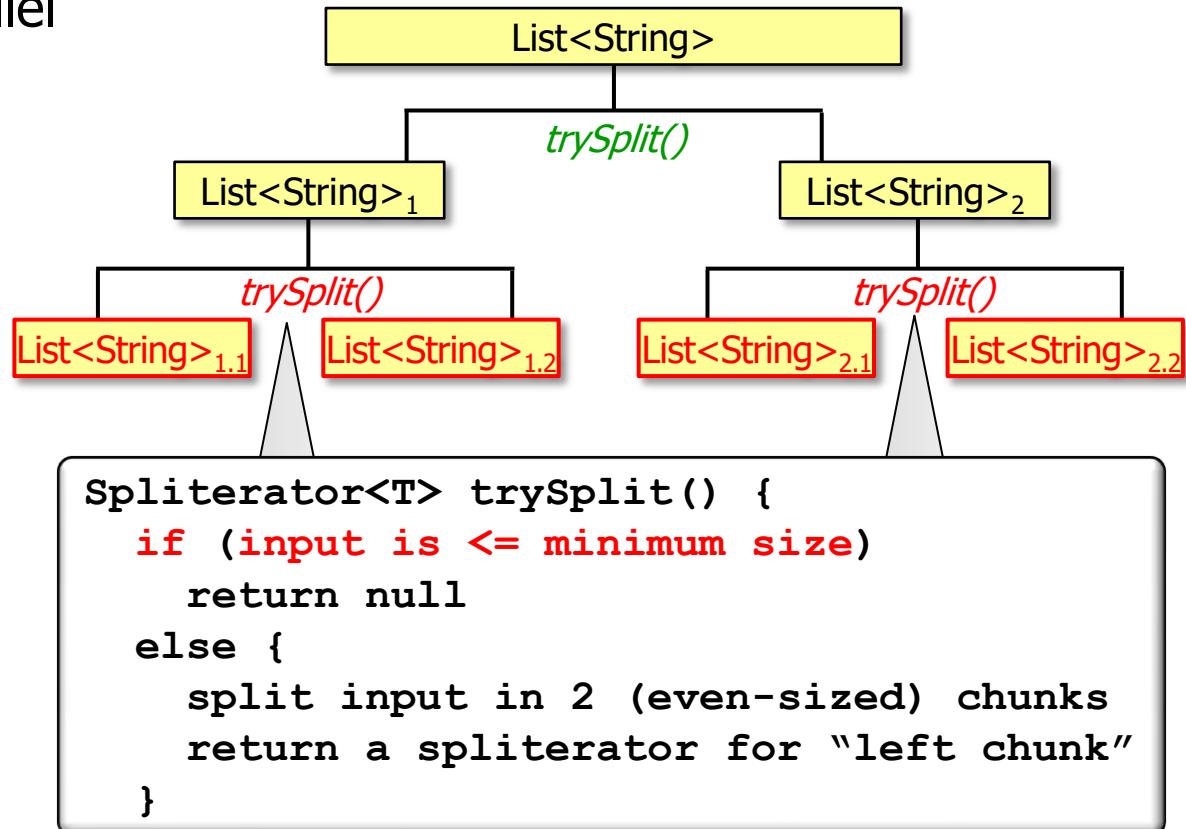
- We now outline how a parallel Spliterator can *partition* all elements in a source.



`trySplit()` is called recursively until all chunks are  $\leq$  to the minimize size

# Partitioning a Parallel Stream

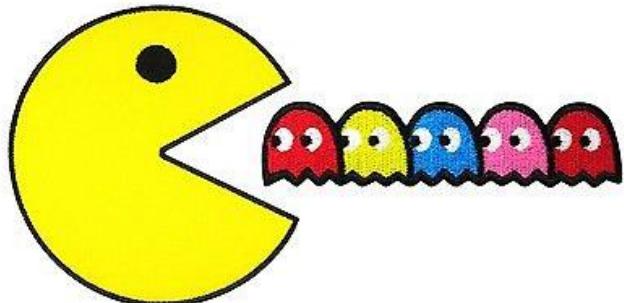
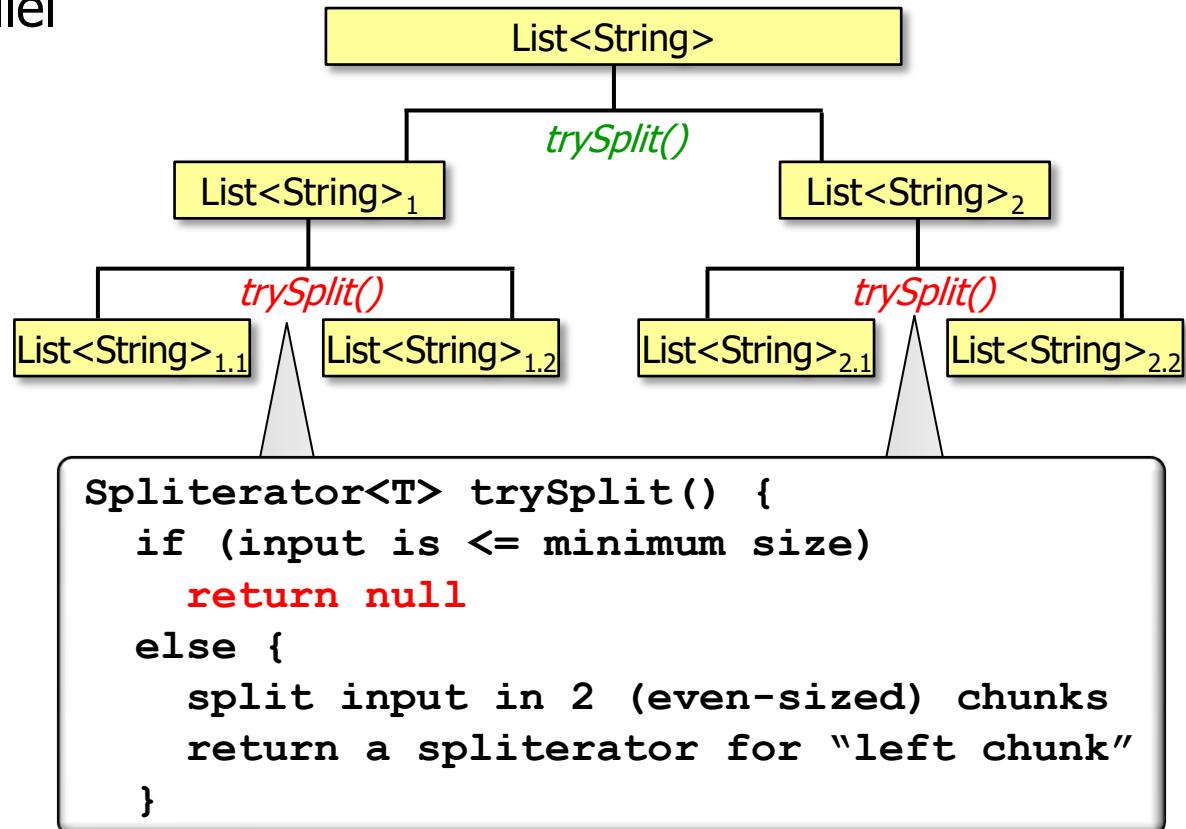
- We now outline how a parallel Spliterator can *partition* all elements in a source.



trySplit() is finished when a chunk is  $\leq$  to the minimize size

# Partitioning a Parallel Stream

- We now outline how a parallel Spliterator can *partition* all elements in a source.



When null is returned, the streams framework processes this chunk sequentially.

# Partitioning a Parallel Stream

- Some Java collections split evenly and efficiently, e.g., ArrayList.

```
ArrayListSpliterator<E> trySplit() {  
    int hi = getFence(), lo = index, mid = (lo + hi) >>> 1;  
    // divide range in half unless too small  
    return lo >= mid ? null : new ArrayListSpliterator<E>  
        (list, lo, index = mid, ...);  
}
```

```
boolean tryAdvance(Consumer<? super E> action) {  
    ...  
    if (index < getFence()) {  
        action.accept((E) list.elementData[i]); ...  
        return true;  
    } return false;  
}
```



See [openjdk/8u40-b25/java/util/ArrayList.java](https://openjdk.java.net/jeps/204)

# Partitioning a Parallel Stream

- Some Java collections split evenly and efficiently, e.g., ArrayList.

```
ArrayListSpliterator<E> trySplit() {  
    int hi = getFence(), lo = index, mid = (lo + hi) >>> 1;  
    // divide range in half unless too small  
    return lo >= mid ? null : new ArrayListSpliterator<E>  
        (list, lo, index = mid, ...);  
}
```

*Split the array evenly each time until there's nothing left to split.*

```
boolean tryAdvance(Consumer<? super E> action) {  
    ...  
    if (index < getFence()) {  
        action.accept( (E) list.elementData[i] ) ; ...  
        return true;  
    } return false;  
}
```

# Partitioning a Parallel Stream

- Some Java collections split evenly and efficiently, e.g., ArrayList.

```
ArrayListSpliterator<E> trySplit() {  
    int hi = getFence(), lo = index, mid = (lo + hi) >>> 1;  
    // divide range in half unless too small  
    return lo >= mid ? null : new ArrayListSpliterator<E>  
        (list, lo, index = mid, ...);  
}
```

*Try to consume a single element on each call.*

```
boolean tryAdvance(Consumer<? super E> action) {  
    ...  
    if (index < getFence()) {  
        action.accept((E) list.elementData[i]); ...  
        return true;  
    } return false;  
}
```

# Partitioning a Parallel Stream

- Other Java collections do *not* split evenly and efficiently, e.g., `LinkedList`.

```
Spliterator<E> trySplit() { ...  
    int n = batch + BATCH_UNIT, j = 0; Object[] a = new Object[n];  
    do { a[j++] = p.item; }  
    while ((p = p.next) != null && j < n); ...  
    return Spliterators.spliterator(a, 0, j, Spliterator.ORDERED);  
}
```

```
boolean tryAdvance(Consumer<? super E> action) { ...  
    Node<E> p;  
    if (getEst() > 0 && (p = current) != null) {  
        --est; E e = p.item; current = p.next;  
        action.accept(e); return true;  
    } return false;  
}
```



See [openjdk/8u40-b25/java/util/LinkedList.java](https://openjdk.java.net/jeps/200/parallel-streams.html)

# Partitioning a Parallel Stream

---

- Other Java collections do *not* split evenly and efficiently, e.g., `LinkedList`.

```
Spliterator<E> trySplit() { ...  
    int n = batch + BATCH_UNIT, j = 0; Object[] a = new Object[n];  
    do { a[j++] = p.item; }  
    while ((p = p.next) != null && j < n); ...  
    return Spliterators.spliterator(a, 0, j, Spliterator.ORDERED);  
}  
    Split the list into "batches," rather than evenly in half.
```

```
boolean tryAdvance(Consumer<? super E> action) { ...  
    Node<E> p;  
    if (getEst() > 0 && (p = current) != null) {  
        --est; E e = p.item; current = p.next;  
        action.accept(e); return true;  
    } return false;  
}
```

# Partitioning a Parallel Stream

- Other Java collections do *not* split evenly and efficiently, e.g., `LinkedList`.

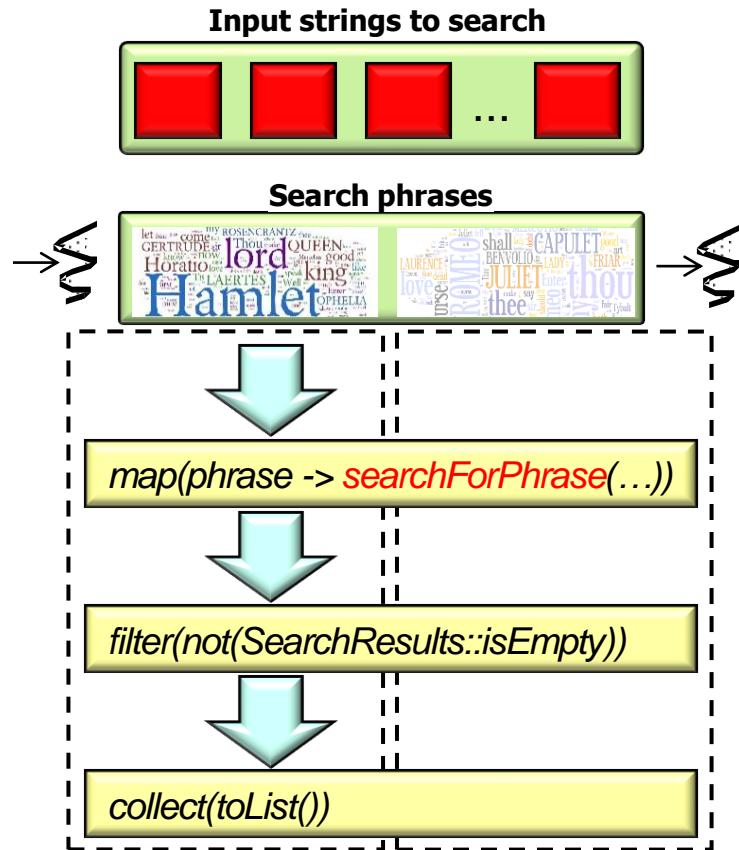
```
Spliterator<E> trySplit() { ...  
    int n = batch + BATCH_UNIT, j = 0; Object[] a = new Object[n];  
    do { a[j++] = p.item; }  
    while ((p = p.next) != null && j < n); ...  
    return Spliterators.spliterator(a, 0, j, Spliterator.ORDERED);  
}
```

*Try to consume a single element on each call.*

```
boolean tryAdvance(Consumer<? super E> action) { ...  
    Node<E> p;  
    if (getEst() > 0 && (p = current) != null) {  
        --est; E e = p.item; current = p.next;  
        action.accept(e); return true;  
    } return false;  
}
```

# Partitioning a Parallel Stream

- We'll cover the implementation details of parallel Spliterators in upcoming lessons.



See "Java SearchWithParallelSpliterator Example: trySplit()"

Java Parallel Stream Internals: Partitioning

---

**The End**

# Java Parallel Stream Internals

---

## Demo'ing Spliterator Performance

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

---

- Understand parallel stream internals,  
e.g.
  - Know what can change and what can't
  - Partition a data source into "chunks"
    - Know the impact of different Java collections on the performance of different spliterators

```
Starting spliterator tests for 1000 words..  
..printing results
```

```
17 msec: ArrayList parallel  
19 msec: LinkedList parallel
```

```
Starting spliterator tests for 10000 words..
```

```
..printing results  
88 msec: LinkedList parallel  
90 msec: ArrayList parallel
```

```
Starting spliterator tests for 100000 words..
```

```
..printing results  
599 msec: ArrayList parallel  
701 msec: LinkedList parallel
```

```
Starting spliterator tests for 883311 words..
```

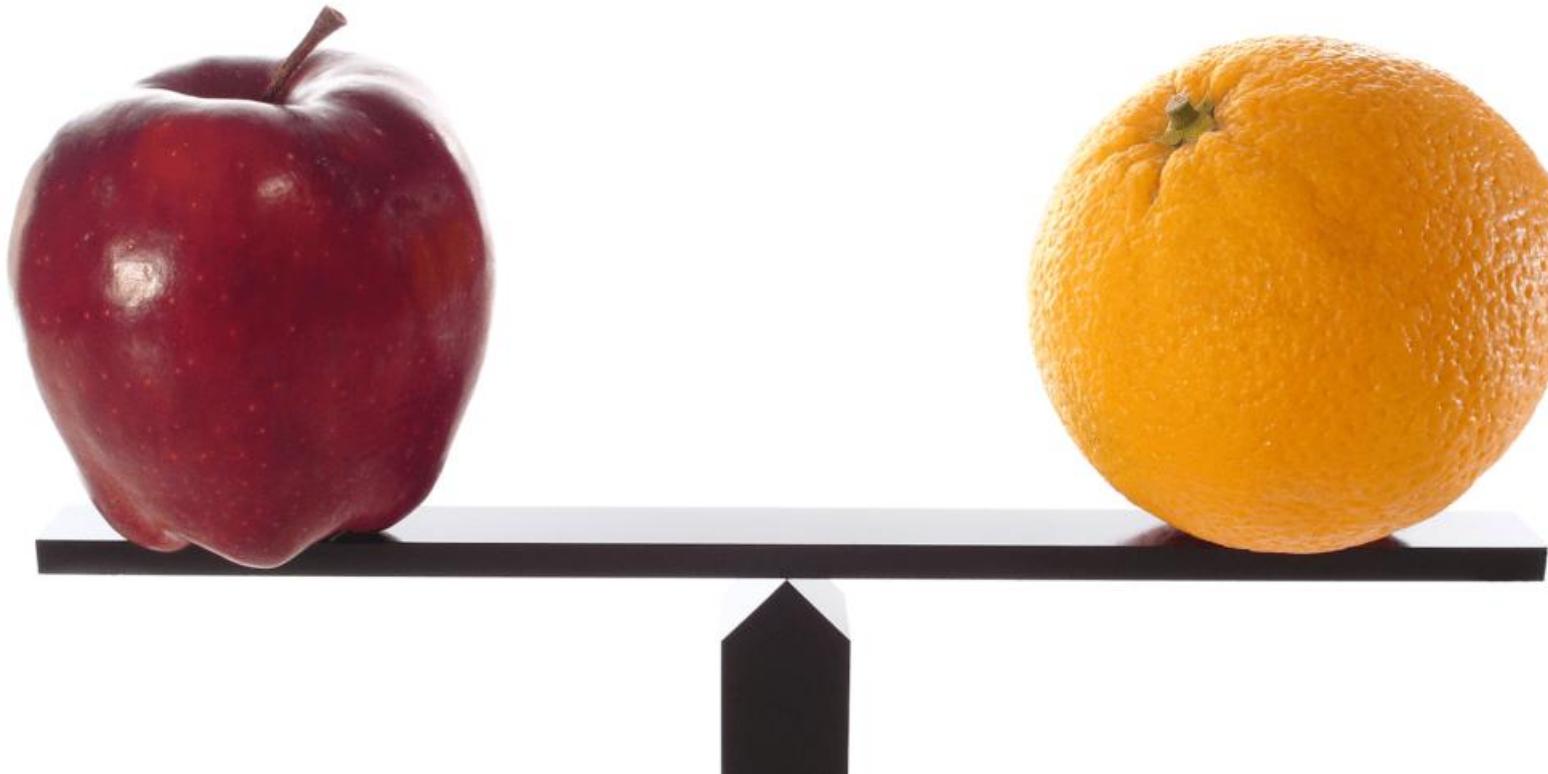
```
..printing results  
5718 msec: ArrayList parallel  
31226 msec: LinkedList parallel
```

---

# Demonstrating Spliterator Performance

# Demonstrating Spliterator Performance

- Spliterators for ArrayList and LinkedList partition data quite differently.



See earlier lesson on “*Java Parallel Streams Internals: Partitioning*”

# Demonstrating Spliterator Performance

- Spliterators for ArrayList and LinkedList partition data quite differently.

```
ArrayListSpliterator<E> trySplit() {  
    int hi = getFence(), lo = index, mid = (lo + hi) >>> 1;  
    // divide range in half unless too small  
    return lo >= mid ? null : new ArrayListSpliterator<E>  
                      (list, lo, index = mid, ...);  
}
```

*ArrayList's spliterator splits evenly & efficiently (e.g., doesn't copy data)*



See [openjdk/8u40-b25/java/util/ArrayList.java](https://openjdk.java.net/jeps/200)

# Demonstrating Spliterator Performance

- Spliterators for ArrayList and LinkedList partition data quite differently.

```
Spliterator<E> trySplit() { ...
```

```
    int n = batch + BATCH_UNIT, j = 0; Object[] a = new Object[n];
    do { a[j++] = p.item; }
    while ((p = p.next) != null && j < n); ...
    return Spliterators.spliterator(a, 0, j, Spliterator.ORDERED);
}
```

*LinkedList's spliterator does not split evenly & efficiently (e.g., it copies data)*



See [openjdk/8u40-b25/java/util/LinkedList.java](https://openjdk.java.net/jeps/200)

# Demonstrating Spliterator Performance

---

- This program shows the performance difference of parallel spliterators for ArrayList and LinkedList when processing the complete works of Shakespeare.

```
void timeParallelStreamUppercase(String testName,
                                  List<CharSequence> words) {
    ...
    List<String> list = new ArrayList<>();
    for (int i = 0; i < sMAX_ITERATIONS; i++) {
        list
            .addAll(words
                    .parallelStream()
                    .map(charSeq ->
                        charSeq.toString().toUpperCase())
                    .collect(toList())));
    ...
}
```

# Demonstrating Spliterator Performance

- This program shows the performance difference of parallel spliterators for ArrayList and LinkedList when processing the complete works of Shakespeare.

```
void timeParallelStreamUppercase(String testName,  
                                List<CharSequence> words) {  
    ...  
    List<String> list = new ArrayList<>();  
  
    for (int i = 0; i < sMAX_ITERATIONS; i++)  
        list  
            .addAll(words  
                    .parallelStream()  
                    .map(charSeq ->  
                        charSeq.toString().toUpperCase())  
                    .collect(toList())); ...  
}
```

*The words param is passed  
an ArrayList & a LinkedList*

# Demonstrating Spliterator Performance

- This program shows the performance difference of parallel spliterators for ArrayList and LinkedList when processing the complete works of Shakespeare.

```
void timeParallelStreamUppercase(String testName,  
                                List<CharSequence> words) {  
    ...  
    List<String> list = new ArrayList<>();  
  
    for (int i = 0; i < sMAX_ITERATIONS; i++)  
        list  
            .addAll(words  
                    .parallelStream()  
                    .map(charSeq ->  
                        charSeq.toString().toUpperCase())  
                    .collect(toList())); ...  
}
```

*Split & uppercase  
the word list via a  
parallel spliterator*

# Demonstrating Spliterator Performance

---

- Results show spliterator differences become more significant as input grows.

Starting spliterator tests for 1000 words....printing results

17 msec: ArrayList parallel

19 msec: LinkedList parallel

Starting spliterator tests for 10000 words....printing results

88 msec: ArrayList parallel

90 msec: LinkedList parallel

Starting spliterator tests for 100000 words....printing results

599 msec: ArrayList parallel

701 msec: LinkedList parallel

Starting spliterator tests for 883311 words....printing results

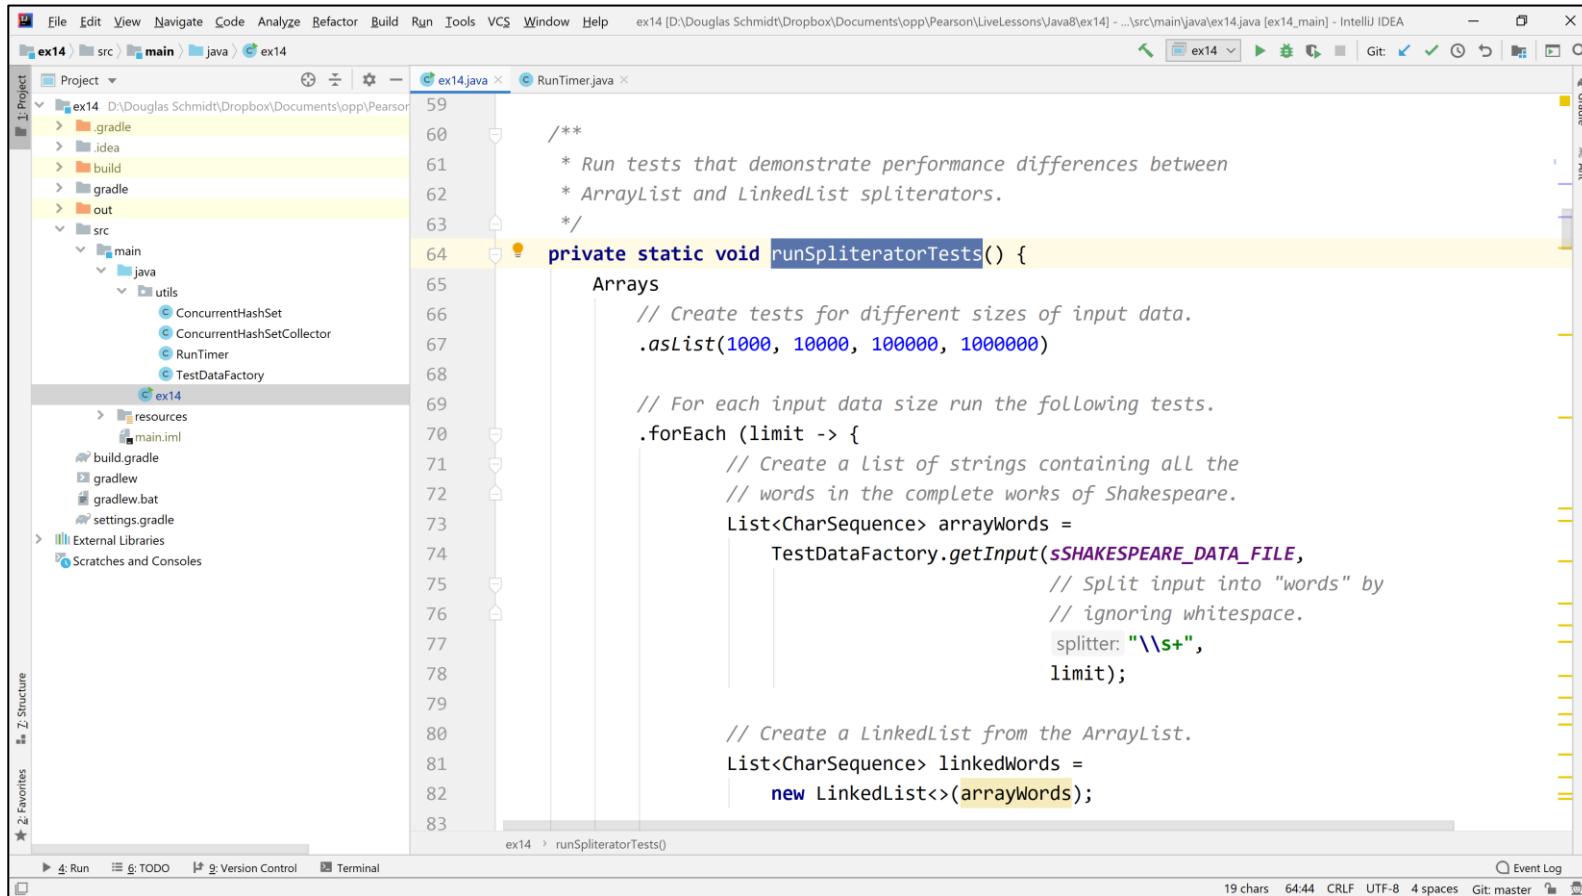
5718 msec: ArrayList parallel

31226 msec: LinkedList parallel

---

See upcoming lessons on “*When [Not] to Use Parallel Streams*”

# Demonstrating Spliterator Performance



The screenshot shows the IntelliJ IDEA IDE interface with the following details:

- File Menu:** File, Edit, View, Navigate, Code, Analyze, Refactor, Build, Run, Tools, VCS, Window, Help.
- Project Bar:** ex14, src, main, java, ex14.
- Toolbars:** Standard, Git, Event Log.
- Code Editor:** The editor shows the file `ex14.java` with the following code:

```
59  /**
60  * Run tests that demonstrate performance differences between
61  * ArrayList and LinkedList spliterators.
62  */
63
64  private static void runSpliteratorTests() {
65      Arrays
66          // Create tests for different sizes of input data.
67          .asList(1000, 10000, 100000, 1000000)
68
69      // For each input data size run the following tests.
70      .forEach (limit -> {
71          // Create a list of strings containing all the
72          // words in the complete works of Shakespeare.
73          List<CharSequence> arrayWords =
74              TestDataFactory.getInput(SHAKESPEARE_DATA_FILE,
75                  // Split input into "words" by
76                  // ignoring whitespace.
77                  splitter: "\\\s+",
78                  limit);
79
80          // Create a LinkedList from the ArrayList.
81          List<CharSequence> linkedWords =
82              new LinkedList<>(arrayWords);
83      })
84  }
```

- Tool Windows:** Project, Structure, Favorites.
- Bottom Bar:** Run, TODO, Version Control, Terminal, Event Log.

See [github.com/douglasraigschmidt/LiveLessons/tree/master/Java8/ex14](https://github.com/douglasraigschmidt/LiveLessons/tree/master/Java8/ex14)

# Demonstrating Spliterator Performance

---

**[Source code analysis  
video goes here!]**

---

Java Parallel Stream Internals: Demo'ing Spliterator Performance

---

# The End

---

# Java Parallel Stream Internals

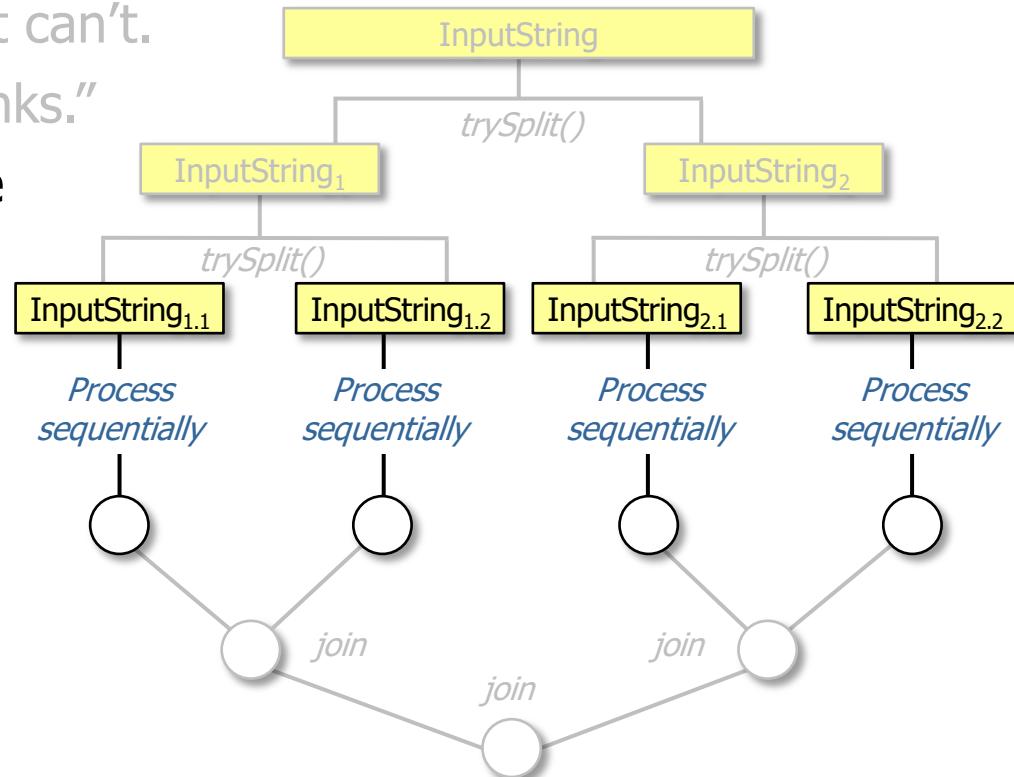
---

Parallel Processing via  
the Common ForkJoinPool

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.

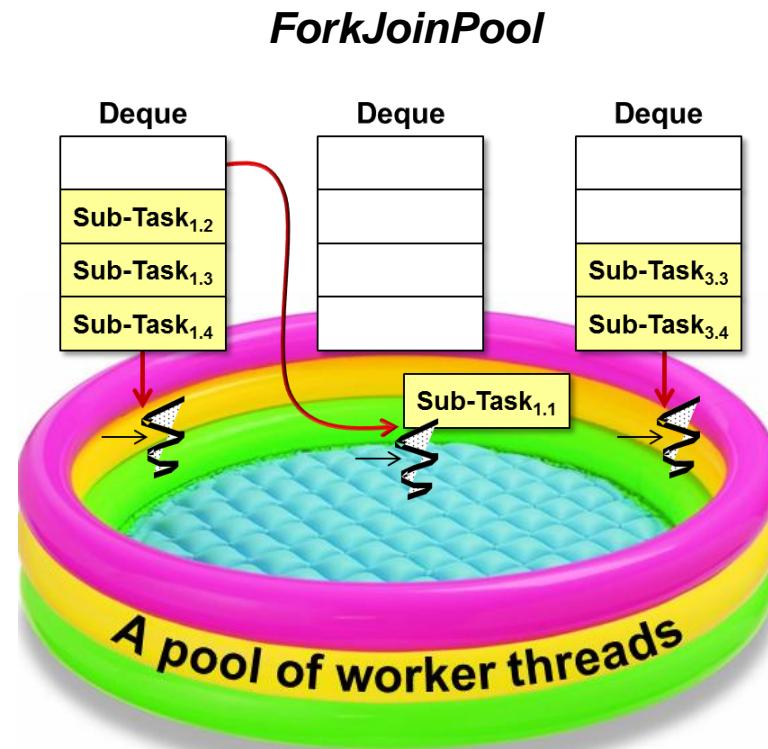


---

# Processing Chunks in Parallel via the Common ForkJoinPool

# Processing Chunks in Parallel via the Common ForkJoinPool

- Chunks created by a Spliterator are processed in the common ForkJoinPool.



# Processing Chunks in Parallel via the Common ForkJoinPool

- A ForkJoinPool provides a high performance, fine-grained task execution framework for Java data parallelism.

## Class ForkJoinPool

```
java.lang.Object
  java.util.concurrent.AbstractExecutorService
    java.util.concurrent.ForkJoinPool
```

### All Implemented Interfaces:

Executor, ExecutorService

```
public class ForkJoinPool
extends AbstractExecutorService
```

An ExecutorService for running ForkJoinTasks. A ForkJoinPool provides the entry point for submissions from non-ForkJoinTask clients, as well as management and monitoring operations.

A ForkJoinPool differs from other kinds of ExecutorService mainly by virtue of employing *work-stealing*: all threads in the pool attempt to find and execute tasks submitted to the pool and/or created by other active tasks (eventually blocking waiting for work if none exist). This enables efficient processing when most tasks spawn other subtasks (as do most ForkJoinTasks), as well as when many small tasks are submitted to the pool from external clients. Especially when setting *asyncMode* to true in constructors, ForkJoinPools may also be appropriate for use with event-style tasks that are never joined.

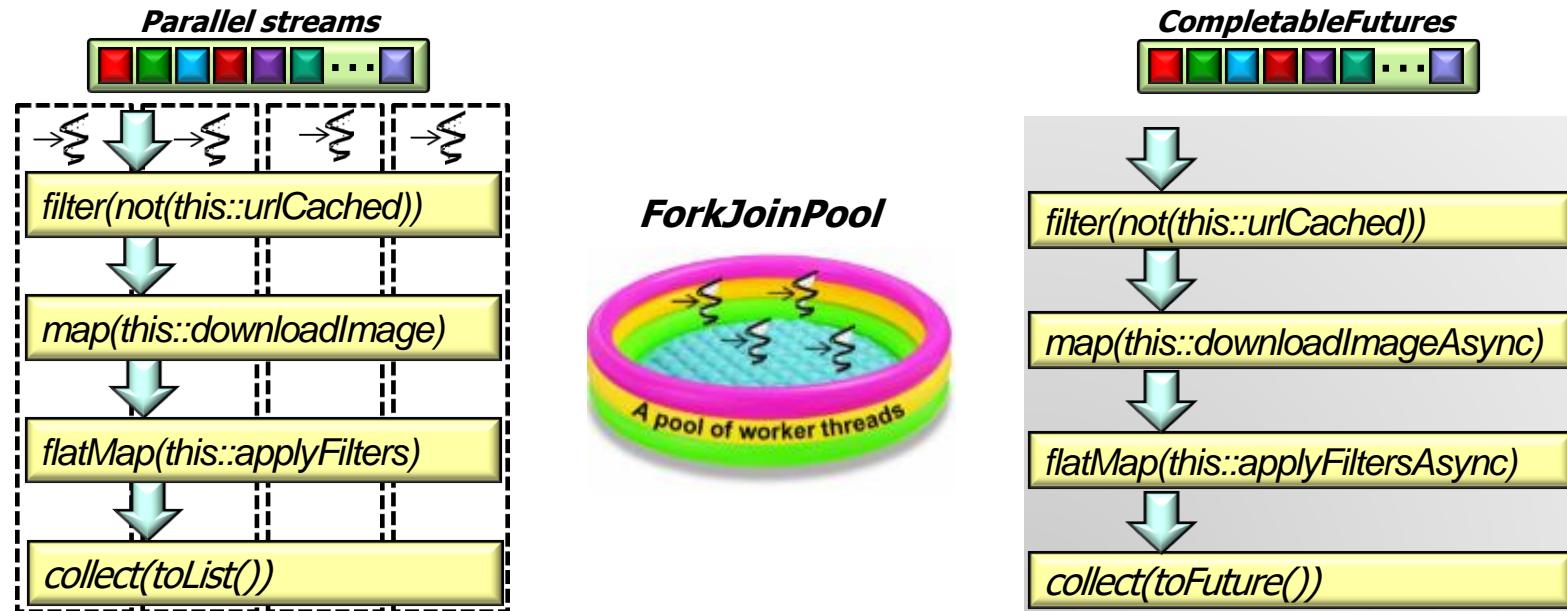
A static *commonPool()* is available and appropriate for most applications. The common pool is used by any ForkJoinTask that is not explicitly submitted to a specified pool. Using the common pool normally reduces resource usage (its threads are slowly reclaimed during periods of non-use, and reinstated upon subsequent use).

For applications that require separate or custom pools, a ForkJoinPool may be constructed with a given target parallelism level; by default, equal to the number of available processors. The pool attempts to maintain enough active (or available) threads by dynamically adding, suspending, or resuming internal worker threads, even if some tasks are stalled waiting to join others. However, no such adjustments are guaranteed in the face of blocked I/O or other unmanaged synchronization. The nested ForkJoinPool.ManagedBlocker interface enables extension of the kinds of synchronization accommodated.



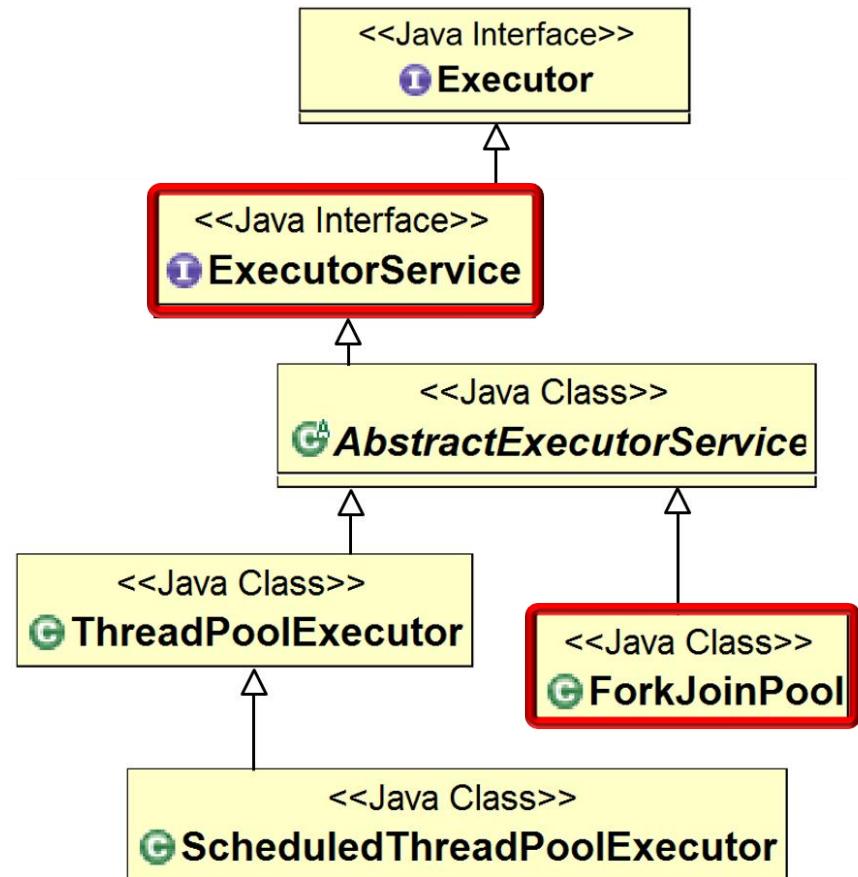
# Processing Chunks in Parallel via the Common ForkJoinPool

- A ForkJoinPool provides a high performance, fine-grained task execution framework for Java data parallelism.
  - It provides a parallel computing engine for many higher-level frameworks.



# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool implements the ExecutorService interface.



# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool implements the ExecutorService interface.
  - A ForkJoinPool executes fork/join tasks.

## Class ForkJoinTask<V>

java.lang.Object  
java.util.concurrent.ForkJoinTask<V>

### All Implemented Interfaces:

Serializable, Future<V>

### Direct Known Subclasses:

CountedCompleter, RecursiveAction,  
RecursiveTask

---

```
public abstract class ForkJoinTask<V>
extends Object
implements Future<V>, Serializable
```

Abstract base class for tasks that run within a ForkJoinPool. A ForkJoinTask is a thread-like entity that is much lighter weight than a normal thread. Huge numbers of tasks and subtasks may be hosted by a small number of actual threads in a ForkJoinPool, at the price of some usage limitations.

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool implements the ExecutorService interface.
  - A ForkJoinPool executes fork/join tasks.
  - A fork/join task associates a chunk of data along with a computation on that data to enable fine-grained parallelism.



## Class ForkJoinTask<V>

java.lang.Object

java.util.concurrent.ForkJoinTask<V>

### All Implemented Interfaces:

Serializable, Future<V>

### Direct Known Subclasses:

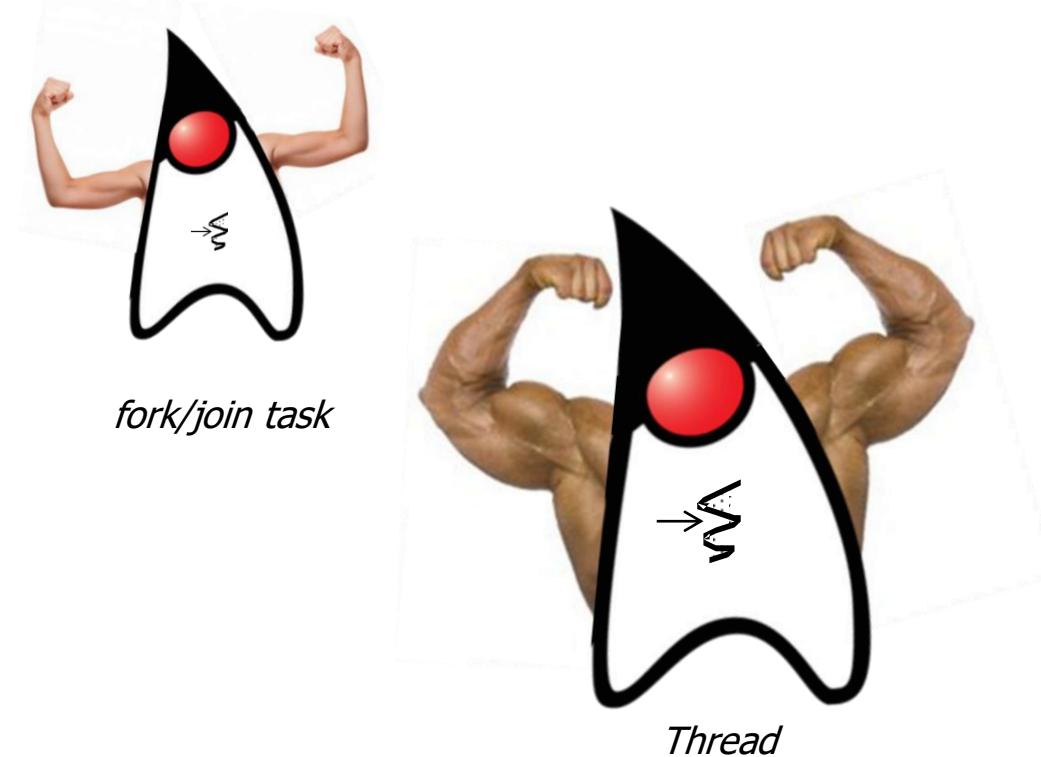
CountedCompleter, RecursiveAction, RecursiveTask

```
public abstract class ForkJoinTask<V>
extends Object
implements Future<V>, Serializable
```

Abstract base class for tasks that run within a ForkJoinPool. A ForkJoinTask is a thread-like entity that is much lighter weight than a normal thread. Huge numbers of tasks and subtasks may be hosted by a small number of actual threads in a ForkJoinPool, at the price of some usage limitations.

# Processing Chunks in Parallel via the Common ForkJoinPool

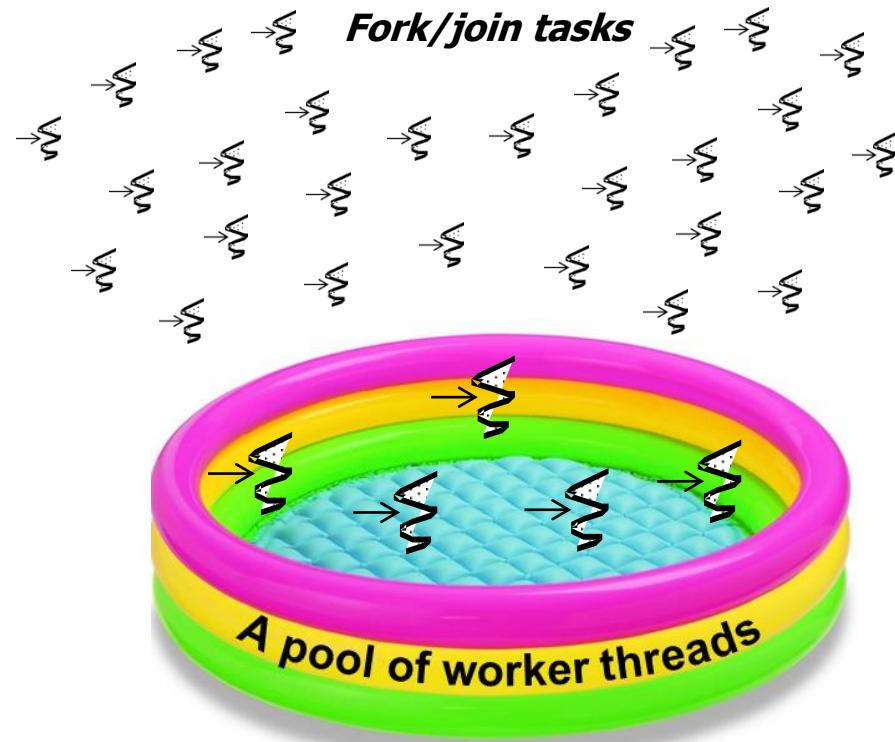
- A fork/join task is similar to—but lighter weight—than a Java thread.



e.g., it omits its own run-time stack, registers, thread-local storage, etc.

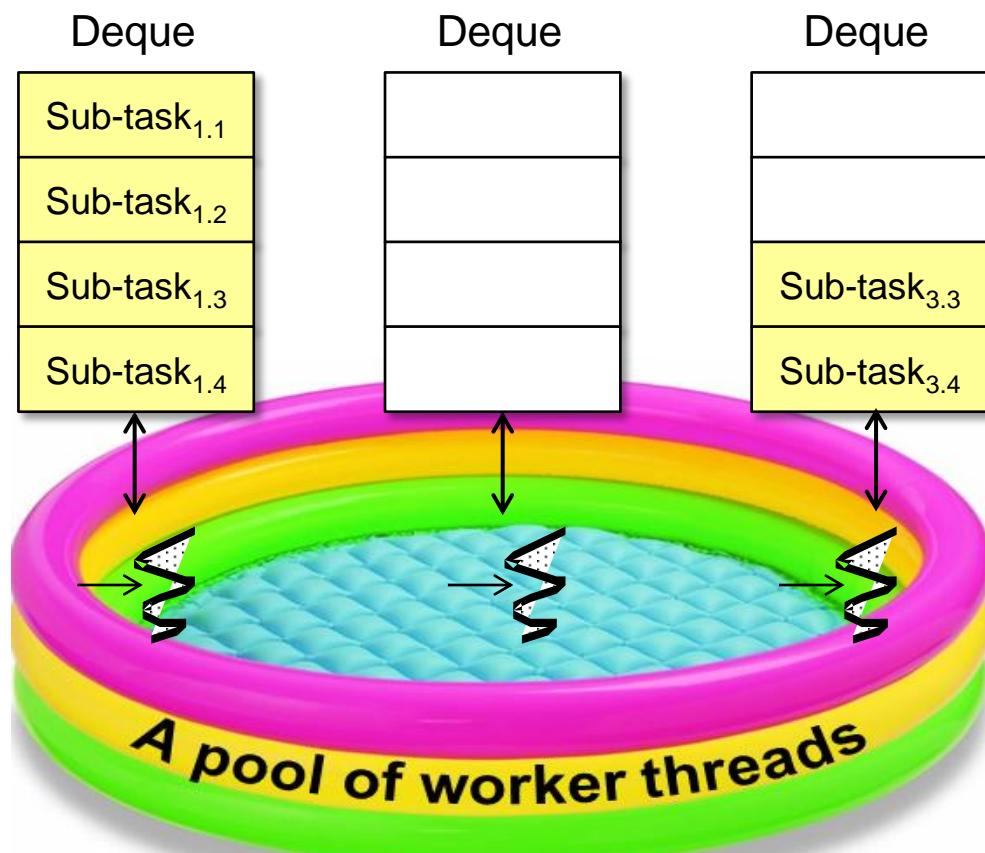
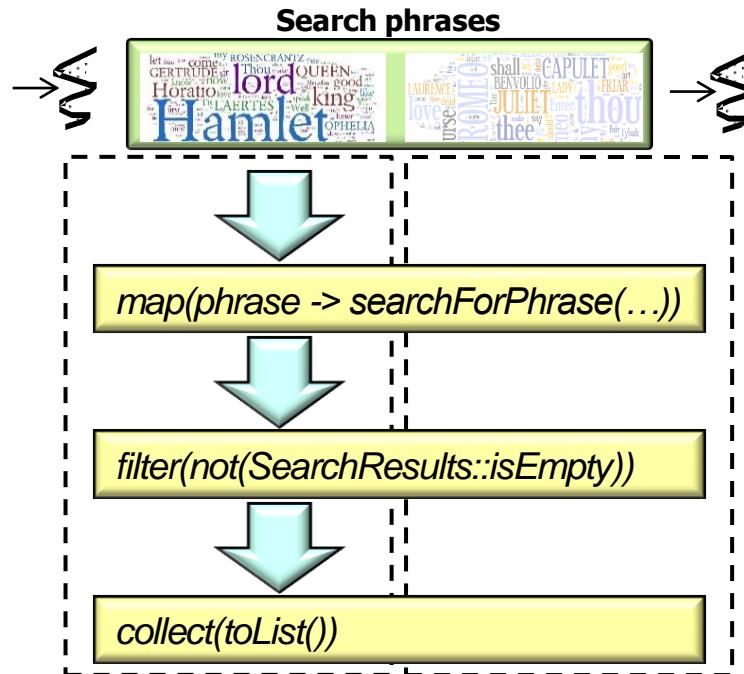
# Processing Chunks in Parallel via the Common ForkJoinPool

- A fork/join task is similar to—but lighter weight—than a Java thread.
  - A large # of fork/join tasks can thus run in a small # of Java worker threads in a ForkJoinPool.



# Processing Chunks in Parallel via the Common ForkJoinPool

- Parallel streams are a “user friendly” ForkJoinPool facade.



See [en.wikipedia.org/wiki/Facade\\_pattern](https://en.wikipedia.org/wiki/Facade_pattern)

# Processing Chunks in Parallel via the Common ForkJoinPool

- You can program directly to the ForkJoinPool API, though it can be somewhat painful!

```
List<List<SearchResults>>
listOfListOfSearchResults =
ForkJoinPool.commonPool()
.invoke(new
SearchWithfork/join task
(inputList,
mPhrasesToFind, . . .));
```

*I gave you the chance of programming Java streams.*



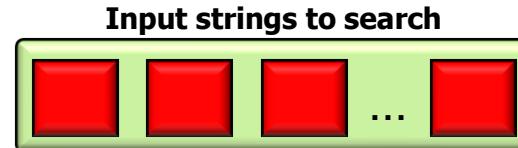
*But you have elected the way of pain!*

# Processing Chunks in Parallel via the Common ForkJoinPool

- You can program directly to the ForkJoinPool API, though it can be somewhat painful!

```
List<List<SearchResults>>
listOfListOfSearchResults =
    ForkJoinPool.commonPool()
        .invoke(new
            SearchWithfork/join task
                (inputList,
                    mPhrasesToFind, ...));
```

*Use the common ForkJoinPool  
to search input strings for  
phrases that match.*



See [livelessons/streamgangs/SearchWithForkJoin.java](#)

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.



```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...  
}
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

*Create a linked list of recursive task objects.*

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

*Create and fork tasks to search folders recursively.*

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

*Create and fork tasks to search documents.*

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

*Join all the tasks together and count the # of search matches.*

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- ForkJoinPool is best used for programs that don't match the parallel streams model.
  - e.g., this program counts the occurrence of a word in document folders.

*Return the final count.*

```
Long compute() {  
    long count = 0L;  
    List<RecursiveTask<Long>> forks =  
        new LinkedList<>();  
    for (Folder sub : mFolder.getSubs()) {  
        FolderSearchTask task = new  
            FolderSearchTask(sub, mWord);  
        forks.add(task); task.fork();  
    }  
    for (Doc doc : mFolder.getDocs()) {  
        DocSearchTask task =  
            new DocSearchTask(doc, mWord);  
        forks.add(task); task.fork();  
    }  
    for (RecursiveTask<Long> task : forks)  
        count += task.join();  
    return count; ...
```

# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.



See [dzone.com/articles/common-fork-join-pool-and-streams](https://dzone.com/articles/common-fork-join-pool-and-streams)

# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.



See [dzone.com/articles/common-fork-join-pool-and-streams](https://dzone.com/articles/common-fork-join-pool-and-streams)

# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.
  - This “global” vs. “local” resource management trade-off is common in computing and other domains.



# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.
  - There are few “knobs” to control this (or any) ForkJoinPool.



# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.
  - There are few “knobs” to control this (or any) ForkJoinPool.
    - This simplicity is intentional.



See [www.youtube.com/watch?v=sq0MX3fHkro](http://www.youtube.com/watch?v=sq0MX3fHkro)

# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.
  - There are few “knobs” to control this (or any) ForkJoinPool.
    - This simplicity is intentional.
    - Contrast ForkJoinPool with ThreadPoolExecutor.



# Processing Chunks in Parallel via the Common ForkJoinPool

- All parallel streams in a process share the common ForkJoinPool.
  - This helps optimize resource utilization by knowing what cores are used globally in a process.
  - There are few “knobs” to control this (or any) ForkJoinPool.
  - However, pool size *can* be configured.

`System.setProperty`

```
("java.util.concurrent"
  + ".ForkJoinPool.common"
  + ".parallelism",
"8");
```

*Set desired # of threads*

## Interface ForkJoinPool.ManagedBlocker

Enclosing class:

`ForkJoinPool`

```
public static interface ForkJoinPool.ManagedBlocker
```

Interface for extending managed parallelism for tasks running in `ForkJoinPools`.



See upcoming lesson on “Java Parallel Stream Internals: Configuration.”

Java Parallel Stream Internals:  
Parallel Processing via the Common ForkJoinPool

---

**The End**

# Java Parallel Stream Internals

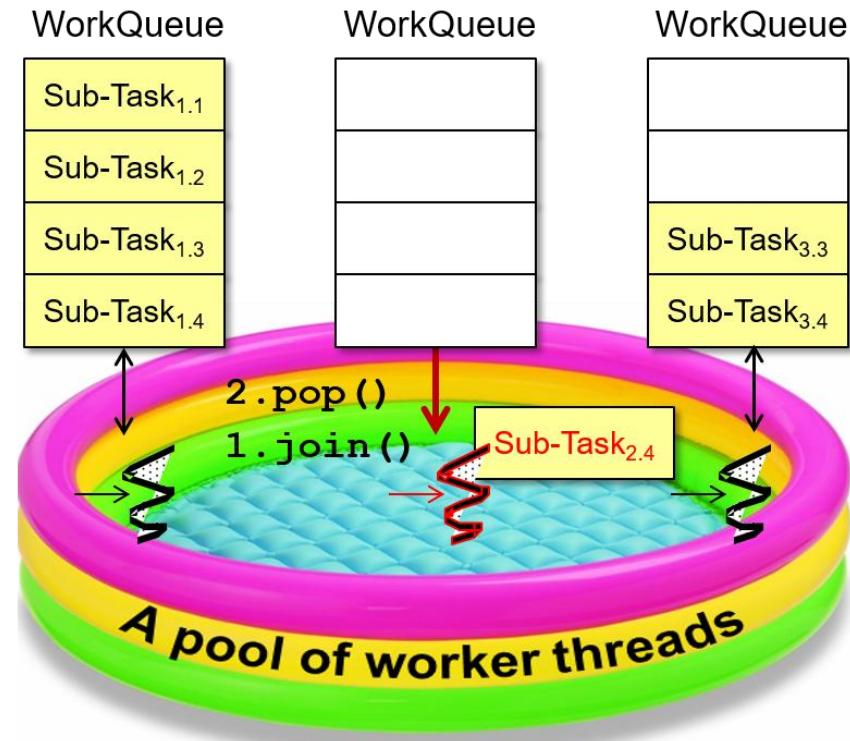
---

## Mapping Onto the Common ForkJoinPool

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.
    - Recognize how parallel streams are mapped onto the common ForkJoinPool framework.



---

# Mapping Parallel Streams Onto the Java ForkJoinPool

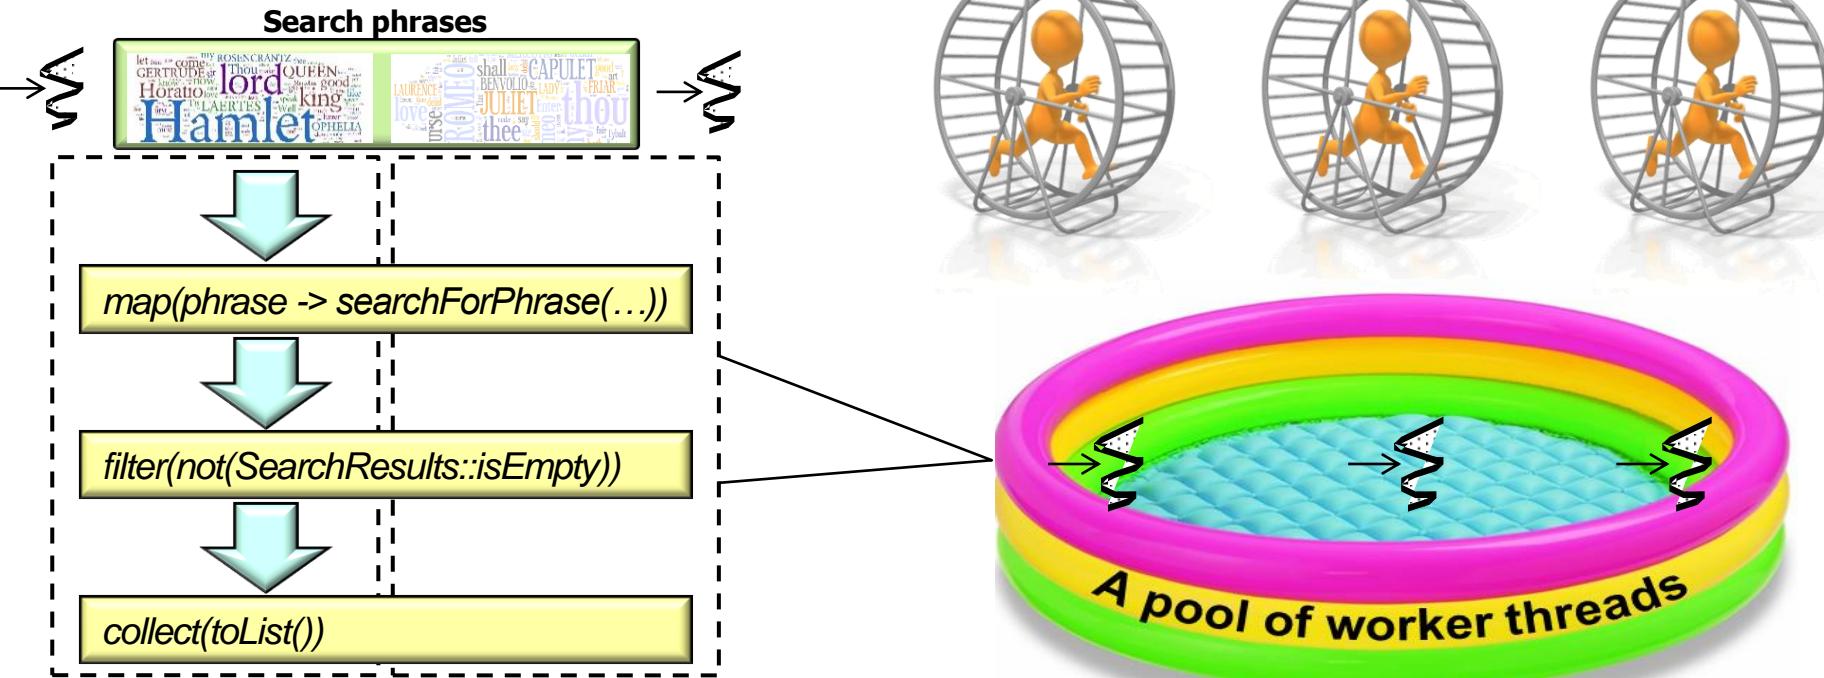
# Mapping Parallel Streams Onto the Common ForkJoinPool

- Each worker thread in the common ForkJoinPool runs a loop scanning for tasks to run.



# Mapping Parallel Streams Onto the Common ForkJoinPool

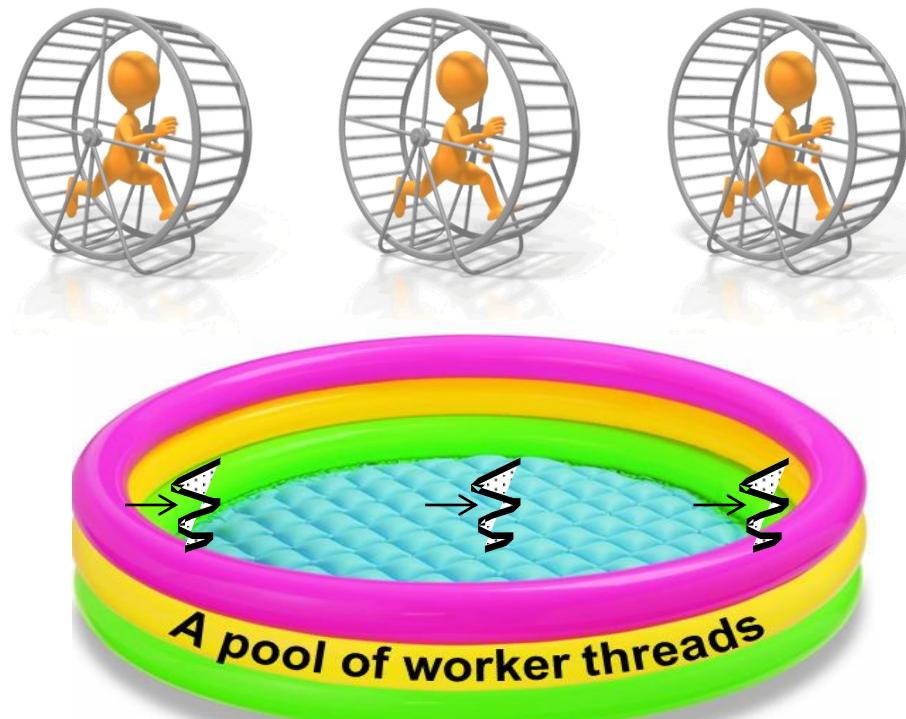
- Each worker thread in the common ForkJoinPool runs a loop scanning for tasks to run.



In this lesson, we just care about tasks associated with parallel streams.

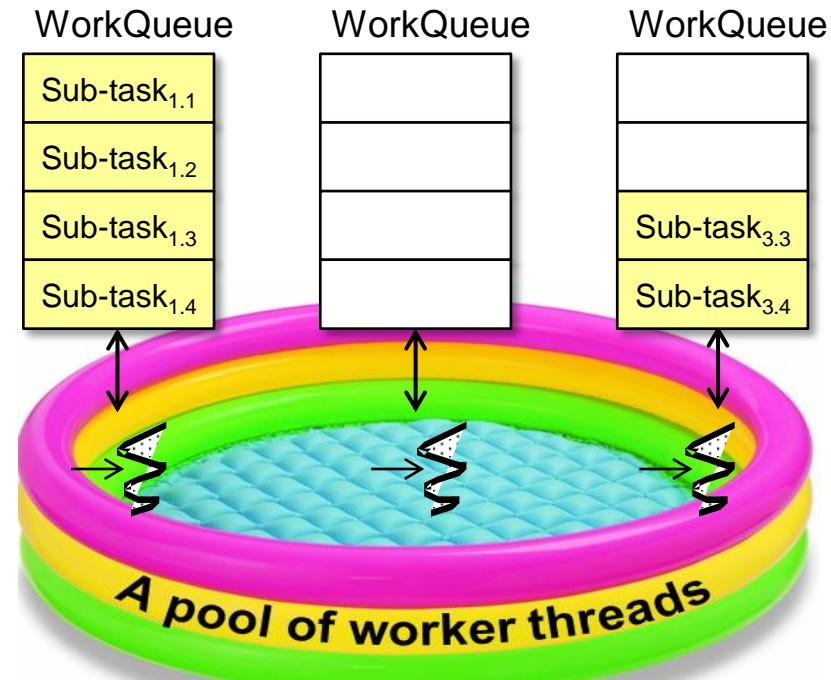
# Mapping Parallel Streams Onto the Common ForkJoinPool

- Each worker thread in the common ForkJoinPool runs a loop scanning for tasks to run.
  - The goal is to keep worker threads and cores as busy as possible!



# Mapping Parallel Streams Onto the Common ForkJoinPool

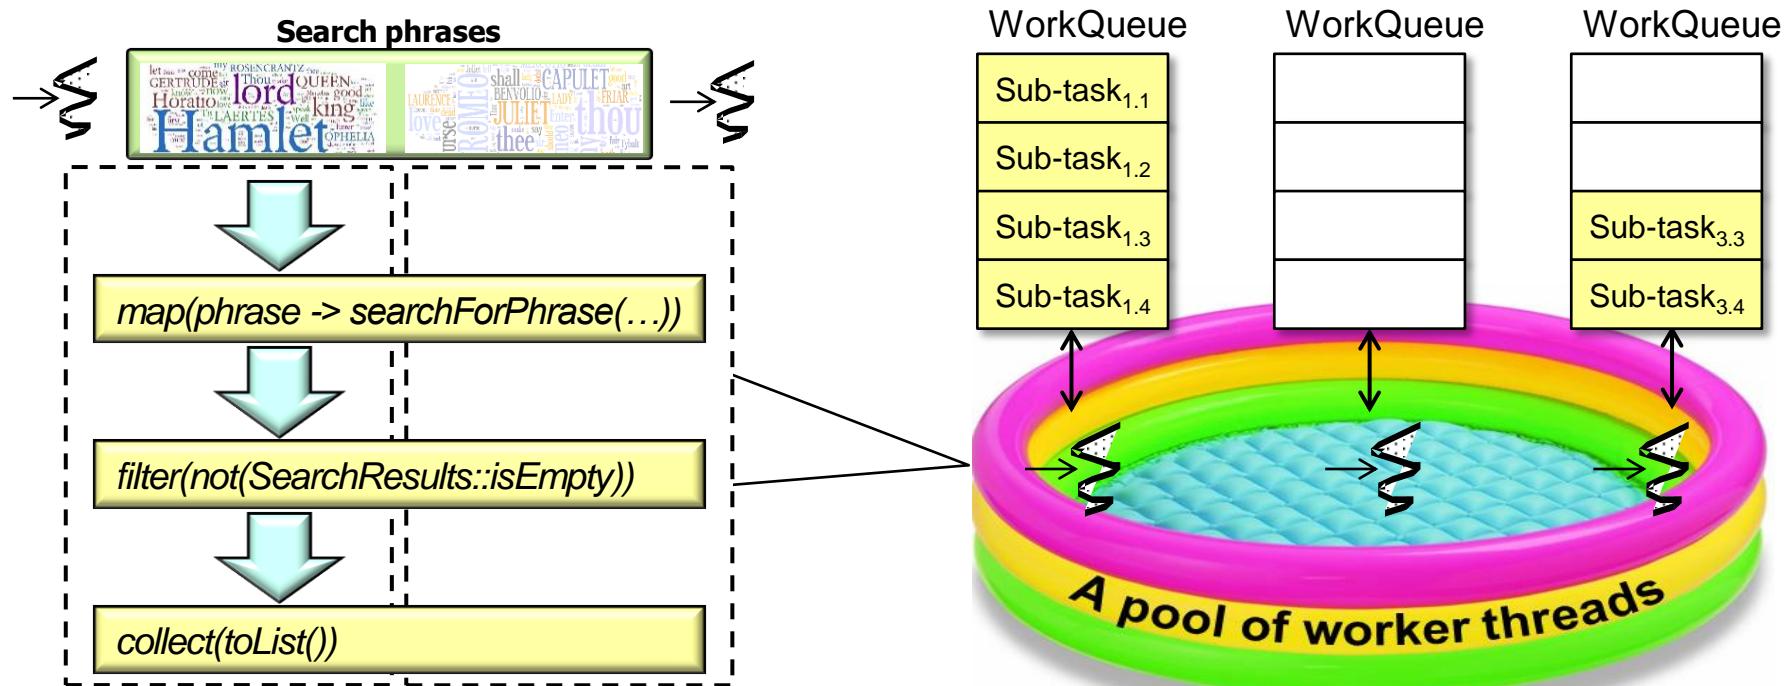
- Each worker thread in the common ForkJoinPool runs a loop scanning for tasks to run.
  - The goal is to keep worker threads and cores as busy as possible!
  - A worker thread has a “double-ended queue” (AKA “deque”) that serves as its main source of tasks.



See [en.wikipedia.org/wiki/Double-ended\\_queue](https://en.wikipedia.org/wiki/Double-ended_queue)

# Mapping Parallel Streams Onto the Common ForkJoinPool

- The parallel streams framework automatically creates fork/join tasks that are run by worker threads in the common ForkJoinPool.



# Mapping Parallel Streams Onto the Common ForkJoinPool

- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

```
abstract class AbstractTask ... {  
    public void compute() {  
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            K taskToFork;  
            if (forkRight)  
                { forkRight = false; ... taskToFork = ...makeChild(rs); }  
            else  
                { forkRight = true; ... taskToFork = ...makeChild(ls); }  
            taskToFork.fork();  
        }  
    } ...
```

*Manages splitting logic, tracking of child tasks, and intermediate results.*

See [openjdk/8-b132/java/util/stream/AbstractTask.java](https://openjdk.java.net/jeps/132/java.util.stream/AbstractTask.java)

# Mapping Parallel Streams Onto the Common ForkJoinPool

- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

```
abstract class AbstractTask ... {  
    public void compute() {
```

*This decides whether to split a task further and/or compute it directly.*

```
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            K taskToFork;  
            if (forkRight)  
                { forkRight = false; ... taskToFork = ...makeChild(rs); }  
            else  
                { forkRight = true; ... taskToFork = ...makeChild(ls); }  
            taskToFork.fork();  
        }  
    } ...
```

# Mapping Parallel Streams Onto the Common ForkJoinPool

- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

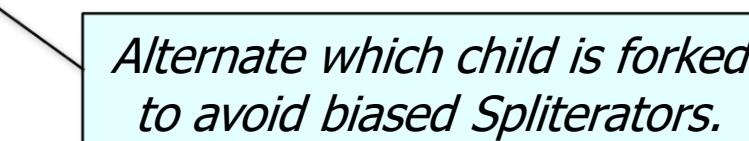
```
abstract class AbstractTask ... { ...  
    public void compute() {  
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            K taskToFork;  
            if (forkRight)  
                { forkRight = false; ... taskToFork = ...makeChild(rs); }  
            else  
                { forkRight = true; ... taskToFork = ...makeChild(ls); }  
            taskToFork.fork();  
        }  
    } ...
```

*Try to partition the input source until trySplit() returns null.*

# Mapping Parallel Streams Onto the Common ForkJoinPool

- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

```
abstract class AbstractTask ... { ...  
    public void compute() {  
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            K taskToFork;  
            if (forkRight)  
                { forkRight = false; ... taskToFork = ...makeChild(rs); }  
            else  
                { forkRight = true; ... taskToFork = ...makeChild(ls); }  
            taskToFork.fork();  
        }  
    } ...
```



*Alternate which child is forked  
to avoid biased Spliterators.*

# Mapping Parallel Streams Onto the Common ForkJoinPool

- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

```
abstract class AbstractTask ... { ...  
    public void compute() {  
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            K taskToFork;  
            if (forkRight)  
                { forkRight = false; ... taskToFork = ...makeChild(rs); }  
            else  
                { forkRight = true; ... taskToFork = ...makeChild(ls); }  
            taskToFork.fork();  
        }  
    } ...
```



*Fork a new sub-task and continue processing the other in the loop.*

# Mapping Parallel Streams Onto the Common ForkJoinPool

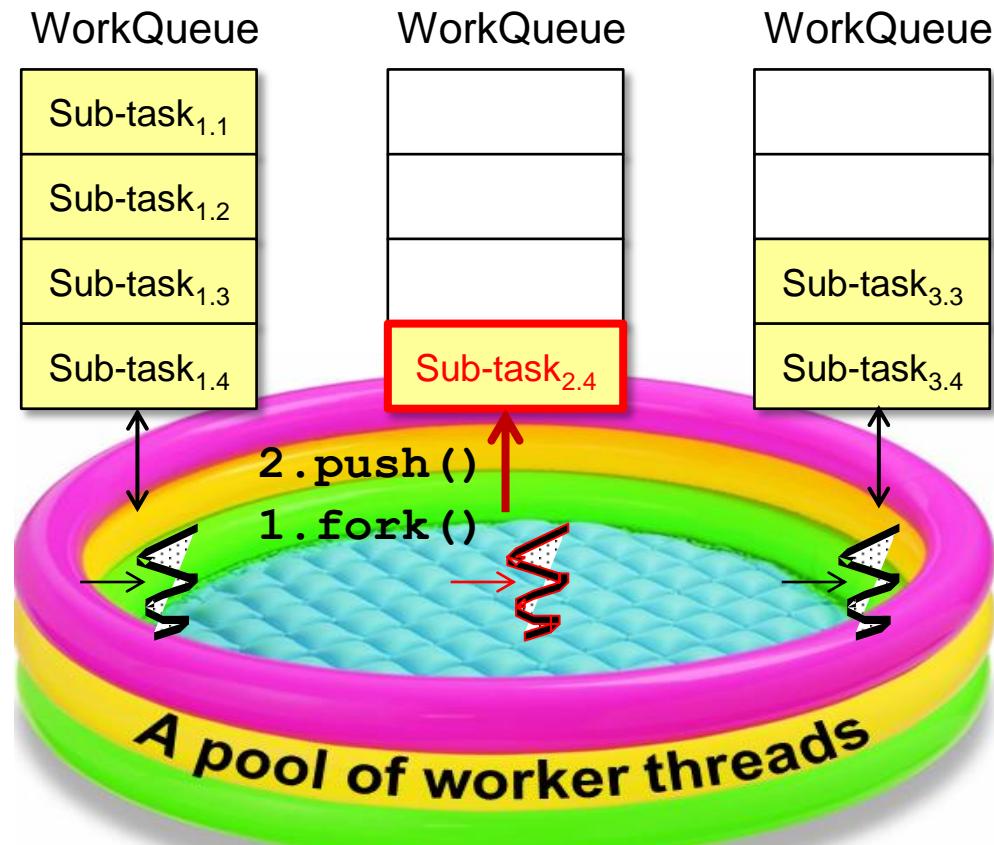
- The AbstractTask super class is used by most fork/join tasks to implement the parallel stream framework.

```
abstract class AbstractTask ... { ...  
    public void compute() {  
        Spliterator<P_IN> rs = spliterator, ls;  
        boolean forkRight = false; ...  
        while(... (ls = rs.trySplit()) != null) {  
            ...  
        }  
        task.setLocalResult(task.doLeaf());  
    } ...
```

*After trySplit() returns null this method typically calls forEachRemaining() to process elements sequentially.*

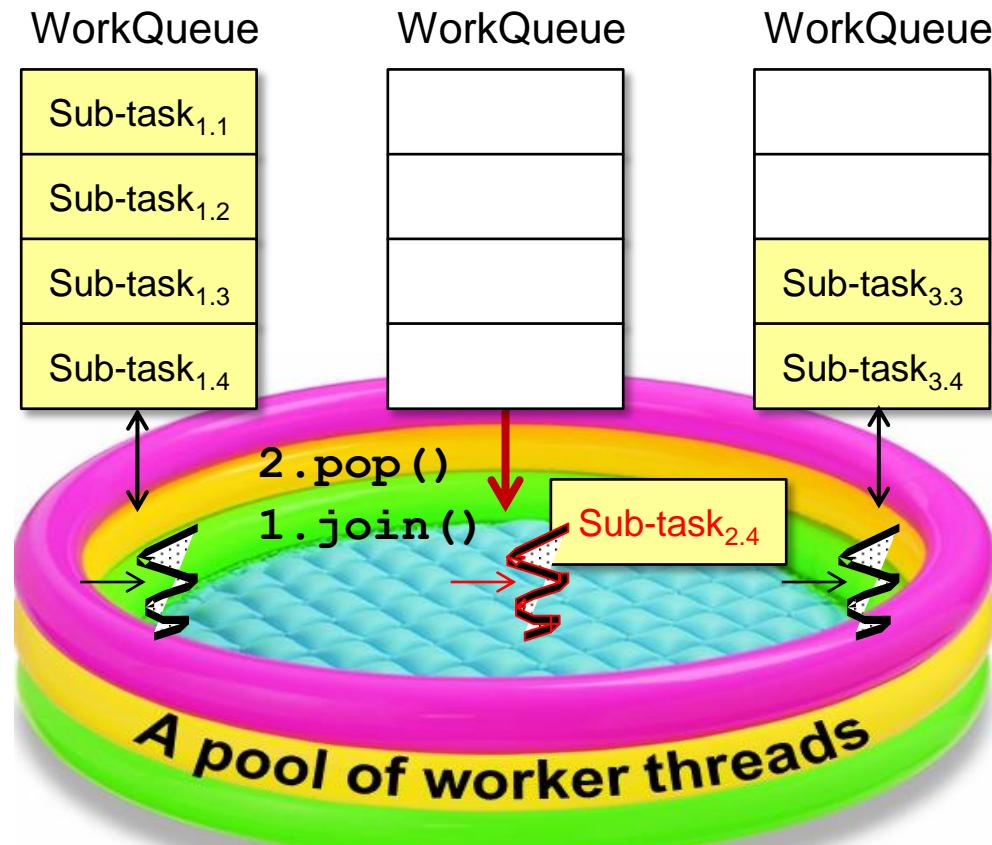
# Mapping Parallel Streams Onto the Common ForkJoinPool

- After the `AbstractTask.compute()` method calls `fork()` on a task, this task is pushed onto the head of its worker thread's deque.



# Mapping Parallel Streams Onto the Common ForkJoinPool

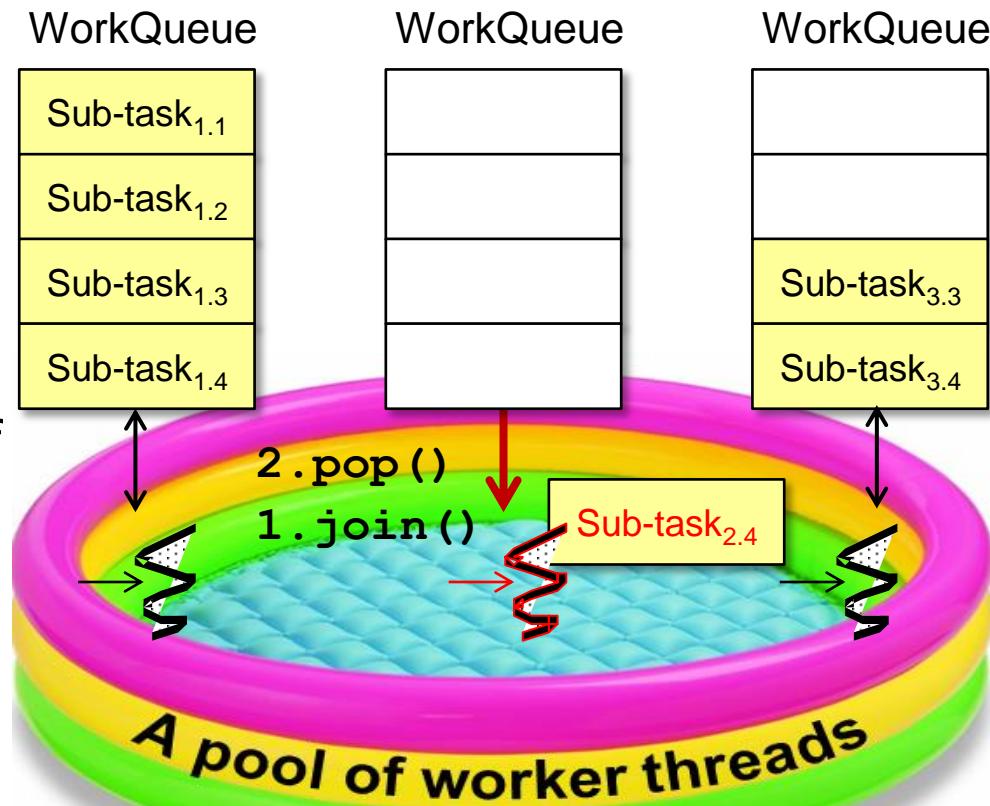
- After the `AbstractTask.compute()` method calls `fork()` on a task, this task is pushed onto the head of its worker thread's deque.
- Each worker thread processes its deque in LIFO order.



See [en.wikipedia.org/wiki/Stack\\_\(abstract\\_data\\_type\)](https://en.wikipedia.org/wiki/Stack_(abstract_data_type))

# Mapping Parallel Streams Onto the Common ForkJoinPool

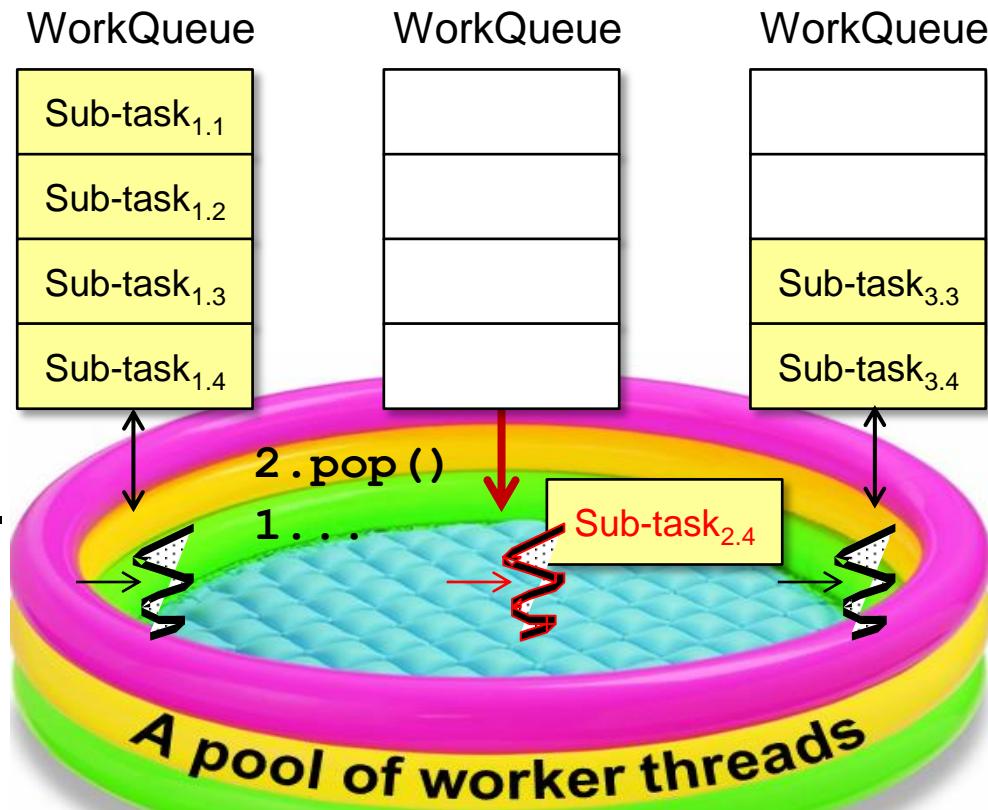
- After the `AbstractTask.compute()` method calls `fork()` on a task, this task is pushed onto the head of its worker thread's deque.
- Each worker thread processes its deque in LIFO order.
- A task popped from the head of a deque is run to completion.



See [en.wikipedia.org/wiki/Run\\_to\\_completion\\_scheduling](https://en.wikipedia.org/wiki/Run_to_completion_scheduling)

# Mapping Parallel Streams Onto the Common ForkJoinPool

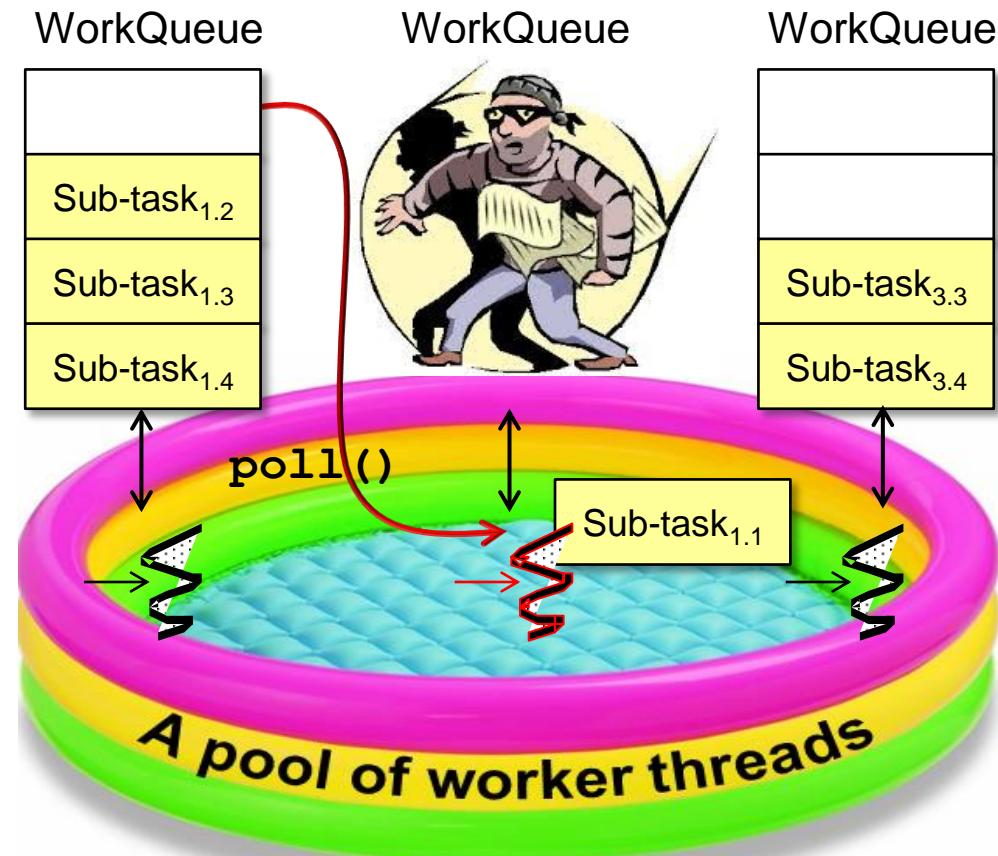
- After the `AbstractTask.compute()` method calls `fork()` on a task, this task is pushed onto the head of its worker thread's deque.
- Each worker thread processes its deque in LIFO order.
- LIFO order improves locality of reference and cache performance.



See [en.wikipedia.org/wiki/Locality\\_of\\_reference](https://en.wikipedia.org/wiki/Locality_of_reference)

# Mapping Parallel Streams Onto the Common ForkJoinPool

- To maximize core utilization, idle worker threads “steal” work from the tail of busy threads’ dequeues.



See upcoming lessons on “*The Java Fork-Join Framework*.”

Java Parallel Stream Internals:  
Mapping Onto the Common ForkJoinPool

---

**The End**

# Java Parallel Stream Internals

---

## Configuring the Common ForkJoinPool

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

---

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.
  - Configure the Java parallel stream common ForkJoinPool.

```
int desiredThreads = "8";  
System.setProperty  
( "java.util.concurrent."  
+ "ForkJoinPool.common."  
+ "parallelism" ,  
desiredThreads );
```



---

# Configuring the Parallel Stream Common ForkJoinPool

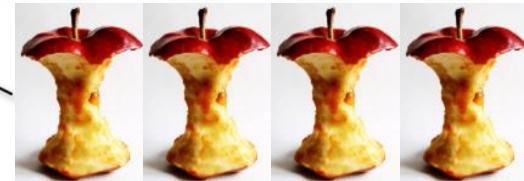
# Configuring the Parallel Stream Common ForkJoinPool

- By default the common ForkJoinPool has one less thread than the # of cores.

```
System.out.println
```

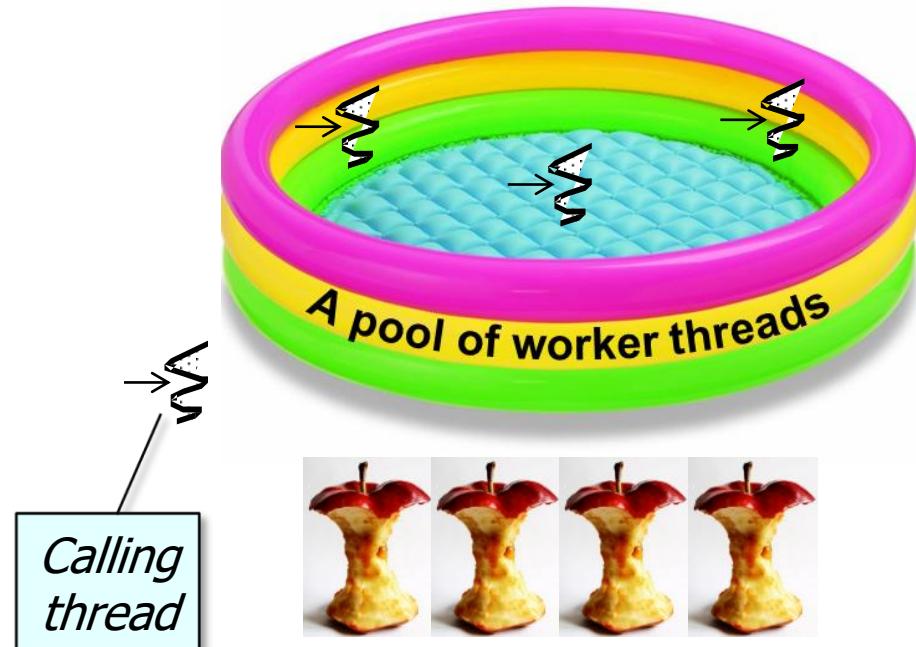
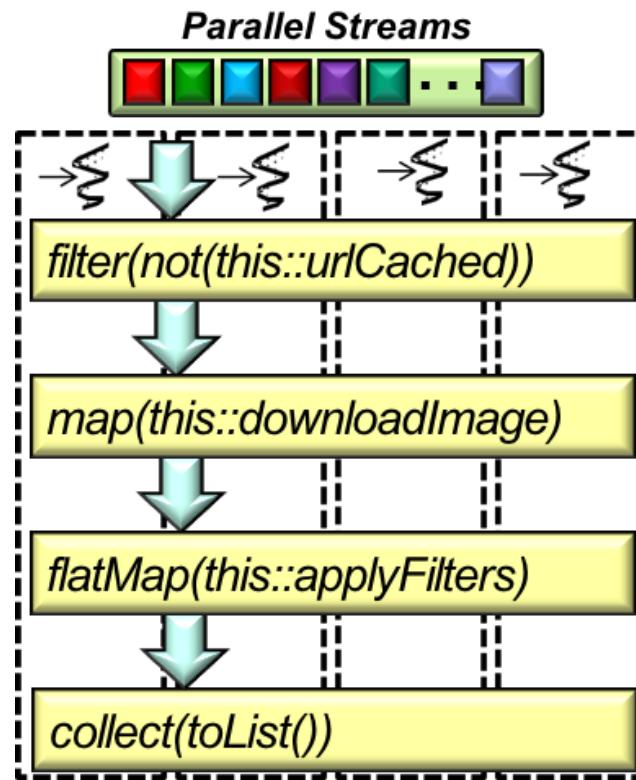
```
  ("The parallelism in the"
   + "common ForkJoinPool is "
   + ForkJoinPool
     .getCommonPoolParallelism());
```

*e.g., it returns three on a quad-core processor.*



# Configuring the Parallel Stream Common ForkJoinPool

- By default the common ForkJoinPool has one less thread than the # of cores.

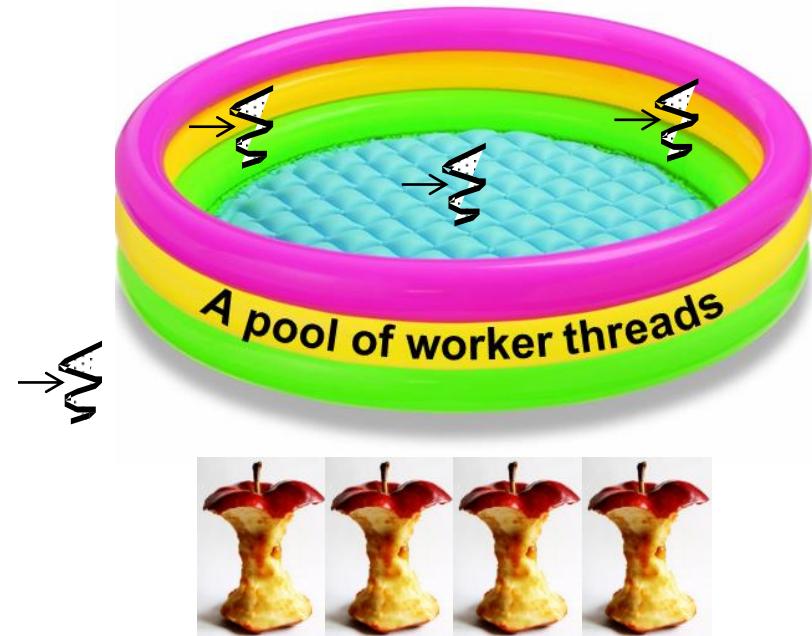
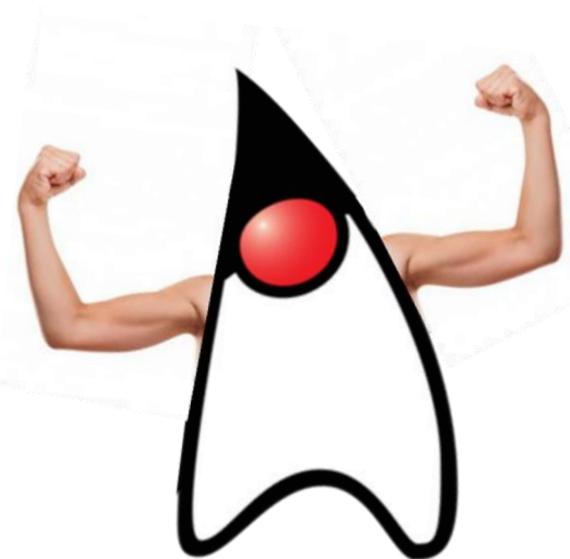


A parallel stream can use all cores since it uses the invoking thread, e.g., main thread.

# Configuring the Parallel Stream Common ForkJoinPool

---

- However, the default # of ForkJoinPool threads may be inadequate.



# Configuring the Parallel Stream Common ForkJoinPool

- However, the default # of ForkJoinPool threads may be inadequate, e.g.
  - Consider a parallel image downloading and processing app

 The image cannot be displayed. Your computer may not have enough memory to open the image, or the image may have been corrupted. Restart your computer, and then open the file again. If the red x still appears, you may have to delete the image and then insert it again.



# Configuring the Parallel Stream Common ForkJoinPool

- However, the default # of ForkJoinPool threads may be inadequate, e.g.
  - Consider a parallel image downloading and processing app

 The image cannot be displayed. Your computer may not have enough memory to open the image, or the image may have been corrupted. Restart your computer, and then open the file again. If the red x still appears, you may have to delete the image and then insert it again.

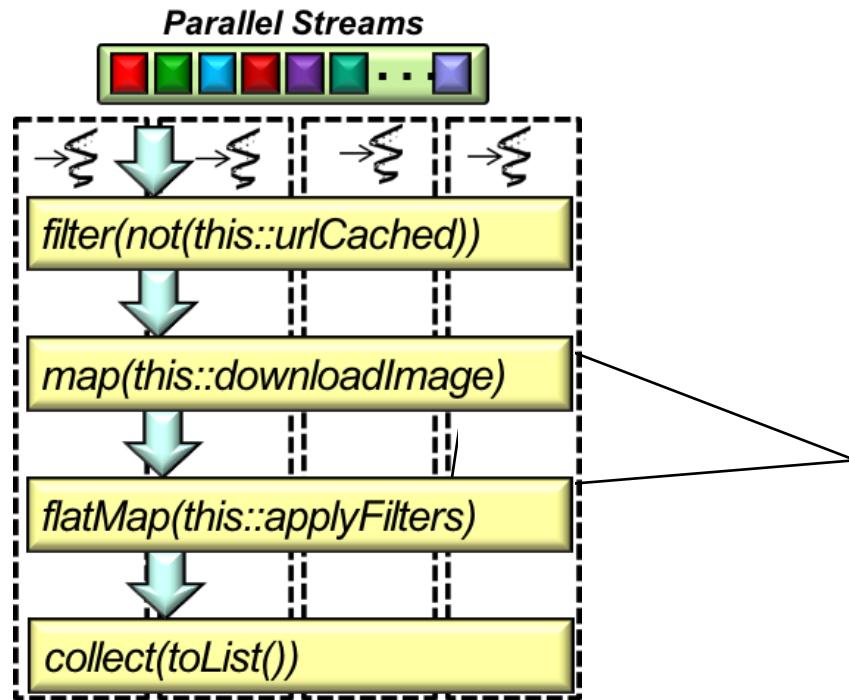


*Problems may occur when trying to download more images than # of cores.*

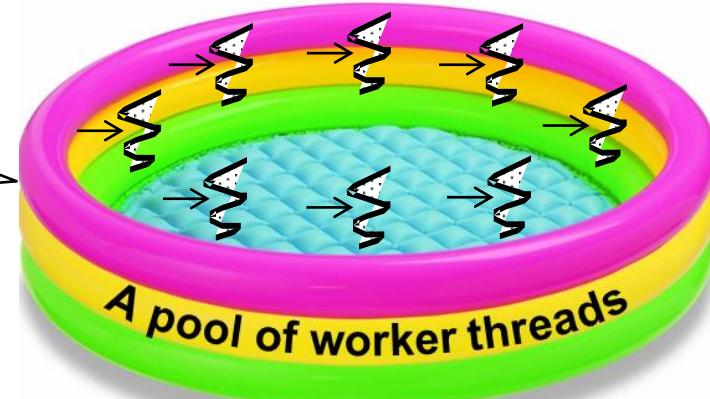
These problems may range from underutilization of processor cores to deadlock.

# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.

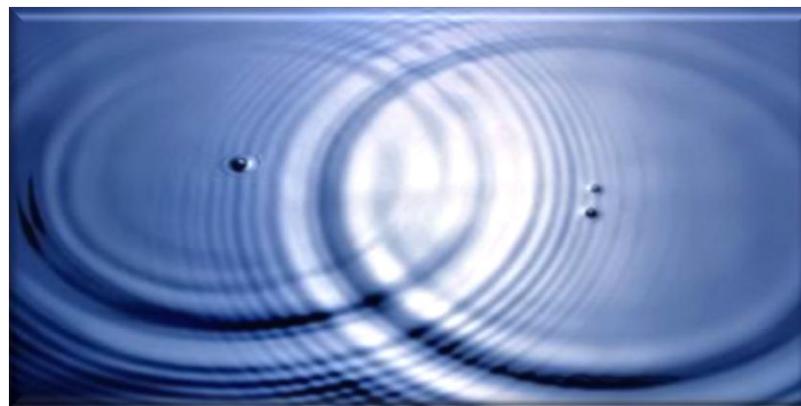


```
int desiredThreads = "8";  
System.setProperty  
  ("java.util.concurrent."  
  + "ForkJoinPool.common."  
  + "parallelism",  
desiredThreads);
```



# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.
  - Setting this property affects all parallel streams in a process.



```
int desiredThreads = "8";  
System.setProperty  
  ("java.util.concurrent."  
  + "ForkJoinPool.common."  
  + "parallelism",  
desiredThreads);
```



It's hard to estimate the total # of threads to set in the common ForkJoinPool.

# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.
  - Setting this property affects all parallel streams in a process.
  - This property can be changed only before the common fork-join pool is initialized.
  - It's initialized "on-demand" the first time it's used.

```
int desiredThreads = "8";  
System.setProperty  
  ("java.util.concurrent."  
  + "ForkJoinPool.common."  
  + "parallelism",  
desiredThreads);
```



# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.
  - Setting this property affects all parallel streams in a process.
  - The ManagedBlocker interface can also be used to add worker threads to common ForkJoinPool temporarily.



```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>(  
        supplier);  
...  
ForkJoinPool.managedBlock(mb);  
...  
return mb.getResult();
```



# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.
  - Setting this property affects all parallel streams in a process.
  - The ManagedBlocker interface can also be used to add worker threads to common ForkJoinPool temporarily.
  - This is useful for behaviors that block on I/O and/or synchronizers.

```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>(  
        supplier);  
    ...  
    ForkJoinPool.managedBlock(mb);  
    ...  
    return mb.getResult();
```



# Configuring the Parallel Stream Common ForkJoinPool

- The common ForkJoinPool size can be controlled programmatically.
  - Setting this property affects all parallel streams in a process.
  - The ManagedBlocker interface can also be used to add worker threads to common ForkJoinPool temporarily.
    - This is useful for behaviors that block on I/O and/or synchronizers.
    - This interface can only be used with the common ForkJoinPool.

```
SupplierManagedBlocker<T> mb =  
    new SupplierManagedBlocker<>(  
        supplier);  
    ...  
    ForkJoinPool.managedBlock(mb);  
    ...  
    return mb.getResult();
```



See lessons on “*The Java ForkJoinPool: the ManagedBlocker Interface.*”

Java Parallel Stream Internals:  
Configuring the Common ForkJoinPool

---

**The End**

# Java Parallel Stream Internals

---

Demo'ing How to Configure the Common  
Fork-Join Pool

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

---

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.
  - Configure the Java parallel stream common ForkJoinPool.
    - Know the performance impact of configuring the common fork-join pool size.

Entering the test program with 12 cores  
`ex20: testDefaultDownloadBehavior()` downloaded and stored 42 images using 12 threads in the pool

`ex20: testAdaptiveMBDownloadBehavior()` downloaded and stored 42 images using 43 threads in the pool

`ex20: testAdaptiveBTDownloadBehavior()` downloaded and stored 42 images using 43 threads in the pool

Printing 3 results from fastest to slowest  
`testAdaptiveBTDownloadBehavior()` executed in 3598 msec

`testAdaptiveMBDownloadBehavior()` executed in 3910 msec

`testDefaultDownloadBehavior()` executed in 4104 msec

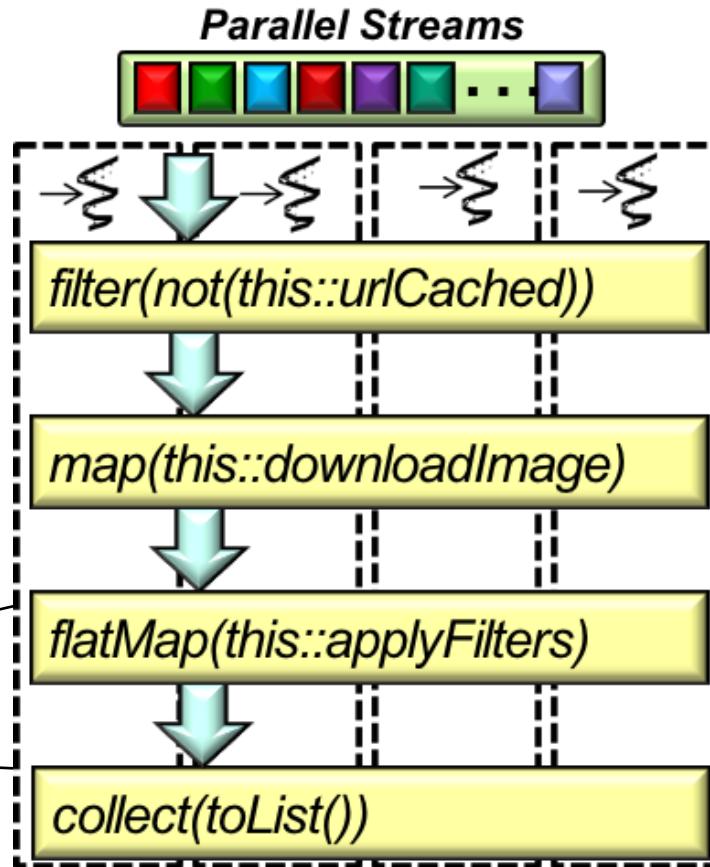
Leaving the test program

---

# Demo'ing Impact of Configuring Common Fork-Join Pool

# Demo'ing Impact of Configuring Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically.



See prior lesson on "Java Parallel Stream Internals: Configuring the Common Fork-Join Pool"

# Demo'ing Impact of Configuring Common Fork-Join Pool

- The common fork-join pool size can be controlled programmatically.
  - This demo applies ManagedBlocker to add new worker threads to the common fork-join pool



```
File downloadAndStoreImageMB  
      (URL url) {  
    final Image[] image =  
        new Image[1];  
    ...  
    ForkJoinPool  
        .managedBlock(new ForkJoinPool  
            .ManagedBlocker() {  
        public boolean block() {  
            image[0] =  
                downloadImage(url);  
            return true;  
        } ... });  
    return image[0].store(); ...  
}
```

# Demo'ing Impact of Configuring Common Fork-Join Pool

---

- This program shows the performance difference of using ManagedBlocker versus not using ManagedBlocker for an I/O-intensive app.

```
void testDownloadBehavior(Function<URL, File>
                           downloadAndStoreImage,
                           String testName) {
    ...
    List<File> imageFiles = Options.instance()
        .getUrlList()
        .parallelStream()

        .map(downloadAndStoreImage)

        .collect(Collectors.toList());
    printStats(testName, imageFiles.size()); ...
}
```

# Demo'ing Impact of Configuring Common Fork-Join Pool

- This program shows the performance difference of using ManagedBlocker versus not using ManagedBlocker for an I/O-intensive app.

```
void testDownloadBehavior(Function<URL, File>
                           downloadAndStoreImage,
                           String testName) {
    ...
    List<File> imageFiles = Options.instance()
        .getUrlList()
        .parallelStream()
        .map(downloadAndStoreImage)
        .collect(Collectors.toList());
    printStats(testName, imageFiles.size()); ...
}
```

*This function param is used to pass different strategies for downloading & storing images from remote websites*

# Demo'ing Impact of Configuring Common Fork-Join Pool

- Results show increasing worker threads in the pool improves performance.

Entering the test program with 12 cores

```
ex20: testDefaultDownloadBehavior() downloaded and stored 42 images
      using 12 threads in the pool
ex20: testAdaptiveMBDownloadBehavior() downloaded and stored 42 images
      using 43 threads in the pool
ex20: testAdaptiveBTDownloadBehavior() downloaded and stored 42 images
      using 43 threads in the pool
```

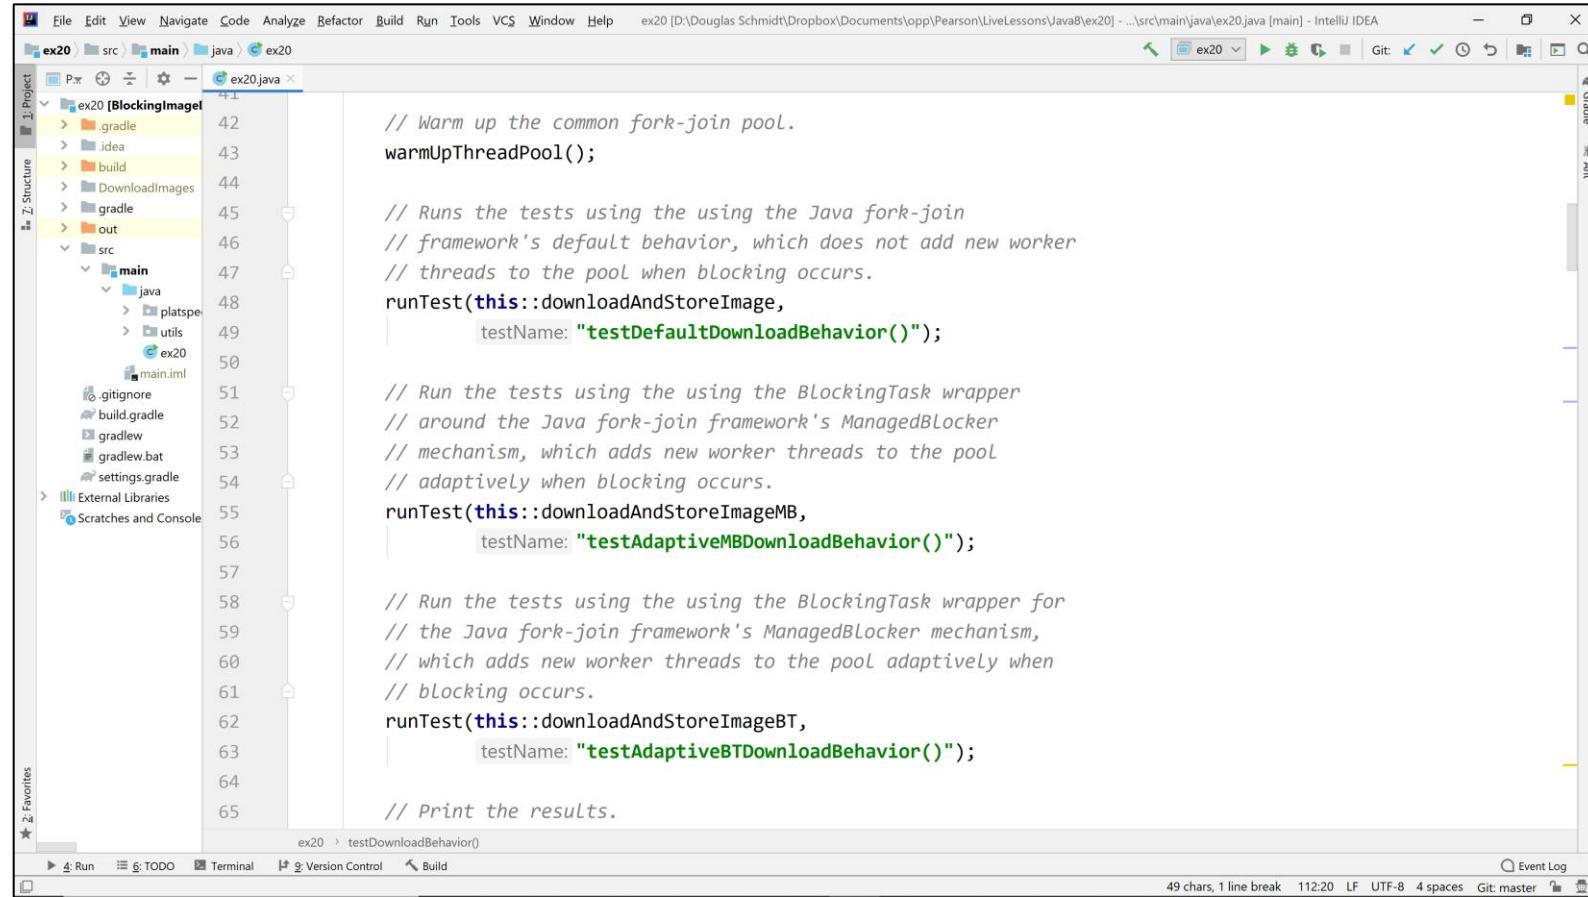
Printing 3 results from fastest to slowest

```
testAdaptiveBTDownloadBehavior() executed in 3598 msecs
testAdaptiveMBDownloadBehavior() executed in 3910 msecs
testDefaultDownloadBehavior() executed in 4104 msecs
```

Leaving the test program

See upcoming lessons on "*The Java Fork-Join Pool: the ManagedBlocker Interface*"

# Demo'ing Impact of Configuring Common Fork-Join Pool



The screenshot shows the IntelliJ IDEA interface with a Java file named `ex20.java` open. The code demonstrates different ways to handle blocking operations using the Java Fork-Join framework's `BlockingImage` class. The code includes comments explaining the behavior of each approach:

```
// Warm up the common fork-join pool.  
warmUpThreadPool();  
  
// Runs the tests using the Java fork-join  
// framework's default behavior, which does not add new worker  
// threads to the pool when blocking occurs.  
runTest(this::downloadAndStoreImage,  
        testName: "testDefaultDownloadBehavior()");  
  
// Run the tests using the BlockingTask wrapper  
// around the Java fork-join framework's ManagedBlocker  
// mechanism, which adds new worker threads to the pool  
// adaptively when blocking occurs.  
runTest(this::downloadAndStoreImageMB,  
        testName: "testAdaptiveMBDownloadBehavior()");  
  
// Run the tests using the BlockingTask wrapper for  
// the Java fork-join framework's ManagedBlocker mechanism,  
// which adds new worker threads to the pool adaptively when  
// blocking occurs.  
runTest(this::downloadAndStoreImageBT,  
        testName: "testAdaptiveBTDownloadBehavior()");  
  
// Print the results.  
ex20 > testDownloadBehavior()
```

The IntelliJ IDEA interface includes a Project tool window on the left, a code editor with line numbers, and various toolbars and status bars at the bottom.

See [github.com/douglasraigschmidt/LiveLessons/tree/master/Java8/ex20](https://github.com/douglasraigschmidt/LiveLessons/tree/master/Java8/ex20)

# Demo'ing Impact of Configuring Common Fork-Join Pool

---

**[Source code analysis  
video goes here!]**

Java Parallel Stream Internals:  
Demo'ing How to Configure the Common Fork-Join Pool

---

**The End**

# Java Parallel Stream Internals

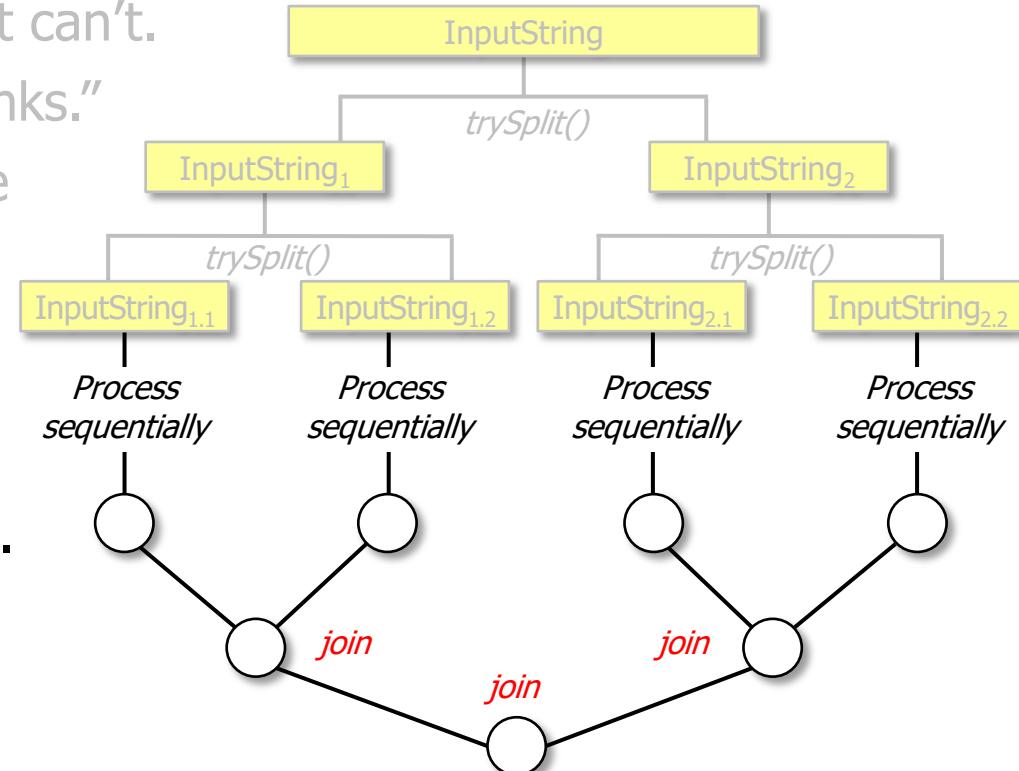
---

## Combining Results (Part I)

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.
  - Configure the Java parallel stream common ForkJoinPool.
  - Perform a reduction to combine partial results into a single result.

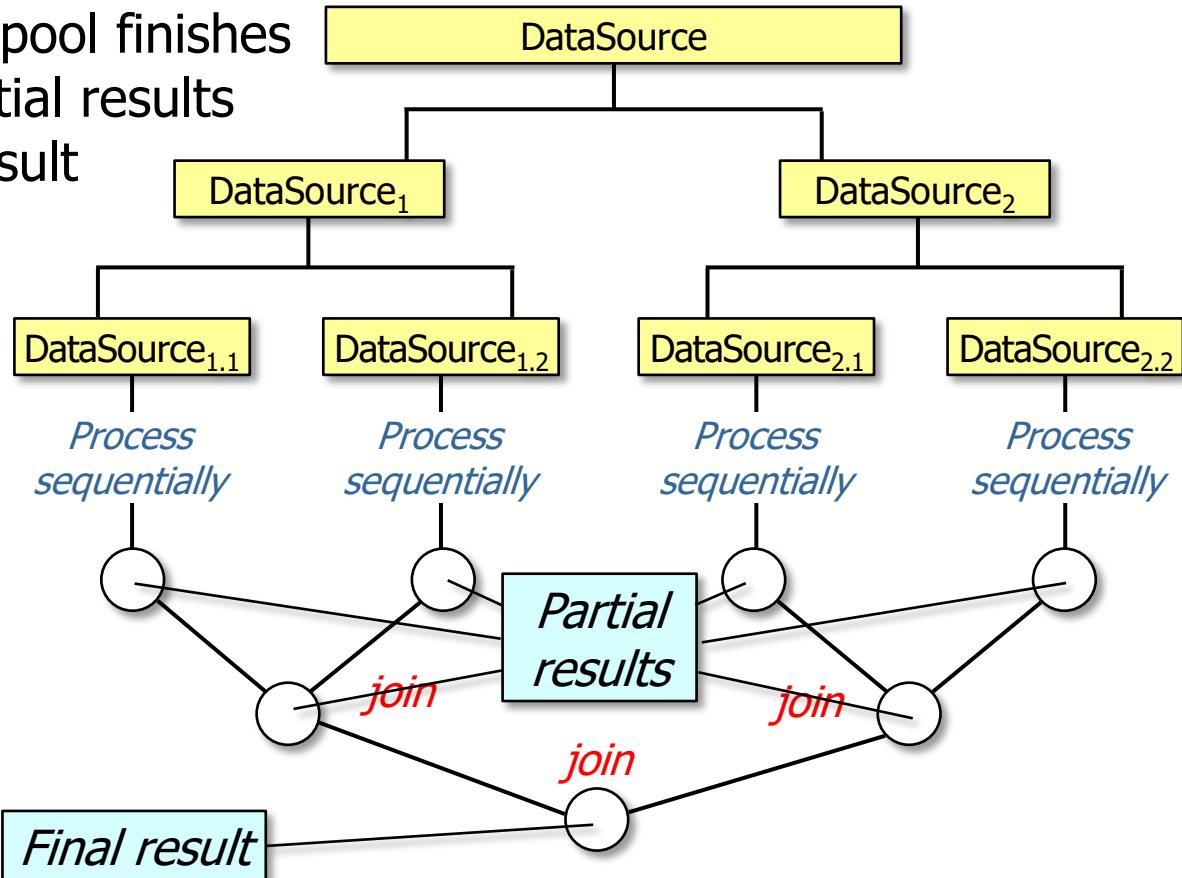


---

# Combining Results in a Parallel Stream

# Combining Results in a Parallel Stream

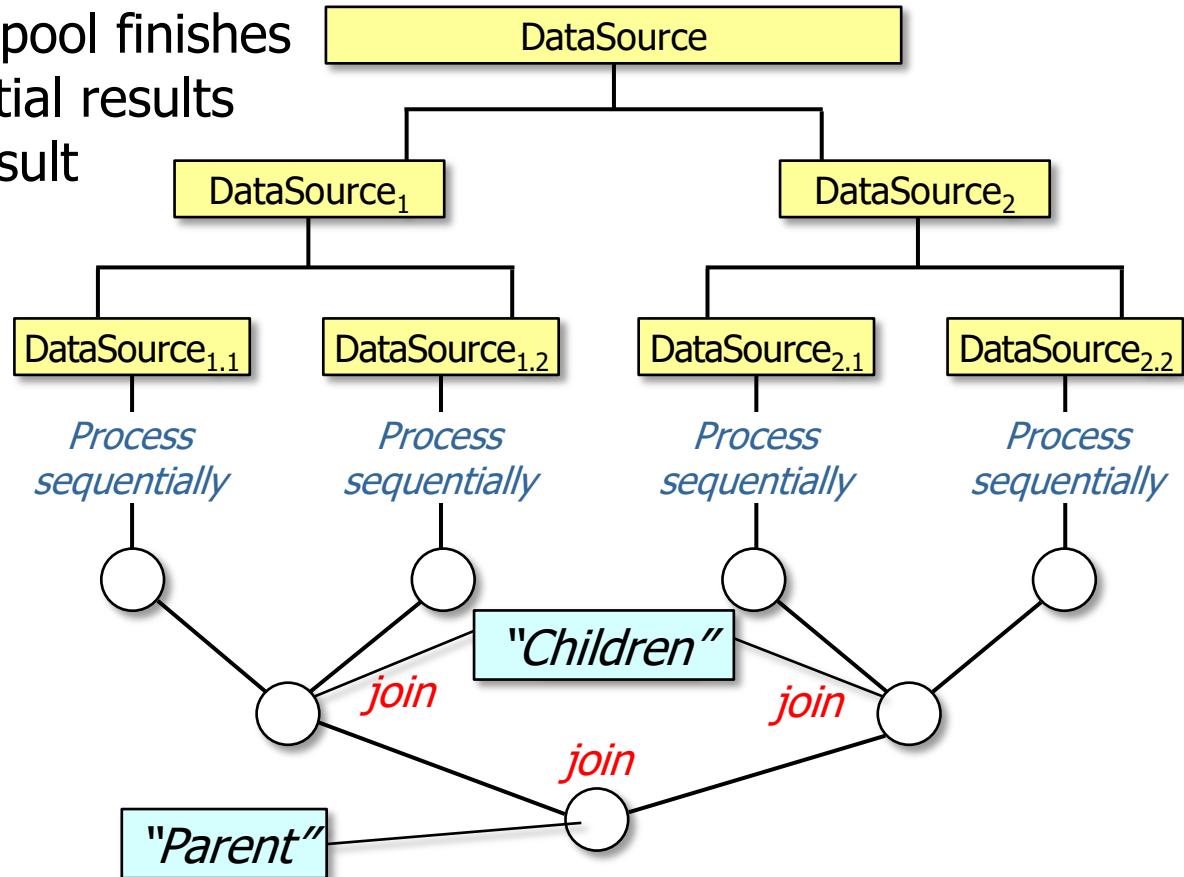
- After the common fork-join pool finishes processing chunks their partial results are combined into a final result



This discussion assumes a non-concurrent collector (other discussions follow)

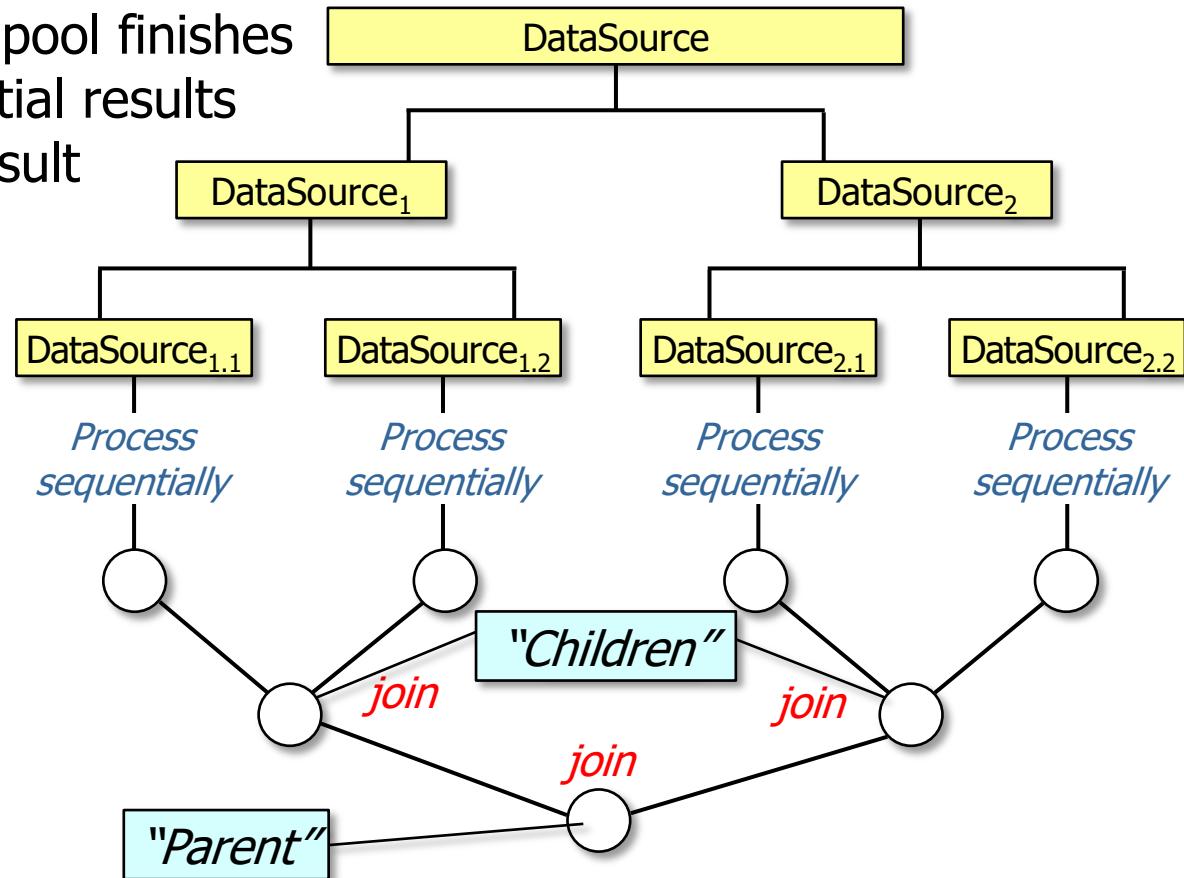
# Combining Results in a Parallel Stream

- After the common fork-join pool finishes processing chunks their partial results are combined into a final result
  - `join()` occurs in a single thread at each level
    - i.e., the “parent”



# Combining Results in a Parallel Stream

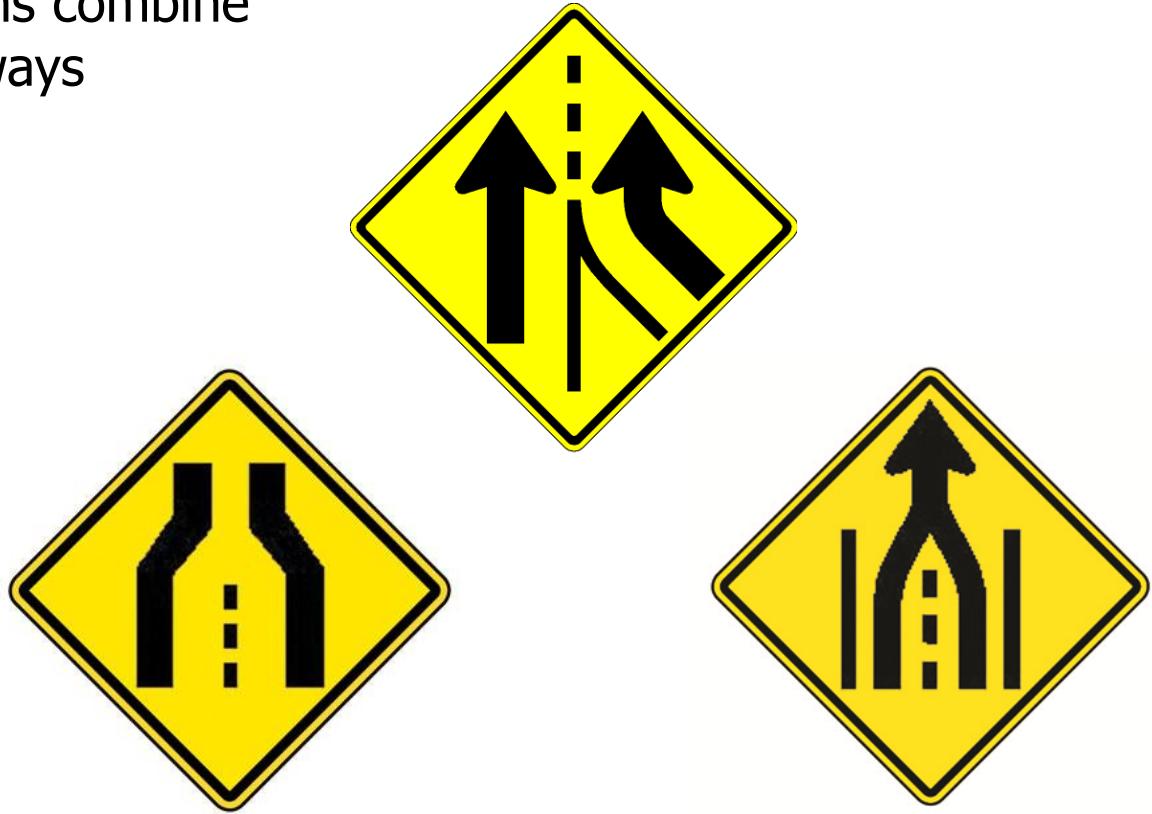
- After the common fork-join pool finishes processing chunks their partial results are combined into a final result
  - `join()` occurs in a single thread at each level
    - i.e., the “parent”



As a result, there's typically no need for synchronizers during the joining

# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways



Understanding these differences is particularly important for parallel streams

# Combining Results in a Parallel Stream

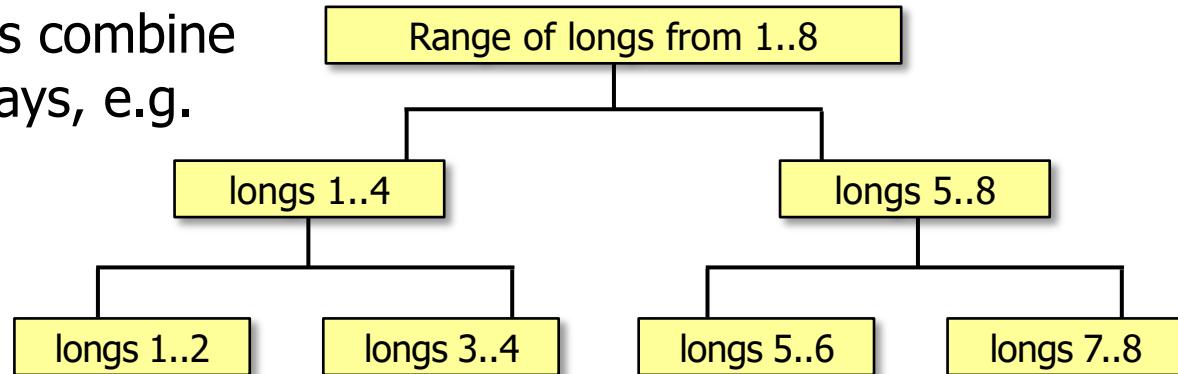
---

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value



# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value

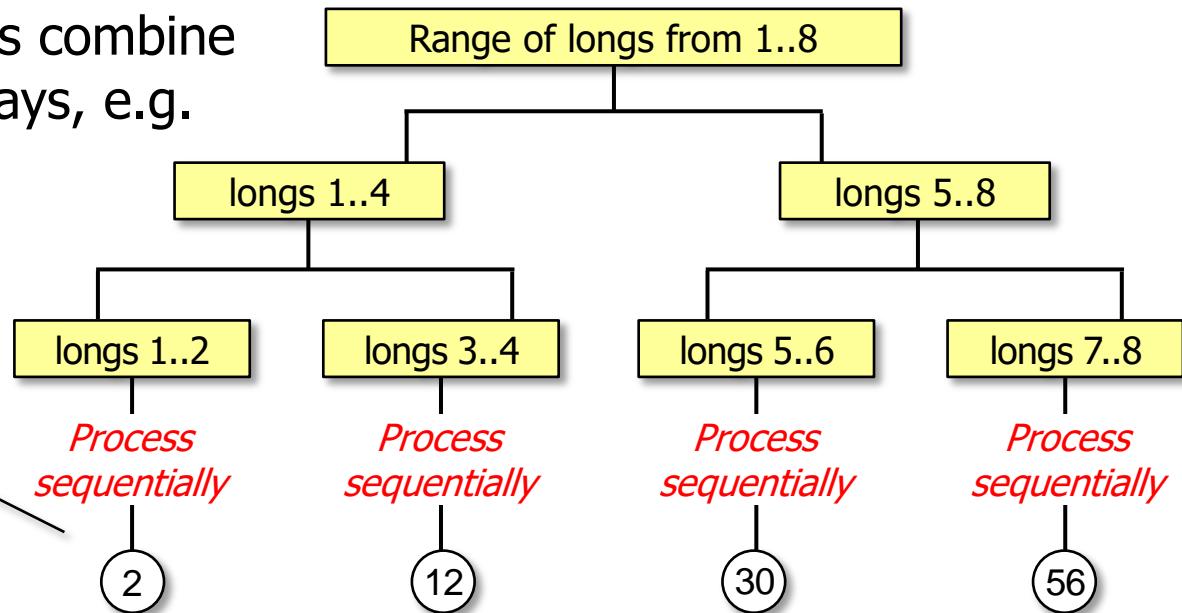


```
long factorial(long n) {  
    return LongStream  
        .rangeClosed(1, n)  
        .parallel()  
        .reduce(1, (a, b) -> a * b,  
               (a, b) -> a * b);  
}
```

*Generate a range of longs  
from 1..8 in parallel*

# Combining Results in a Parallel Stream

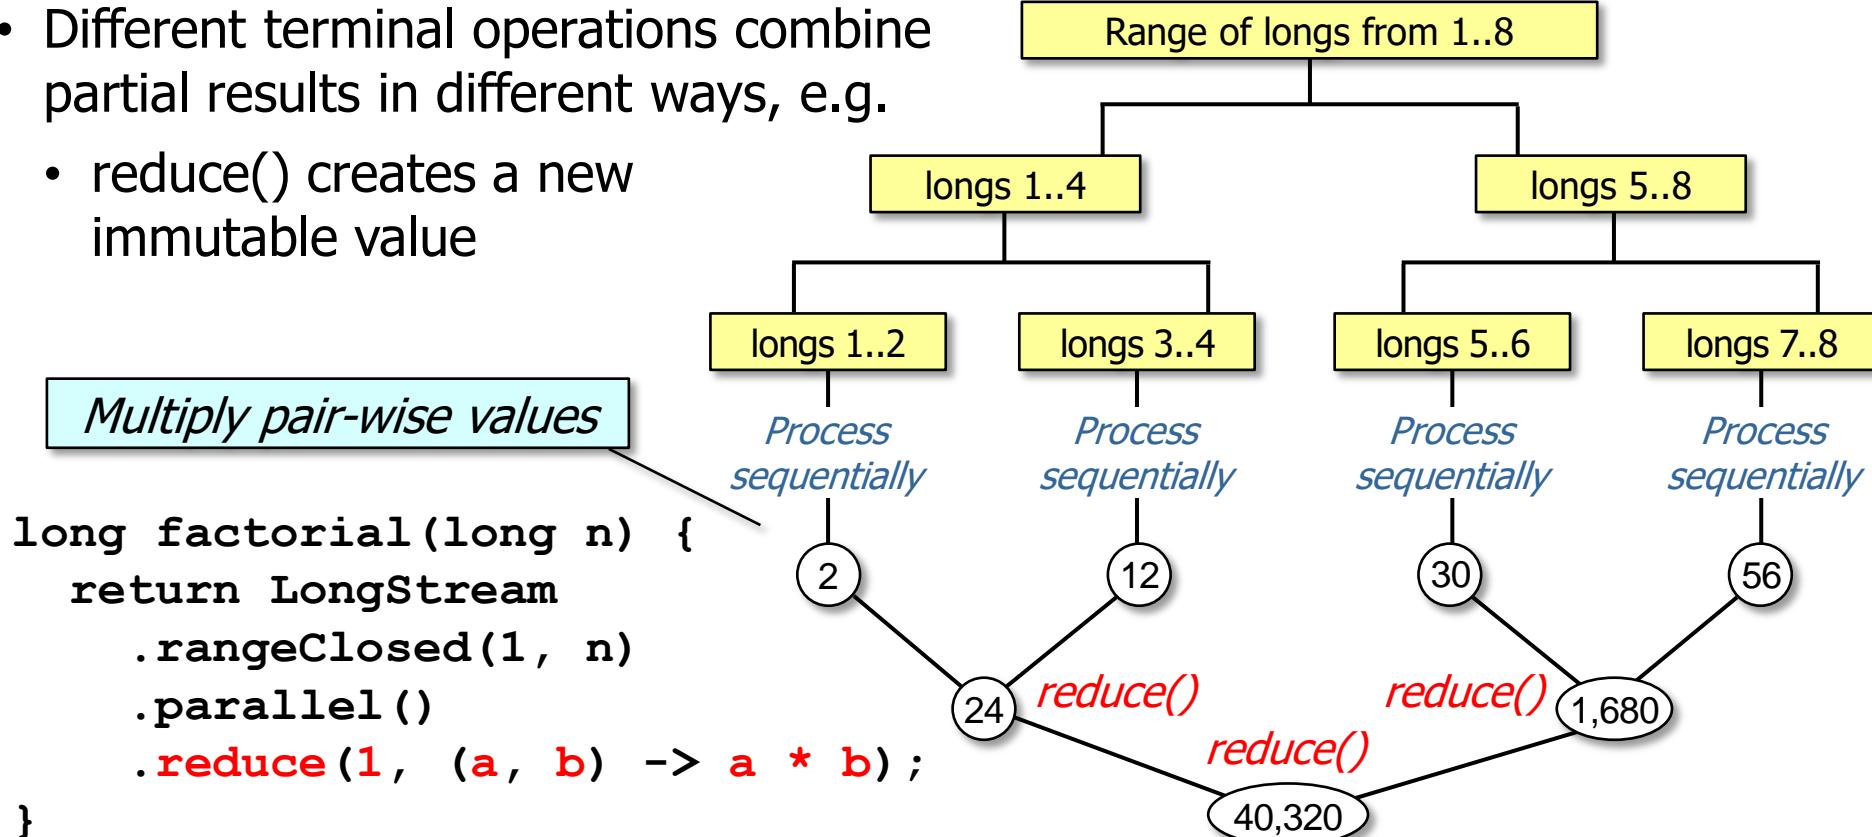
- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value



```
long factorial(long n) {  
    return LongStream  
        .rangeClosed(1, n)  
        .parallel()  
        .reduce(1, (a, b) -> a * b);  
}
```

# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value



`reduce()` combines two immutable values (e.g., long) and produces a new one

# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value
  - `collect()` mutates an existing results container

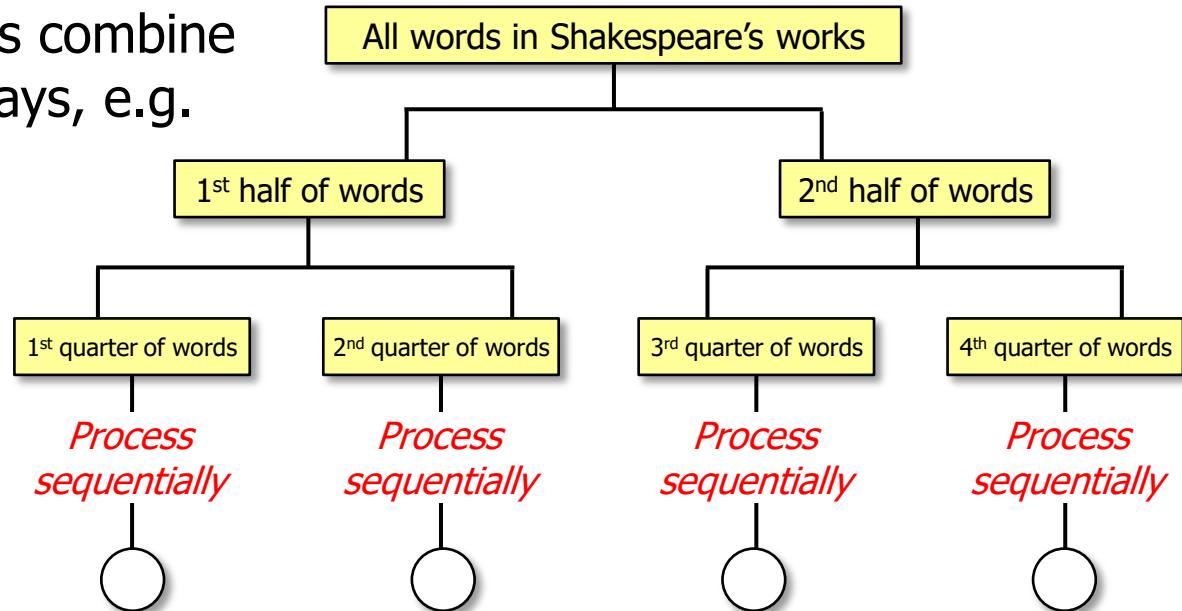


See [greenteapress.com/thinkapjava/html/thinkjava011.html](http://greenteapress.com/thinkapjava/html/thinkjava011.html)

# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value
  - `collect()` mutates an existing results container

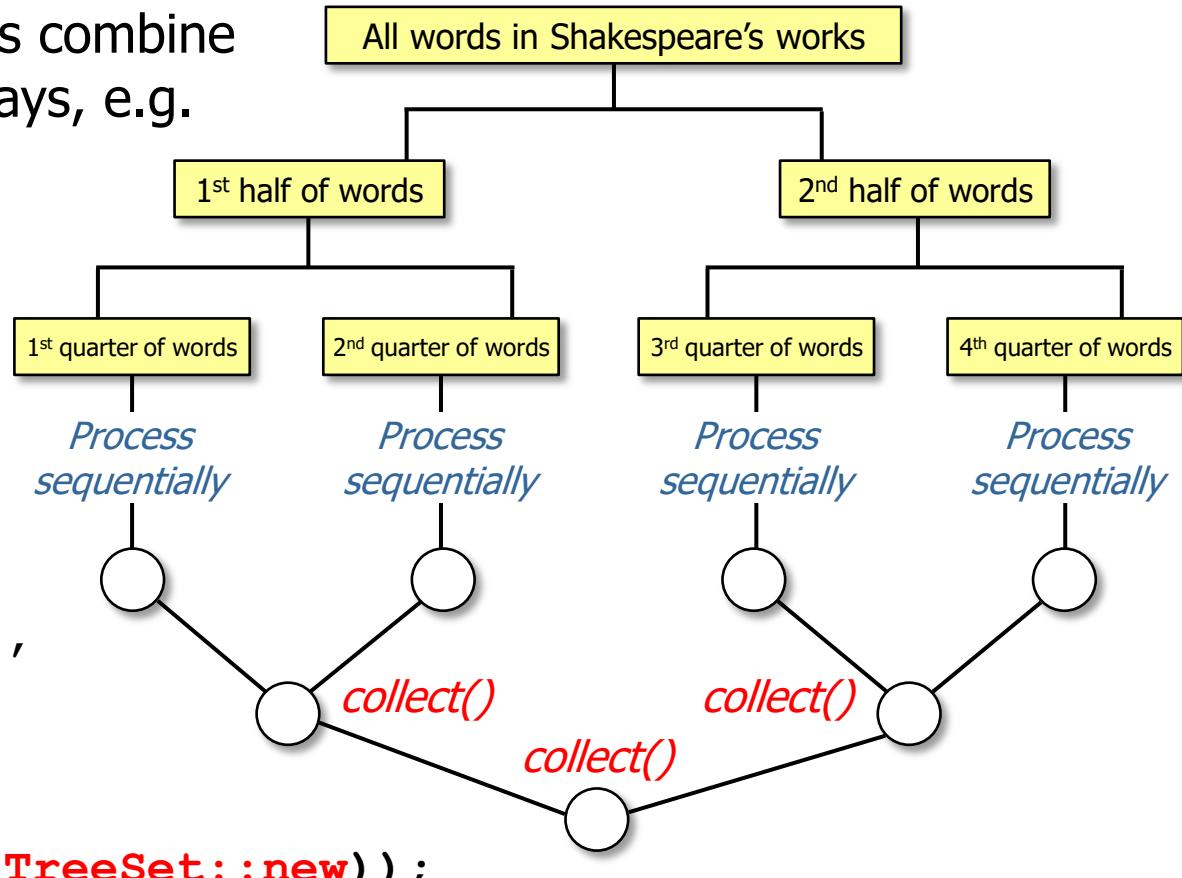
```
Set<CharSequence>
    uniqueWords =
        getInput(sSHAKESPEARE),
        "\s+")
    .parallelStream()
    ...
    .collect(toCollection(TreeSet::new));
```



# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value
  - `collect()` mutates an existing results container

```
Set<CharSequence>
    uniqueWords =
getInput(sSHAKESPEARE),
    "\s+")
.parallelStream()
...
.collect(toCollection(TreeSet::new));
```

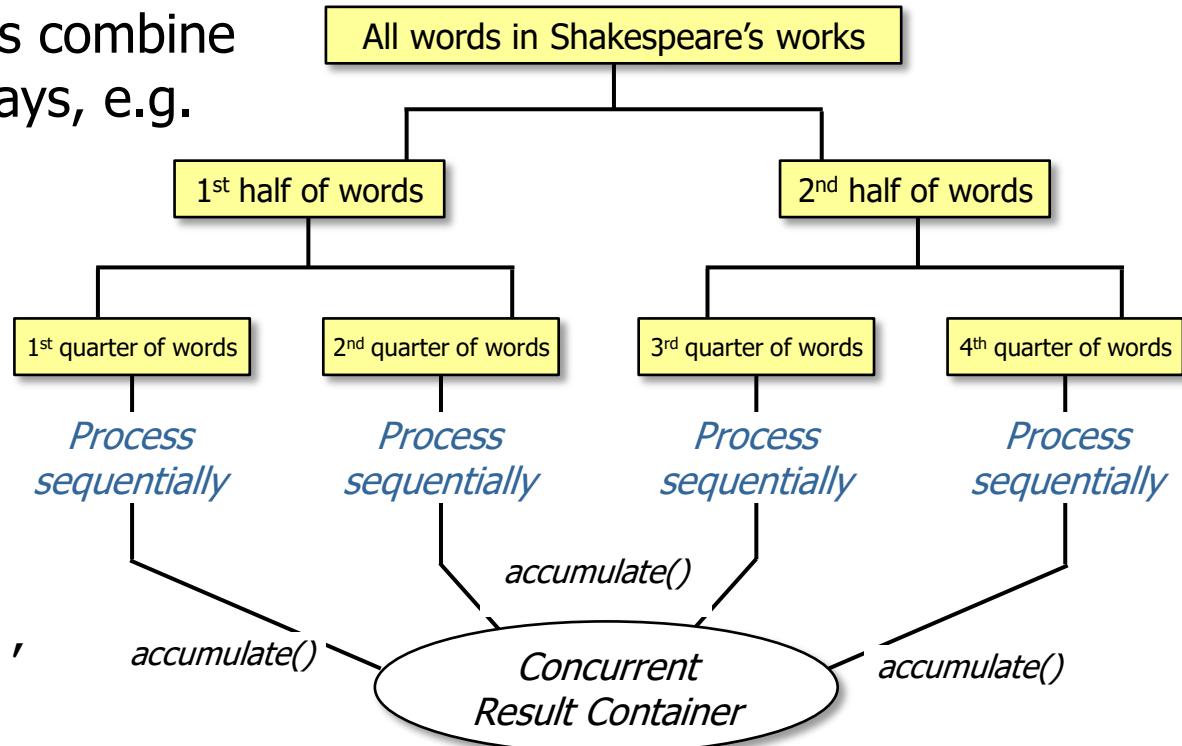


collect() mutates a container to accumulate the result it's producing

# Combining Results in a Parallel Stream

- Different terminal operations combine partial results in different ways, e.g.
  - `reduce()` creates a new immutable value
  - `collect()` mutates an existing results container

```
Set<CharSequence>
    uniqueWords =
getInput(sSHAKESPEARE),
    "\s+")
.parallelStream()
...
.collect(ConcurrentHashSetCollector.toSet());
```



Concurrent collectors (covered later) are different than non-concurrent collectors

Java Parallel Stream Internals:  
Combining Results (Part I)

---

**The End**

# Java Parallel Stream Internals

---

## Combining Results (Part II)

Douglas C. Schmidt

# Learning Objectives in This Part of the Lesson

---

- Understand parallel stream internals, e.g.
  - Know what can change and what can't.
  - Partition a data source into "chunks."
  - Process chunks in parallel via the common ForkJoinPool.
  - Configure the Java parallel stream common ForkJoinPool.
  - Perform a reduction to combine partial results into a single result.
    - Be aware of common traps and pitfalls with parallel streams.



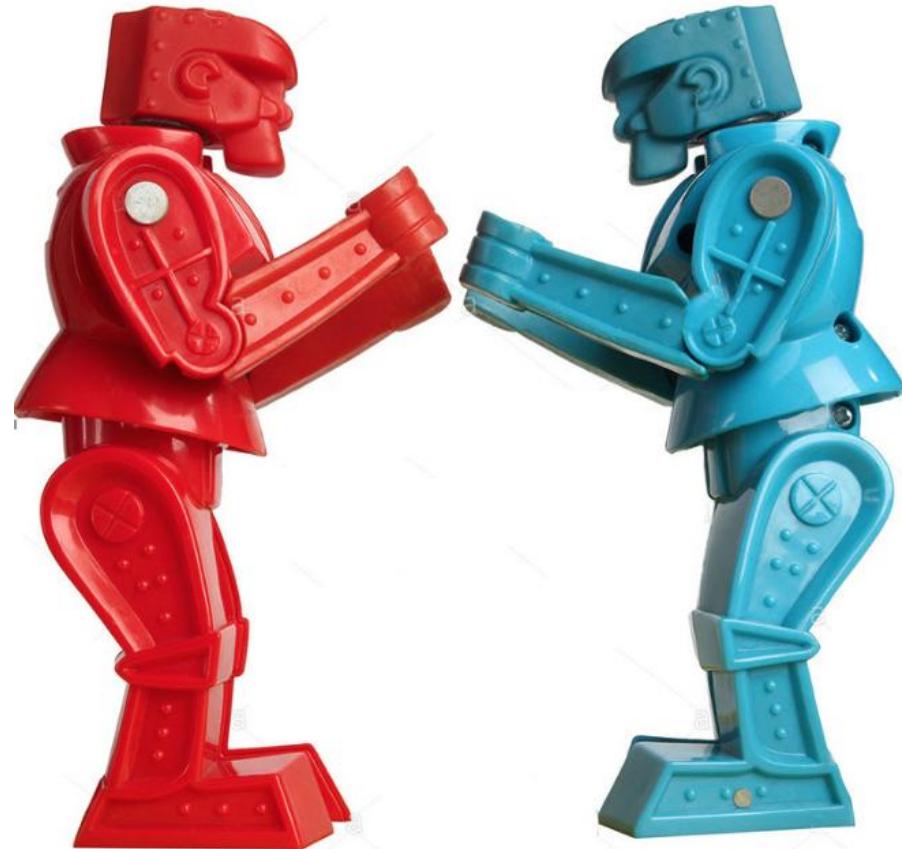
---

# Differences for collect() and reduce() in a Parallel Stream

# Differences for collect() and reduce() in a Parallel Stream

---

- It's important to understand the semantic differences between collect() and reduce()



# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

```
void buggyStreamReduce3
(boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new StringBuilder(),
        StringBuilder::append,
        StringBuilder::append)
    .toString();
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*Convert a list of words into a stream of words*

```
void buggyStreamReduce3
(boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();
if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new StringBuilder(),
        StringBuilder::append,
        StringBuilder::append)
    .toString();
```

Naturally, this call doesn't really do any work since streams are "lazy"

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*A stream can be dynamically switched to "parallel" mode!*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*The "last" call to .parallel() or .sequential() in a stream "wins"*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new StringBuilder(),
        StringBuilder::append,
        StringBuilder::append)
    .toString();
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*This code works when parallel is false since the StringBuilder is only called in a single thread*

```
void buggyStreamReduce3
(boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new StringBuilder(),
        StringBuilder::append,
        StringBuilder::append)
    .toString();
```



# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions

*This code fails when parallel is true since reduce() expects to do an "immutable" reduction*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new StringBuilder(),
        StringBuilder::append,
        StringBuilder::append)
    .toString();
```



# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions



*There's a race condition here since  
StringBuilder is not thread-safe..*

```
void buggyStreamReduce3
    (boolean parallel) {
    ...
    Stream<String> wordStream =
        allWords.stream();

    if (parallel)
        wordStream.parallel();

    String words = wordStream
        .reduce(new StringBuilder(),
            StringBuilder::append,
            StringBuilder::append)
        .toString();
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
    - One solution use reduce() with string concatenation

```
void streamReduceConcat
(boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new String(),
        (x, y) -> x + y);
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
  - One solution use reduce() with string concatenation

```
void streamReduceConcat
(boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .reduce(new String(),
        (x, y) -> x + y);
```

*This simple fix is inefficient due to string concatenation overhead*

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
    - One solution use reduce() with string concatenation
    - Another solution uses collect() with the joining collector

```
void streamCollectJoining
  (boolean parallel) {
  ...
Stream<String> wordStream =
  allWords.stream();

if (parallel)
  wordStream.parallel();

String words = wordStream
  .collect(joining());
```

# Differences for collect() and reduce() in a Parallel Stream

- It's important to understand the semantic differences between collect() and reduce(), e.g.
  - Always test w/a parallel stream to detect mistakes wrt mutable vs. immutable reductions
    - One solution use reduce() with string concatenation
    - Another solution uses collect() with the joining collector

```
void streamCollectJoining
    (boolean parallel) {
    ...
Stream<String> wordStream =
    allWords.stream();

if (parallel)
    wordStream.parallel();

String words = wordStream
    .collect(joining());
```

*This is a much better solution!!*



# Differences for collect() and reduce() in a Parallel Stream

- Also beware of issues related to associativity and identity with reduce()

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
               (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
               // Could use (x, y) -> x + y  
               Math::addExact);
```

# Differences for collect() and reduce() in a Parallel Stream

- Also beware of issues related to associativity and identity with reduce()

*This code fails for a parallel stream since subtraction is not associative*

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
               (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
               // Could use (x, y) -> x + y  
               Math::addExact);
```

# Differences for collect() and reduce() in a Parallel Stream

- Also beware of issues related to associativity and identity with reduce()

```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
               (x, y) -> x - y);  
}
```

```
void testSum(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
               // Could use (x, y) -> x + y  
               Math::addExact);
```

*This code fails if identity is not 0L*

The “identity” of an OP is defined as “identity OP value == value” (and inverse)

# Differences for collect() and reduce() in a Parallel Stream

- Also beware of issues related to associativity and identity with reduce()

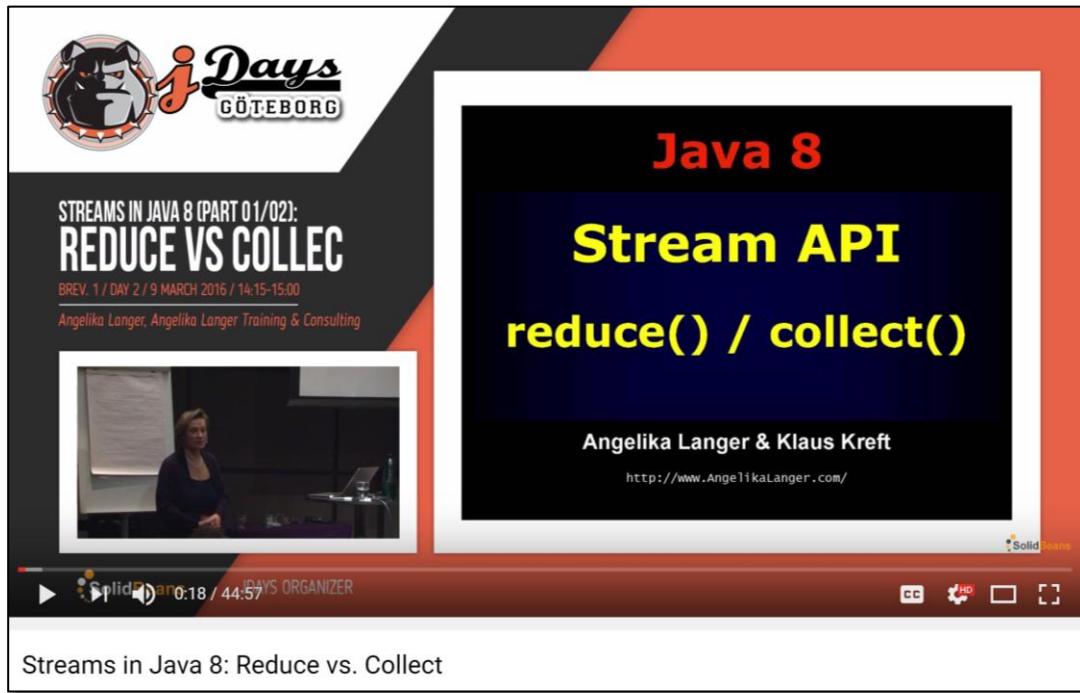
```
void testDifferenceReduce(...) {  
    long difference = LongStream  
        .rangeClosed(1, 100)  
        .parallel()  
        .reduce(0L,  
               (x, y) -> x - y);  
}
```

```
void testProd(long identity, ...) {  
    long sum = LongStream  
        .rangeClosed(1, 100)  
        .reduce(identity,  
               (x, y) -> x * y);  
}
```

*This code fails if identity is not 1L*

# Differences for collect() and reduce() in a Parallel Stream

- More good discussions about reduce() vs. collect() appear online



See [www.youtube.com/watch?v=oWIWEKNM5Aw](https://www.youtube.com/watch?v=oWIWEKNM5Aw)

Java Parallel Stream Internals:  
Combining Results (Part II)

---

**The End**