

When to Use Java Parallel Streams

Douglas C. Schmidt

Learning Objectives in This Lesson

- Know when to use parallel streams.



Parallelism Is Not a Panacea

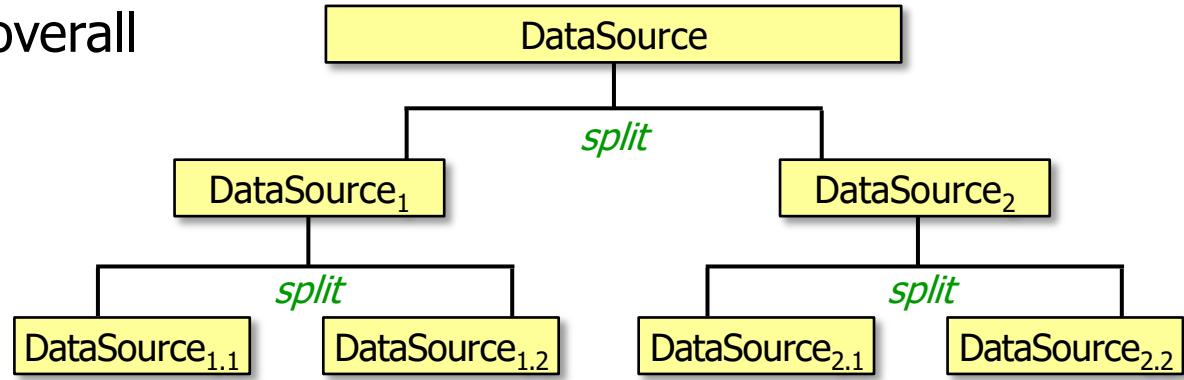
Parallelism Is Not a Panacea

- A parallel program *always* does more work than a nonparallel program.



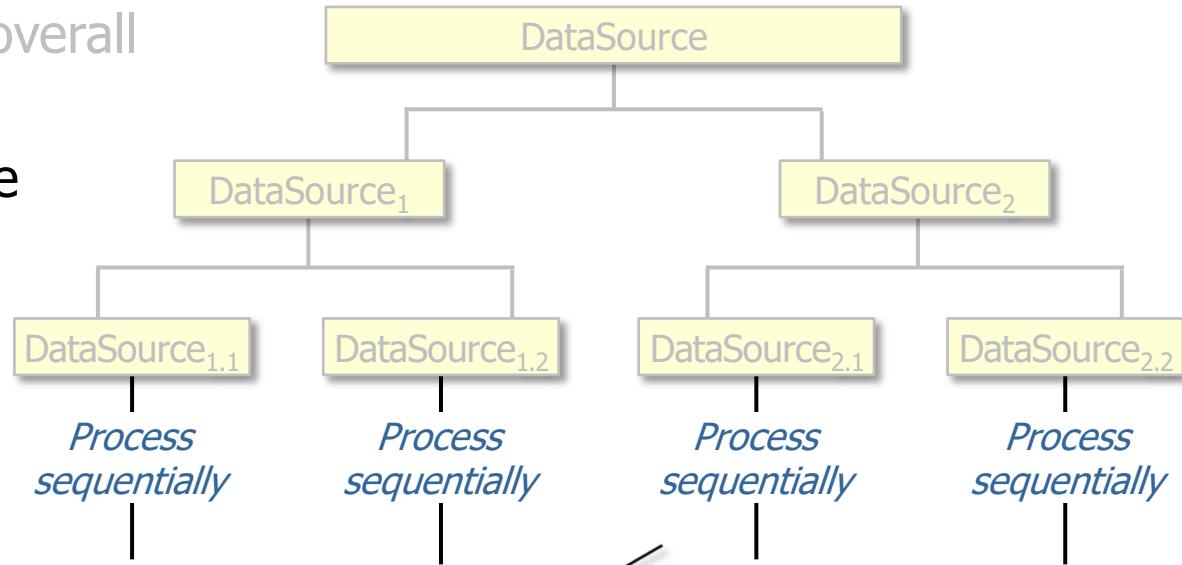
Parallelism Is Not a Panacea

- A parallel program *always* does more work than a nonparallel program, e.g.
 1. It needs to partition the overall task into sub-tasks.



Parallelism Is Not a Panacea

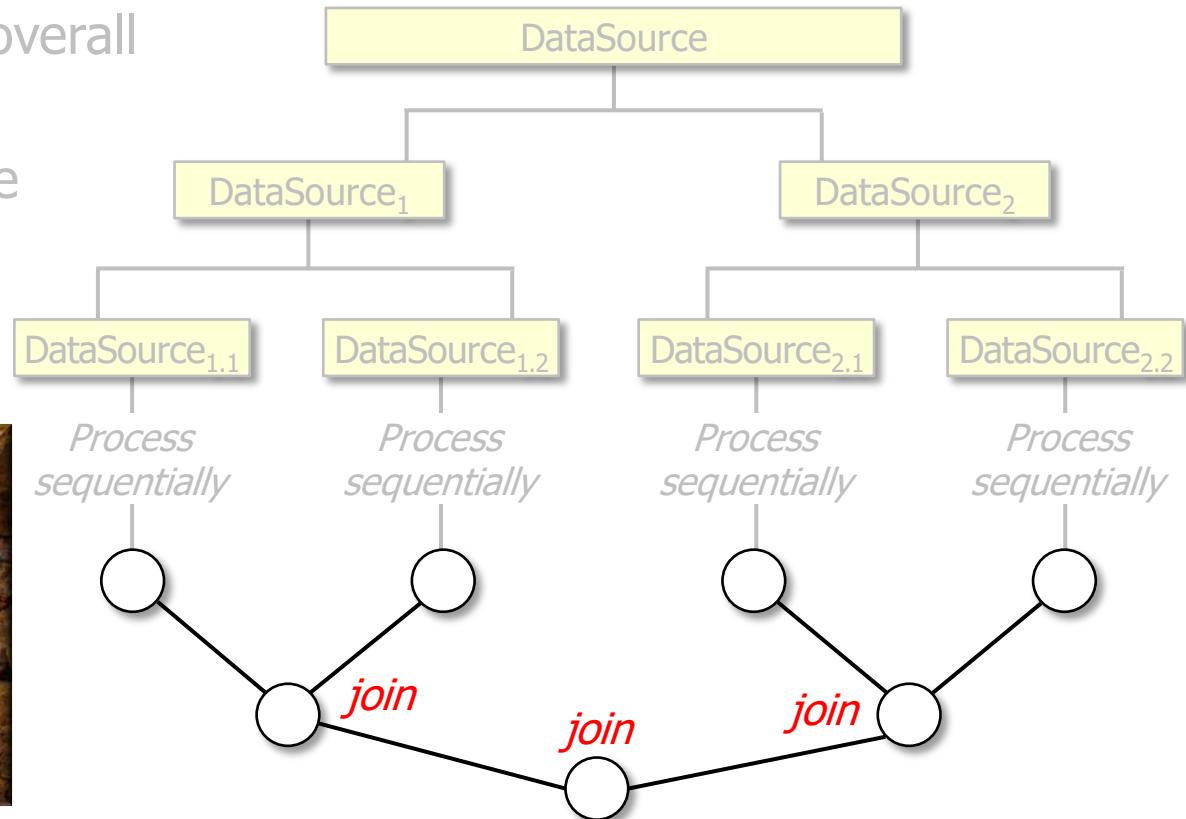
- A parallel program *always* does more work than a nonparallel program, e.g.
 1. It needs to partition the overall task into sub-tasks.
 2. It needs to process all the sub-tasks.



This step is typically *all* that a sequential program does!

Parallelism Is Not a Panacea

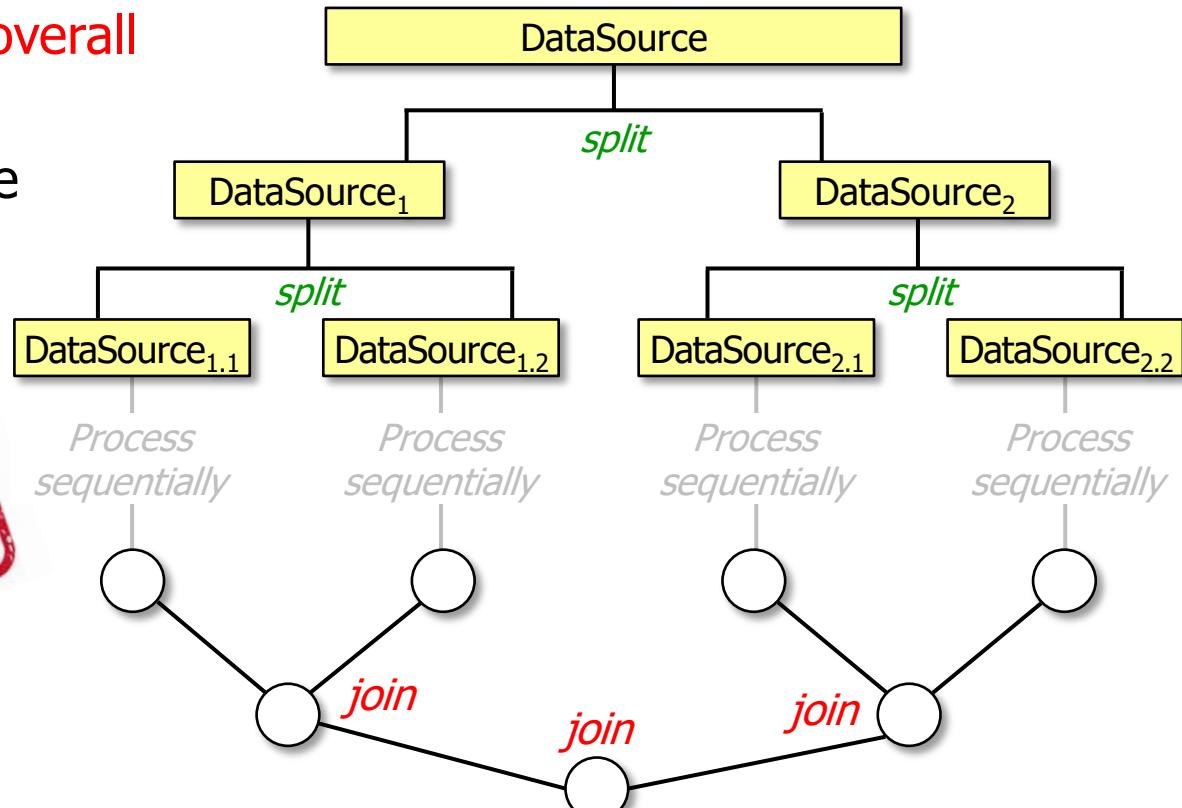
- A parallel program *always* does more work than a nonparallel program, e.g.
 - It needs to partition the overall task into sub-tasks.
 - It needs to process all the sub-tasks.
 - It needs to combine the sub-task results.



Parallelism Is Not a Panacea

- A parallel program *always* does more work than a nonparallel program, e.g.
 1. It needs to partition the overall task into sub-tasks.
 2. It needs to process all the sub-tasks.
 3. It needs to combine the sub-task results.

EXTRA COST



A sequential program needn't do steps one and three...

Parallelism Is Not a Panacea

- Java parallel streams are thus useful in some (but not all) conditions.



When to Use Java Parallel Streams

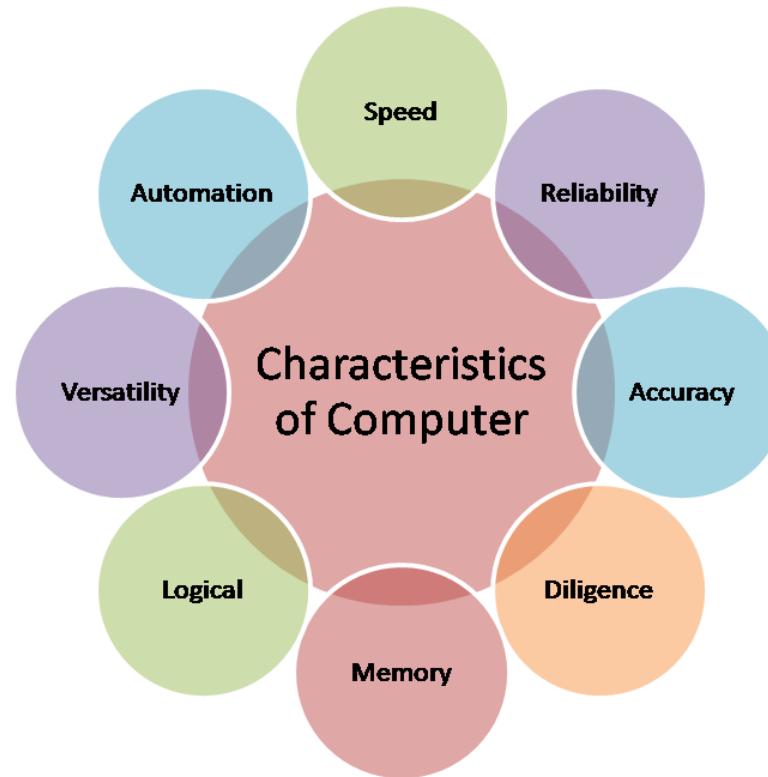
When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions.



When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics



When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent

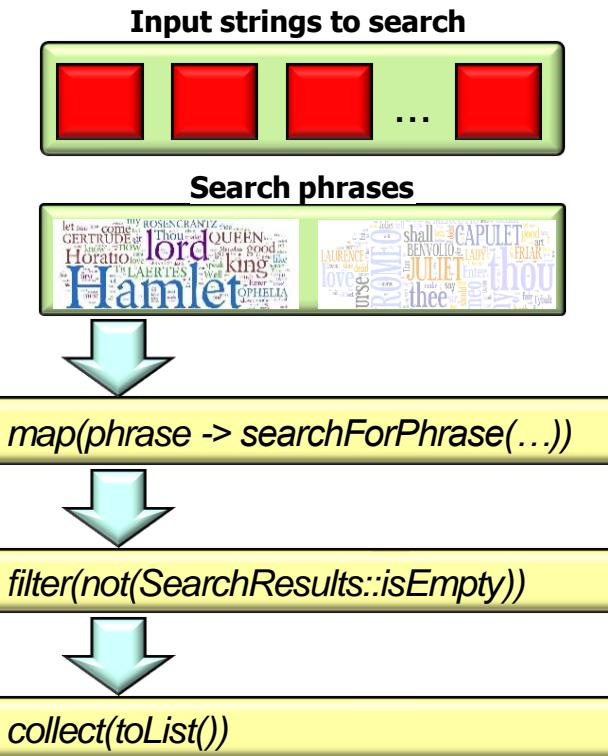
"Embarrassingly parallel" tasks have little/no dependency or need for communication between tasks or for sharing results between them.



See en.wikipedia.org/wiki/Embarrassingly_parallel

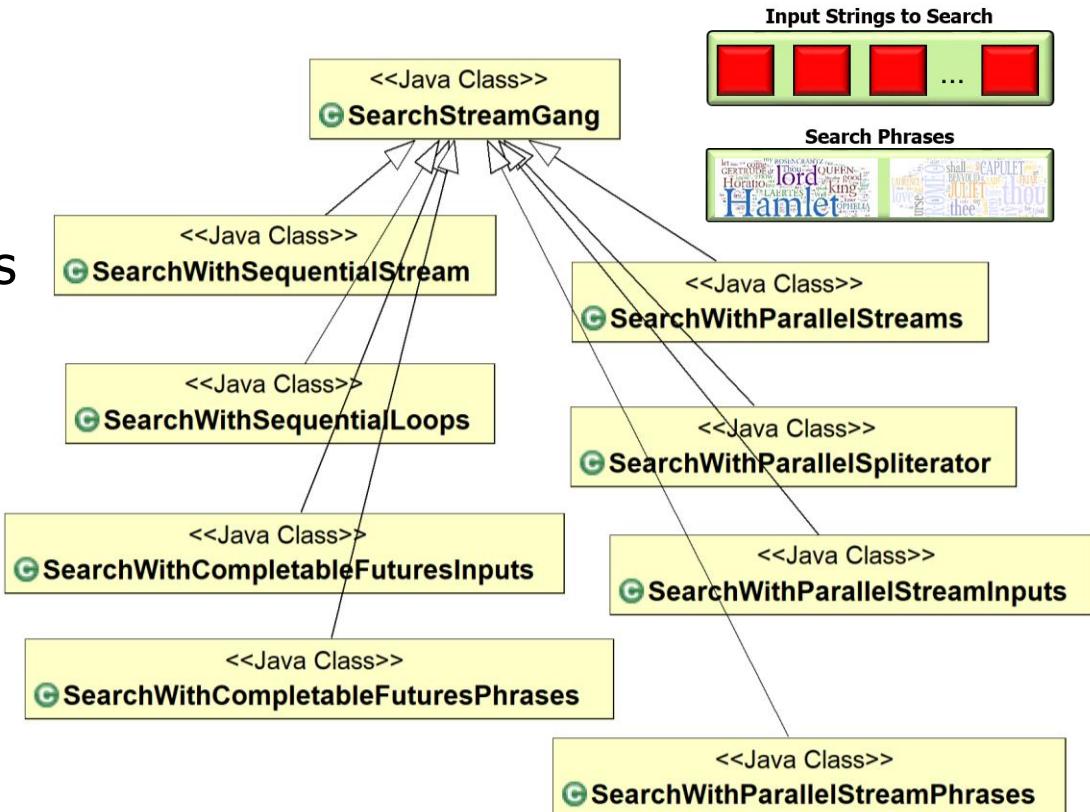
When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings



When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings

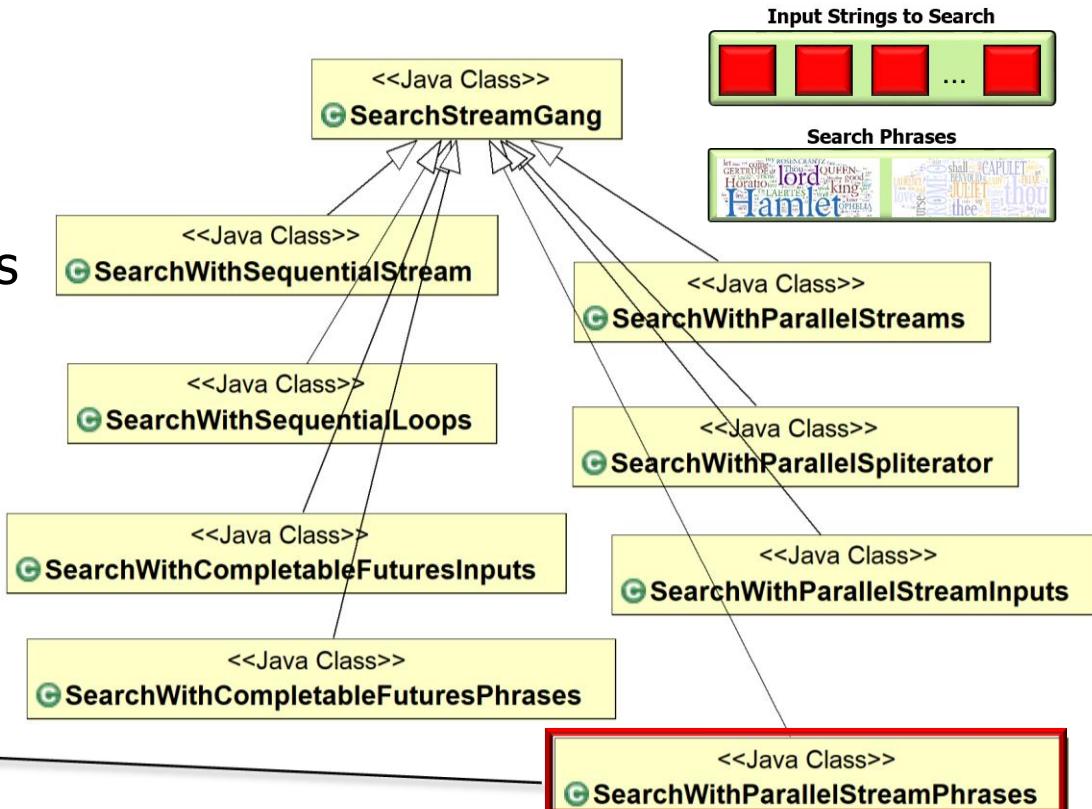


When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can

- Search each phrase in parallel*
- Search each input string in parallel*
- Search chunks of each input string in parallel*

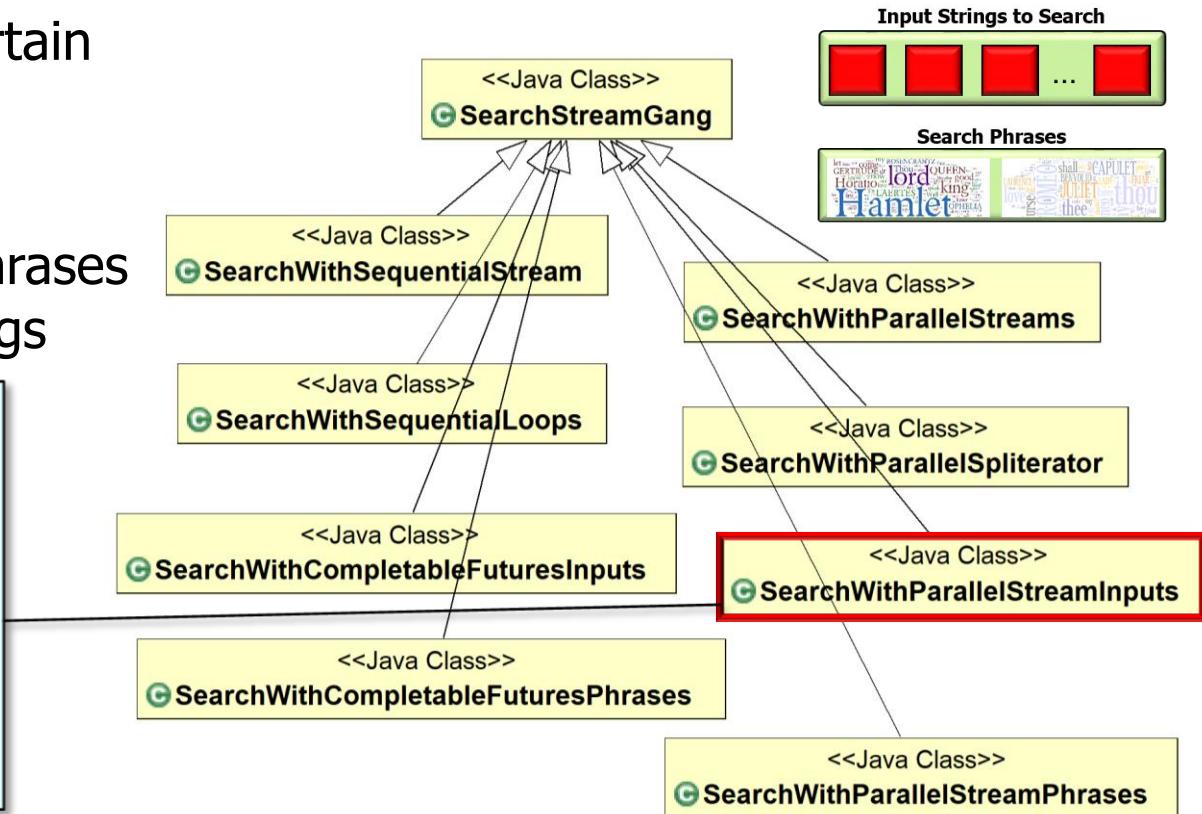


When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can

- Search each phrase in parallel*
- Search each input string in parallel*
- Search chunks of each input string in parallel*

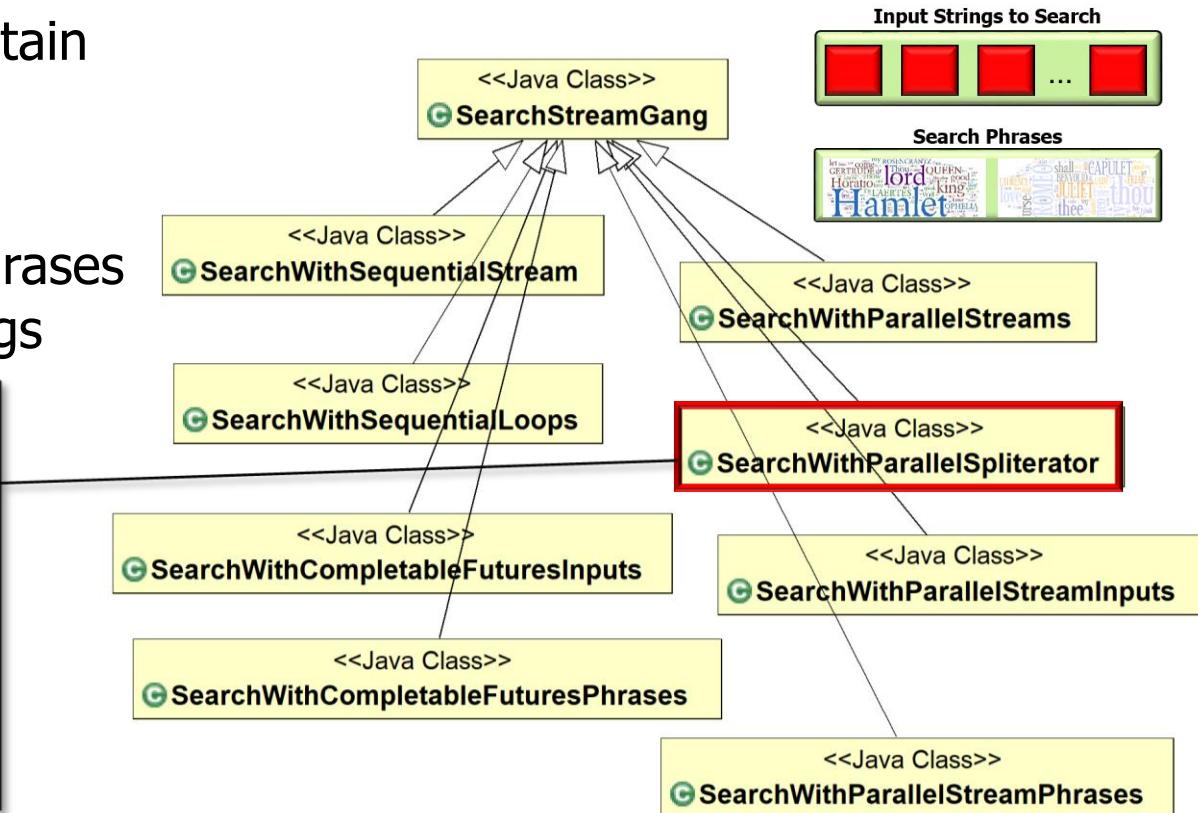


When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can

- Search each phrase in parallel*
- Search each input string in parallel*
- Search chunks of each input string in parallel*

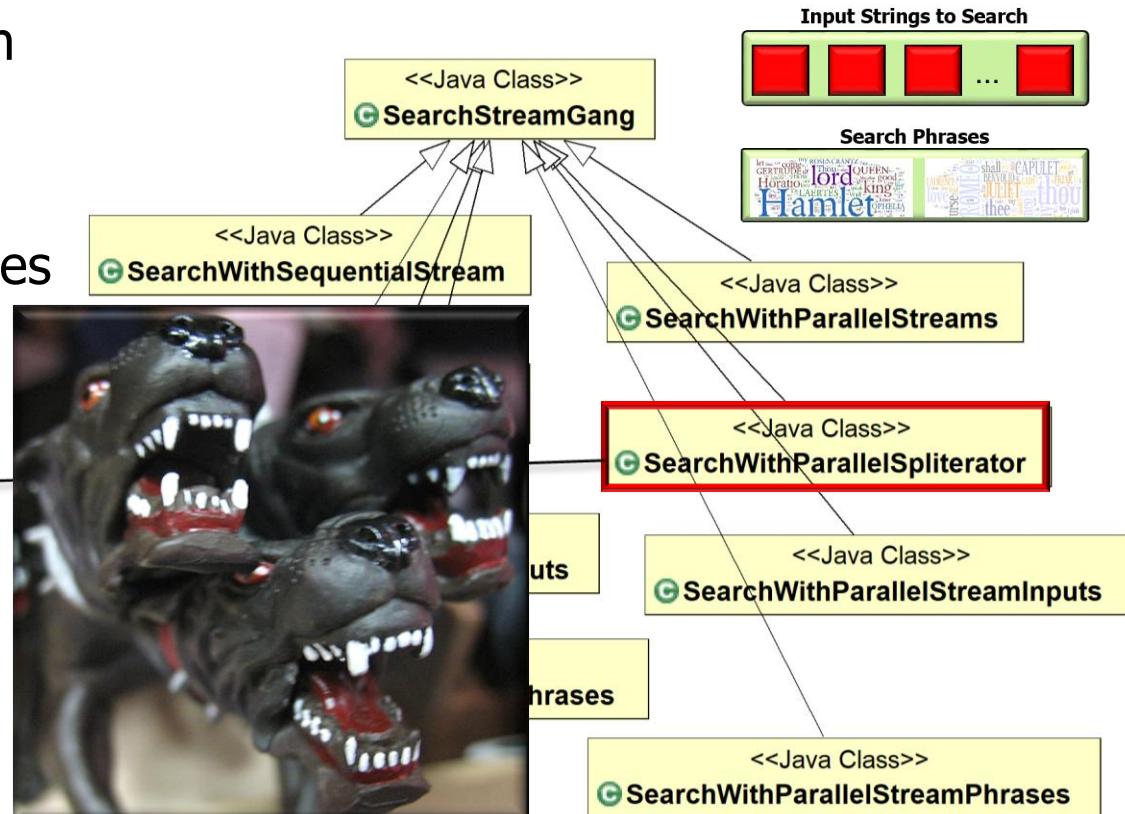


When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - e.g., searching for phrases in a list of input strings

Parallel streams can

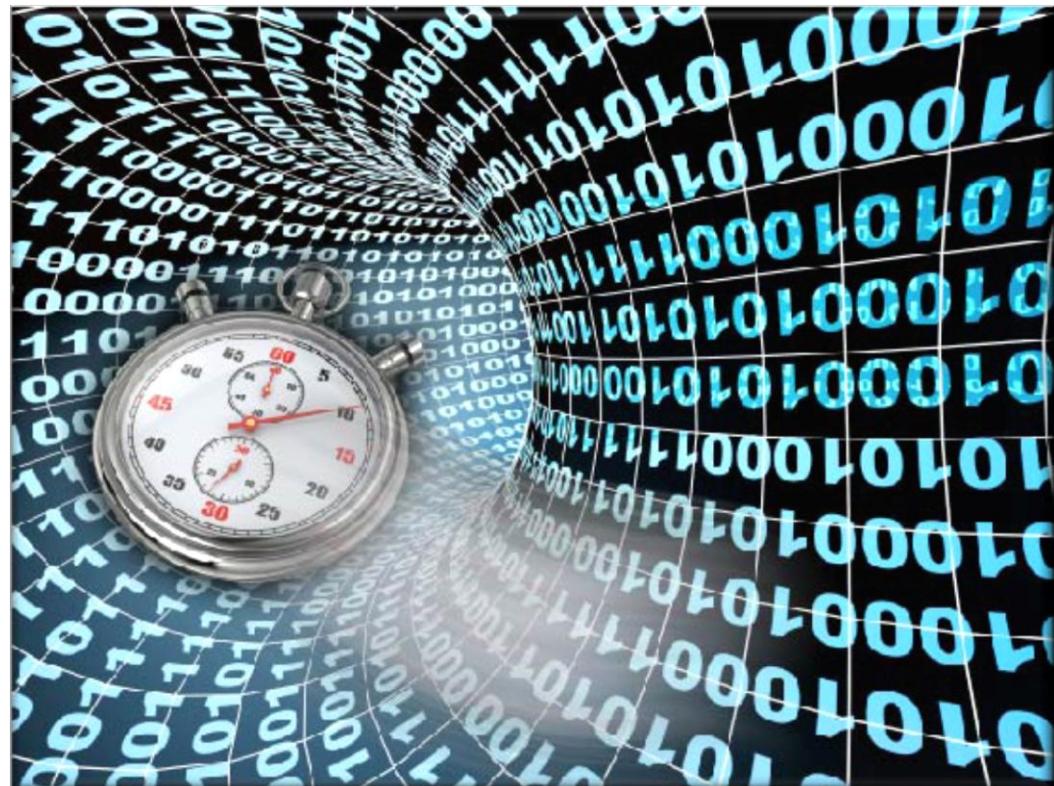
- Search chunks of phrases in parallel
- Search chunks of input in parallel
- Search chunks of each input string in parallel



SearchWithParallelSpliterator is the most aggressive parallelism strategy!

When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - Computationally expensive
 - e.g., when behavior(s) applied to each element take a “long time” to run



When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - Computationally expensive
 - Applied to many elements of data sources



See www.ibm.com/developerworks/library/j-java-streams-5-brian-goetz

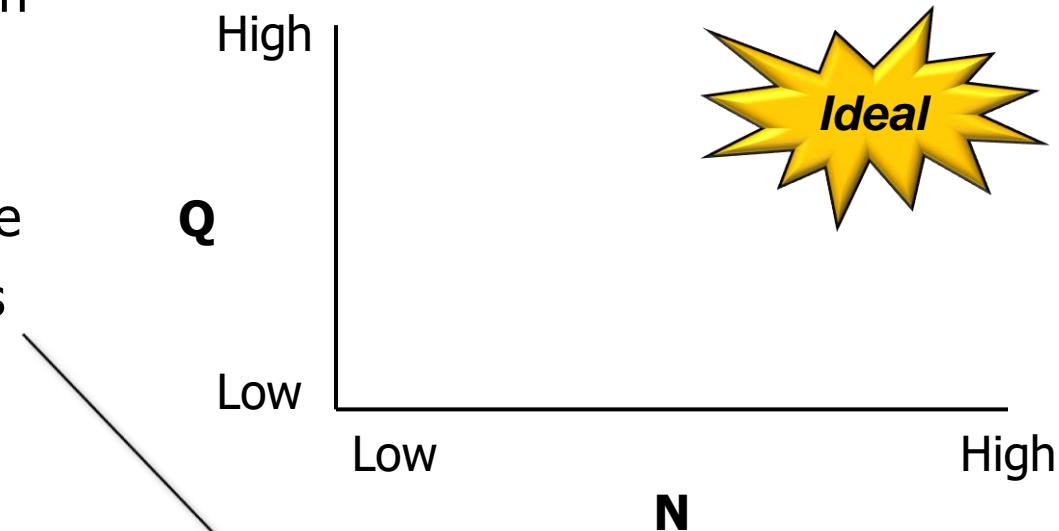
When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - Computationally expensive
 - Applied to many elements of data sources
 - Where these sources can be split efficiently/evenly



When to Use Java Parallel Streams

- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - Computationally expensive
 - Applied to many elements of data sources



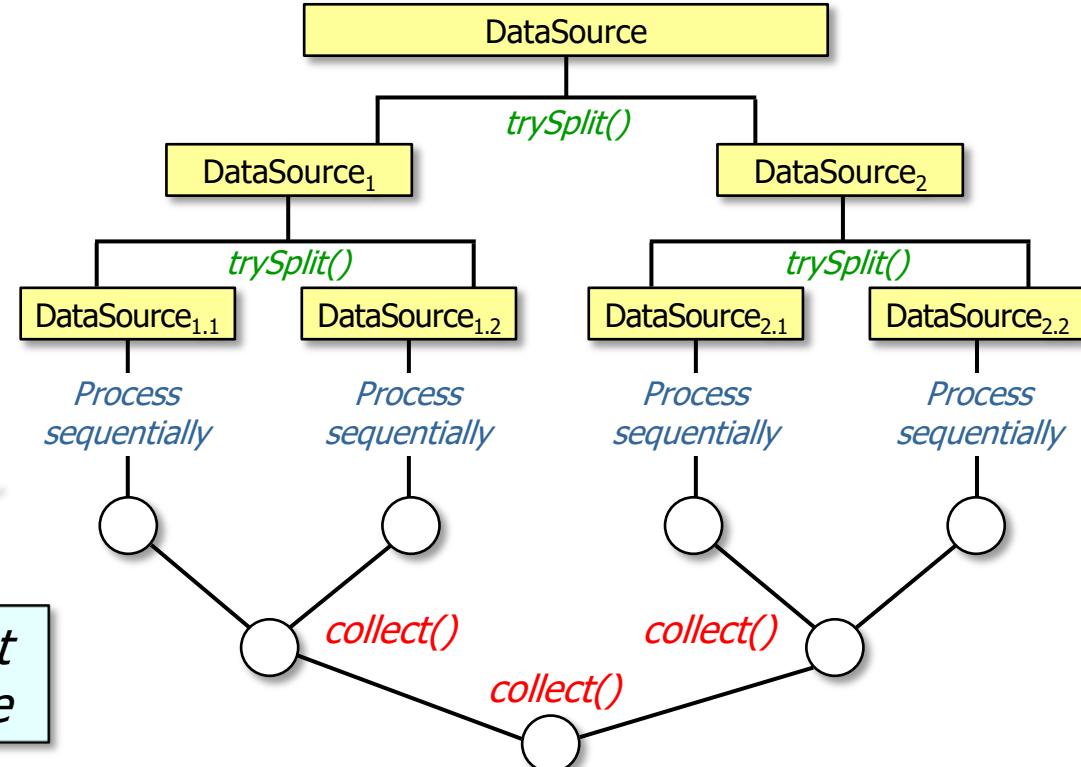
The "NQ" model

- N is the number of data elements to process per thread.*
- Q quantifies how CPU-intensive the processing is.*

When to Use Java Parallel Streams

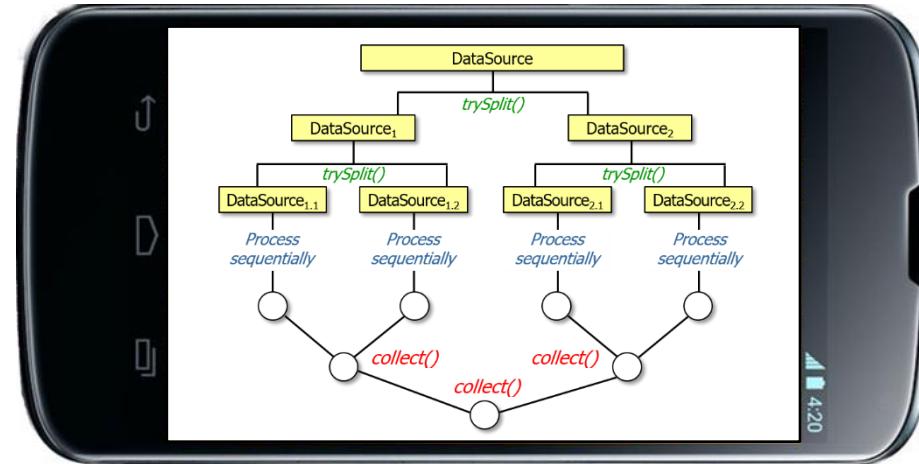
- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - Independent
 - Computationally expensive
 - Applied to many elements of data sources

e.g., searching for phrases that match in works of Shakespeare



When to Use Java Parallel Streams

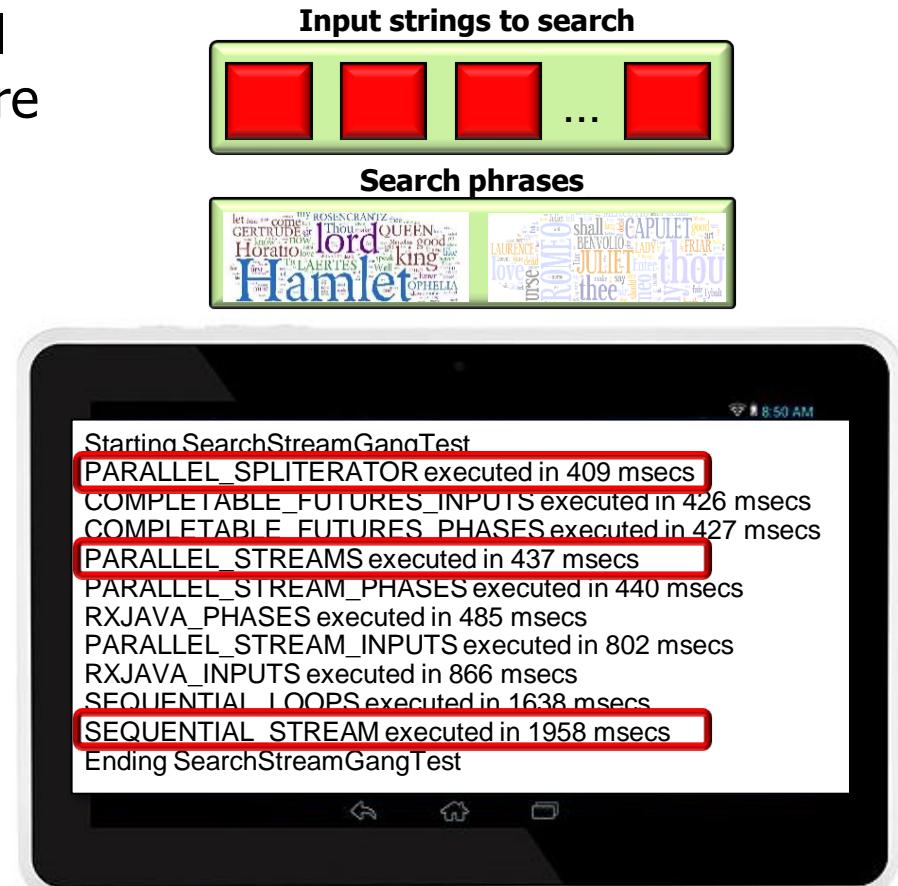
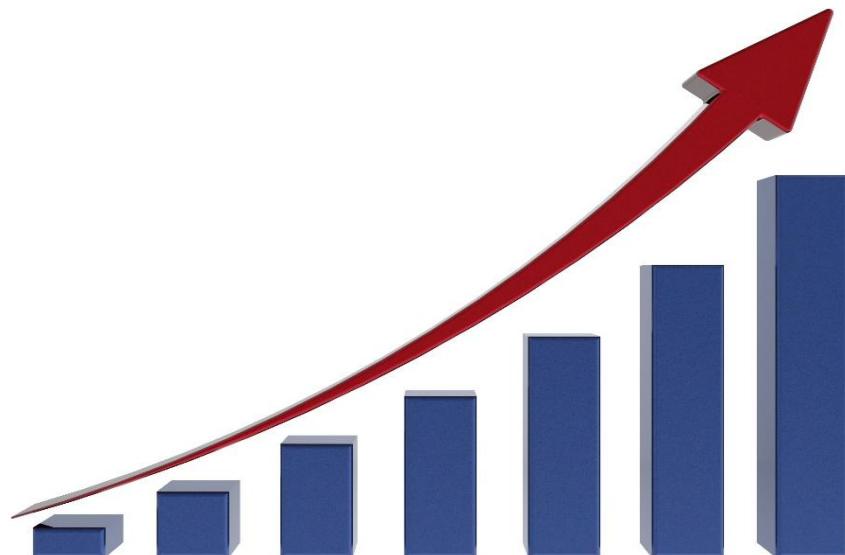
- Java parallel streams are most useful under certain conditions, e.g.
 - When behaviors have certain characteristics
 - If there are multiple cores



See blog.oio.de/2016/01/22/parallel-stream-processing-in-java-8-performance-of-sequential-vs-parallel-stream-processing

When to Use Java Parallel Streams

- Under the right conditions Java parallel streams can scale up nicely on multicore and many-core processors.



When to Use Java Parallel Streams

The End

When Not to Use Java Parallel Streams

Douglas C. Schmidt

Learning Objectives in This Lesson

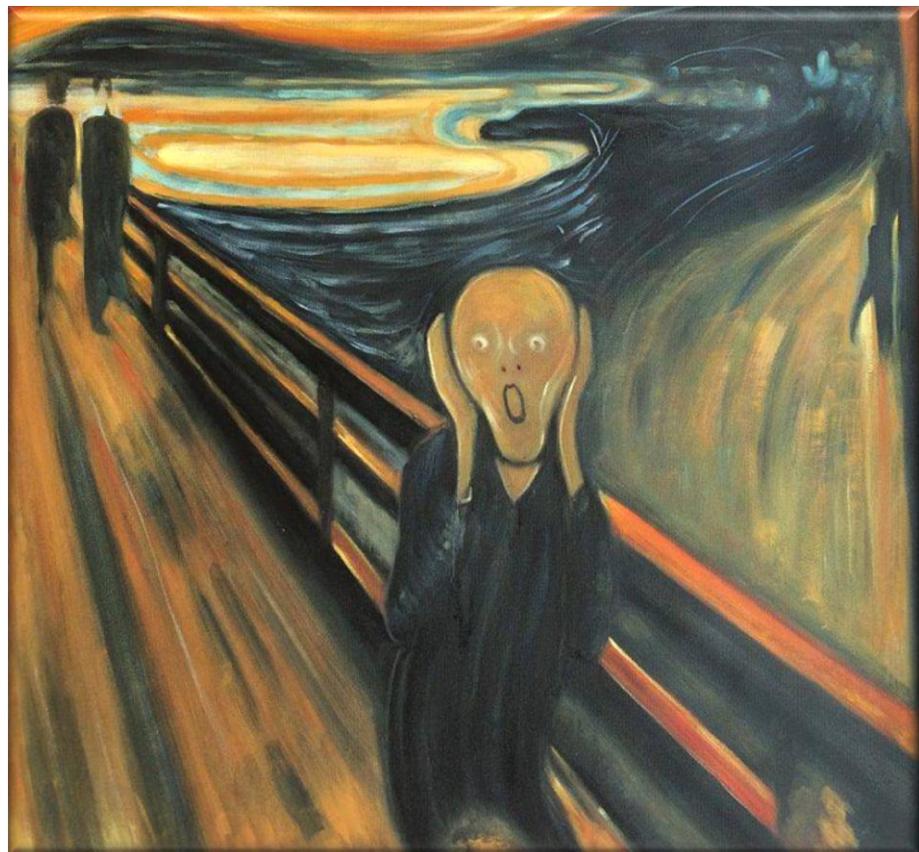
- Know when to use parallel streams.
 - And when *not* to use parallel streams.



When Not to Use Java Parallel Streams

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs.



See www.ibm.com/developerworks/library/j-java-streams-5-brian-goetz

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.



```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (sSHAKESPEARE_WORKS, "\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

```
arrayAllWords.parallelStream()  
    ...;
```

```
listAllWords.parallelStream()  
    ...;
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Make an ArrayList that contains all words in the works of Shakespeare.

```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (sSHAKESPEARE_WORKS, "\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

```
arrayAllWords.parallelStream()  
    ...;
```

```
listAllWords.parallelStream()  
    ...;
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Make a LinkedList that contains all words in the works of Shakespeare.

```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (sSHAKESPEARE_WORKS, "\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

```
arrayAllWords.parallelStream()  
    ...;
```

```
listAllWords.parallelStream()  
    ...;
```

LinkedList doesn't split evenly or efficiently compared with ArrayList.

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

```
Starting spliterator tests for 100000  
words....printing results  
599 msecs: ArrayList parallel  
701 msecs: LinkedList parallel  
  
Starting spliterator tests for 883311  
words....printing results  
5718 msecs: ArrayList parallel  
31226 msecs: LinkedList parallel
```

```
List<CharSequence> arrayAllWords =  
    TestDataFactory.getInput  
        (sSHAKESPEARE_WORKS, "\s+");
```

```
List<CharSequence> listAllWords =  
    new LinkedList<>(arrayAllWords);
```

```
arrayAllWords.parallelStream()
```

```
...;
```

```
listAllWords.parallelStream()
```

```
...;
```

The ArrayList parallel stream is much faster than the LinkedList parallel stream.

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

The ArrayList Spliterator runs in O(1) constant time.

```
class ArrayListSpliterator {  
    ...  
    ArrayListSpliterator<E>  
    trySplit() {  
        int hi = getFence(), lo =  
            index, mid = (lo + hi) >>> 1;  
        return lo >= mid  
            ? null  
            : new  
                ArrayListSpliterator<E>  
                (list, lo, index = mid,  
                 expectedModCount);  
    }  
    ...
```

See [openjdk/8u40-b25/java/util/ArrayList.java](https://openjdk.java.net/jeps/204)

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Compute the midpoint efficiently.

```
class ArrayListSpliterator {  
    ...  
    ArrayListSpliterator<E>  
        trySplit() {  
            int hi = getFence(), lo =  
                index, mid = (lo + hi) >>> 1;  
            return lo >= mid  
                ? null  
                : new  
                    ArrayListSpliterator<E>  
                        (list, lo, index = mid,  
                         expectedModCount);  
        }  
        ...  
}
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Split the array list evenly without copying the data.

```
class ArrayListSpliterator {  
    ...  
    ArrayListSpliterator<E>  
        trySplit() {  
            int hi = getFence(), lo =  
                index, mid = (lo + hi) >>> 1;  
            return lo >= mid  
                ? null  
                : new  
                    ArrayListSpliterator<E>  
                        (list, lo, index = mid,  
                         expectedModCount);  
        }  
        ...  
}
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

The LinkedList Splitterator runs in $O(n)$ linear time.

```
class LLSpliterator {  
    ...  
    public Spliterator<E> trySplit() {  
        ...  
        int n = batch + BATCH_UNIT;  
        ...  
        Object[] a = new Object[n];  
        int j = 0;  
        do { a[j++] = p.item; }  
        while ((p = p.next) != null  
               && j < n);  
        ...  
        return Spliterators  
            .spliterator(a, 0, j,  
                         Spliterator.ORDERED);
```

See [openjdk/8-b132/java/util/LinkedList.java](https://openjdk.java.net/jeps/123)

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Create a fixed-size chunk.

```
class LLSpliterator {  
    ...  
    public Spliterator<E> trySplit() {  
        ...  
        int n = batch + BATCH_UNIT;  
        ...  
        Object[] a = new Object[n];  
        int j = 0;  
        do { a[j++] = p.item; }  
        while ((p = p.next) != null  
               && j < n);  
        ...  
        return Spliterators  
            .spliterator(a, 0, j,  
                         Spliterator.ORDERED);
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

Copy data into the chunk.

```
class LLSpliterator {  
    ...  
    public Spliterator<E> trySplit() {  
        ...  
        int n = batch + BATCH_UNIT;  
        ...  
        Object[] a = new Object[n];  
        int j = 0;  
        do { a[j++] = p.item; }  
        while ((p = p.next) != null  
               && j < n);  
        ...  
        return Spliterators  
            .spliterator(a, 0, j,  
                        Spliterator.ORDERED);
```

When Not to Use Java Parallel Streams

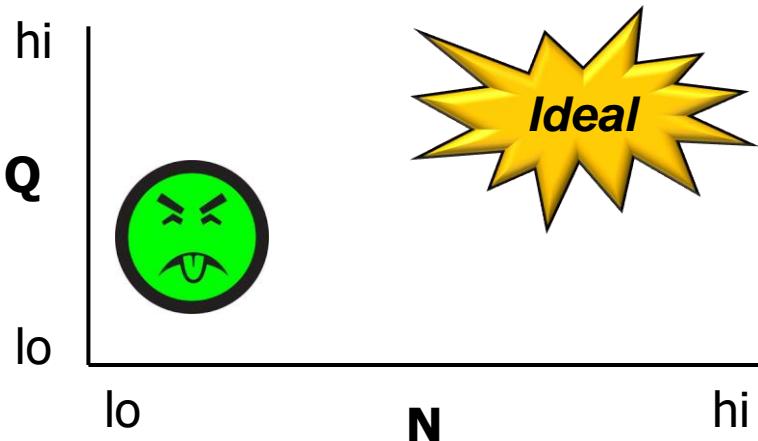
- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.

```
class LLSpliterator {  
    ...  
    public Spliterator<E> trySplit() {  
        ...  
        int n = batch + BATCH_UNIT;  
        ...  
        Object[] a = new Object[n];  
        int j = 0;  
        do { a[j++] = p.item; }  
        while ((p = p.next) != null  
               && j < n);  
        ...  
        return Spliterators  
            .spliterator(a, 0, j,  
                         Spliterator.ORDERED);
```

*Create a new Spliterator
that covers the chunk.*

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.



```
class ParallelStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n)  
            .parallel() ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    ...  
}  
  
class SequentialStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n) ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    ...  
}
```

See previous lesson on "When to Use Parallel Streams"

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.

The overhead of creating a parallel stream is greater than the benefits of parallelism for small values of "n."

```
class ParallelStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n)  
            .parallel() ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
}
```

```
class SequentialStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n) ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    }  
}
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.

If "n" is small, then this parallel solution will be inefficient.

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.

```
class ParallelStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n)  
            .parallel() ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    ...  
}
```

```
class SequentialStreamFactorial {  
    BigInteger factorial(long n) {  
        return LongStream  
            .rangeClosed(1, n) ...  
            .reduce(BigInteger.ONE,  
                    BigInteger::multiply);  
    ...  
}
```

If "n" is small, then this sequential solution will be more efficient.

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.



```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (sSHAKESPEARE_DATA_FILE,  
                "\\\\s+")) ;  
    ...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
    ...  
        .collect(toCollection  
            (TreeSet::new)) ;
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.

A linked list of all words in the complete works of Shakespeare.

```
List<CharSequence> allWords =  
    new LinkedList<>(  
        TestDataFactory.getInput  
            (sSHAKESPEARE_DATA_FILE,  
             "\s+"));  
    ...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
    ...  
        .collect(toCollection  
            (TreeSet::new));
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.

Performance will be poor due to the overhead of combining partial results for a set in a parallel stream.

```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (sSHAKESPEARE_DATA_FILE,  
                "\\\\s+")) ;  
    ...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
    ...  
        .collect(toCollection  
            (TreeSet::new)) ;
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.

Combining costs can be alleviated if the amount of work performed per element is rather large (i.e., the "NQ model").

```
List<CharSequence> allWords =  
    new LinkedList<>  
        (TestDataFactory.getInput  
            (sSHAKESPEARE_DATA_FILE,  
                "\\\s+")) ;  
    ...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
    ...  
        .collect(toCollection  
            (TreeSet::new)) ;
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.

```
List<CharSequence> allWords =  
    new LinkedList<>(  
        TestDataFactory.getInput(  
            sSHAKESPEARE_DATA_FILE,  
            "\\\s+"));  
    ...  
  
Set<CharSequence> uniqueWords =  
    allWords  
        .parallelStream()  
        ...  
        .collect(toSet()));
```

A concurrent collector can also be used to optimize the reduction phase.

See [Java8/ex14/src/main/java/utils/ConcurrentHashSetCollector.java](#)

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.

```
Starting collector tests for 100000 words..printing results
 219 msecs: parallel timeStreamCollectToConcurrentSet()
 364 msecs: parallel timeStreamCollectToSet()
 657 msecs: sequential timeStreamCollectToSet()
 804 msecs: sequential timeStreamCollectToConcurrentSet()

Starting collector tests for 883311 words..printing results
1782 msecs: parallel timeStreamCollectToConcurrentSet()
3010 msecs: parallel timeStreamCollectToSet()
6169 msecs: sequential timeStreamCollectToSet()
7652 msecs: sequential timeStreamCollectToConcurrentSet()
```

```
List<CharSequence> allWords =
    new LinkedList<>
        (TestDataFactory.getInput(
            sSHAKESPEARE_DATA_FILE,
            "\\\s+"));
```

...

```
Set<CharSequence> uniqueWords =
    allWords
        .parallelStream()
        ...
        .collect(toSet()));
```

Concurrent collector scales much better than non-concurrent collector

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.

```
List<Double> result = Stream
    .iterate(2, i -> i + 1)
    .parallel()
    .filter(this::isEven)
    .limit(n)
    .map(this::findSQRT)
    .collect(toList());
```

```
List<Double> result = LongStream
    .range(2, (n * 2) + 1)
    .parallel()
    .filter(this::isEven)
    .mapToObj(this::findSQRT)
    .collect(toList());
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.

Create a list containing square root of the first "n" even numbers.

```
List<Double> result = Stream  
    .iterate(2, i -> i + 1)  
    .parallel()  
    .filter(this::isEven)  
    .limit(n)  
    .map(this::findSQRT)  
    .collect(toList());
```

```
List<Double> result = LongStream  
    .range(2, (n * 2) + 1)  
    .parallel()  
    .filter(this::isEven)  
    .mapToObj(this::findSQRT)  
    .collect(toList());
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.

Stream.iterate() and limit() split and parallelize poorly since iterate creates an ordered stream . . .

```
List<Double> result = Stream
    .iterate(2, i -> i + 1)
    .parallel()
    .filter(this::isEven)
    .limit(n)
    .map(this::findSQRT)
    .collect(toList());
```

```
List<Double> result = LongStream
    .range(2, (n * 2) + 1)
    .parallel()
    .filter(this::isEven)
    .mapToObj(this::findSQRT)
    .collect(toList());
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.

Create a list containing square root of the first "n" even numbers.

```
List<Double> result = Stream  
    .iterate(2, i -> i + 1)  
    .parallel()  
    .filter(this::isEven)  
    .limit(n)  
    .map(this::findSQRT)  
    .collect(toList());
```

```
List<Double> result = LongStream  
    .range(2, (n * 2) + 1)  
    .parallel()  
    .filter(this::isEven)  
    .mapToObj(this::findSQRT)  
    .collect(toList());
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.

LongStream.range() splits nicely and thus runs efficiently in parallel.

```
List<Double> result = Stream  
    .iterate(2, i -> i + 1)  
    .parallel()  
    .filter(this::isEven)  
    .limit(n)  
    .map(this::findSQRT)  
    .collect(toList());
```

```
List<Double> result = LongStream  
    .range(2, (n * 2) + 1)  
    .parallel()  
    .filter(this::isEven)  
    .mapToObj(this::findSQRT)  
    .collect(toList());
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.
 - There aren't many/any cores.



Older computing devices just have a single core, which limits available parallelism.

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.
 - There aren't many/any cores.
 - No built-in means to shut down processing of a parallel stream.



When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.
 - There aren't many/any cores.
 - No built-in means to shut down processing of a parallel stream.

```
private static volatile  
boolean mCancelled;
```

Define a static volatile flag.

```
Image downloadImage(Cache.Item  
item) {  
    if (mCancelled)  
        throw new  
        CancellationException  
        ("Canceling crawl.");  
    ...
```

When Not to Use Java Parallel Streams

- Parallel streams aren't suitable for certain types of programs, e.g.
 - The source is expensive to split or splits unevenly.
 - The startup costs of parallelism overwhelm the amount of data.
 - Combining partial results is costly.
 - Some streams operations don't sufficiently exploit parallelism.
 - There aren't many/any cores.
 - No built-in means to shut down processing of a parallel stream.

```
private static volatile  
boolean mCancelled;  
  
Image downloadImage(Cache.Item  
item) {  
    if (mCancelled)  
        throw new  
        CancellationException  
        ("Canceling crawl.");  
    ...  
}
```

Before downloading the next image, check for cancellation and throw an exception if cancelled.

When Not to Use Java Parallel Streams

The End

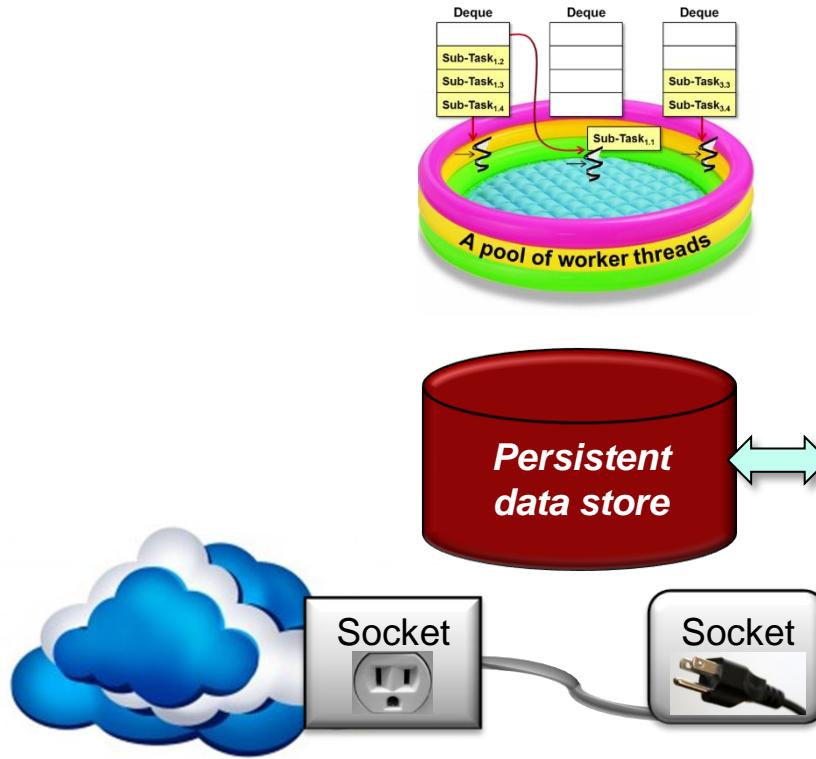
Java Parallel ImageStreamGang Example

Introduction

Douglas C. Schmidt

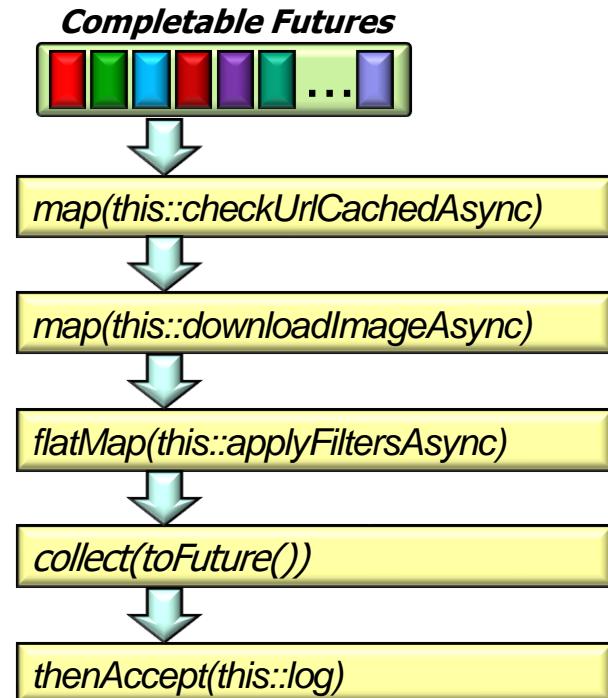
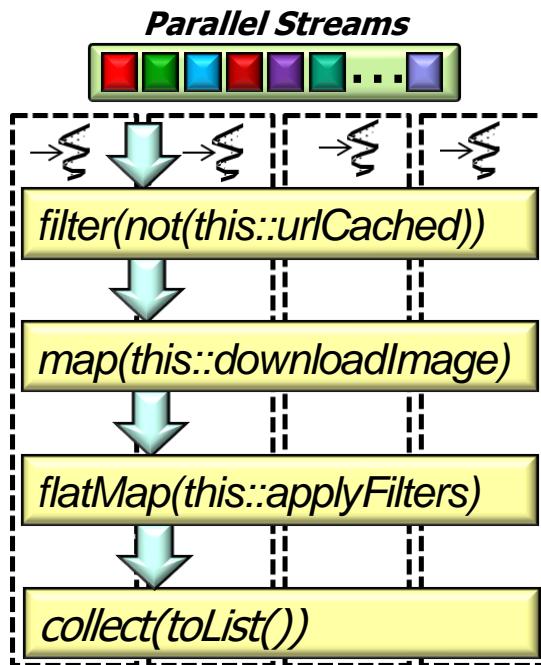
Learning Objectives in This Part of the Lesson

- Recognize the structure and functionality of the ImageStreamGang app.



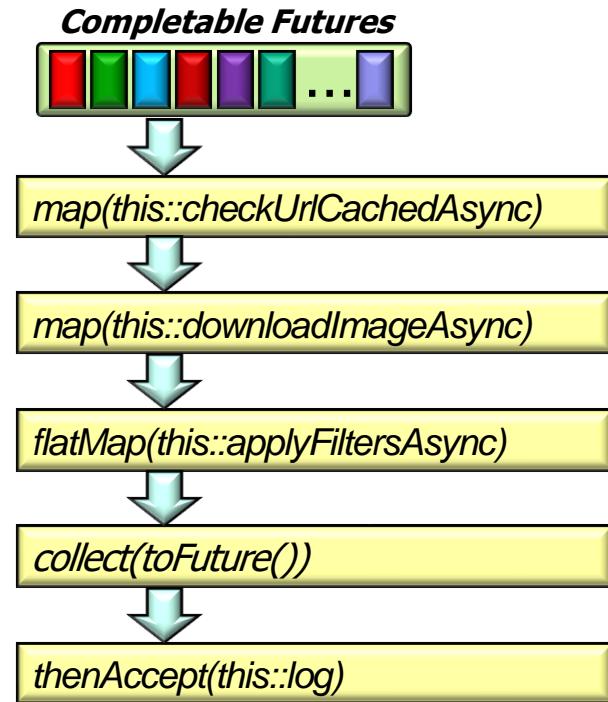
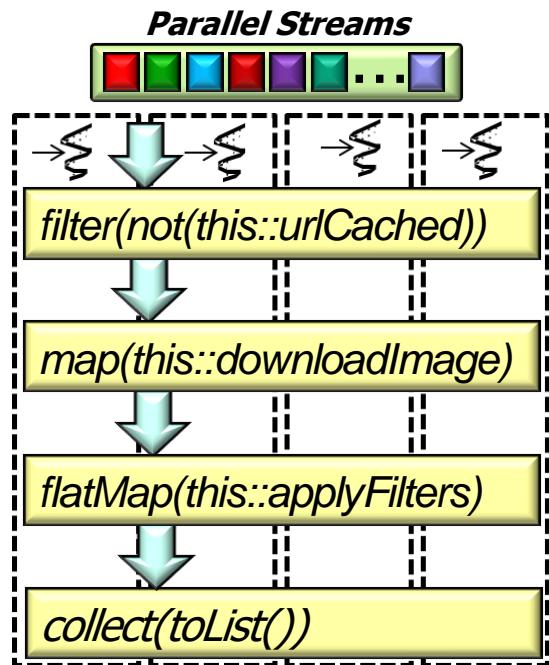
Learning Objectives in This Part of the Lesson

- Recognize the structure and functionality of the ImageStreamGang app.
 - It applies several Java parallelism frameworks.



Learning Objectives in This Part of the Lesson

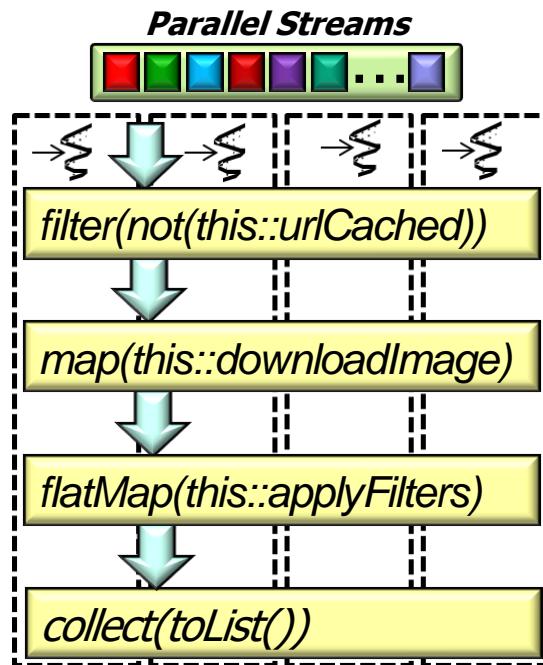
- Recognize the structure and functionality of the ImageStreamGang app.
 - It applies several Java parallelism frameworks.



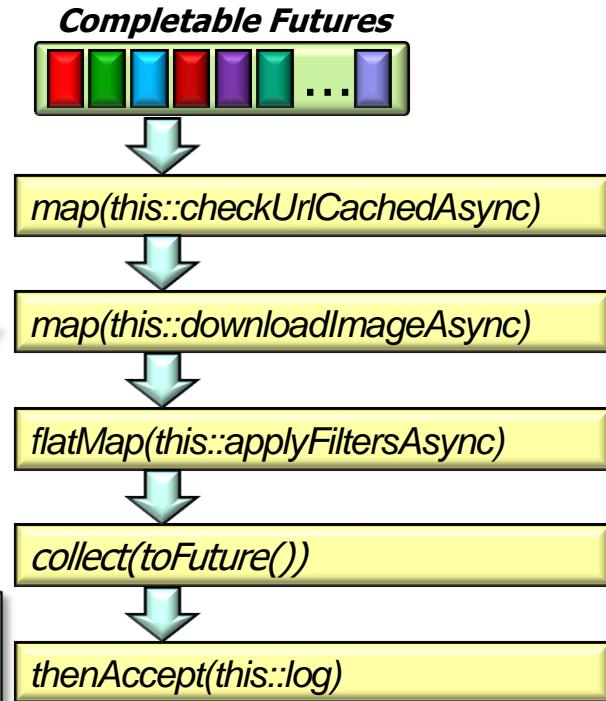
See docs.oracle.com/javase/tutorial/collections/stream/parallelism.html

Learning Objectives in This Part of the Lesson

- Recognize the structure and functionality of the ImageStreamGang app.
 - It applies several Java parallelism frameworks.

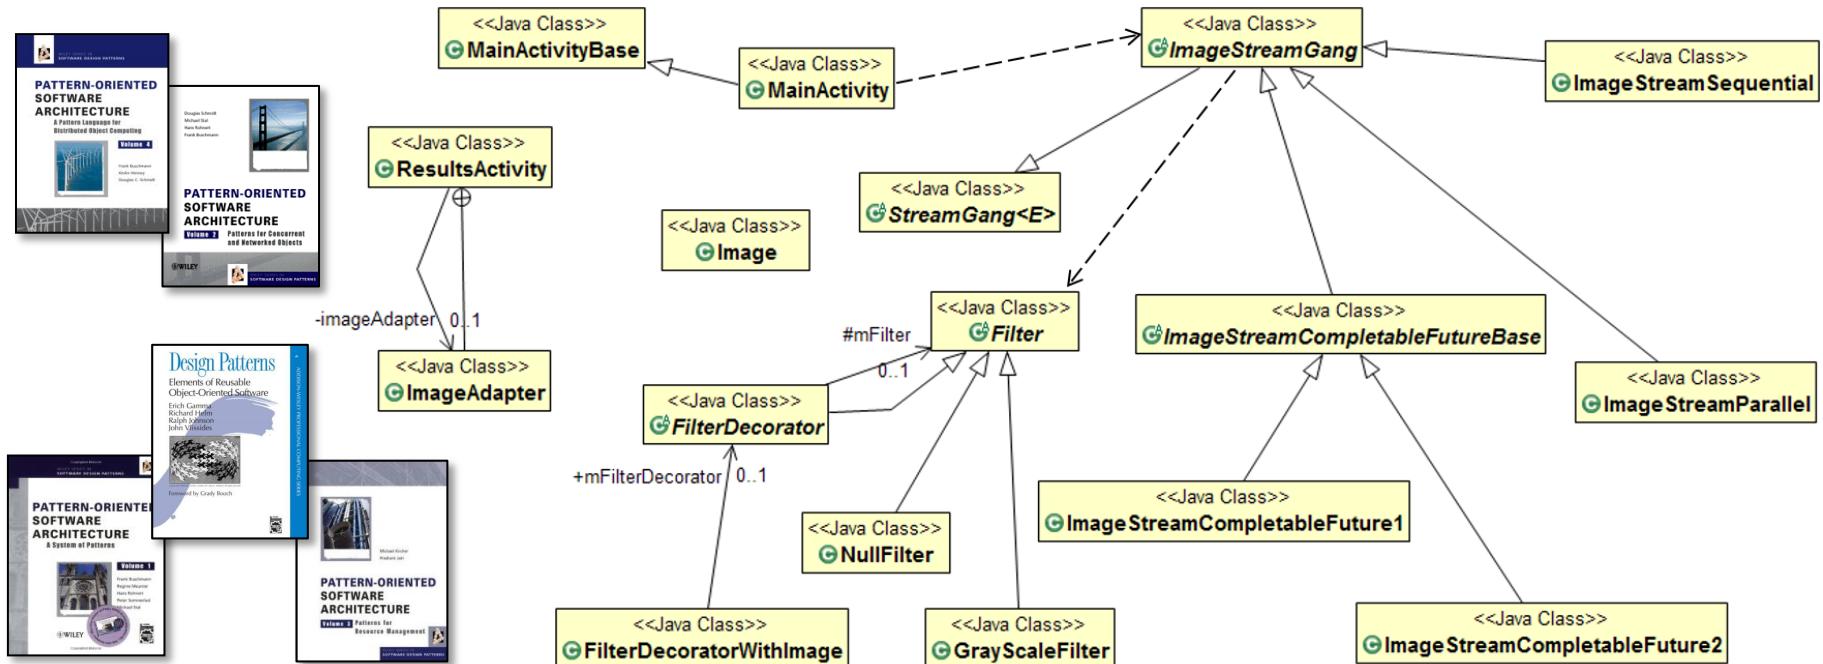


CompletableFuture may use ForkJoinPool framework.



Learning Objectives in This Part of the Lesson

- Recognize the structure and functionality of the ImageStreamGang app.
 - It applies several Java parallelism frameworks.
 - The focus is on integrating object-oriented and functional programming paradigms.

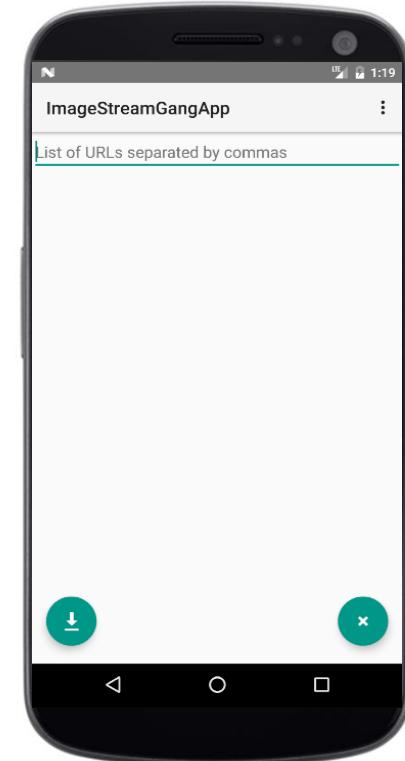
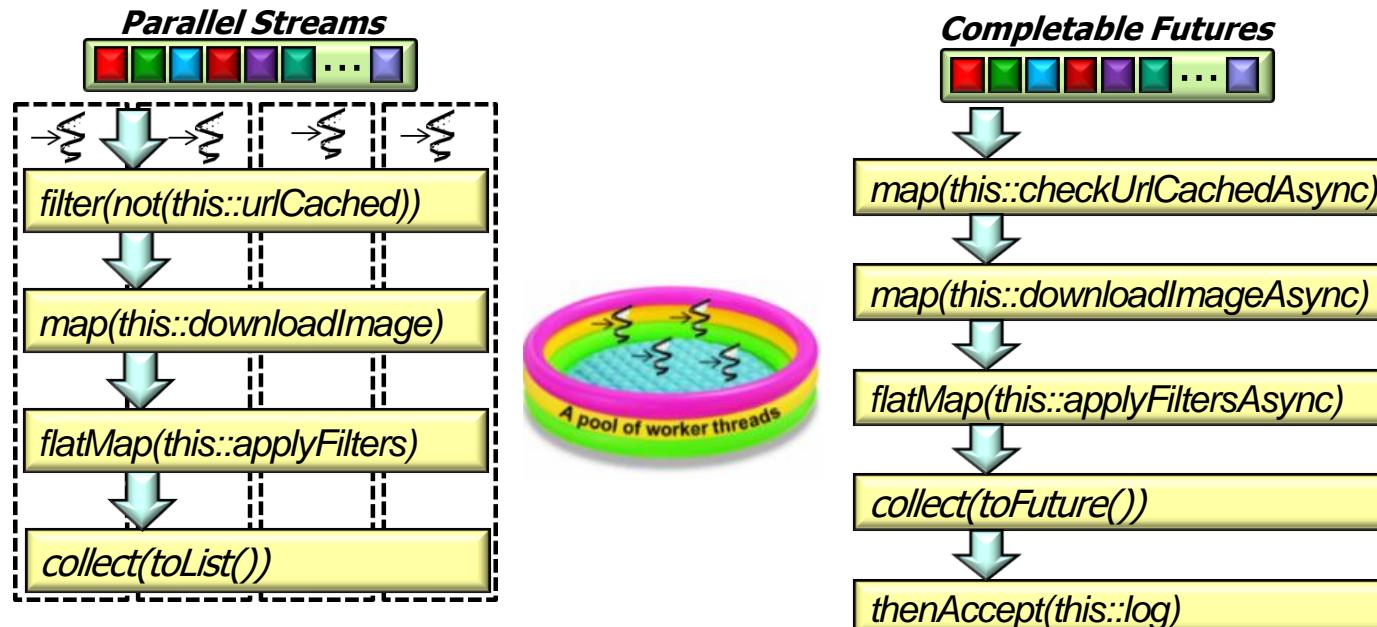


Patterns are used to emphasize key roles and responsibilities in the app's design.

Overview of the Pattern-Oriented ImageStreamGang App

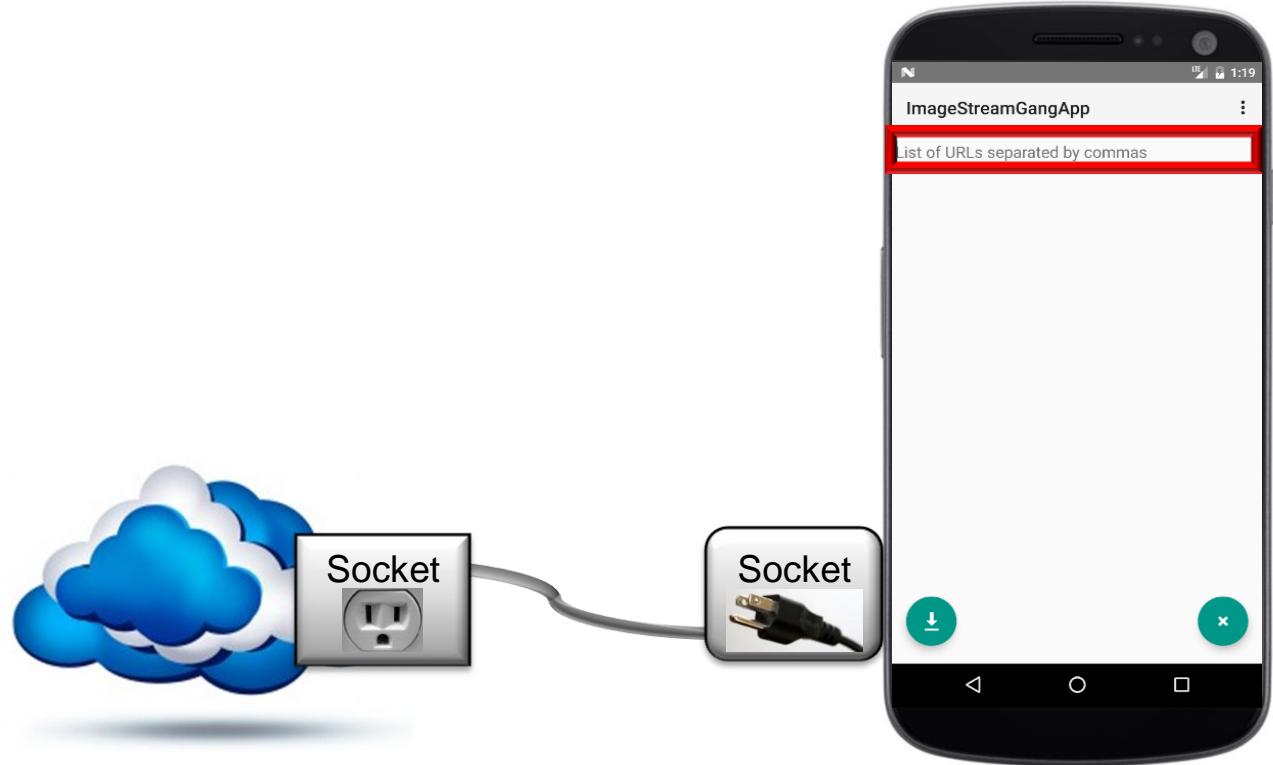
Overview of the Pattern-Oriented ImageStreamGang App

- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images.



Overview of the Pattern-Oriented ImageStreamGang App

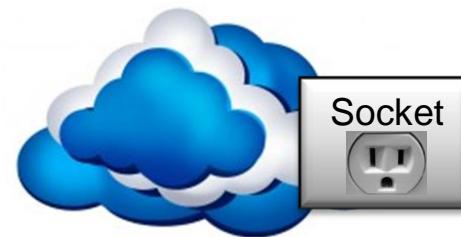
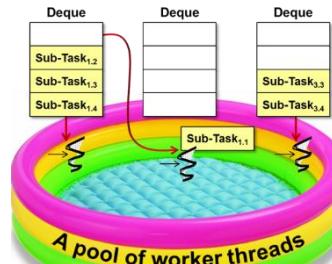
- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images, for example.



Prompt user for list of URLs to download.

Overview of the Pattern-Oriented ImageStreamGang App

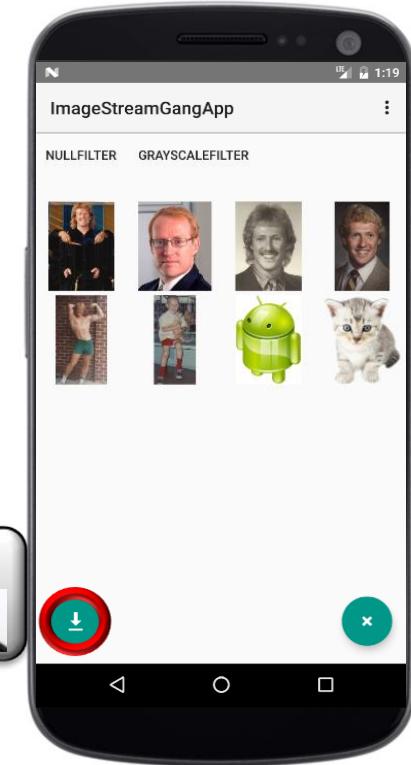
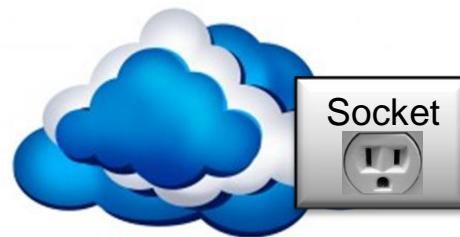
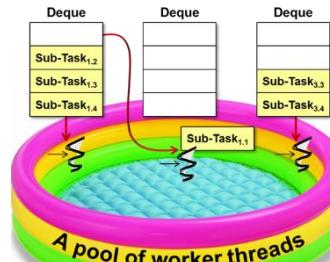
- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images, for example.



User supplies the list of URLs to download.

Overview of the Pattern-Oriented ImageStreamGang App

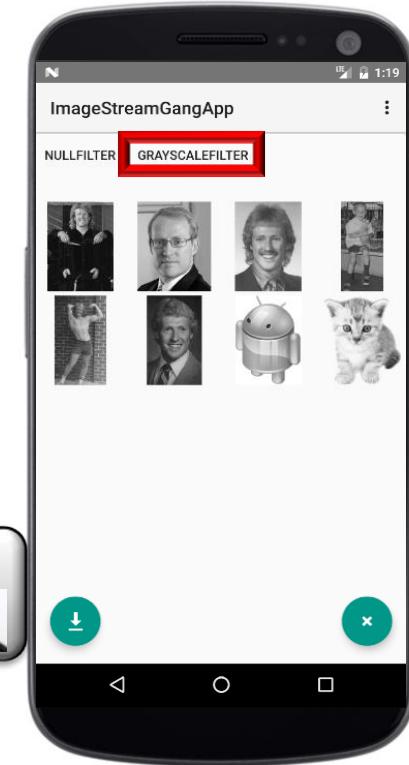
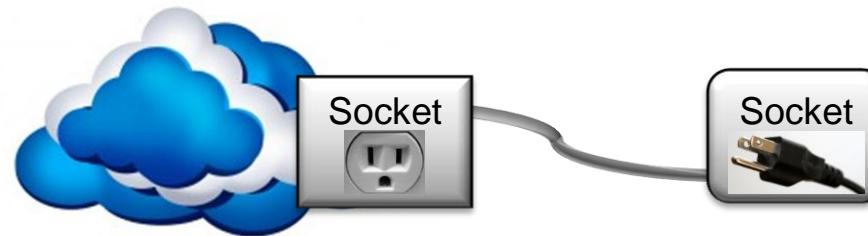
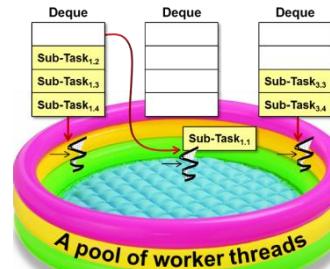
- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images, for example.



Download images via one or more threads.

Overview of the Pattern-Oriented ImageStreamGang App

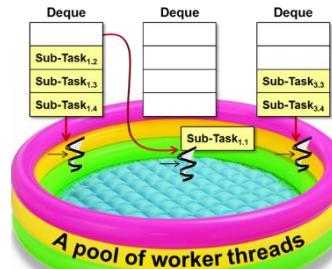
- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images, for example.



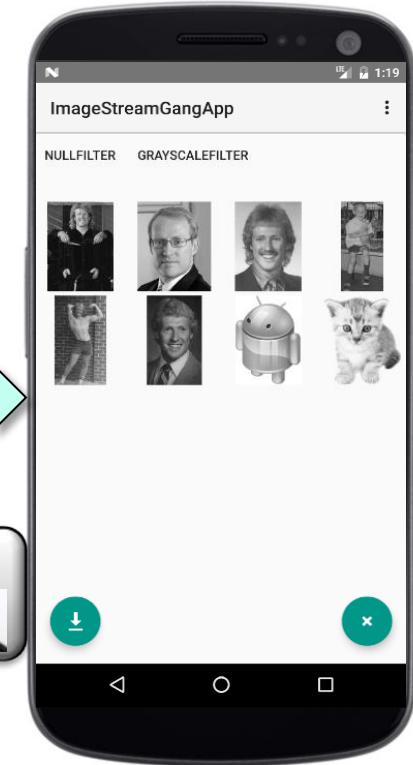
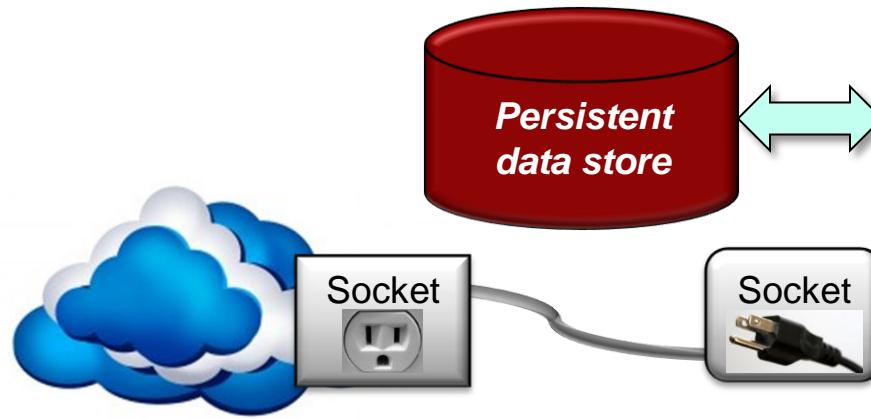
Apply filters to transform downloaded images.

Overview of the Pattern-Oriented ImageStreamGang App

- This app combines streams and CompletableFutures with the StreamGang framework to download, transform, store, and display images, for example.



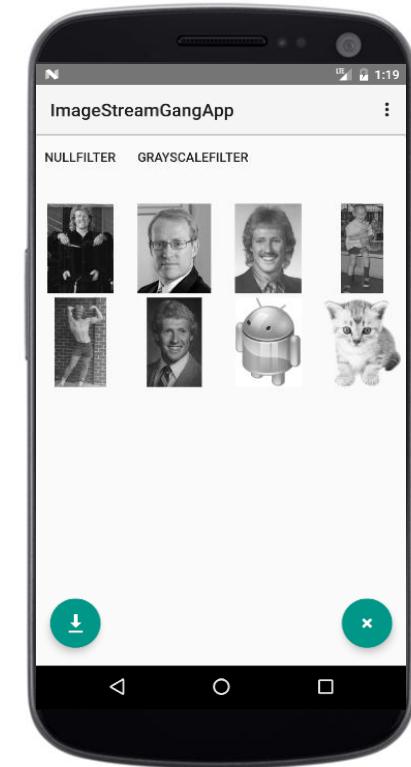
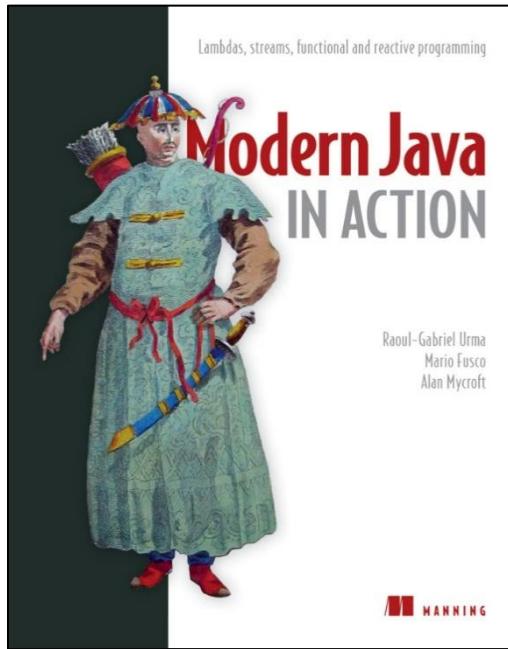
List of filters to apply



Output filtered images to persistent storage.

Overview of the Pattern-Oriented ImageStreamGang App

- The ImageStreamGang app applies a range of modern Java features.

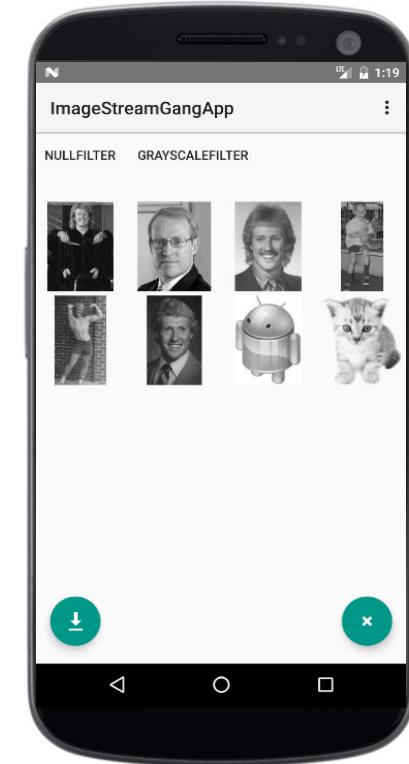


See www.manning.com/books/modern-java-in-action

Overview of the Pattern-Oriented ImageStreamGang App

- The ImageStreamGang app applies a range of modern Java features, e.g.
 - Sequential and parallel streams

```
List<Image> filteredImages =  
    getInput()  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::downloadImage)  
        .flatMap(this::applyFilters)  
        .collect(toList());
```

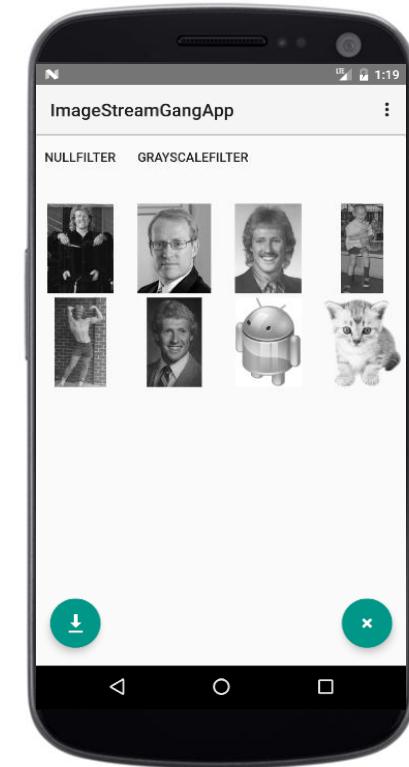


We'll cover parallel streams now.

Overview of the Pattern-Oriented ImageStreamGang App

- The ImageStreamGang app applies a range of modern Java features, e.g.
 - Sequential and parallel streams
 - CompletableFuture

```
getInput()
    .stream()
    .map(this::checkUrlCachedAsync)
    .map(this::downloadImageAsync)
    .flatMap(this::applyFiltersAsync)
    .collect(toFuture())
    .thenAccept
        (stream ->
            log(stream.flatMap(Optional::stream) ,
                urls.size()))
    .join();
```

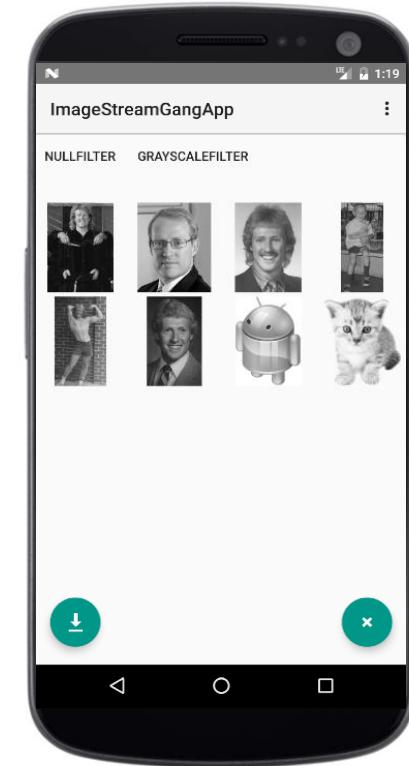


We'll cover CompletableFuture later.

Overview of the Pattern-Oriented ImageStreamGang App

- The ImageStreamGang app applies a range of modern Java features, e.g.
 - Sequential and parallel streams
 - CompletableFuture
 - Lambda expressions and method references

```
Runnable mCompletionHook =  
    () -> MainActivity.this.runOnUiThread  
        (this::goToResultActivity);
```



We covered these foundational Java features earlier.

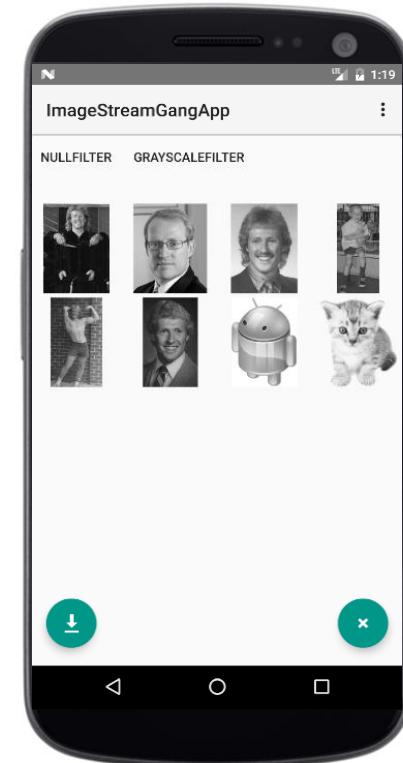
Overview of the Pattern-Oriented ImageStreamGang App

- The ImageStreamGang app applies a range of modern Java features, e.g.
 - Sequential and parallel streams
 - CompletableFuture
 - Lambda expressions and method references

```
Runnable mCompletionHook =  
    () -> MainActivity.this.runOnUiThread  
        (this::goToResultActivity);
```

versus

```
Runnable mCompletionHook = new Runnable() {  
    public void run() {  
        MainActivity.this.runOnUiThread  
            (new Runnable() { public void run()  
                { goToResultActivity(); } } ); } };
```



Overview of Patterns Applied in the ImageStreamGang App

Overview of Patterns Applied in the ImageStreamGang App

- “Gang-of-Four” and POSA patterns are applied to enhance its framework-based architecture.



See en.wikipedia.org/wiki/Design_Patterns & www.dre.vanderbilt.edu/~schmidt/POSA

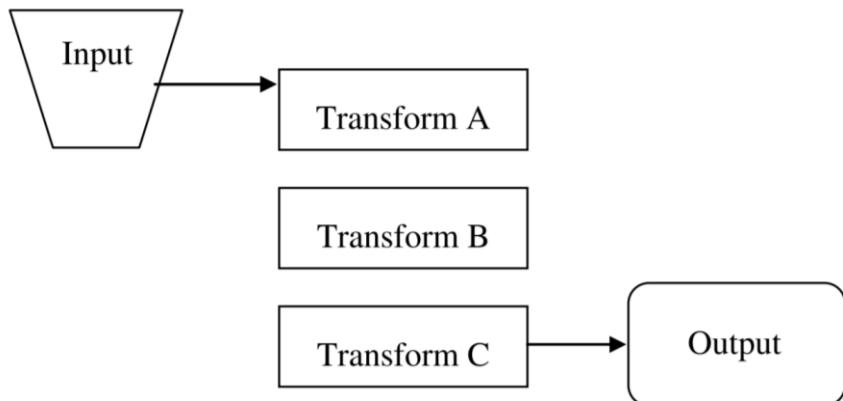
Overview of Patterns Applied in the ImageStreamGang App

- Some patterns are essential to its design.



Overview of Patterns Applied in the ImageStreamGang App

- Some patterns are essential to its design.
 - *Pipes and filters*
 - Divide application's tasks into several self-contained data processing steps and connect these steps via intermediate data buffers to a data processing pipeline



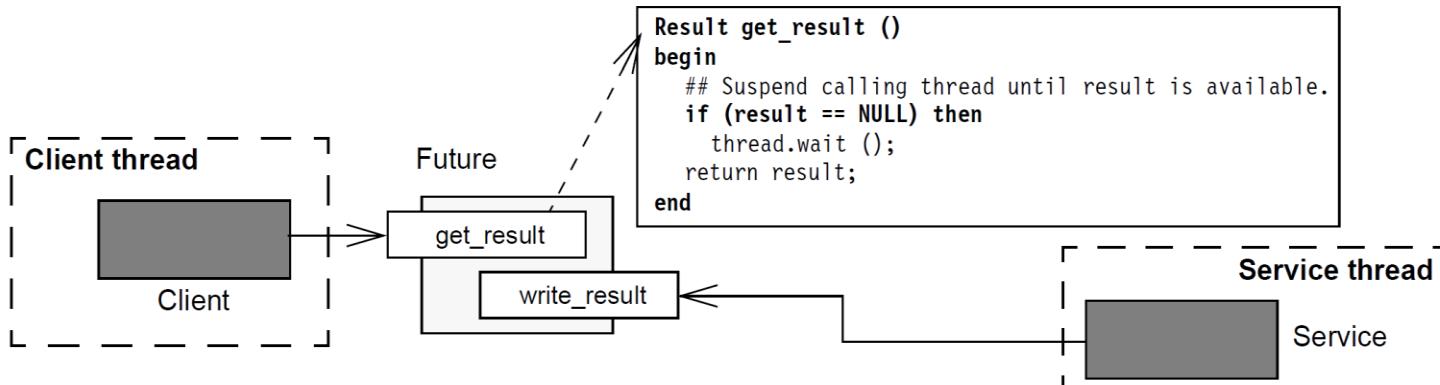
See www.hillside.net/plop/2011/papers/B-10-Hanmer.pdf

Overview of Patterns Applied in the ImageStreamGang App

- Some patterns are essential to its design.

- *Future*

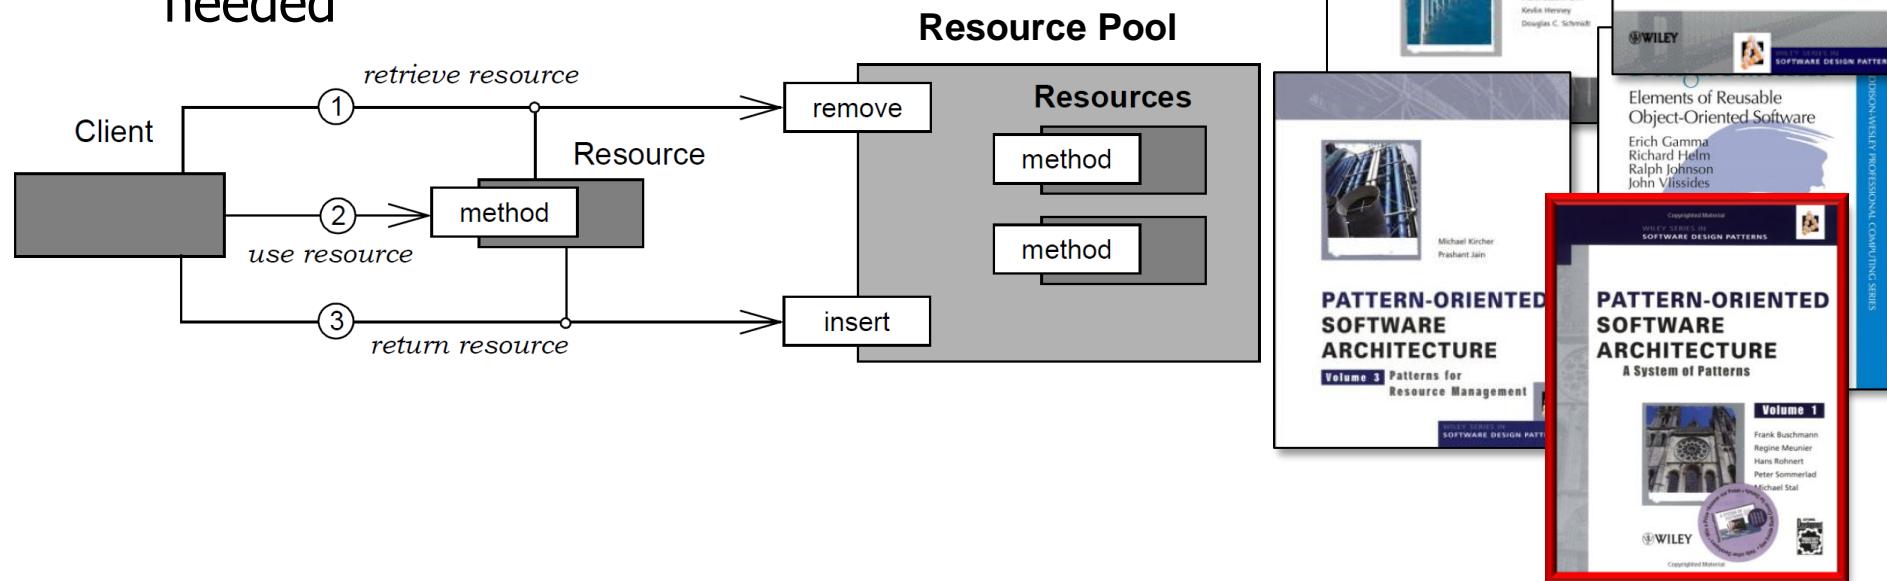
- Provides a proxy to a client when it invokes a service to keep track of the state of the service's concurrent computation and only returns a value to the client when the computation completes



See en.wikipedia.org/wiki/Futures_and_promises

Overview of Patterns Applied in the ImageStreamGang App

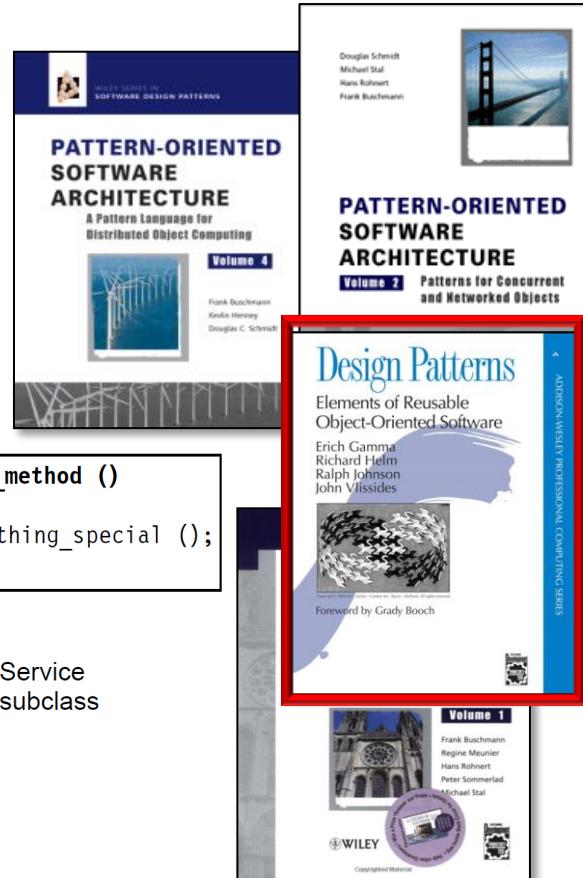
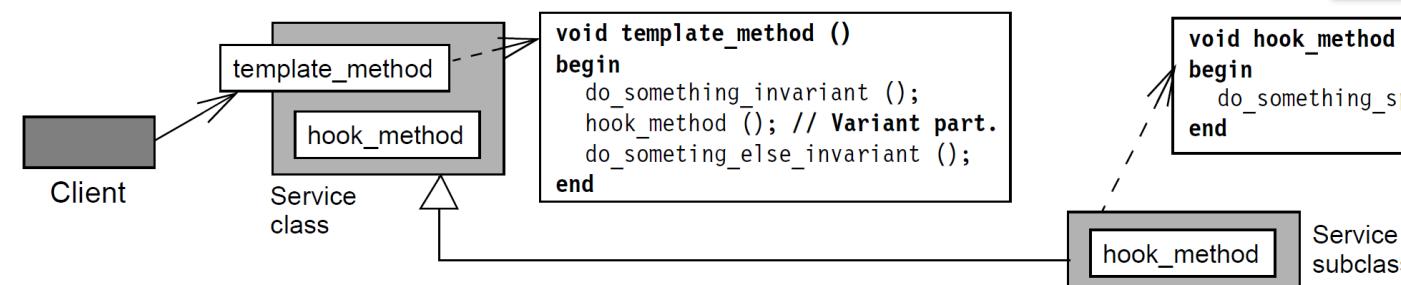
- Some patterns are essential to its design.
 - *ResourcePool*
 - Prevents expensive acquisition and release of resources by recycling resources no longer needed



See kircher-schwanninger.de/michael/publications/Pooling.pdf

Overview of Patterns Applied in the ImageStreamGang App

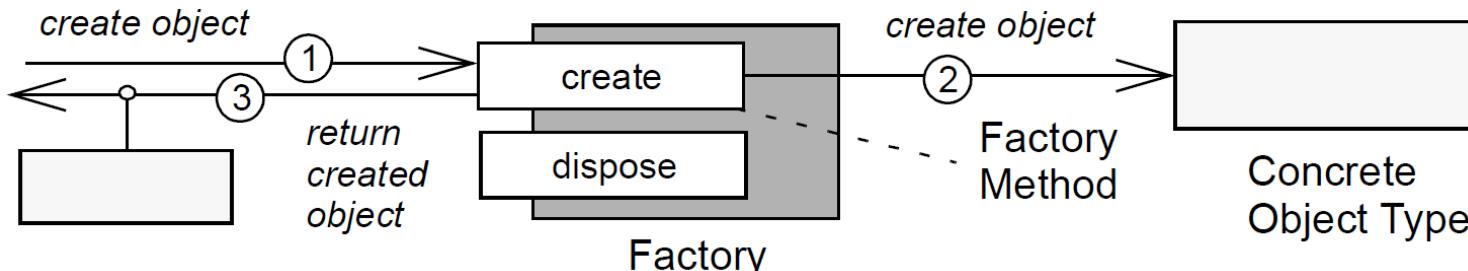
- Some patterns are essential to its design.
 - *Template method*
 - Defines the overall structure of a method, while allowing subclasses to refine, or redefine, certain steps



See en.wikipedia.org/wiki/Template_method_pattern

Overview of Patterns Applied in the ImageStreamGang App

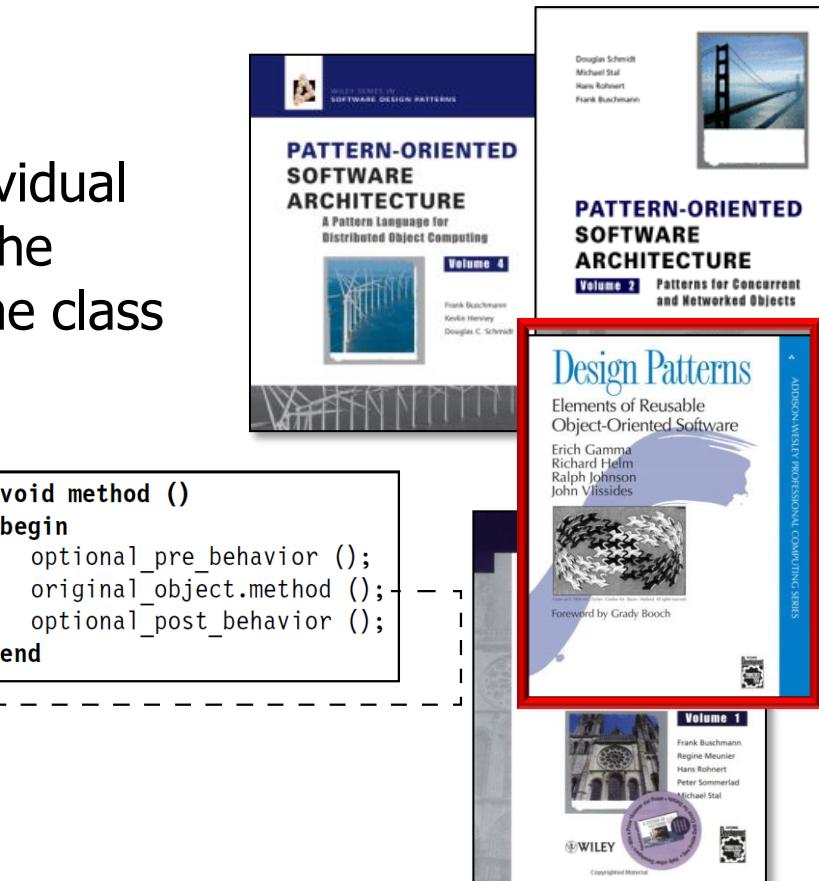
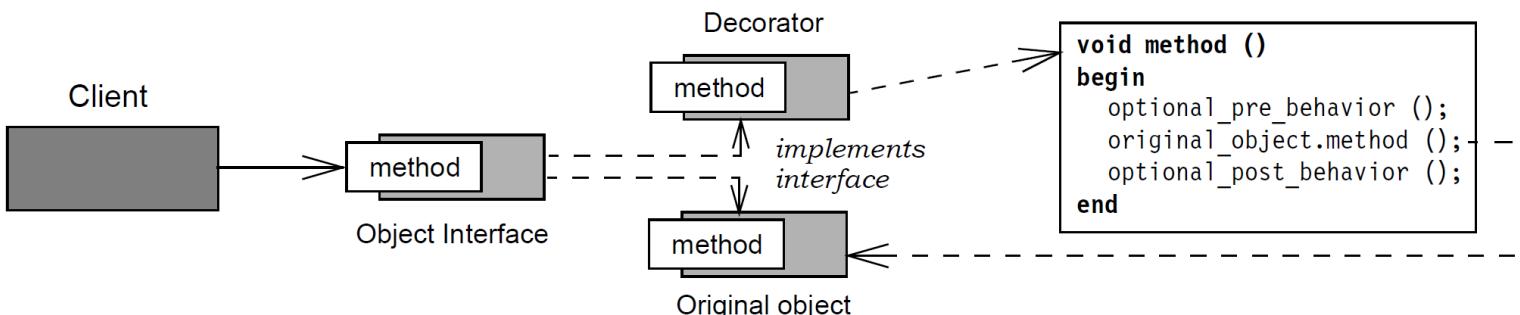
- Some patterns are essential to its design.
 - *Factory method*
 - Encapsulate the concrete details of object creation inside a factory method, rather than letting clients create the object themselves



See en.wikipedia.org/wiki/Factory_method_pattern

Overview of Patterns Applied in the ImageStreamGang App

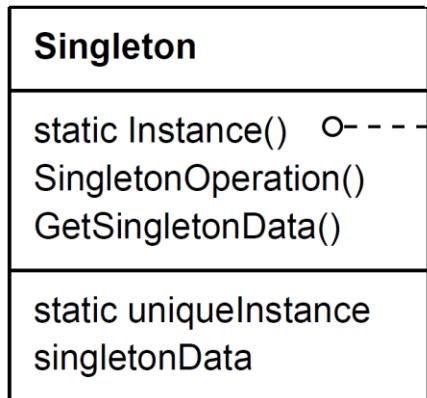
- Some patterns are essential to its design.
 - *Decorator*
 - Allows behavior to be added to an individual object, dynamically, without affecting the behavior of other objects from the same class



See en.wikipedia.org/wiki/Decorator_pattern

Overview of Patterns Applied in the ImageStreamGang App

- Other patterns are also applied.
 - *Singleton*
 - Ensures a class has only one instance and provide a global point of access to it



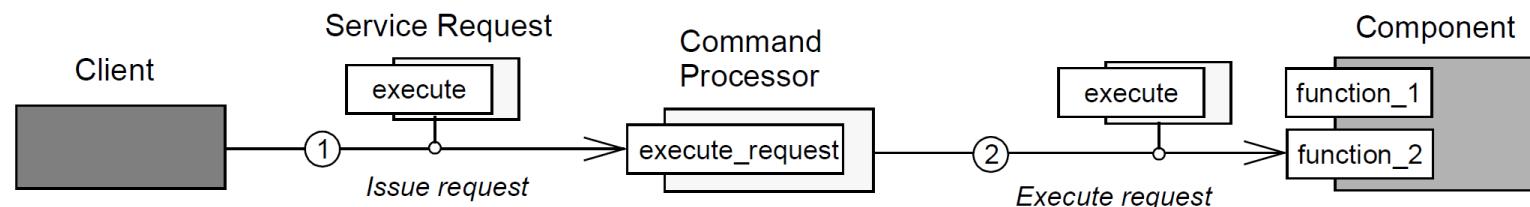
return uniqueInstance



See en.wikipedia.org/wiki/Singleton_pattern

Overview of Patterns Applied in the ImageStreamGang App

- Other patterns are also applied.
 - *CommandProcessor*
 - Packages a piece of application functionality—as well as its parameterization in an object—to make it usable in another context, such as later in time or in a different thread

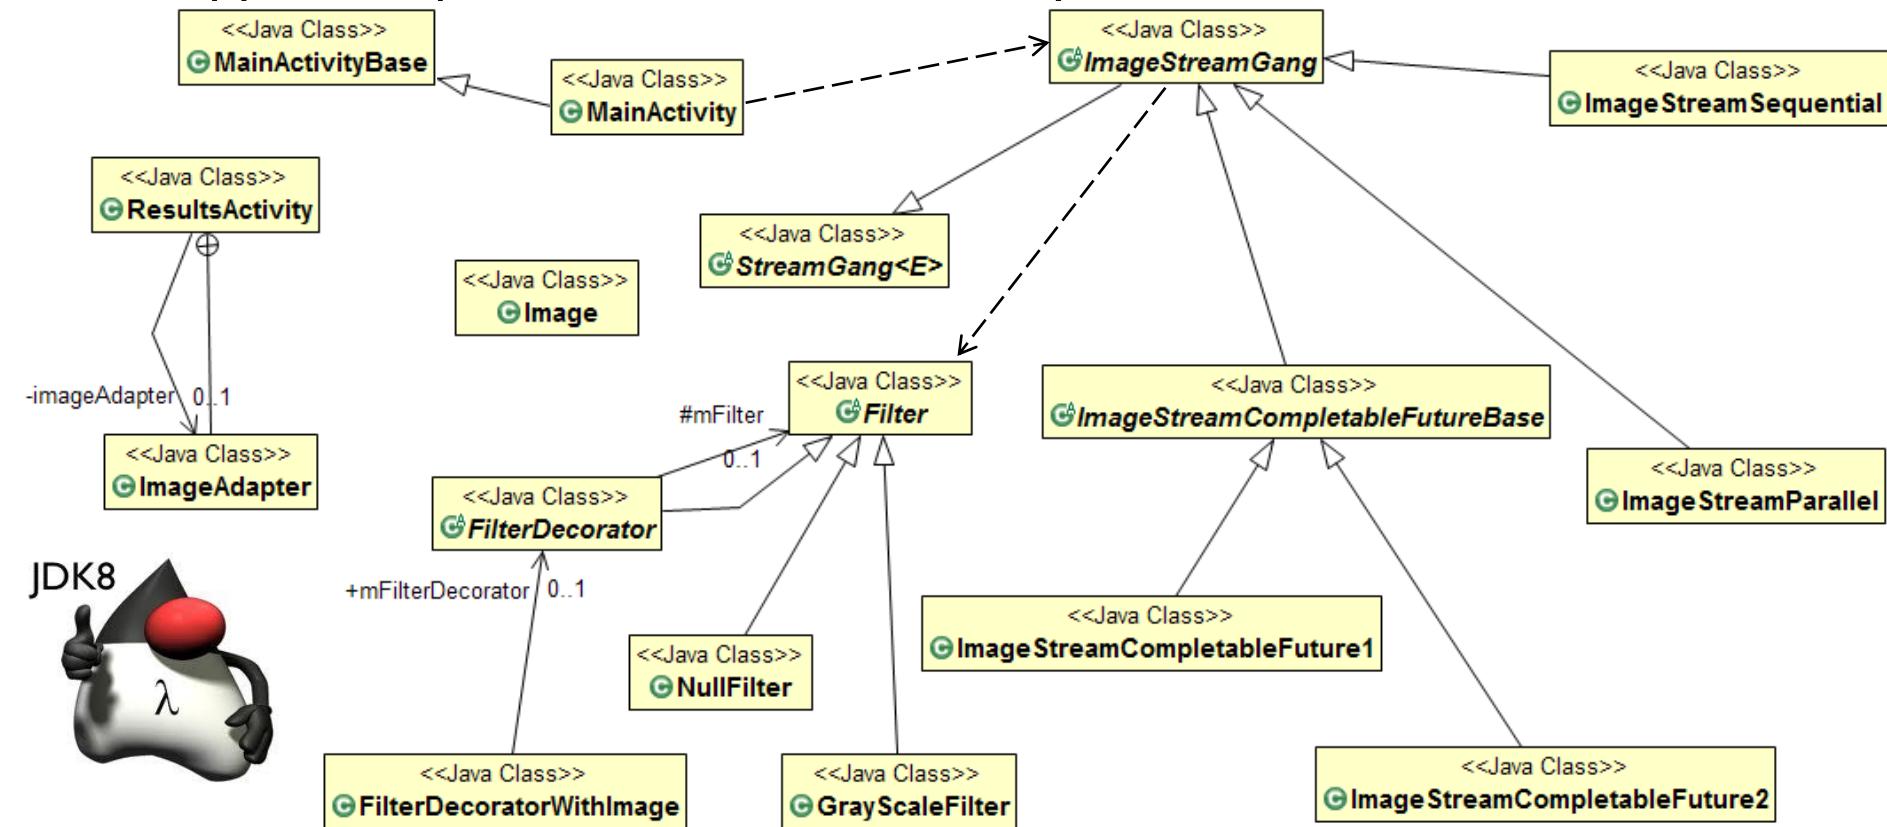


See www.dre.vanderbilt.edu/~schmidt/CommandProcessor.pdf

Strategy for Understanding the ImageStreamGang App

Strategy for Understanding the ImageStreamGang App

- This app is complicated and contains many classes.



Strategy for Understanding the ImageStreamGang App

- This app is complicated and contains many classes.
 - We therefore analyze it from various perspectives.



Including pattern-oriented design, data flows, and detailed code walk-throughs

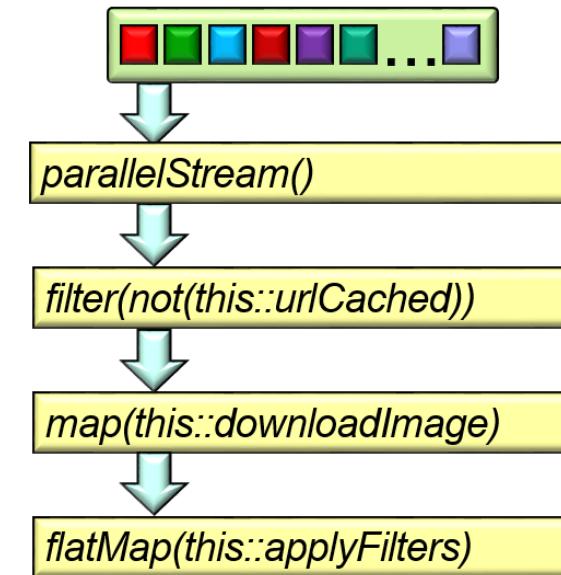
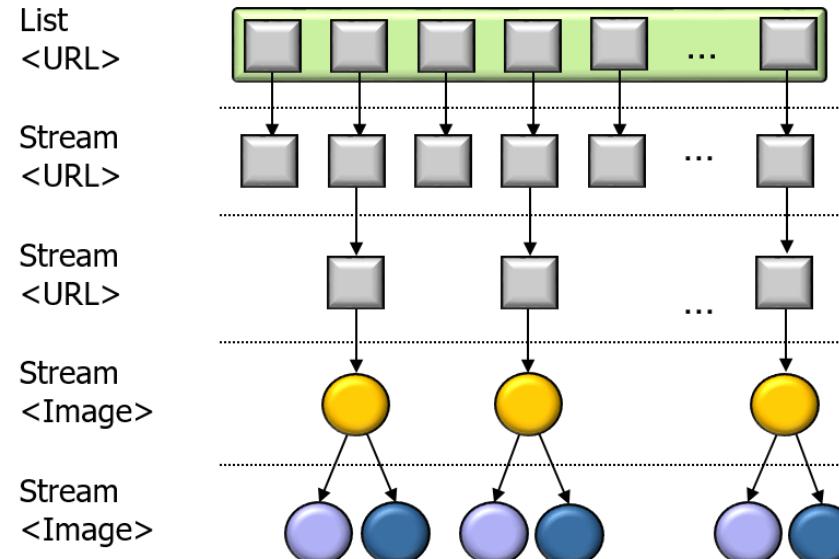
Strategy for Understanding the ImageStreamGang App

- This app is complicated and contains many classes.
 - We therefore analyze it from various perspectives.
 - Watch this entire lesson carefully to understand how it all works.



Strategy for Understanding the ImageStreamGang App

- This app is complicated and contains many classes.
 - We therefore analyze it from various perspectives.
 - Watch this entire lesson carefully to understand how it all works.
 - Visualize the data flow in a parallel stream.



Strategy for Understanding the ImageStreamGang App

- This app is complicated and contains many classes.
 - We therefore analyze it from various perspectives.
 - Watch this entire lesson carefully to understand how it all works.
 - Visualize the data flow in a parallel stream.
 - Run/read the code to see how it all works.

USE THE
SOURCE LUKE!
SOURCE FORGE!



Java Parallel ImageStreamGang Example: Introduction

The End

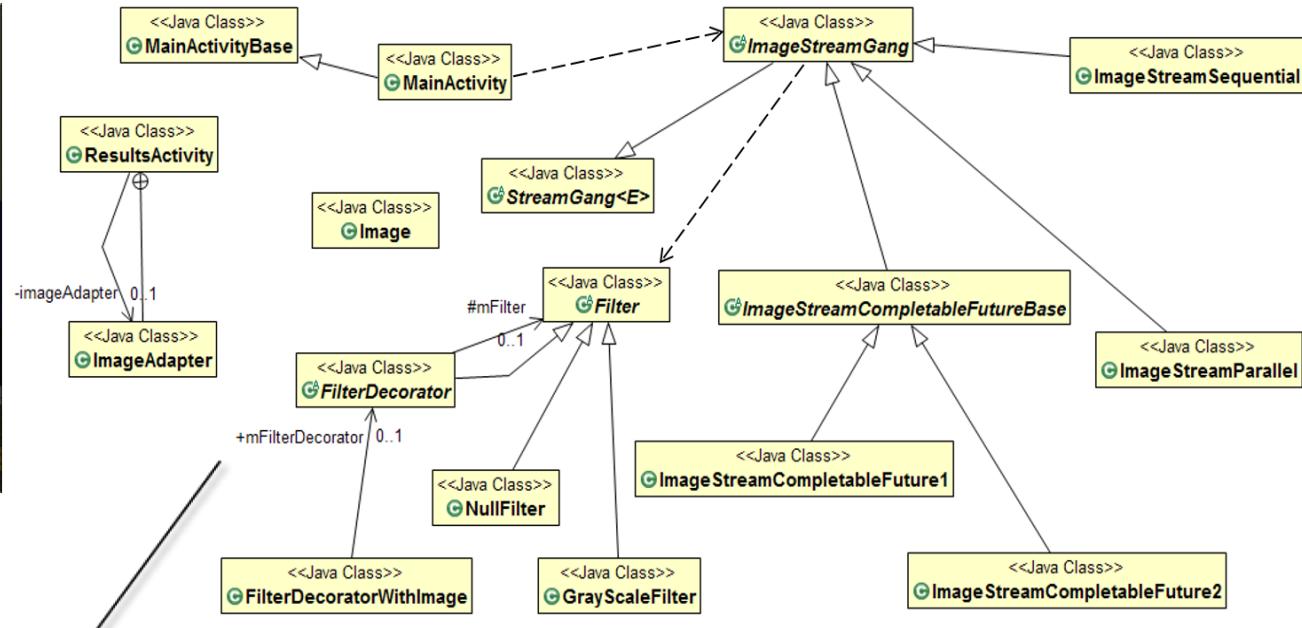
Java Parallel ImageStreamGang Example

Structure and Functionality

Douglas C. Schmidt

Learning Objectives in This Part of the Lesson

- Recognize the structure and functionality of the ImageStreamGang app.
 - It applies several Java parallelism frameworks.
 - The focus is on integrating object-oriented and functional programming paradigms.

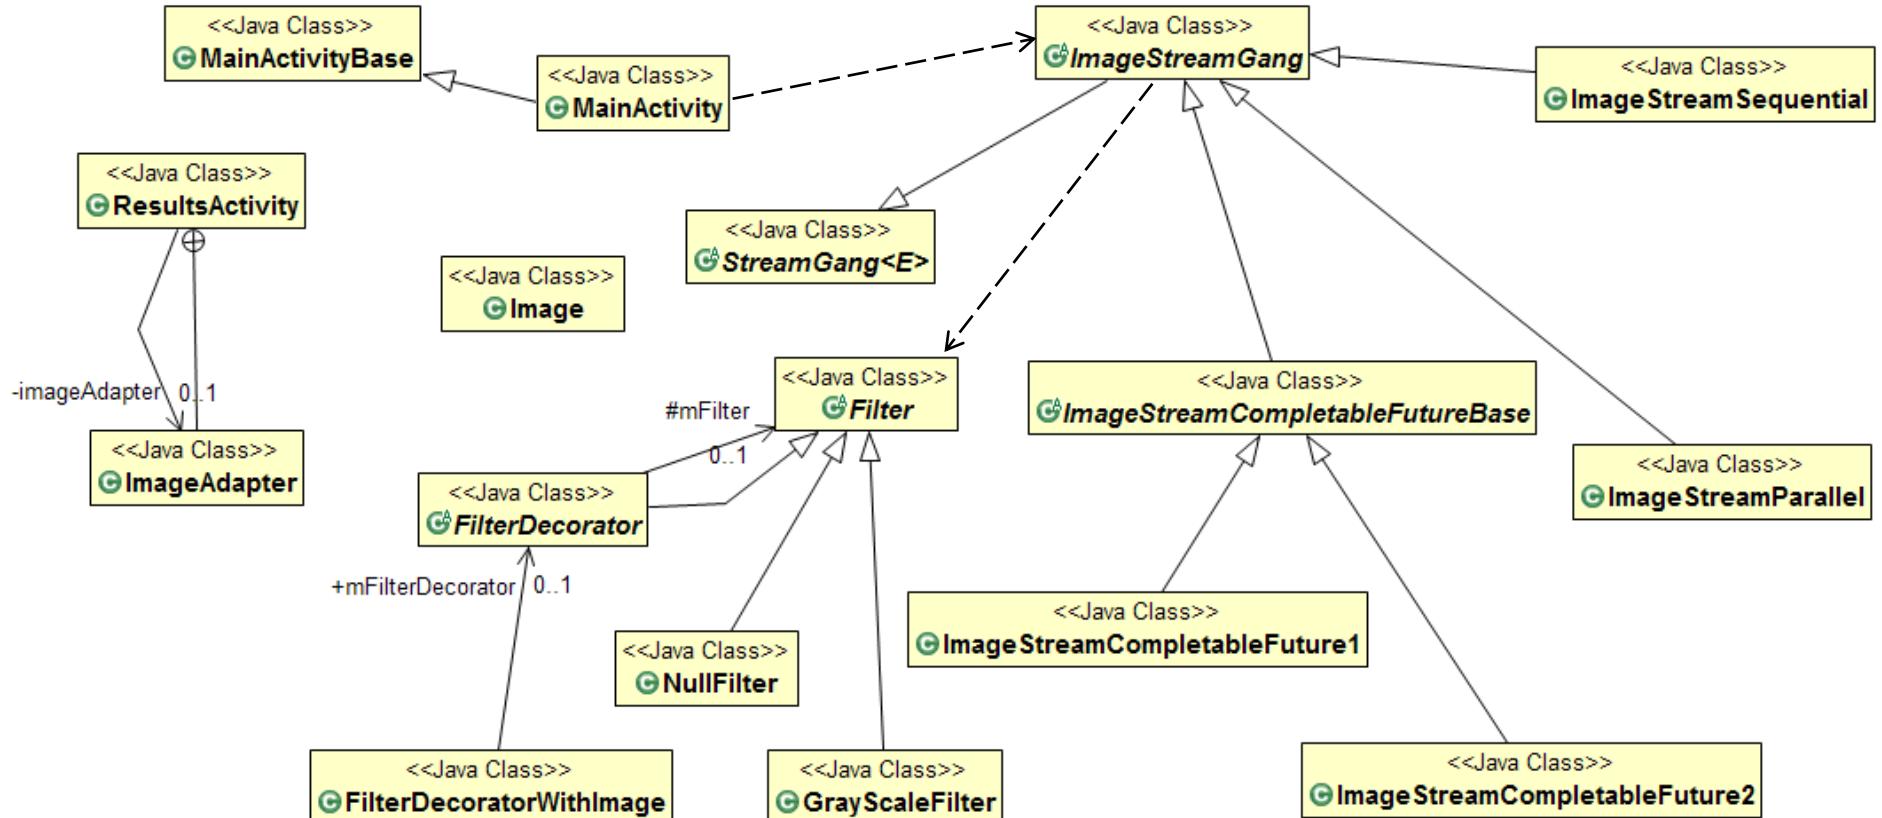


This design shows the synergy between object-oriented and functional programming.

The Structure of the ImageStreamGang App

The Structure of the ImageStreamGang App

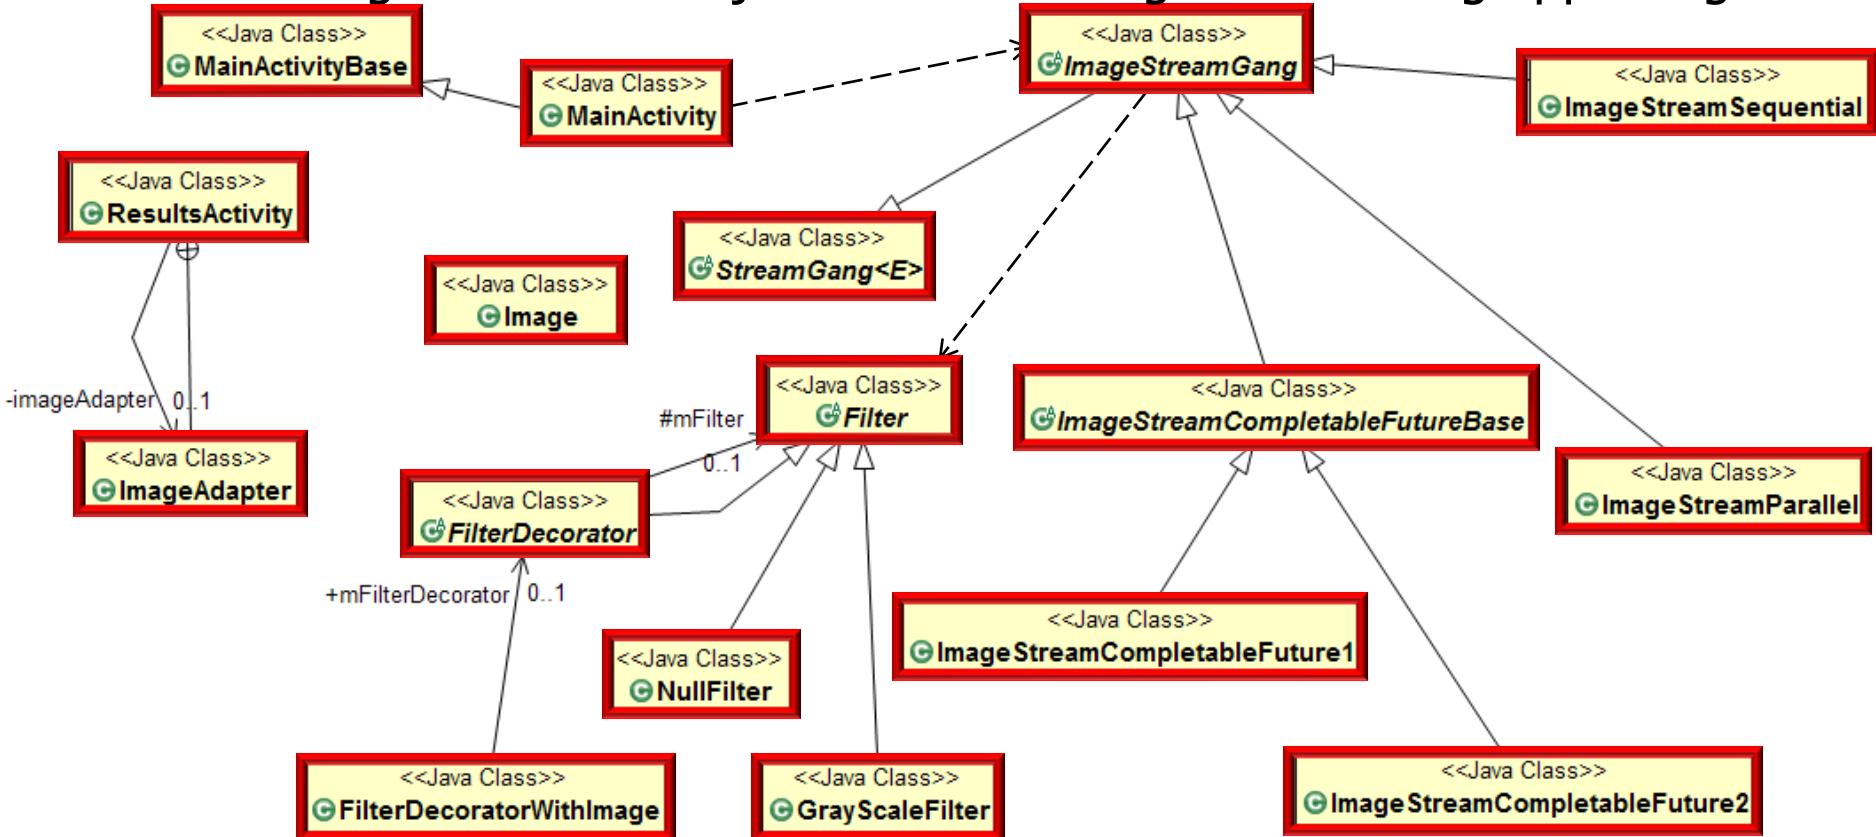
- UML class diagram for the object-oriented ImageStreamGang app design



See en.wikipedia.org/wiki/Unified_Modeling_Language

The Structure of the ImageStreamGang App

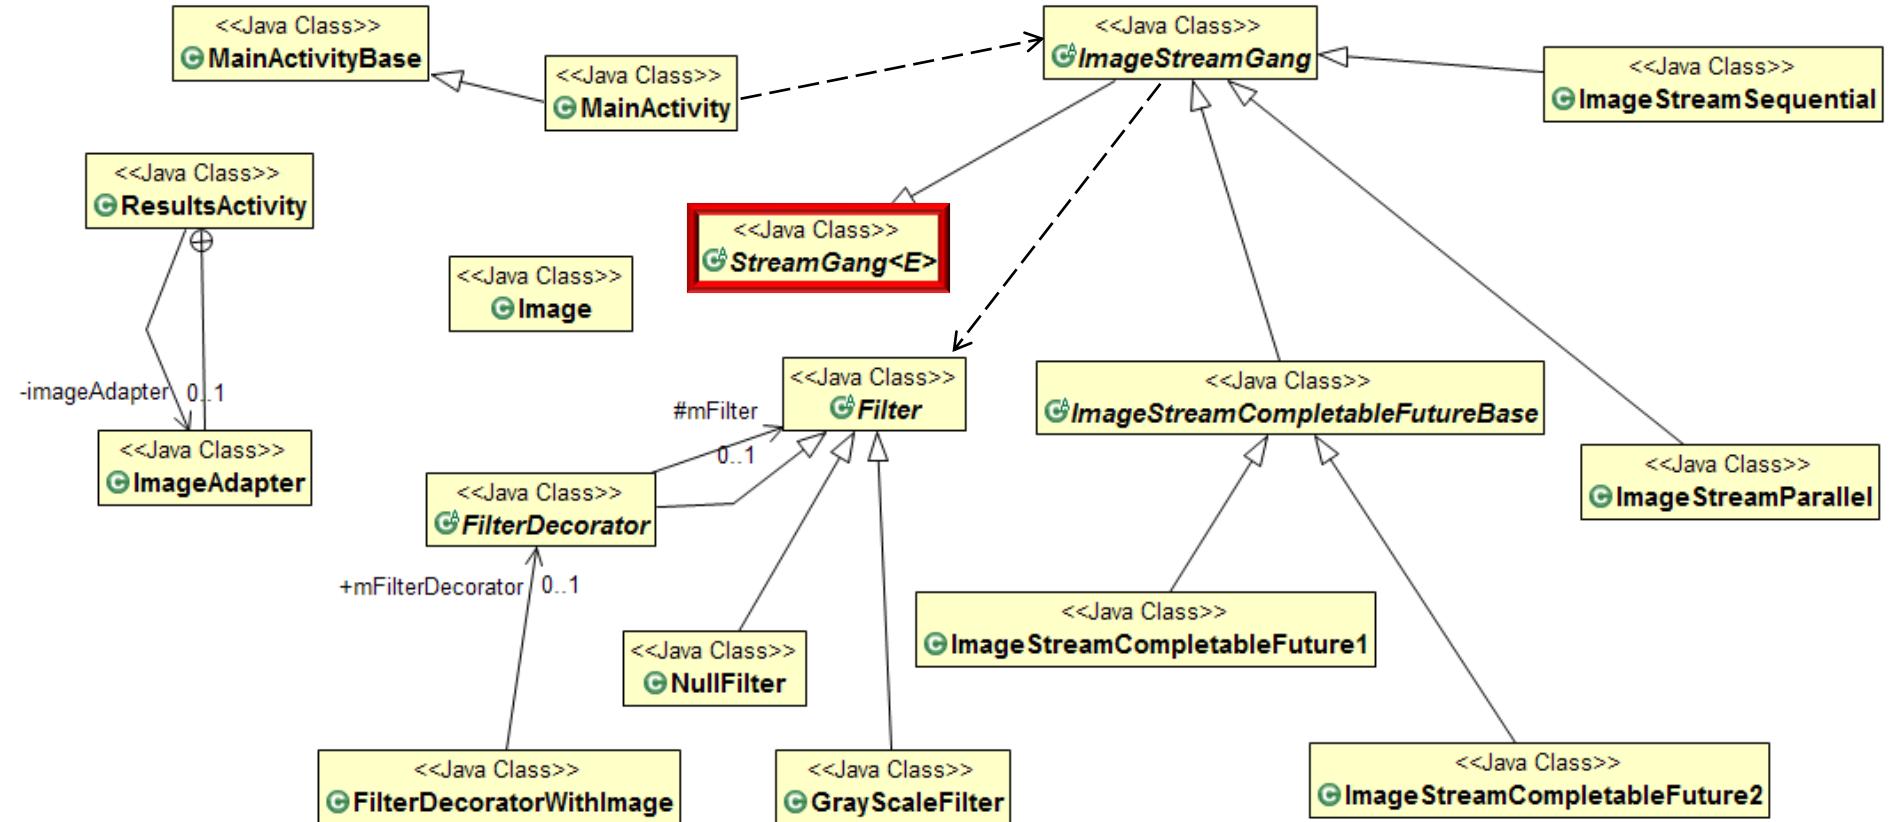
- UML class diagram for the object-oriented ImageStreamGang app design



These classes apply Java features to image downloading and processing.

The Structure of the ImageStreamGang App

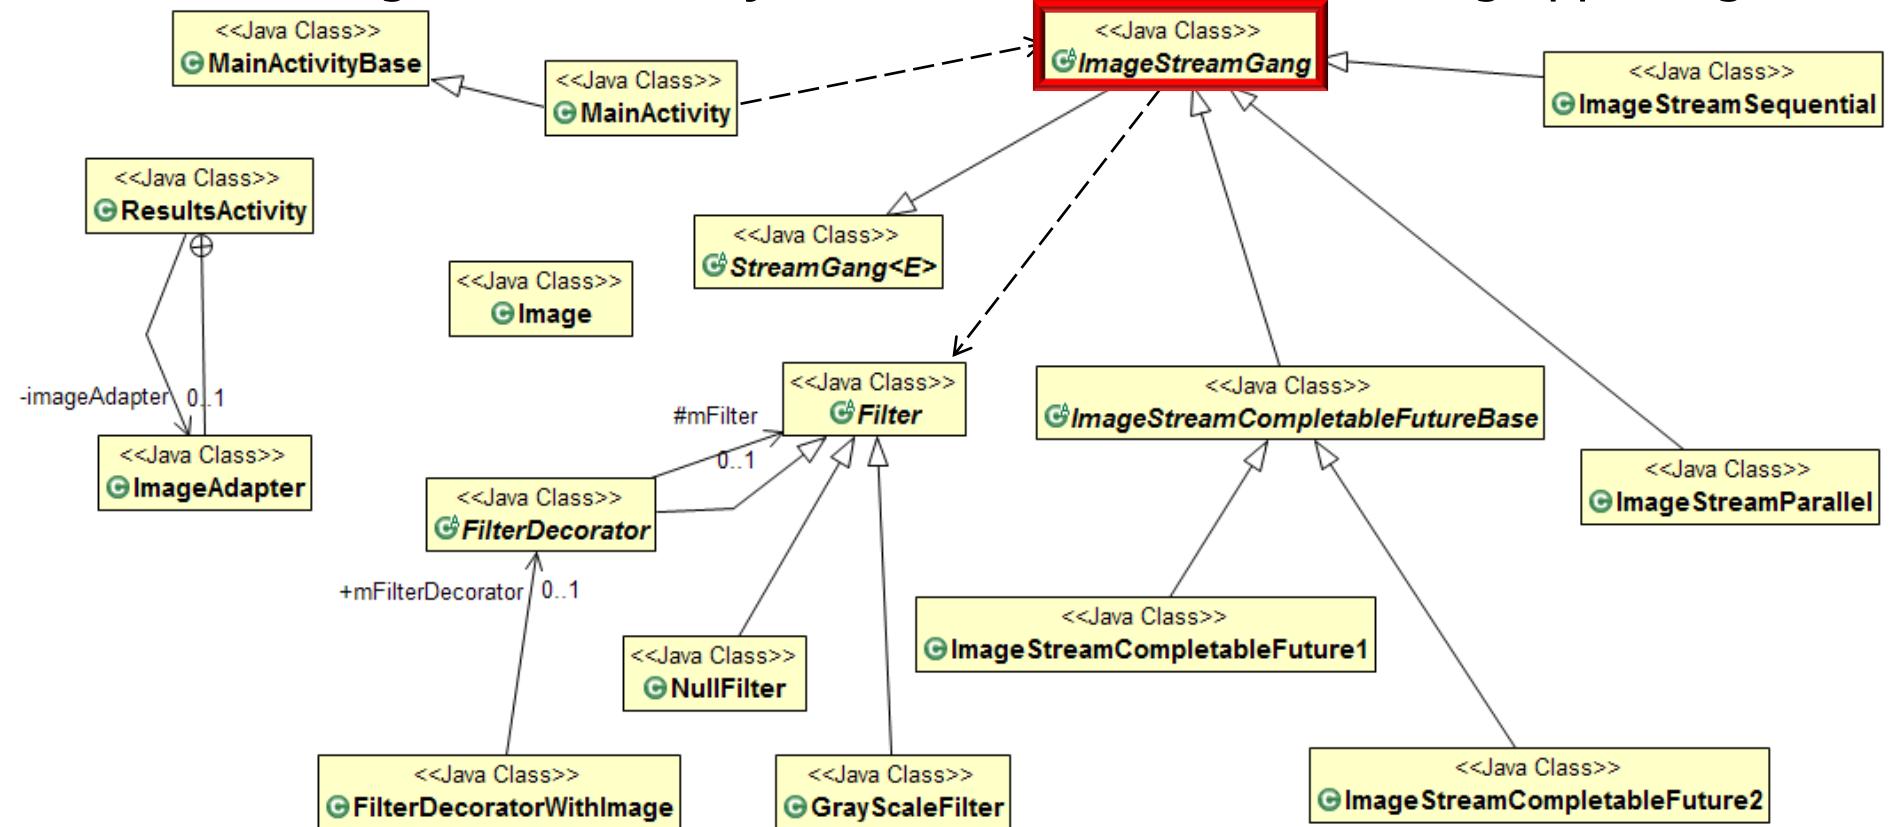
- UML class diagram for the object-oriented ImageStreamGang app design



A framework for initiating streams that process input from a list of elements

The Structure of the ImageStreamGang App

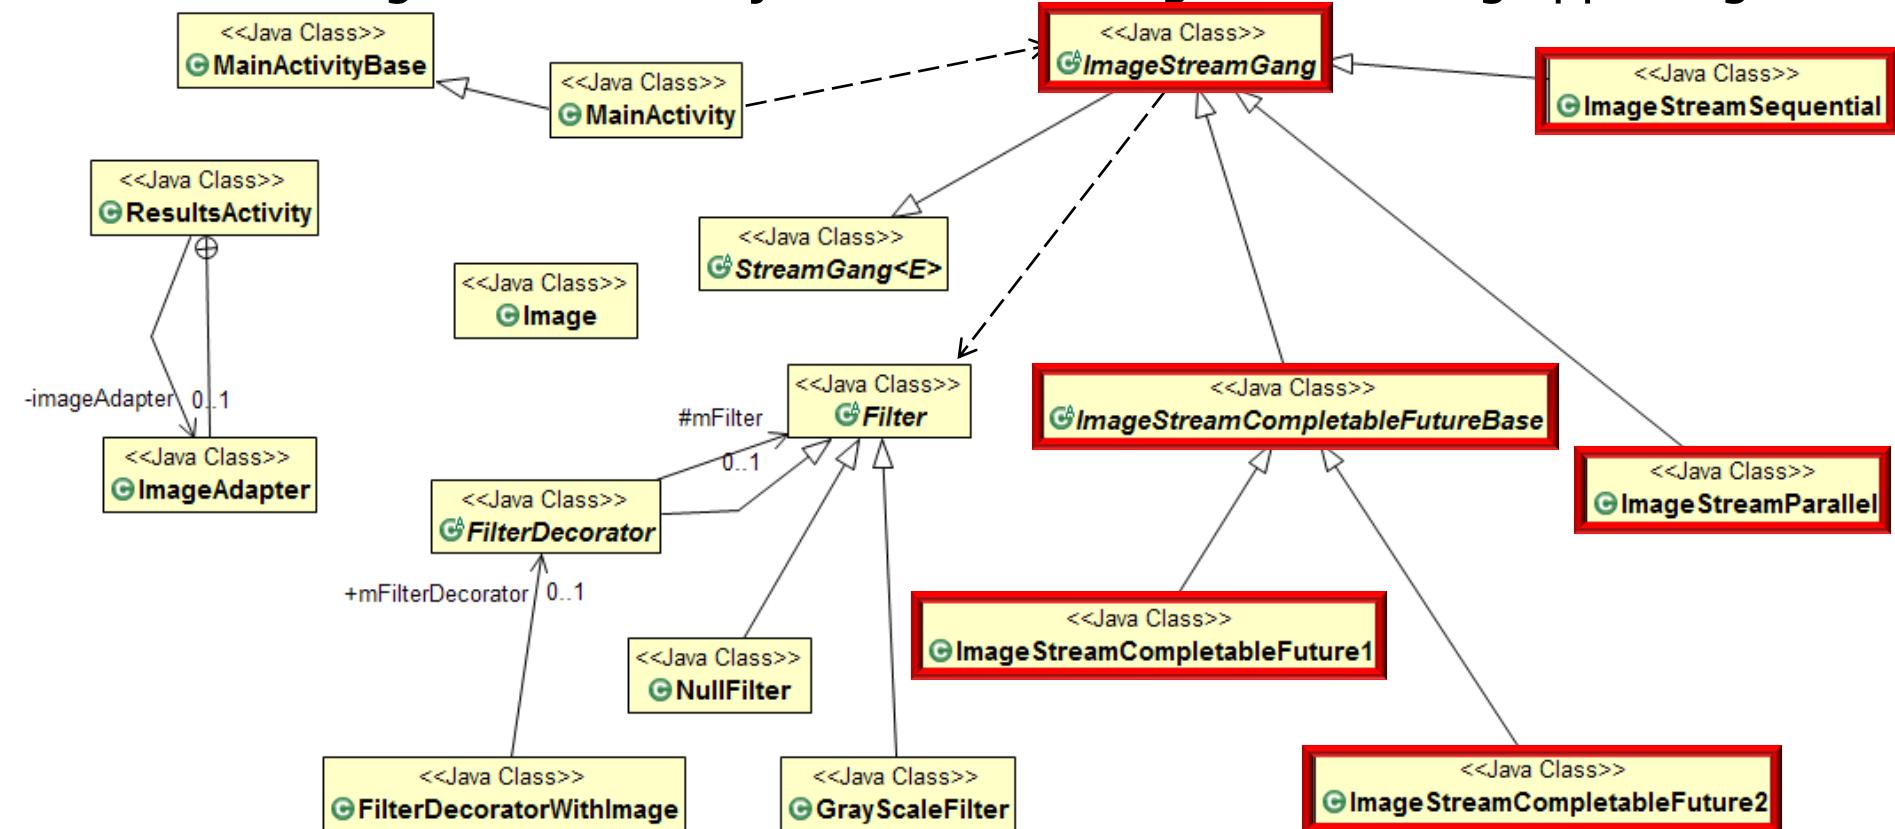
- UML class diagram for the object-oriented ImageStreamGang app design



Customizes the StreamGang framework to download and process images . . .

The Structure of the ImageStreamGang App

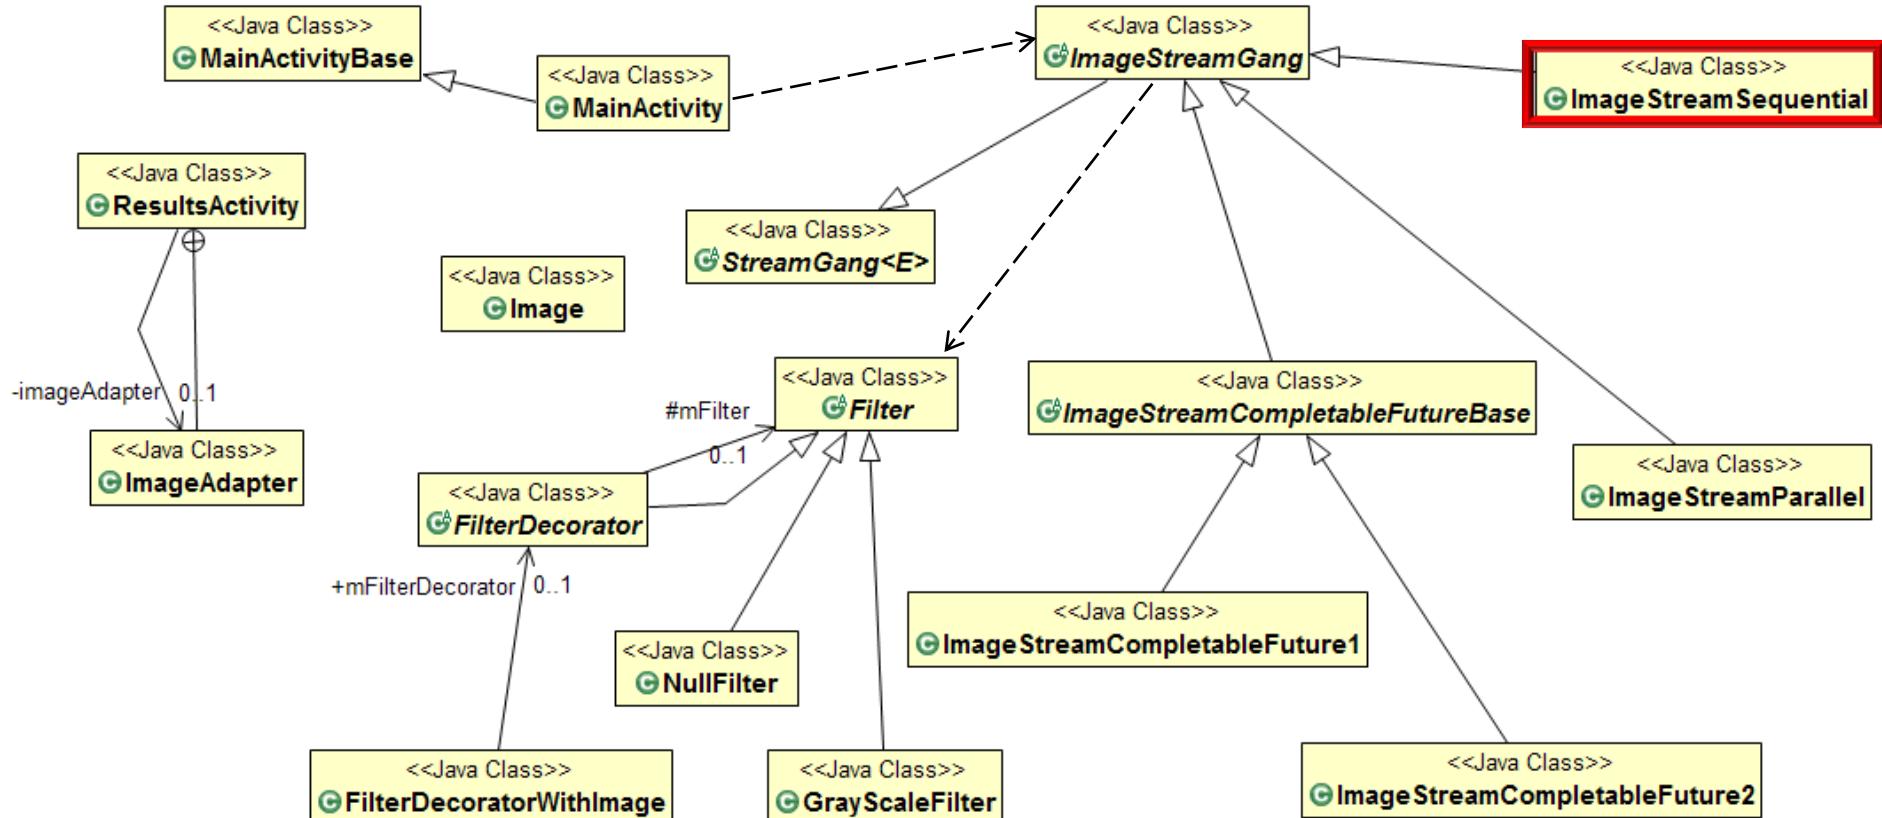
- UML class diagram for the object-oriented ImageStreamGang app design



... based on different Java concurrency and parallelism frameworks

The Structure of the ImageStreamGang App

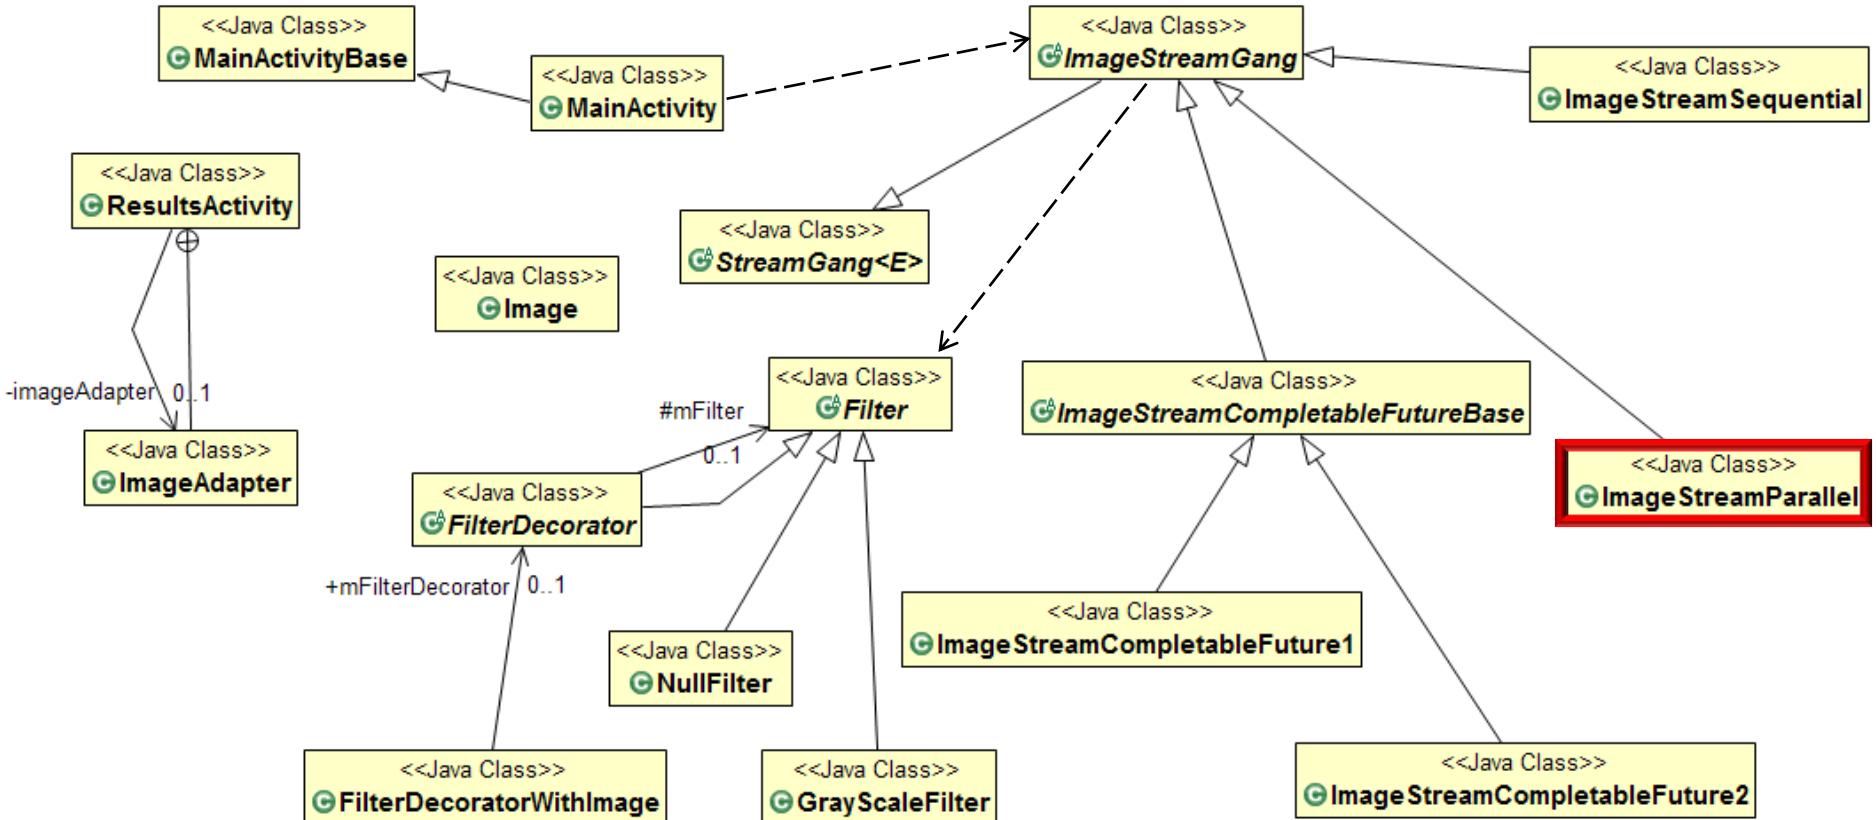
- UML class diagram for the object-oriented ImageStreamGang app design



Uses Java streams to download and filter images sequentially

The Structure of the ImageStreamGang App

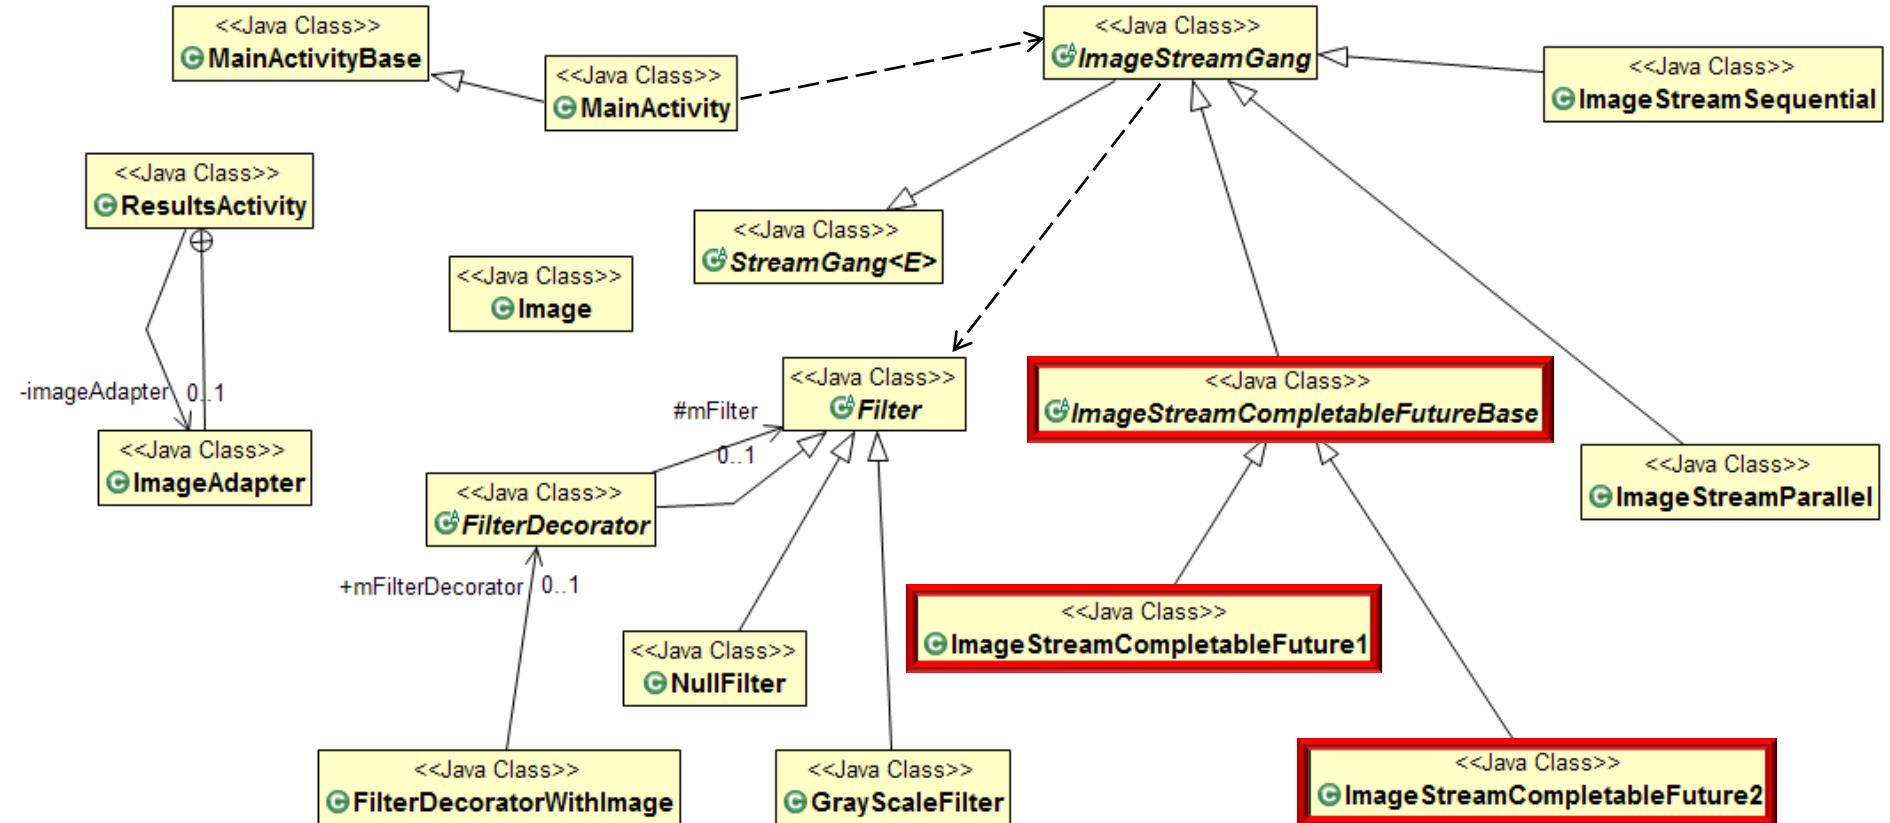
- UML class diagram for the object-oriented ImageStreamGang app design



Uses Java parallel streams to download and filter images concurrently

The Structure of the ImageStreamGang App

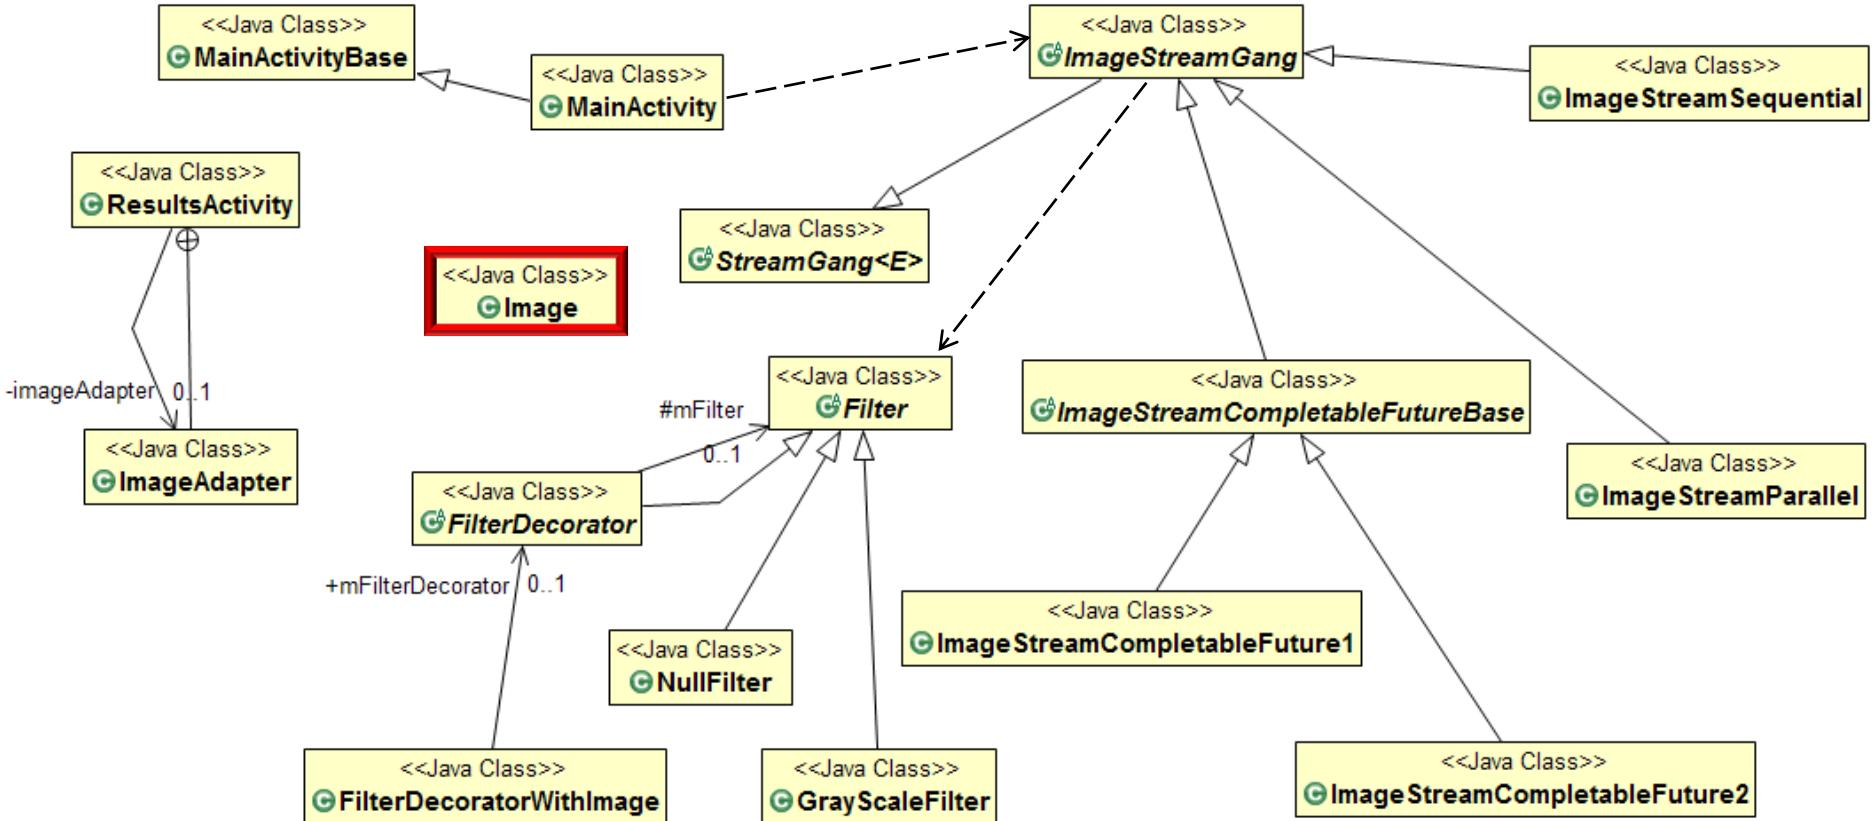
- UML class diagram for the object-oriented ImageStreamGang app design



Uses Java CompletableFutures to download and filter images asynchronously

The Structure of the ImageStreamGang App

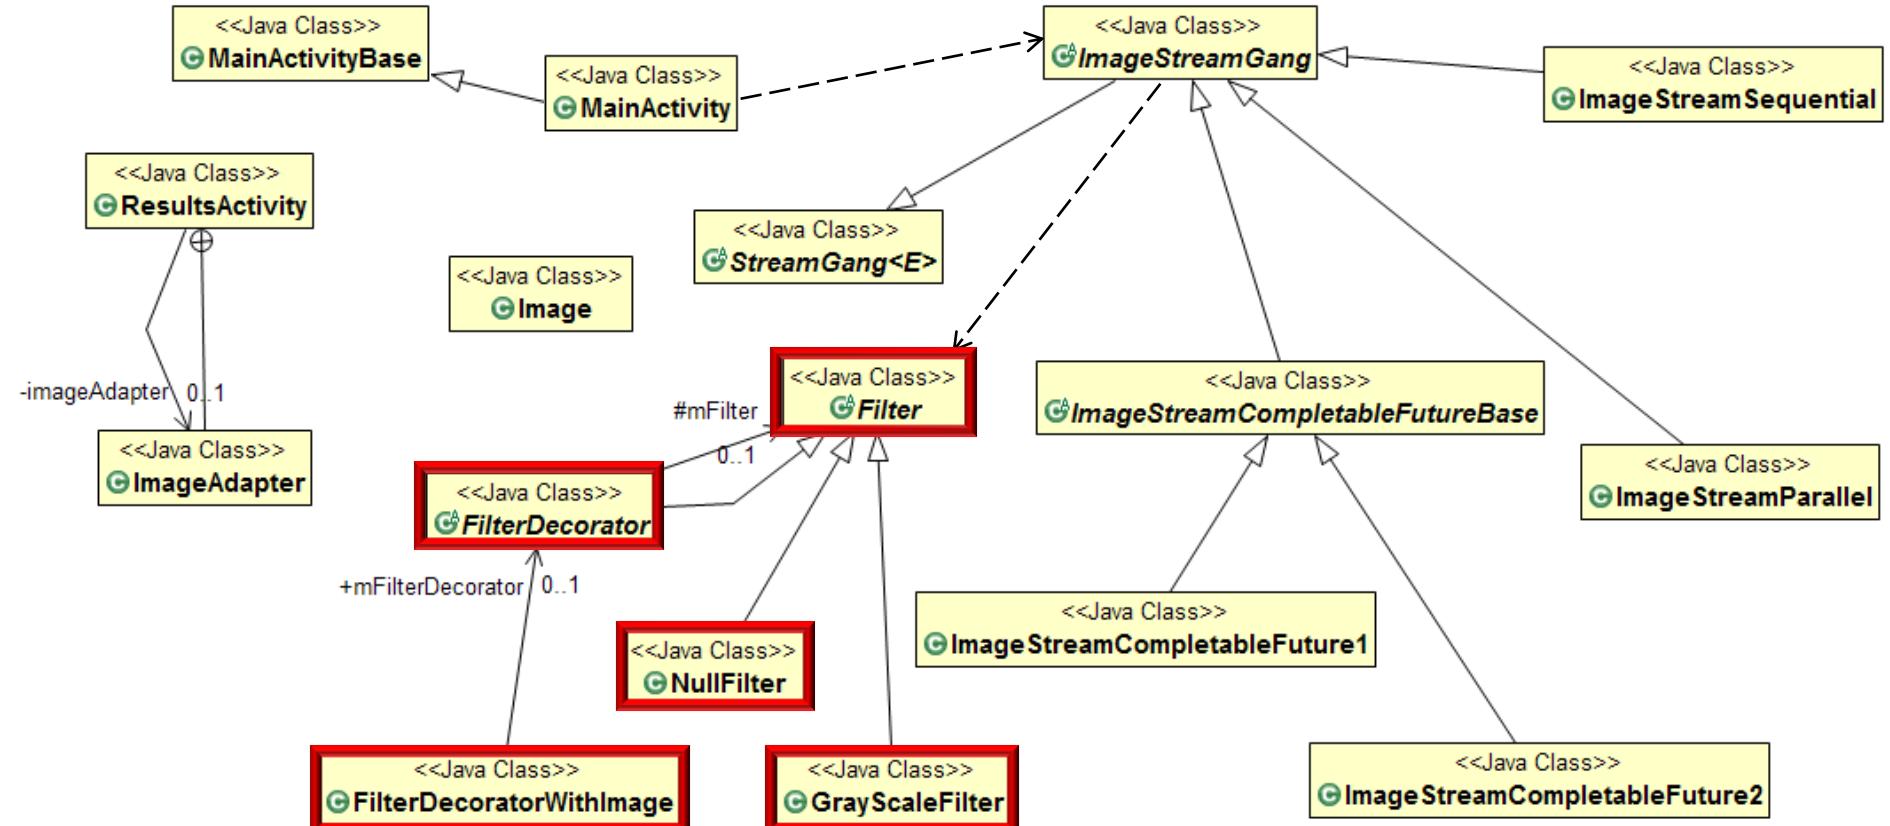
- UML class diagram for the object-oriented ImageStreamGang app design



Stores image metadata and provides methods for common image-/file-related tasks

The Structure of the ImageStreamGang App

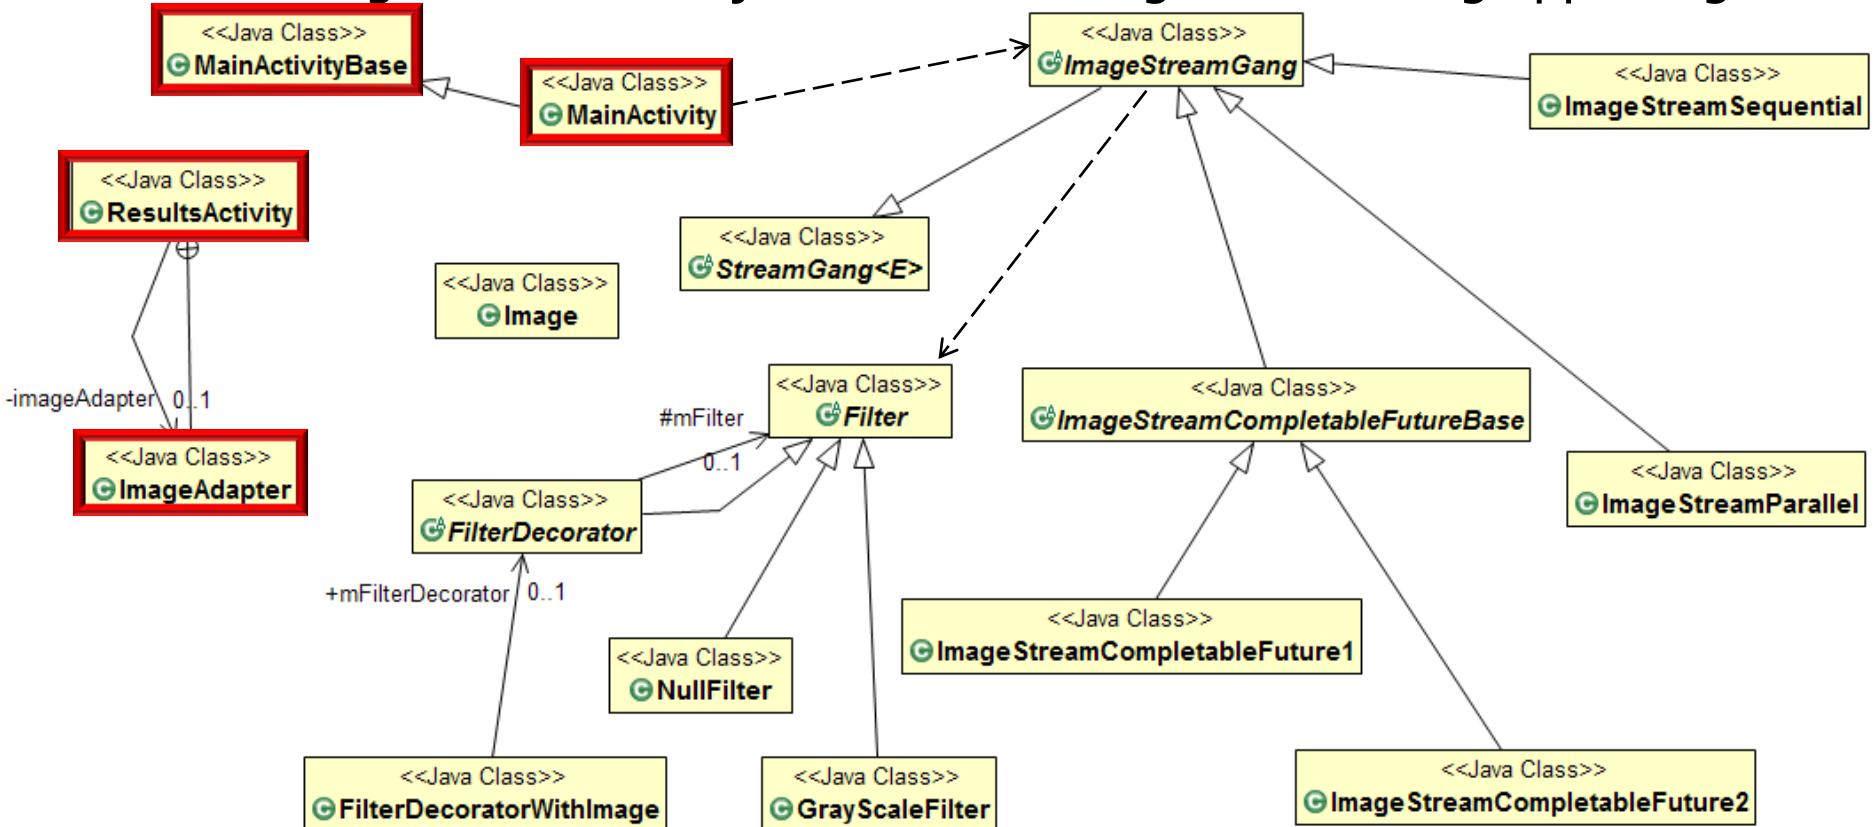
- UML class diagram for the object-oriented ImageStreamGang app design



This class hierarchy applies operations to filter and store images.

The Structure of the ImageStreamGang App

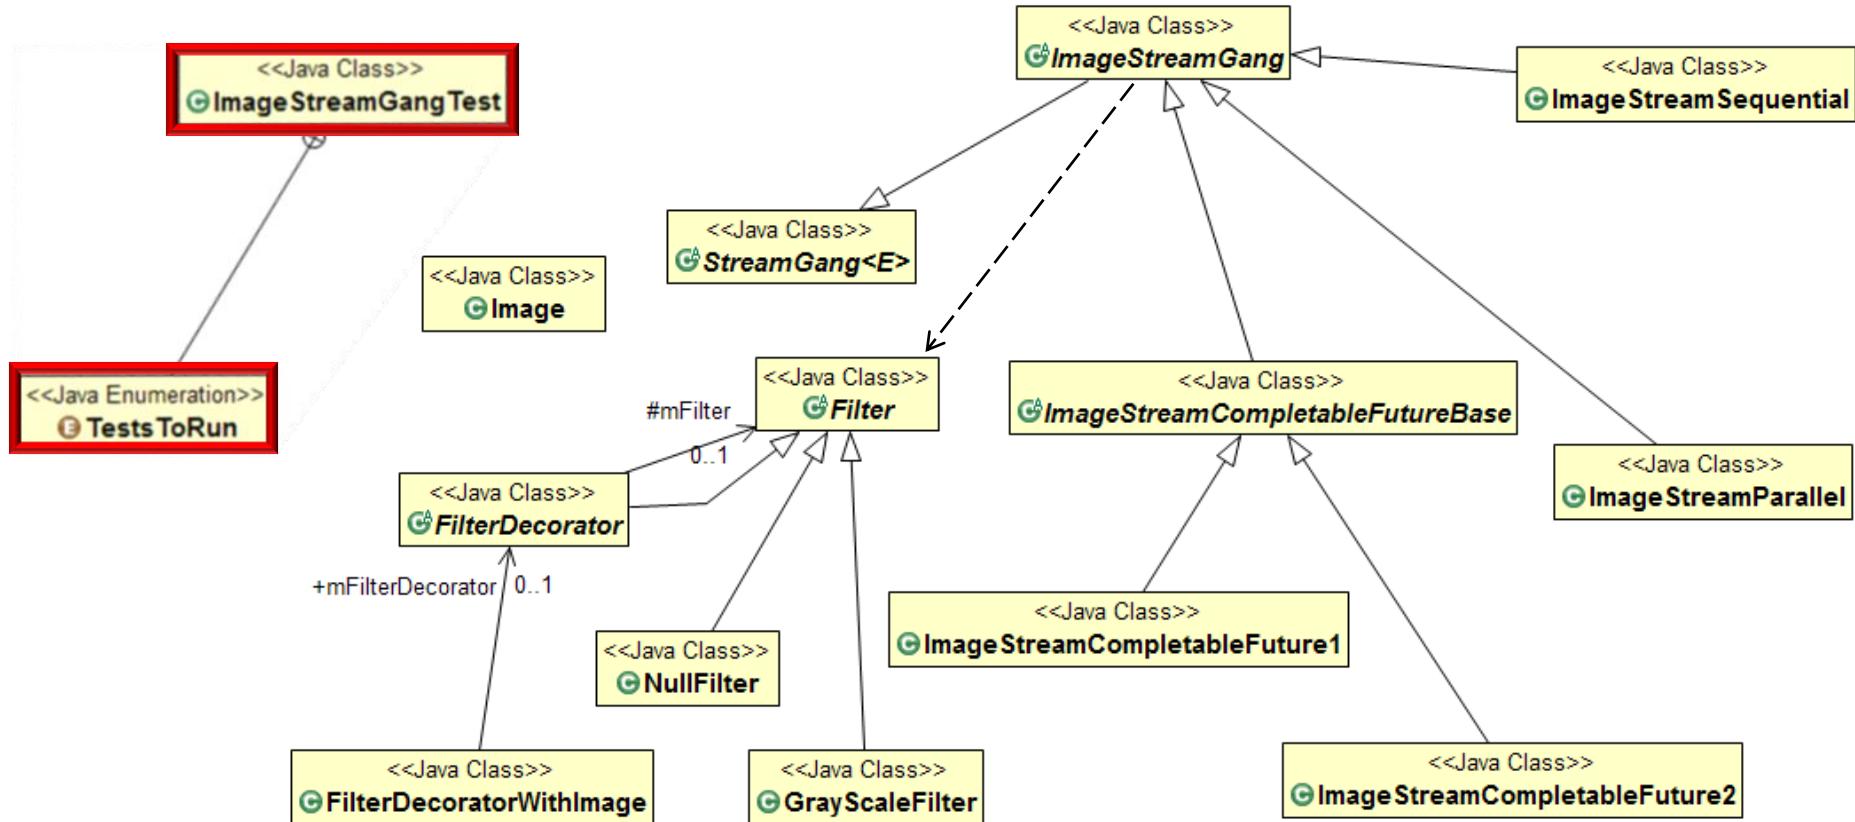
- UML class diagram for the object-oriented ImageStreamGang app design



Provides the user interface for an Android app

The Structure of the ImageStreamGang App

- UML class diagram for the object-oriented ImageStreamGang app design



There's a Java console version of ImageStreamGang that shares most of the code.

Running the ImageStreamGang App

Running the ImageStreamGang App

Starting ImageStreamGangTest

Printing 4 results for input file 1 from fastest to slowest

COMPLETABLE_FUTURES_1 executed in 312 msecs

COMPLETABLE_FUTURES_2 executed in 335 msecs

PARALLEL_STREAM executed in 428 msecs

SEQUENTIAL_STREAM executed in 981 msecs

Printing 4 results for input file 2 from fastest to slowest

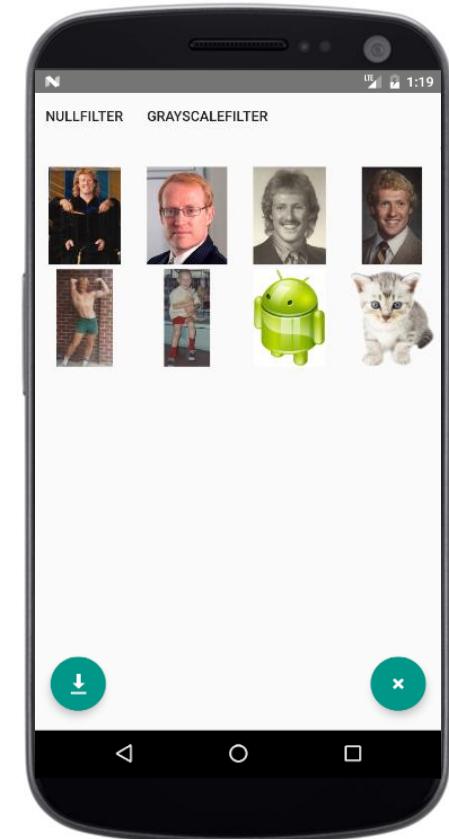
COMPLETABLE_FUTURES_2 executed in 82 msecs

COMPLETABLE_FUTURES_1 executed in 83 msecs

PARALLEL_STREAM executed in 102 msecs

SEQUENTIAL_STREAM executed in 251 msecs

Ending ImageStreamGangTest



Tests conducted on a 2.6 GHz six-core Lenovo P52 with 64 GB of RAM

Java Parallel ImageStreamGang Example:
Structure and Functionality

The End

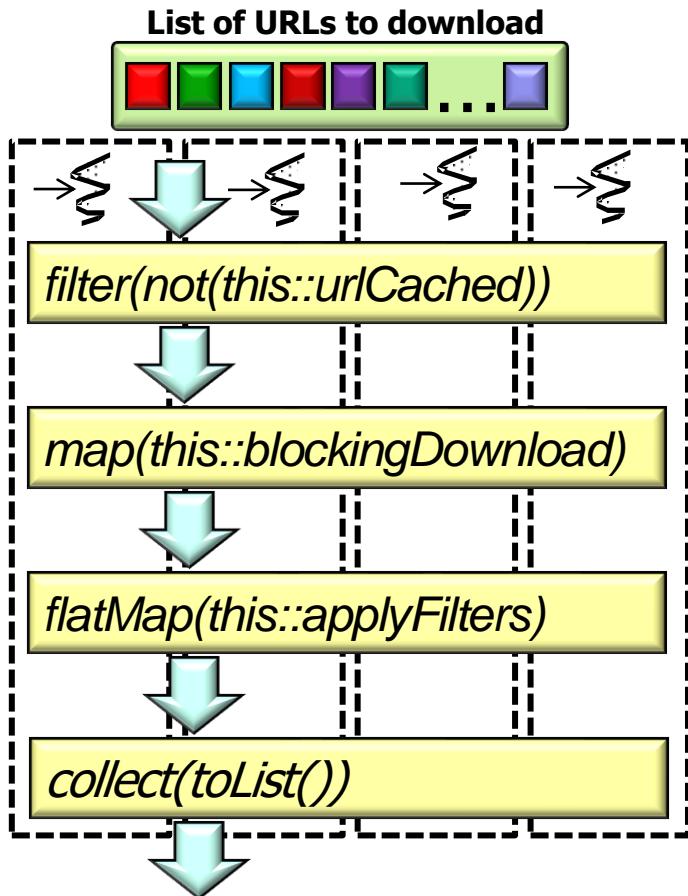
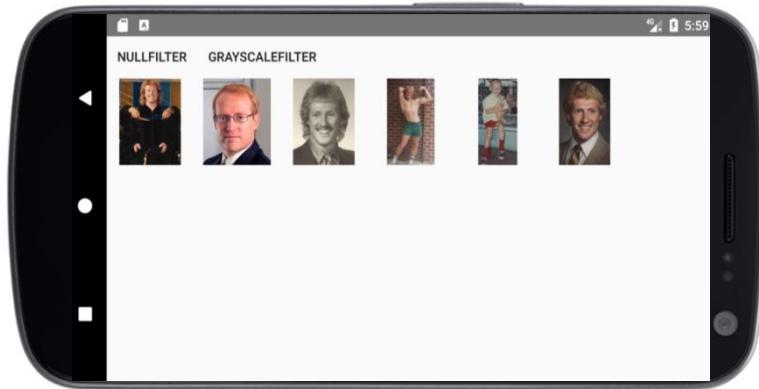
Java Parallel ImageStreamGang Example

Visualizing Behaviors

Douglas C. Schmidt

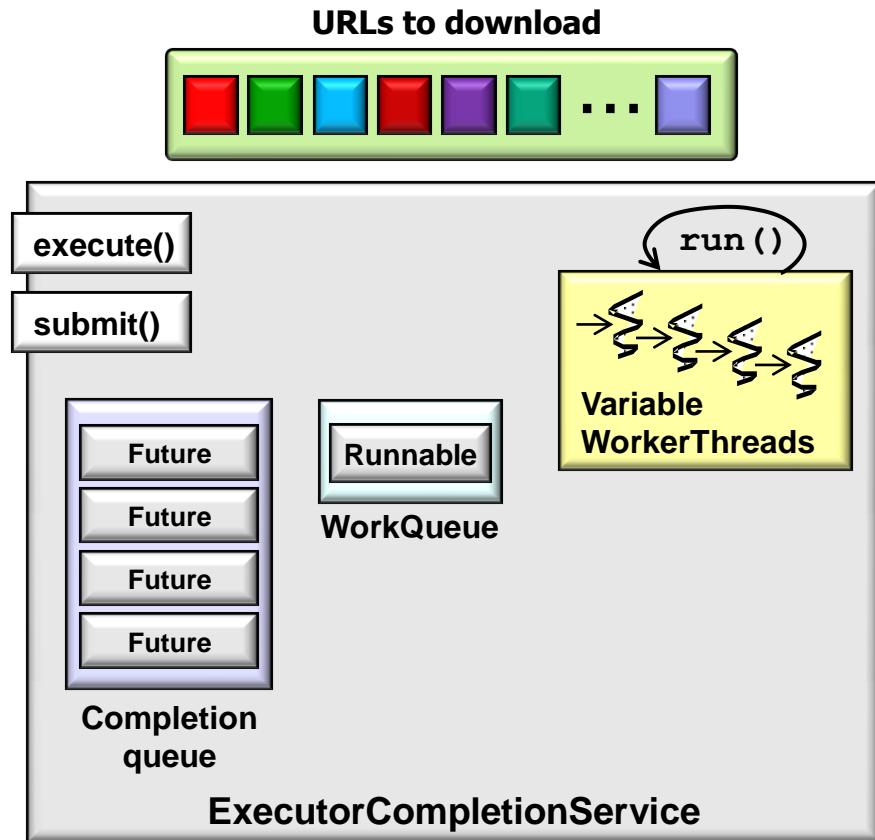
Learning Objectives in This Part of the Lesson

- Recognize the structure/functionality of the ImageStreamGang app.
- Know how Java parallel streams are applied to the ImageStreamGang app.



Learning Objectives in This Part of the Lesson

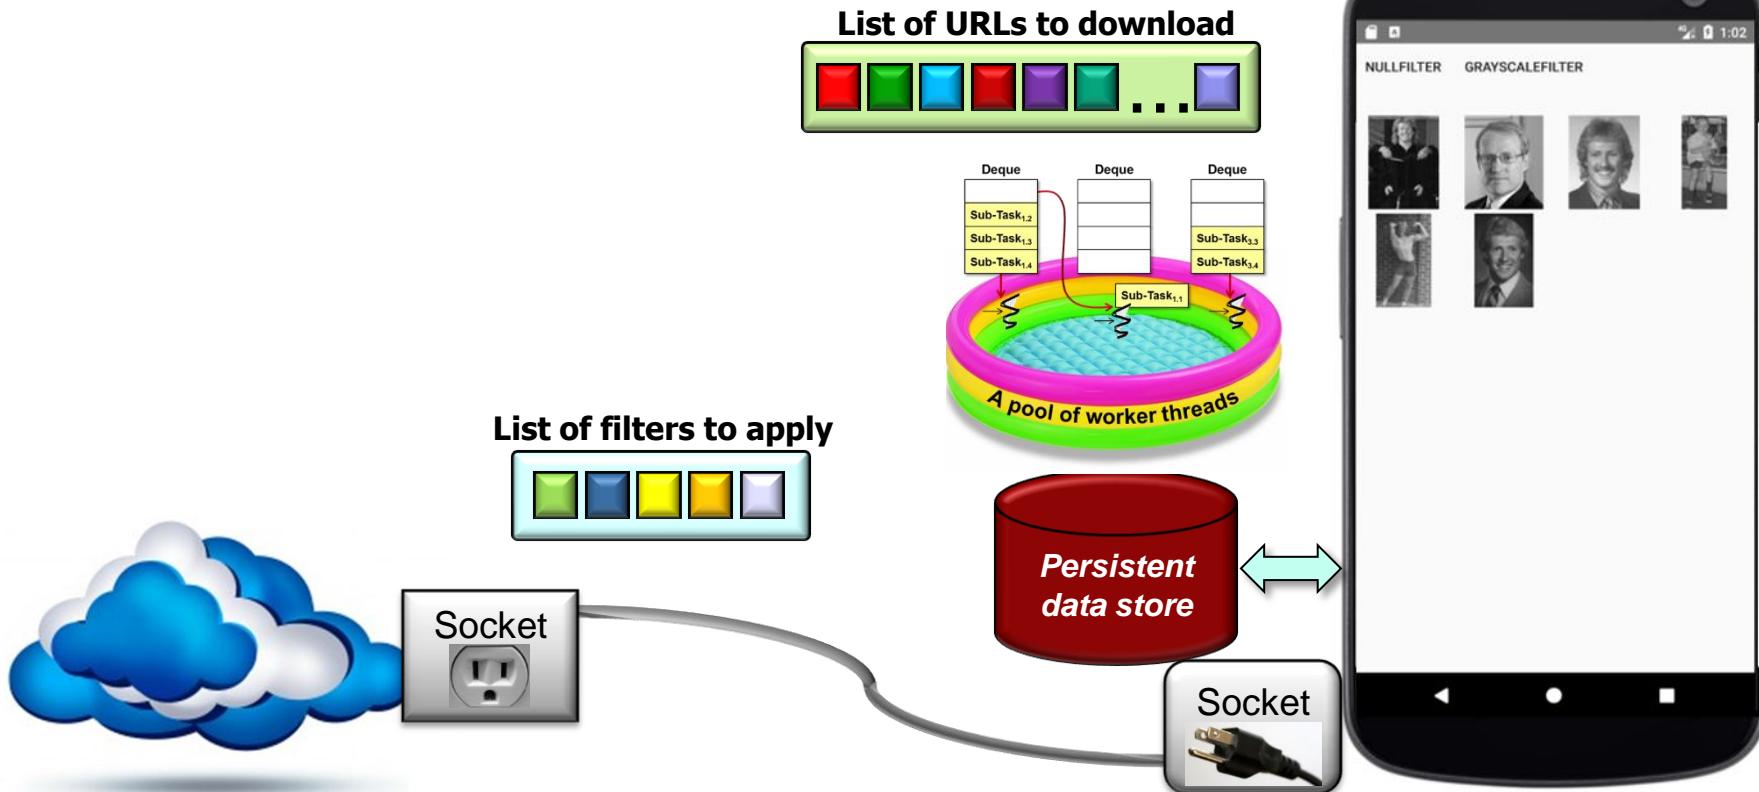
- Recognize the structure/functionality of the ImageStreamGang app.
- Know how Java parallel streams are applied to the ImageStreamGang app.
 - This app enhances ImageTaskGang.



Overview of Parallel Streams in ImageStreamGang

Overview of Parallel Streams in ImageStreamGang

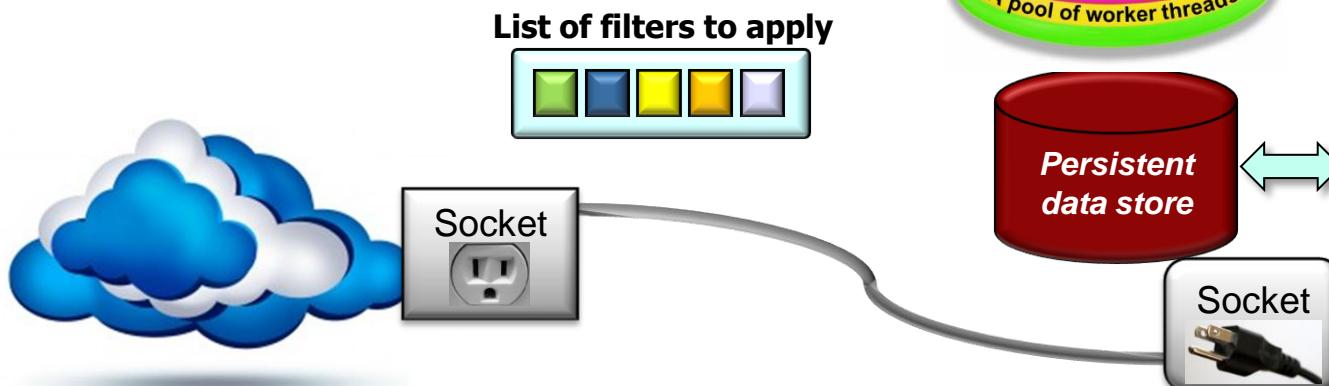
- This app uses a parallel stream with blocking I/O.



See github.com/douglasraigschmidt/LiveLessons/blob/master/ImageStreamGang/AndroidGUI

Overview of Parallel Streams in ImageStreamGang

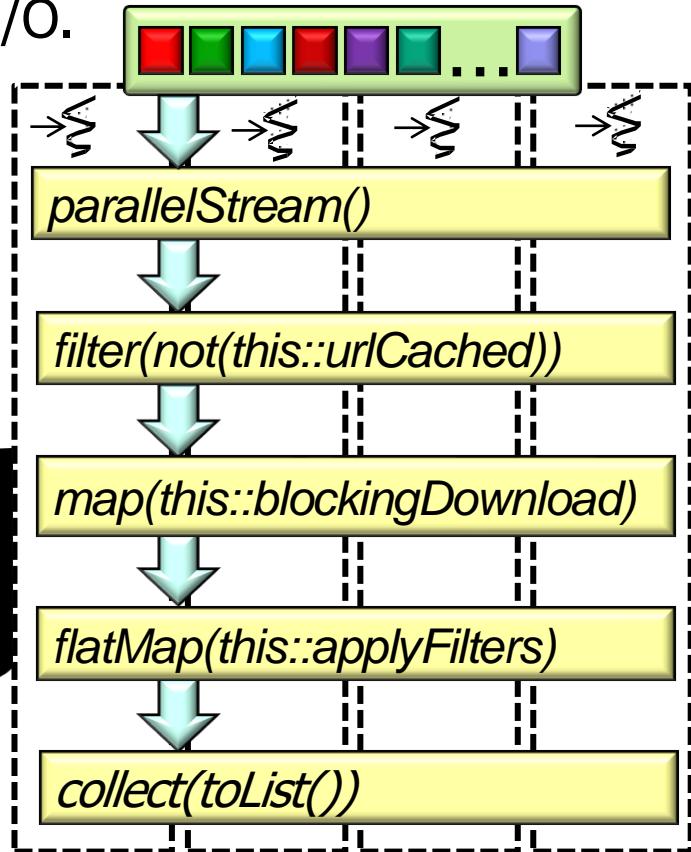
- This app uses a parallel stream with blocking I/O.
 - Ignore cached images
 - Download non-cached images
 - Apply list of filters to each image
 - Store filtered images in the file system
 - Display images to the user



Combines Java object-oriented and functional programming features

Overview of Parallel Streams in ImageStreamGang

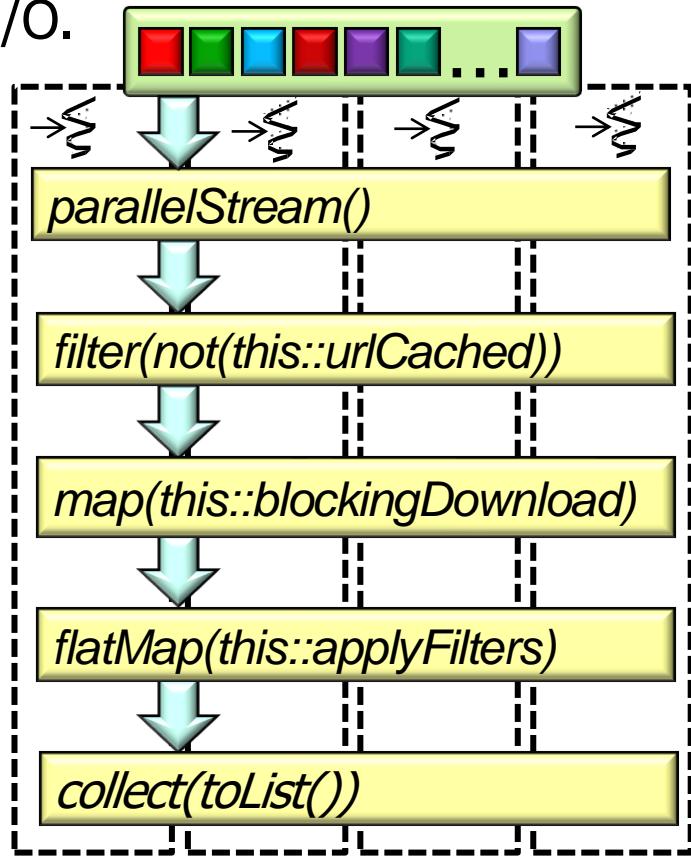
- This app uses a parallel stream with blocking I/O.
 - Ignore cached images
 - Download non-cached images
 - Apply list of filters to each image
 - Store filtered images in the file system
 - Display images to the user



Declarative stream pipeline closely aligns with the app description.

Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.



Closes gap between domain intent and computations that implement the intent

Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

List
<URL>



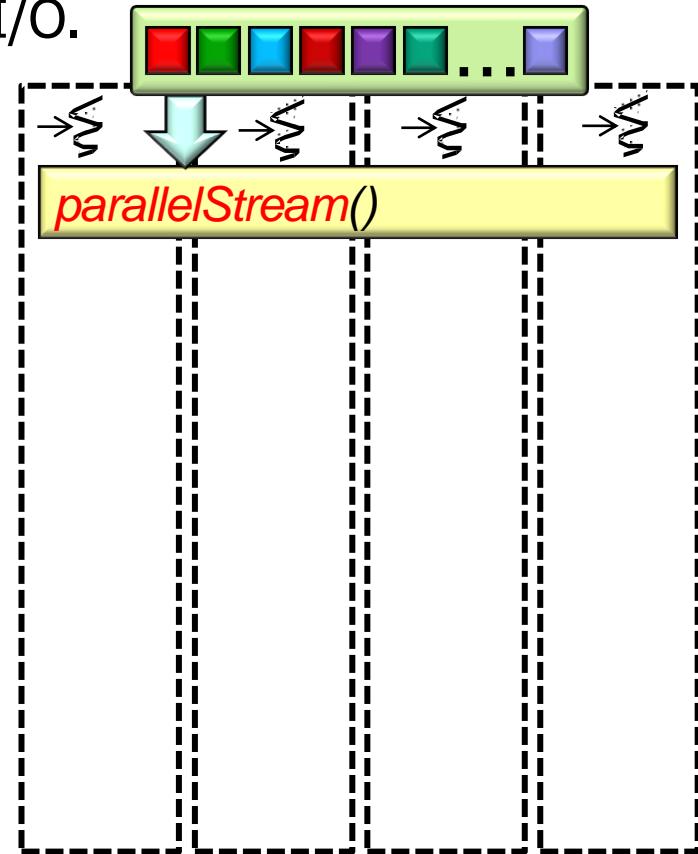
`parallelStream()`

Input a list of image URLs

Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

List
<URL>

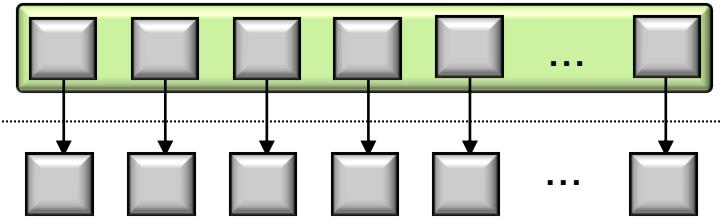


Convert collection to a parallel stream

Overview of Parallel Streams in ImageStreamGang

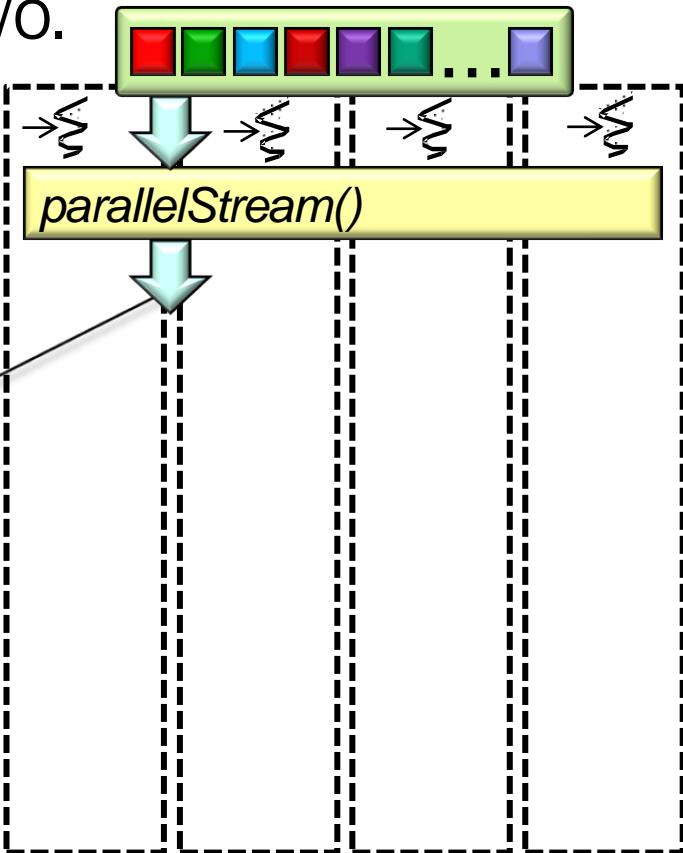
- This app uses a parallel stream with blocking I/O.

List
<URL>



Stream
<URL>

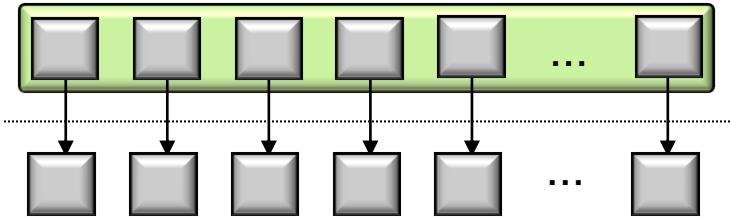
Output a stream of image URLs



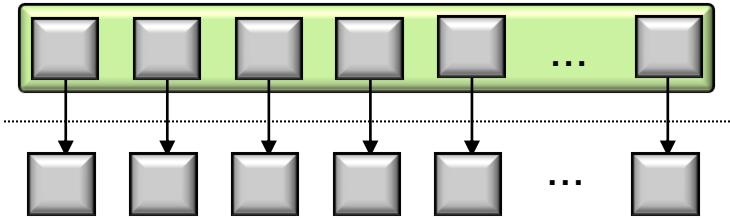
Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

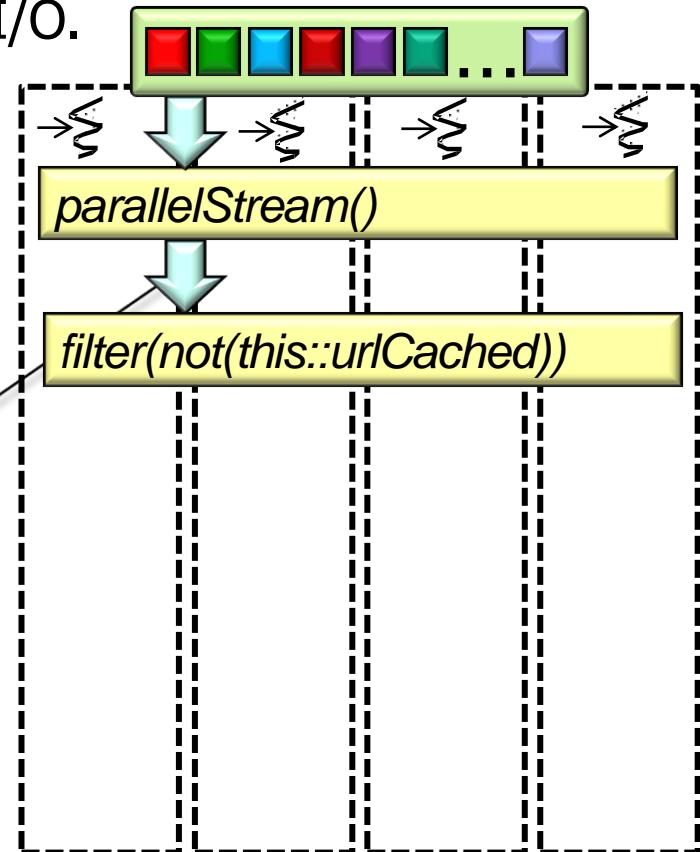
List
<URL>



Stream
<URL>



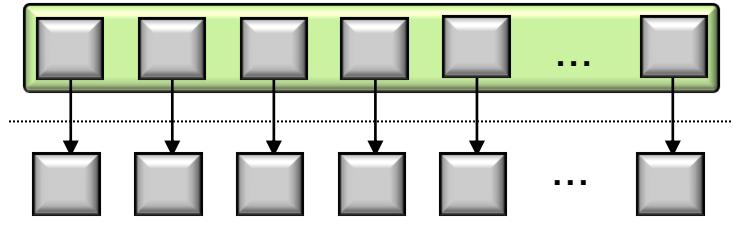
Input a stream of image URLs



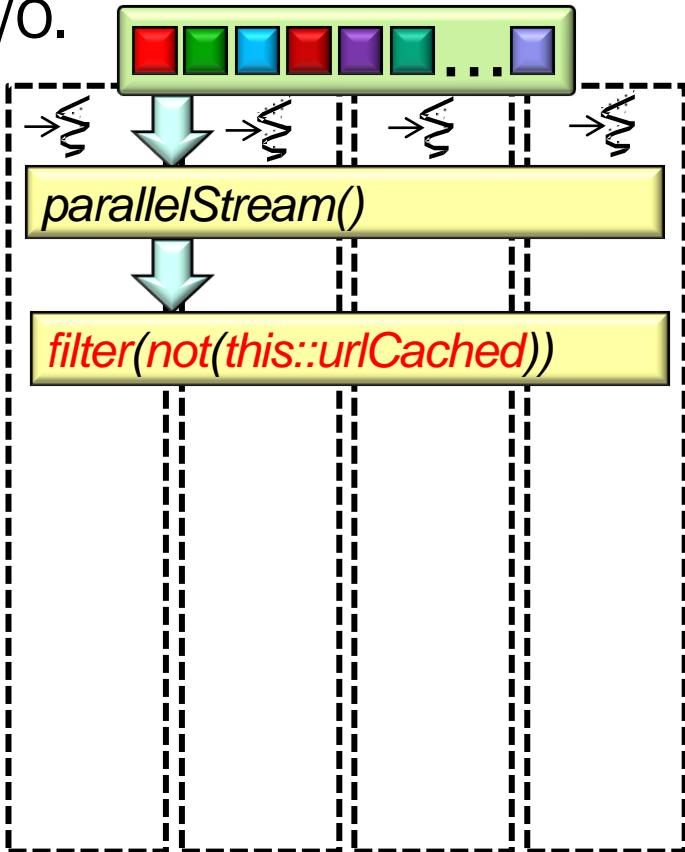
Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

List
<URL>



Stream
<URL>



Filter() ignores cached images.

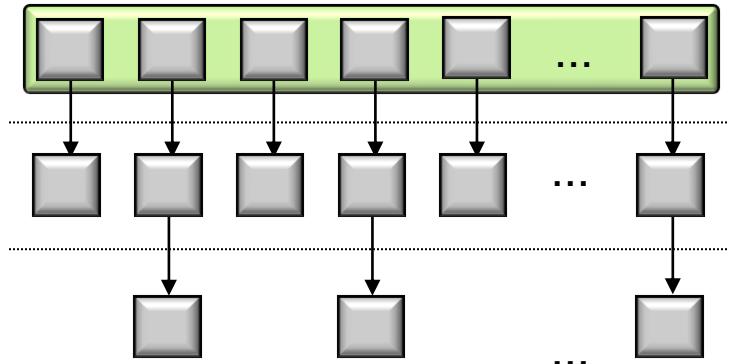
Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

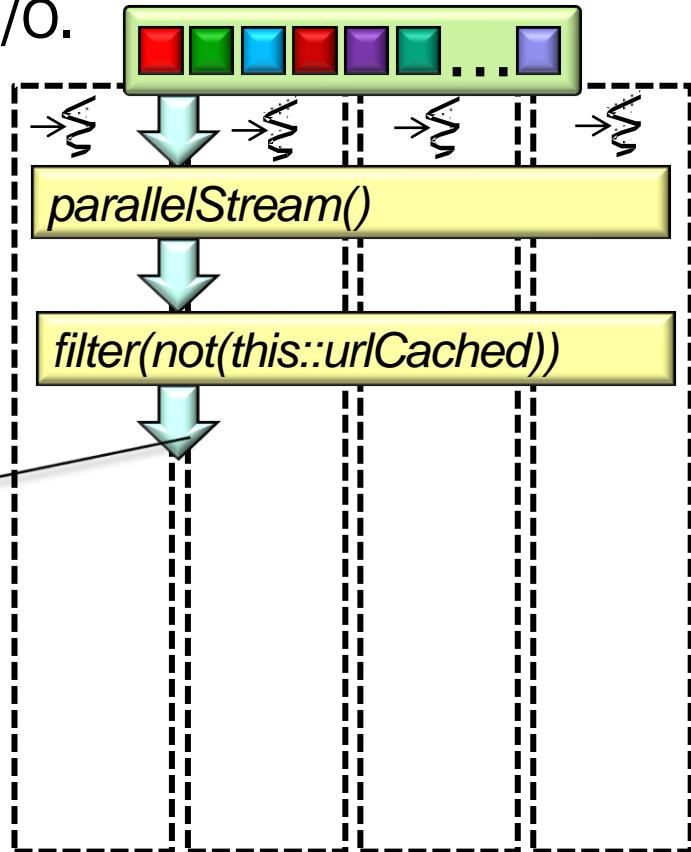
List
<URL>

Stream
<URL>

Stream
<URL>



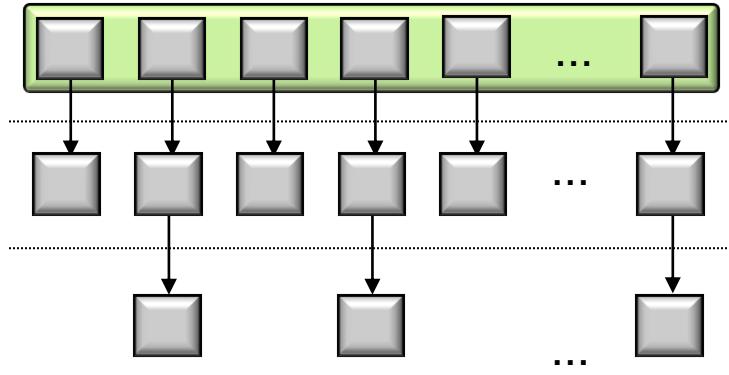
Output a stream of non-cached image URLs



Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

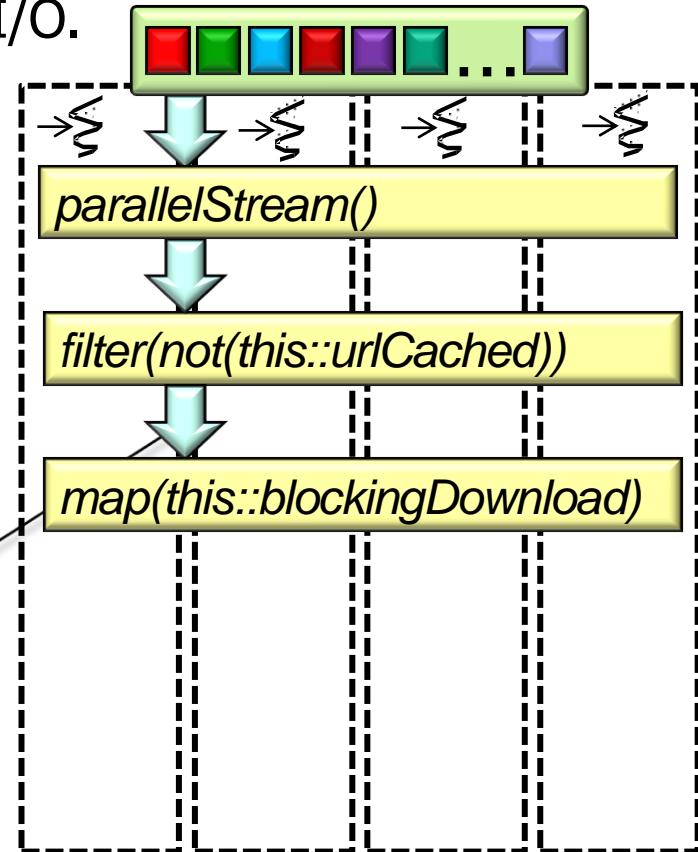
List
<URL>



Stream
<URL>

Stream
<URL>

Input a stream of non-cached image URLs



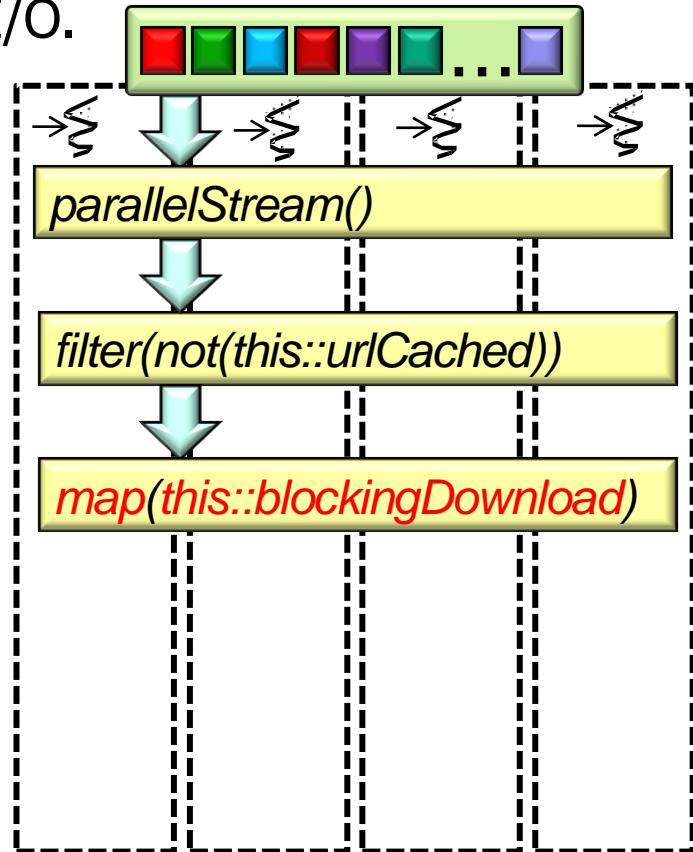
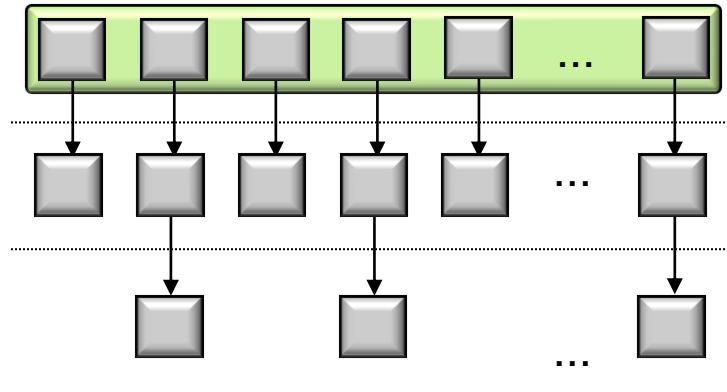
Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.

List
<URL>

Stream
<URL>

Stream
<URL>



Download non-cached images in parallel

Overview of Parallel Streams in ImageStreamGang

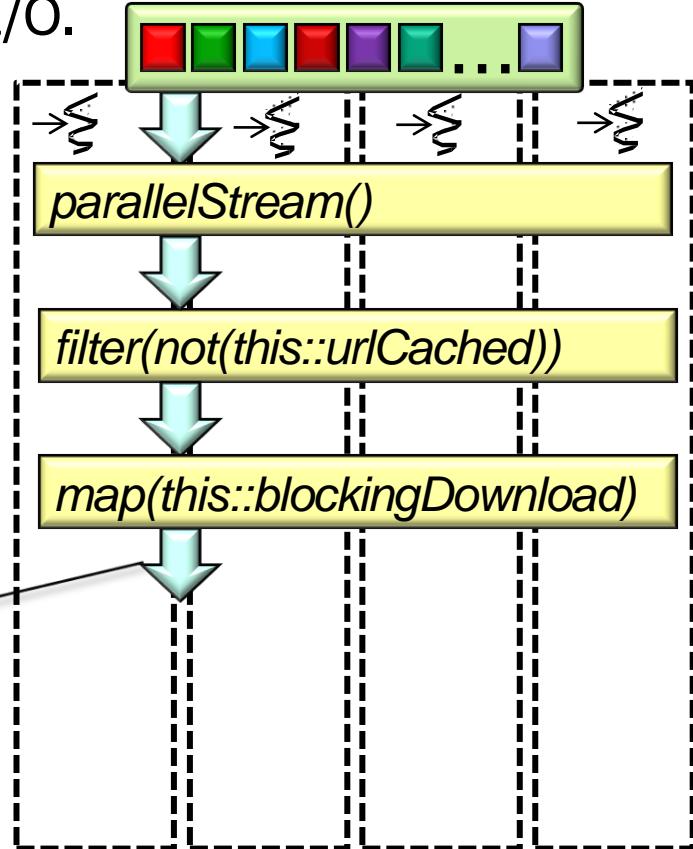
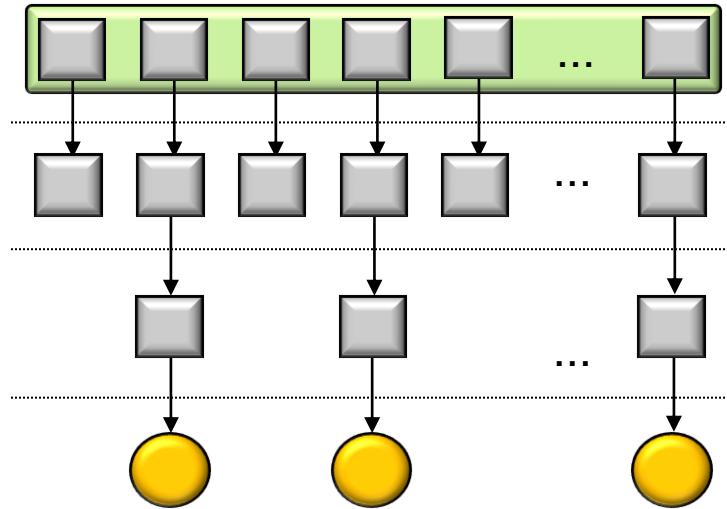
- This app uses a parallel stream with blocking I/O.

List
<URL>

Stream
<URL>

Stream
<URL>

Stream
<Image>



Overview of Parallel Streams in ImageStreamGang

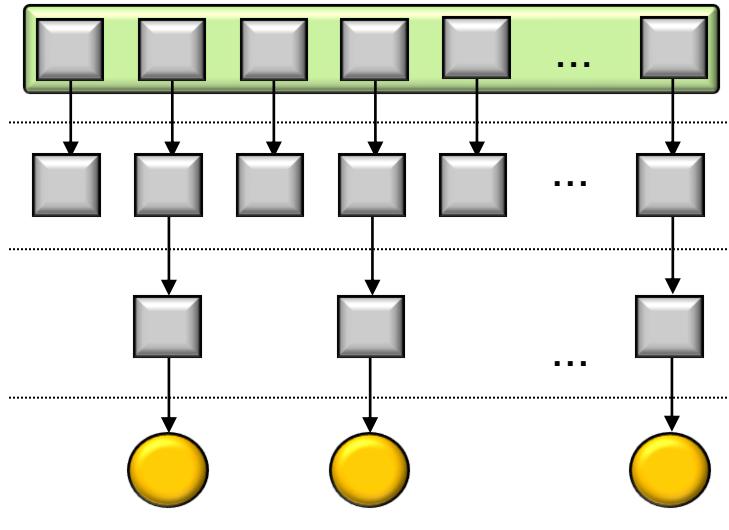
- This app uses a parallel stream with blocking I/O.

List
<URL>

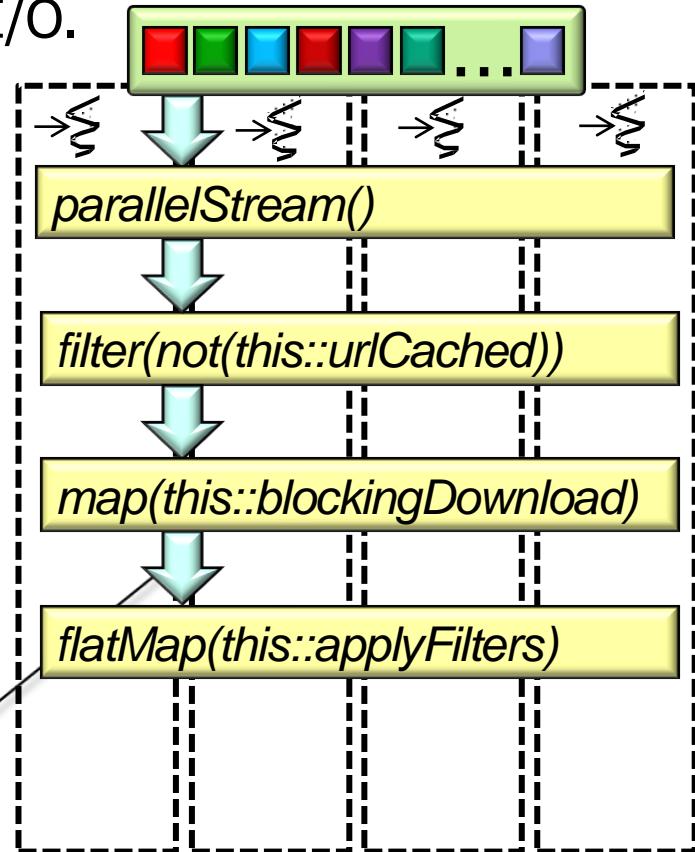
Stream
<URL>

Stream
<URL>

Stream
<Image>

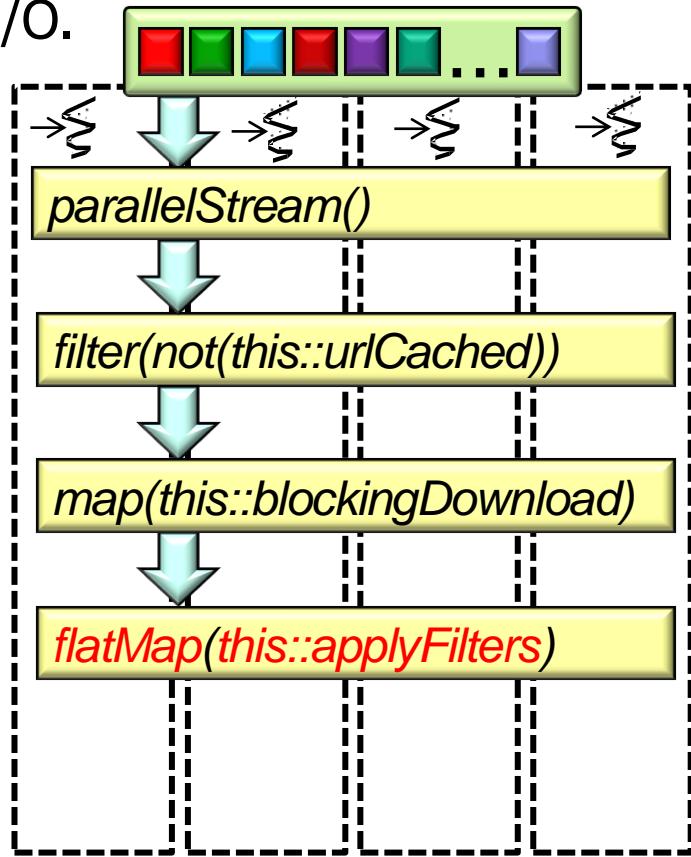
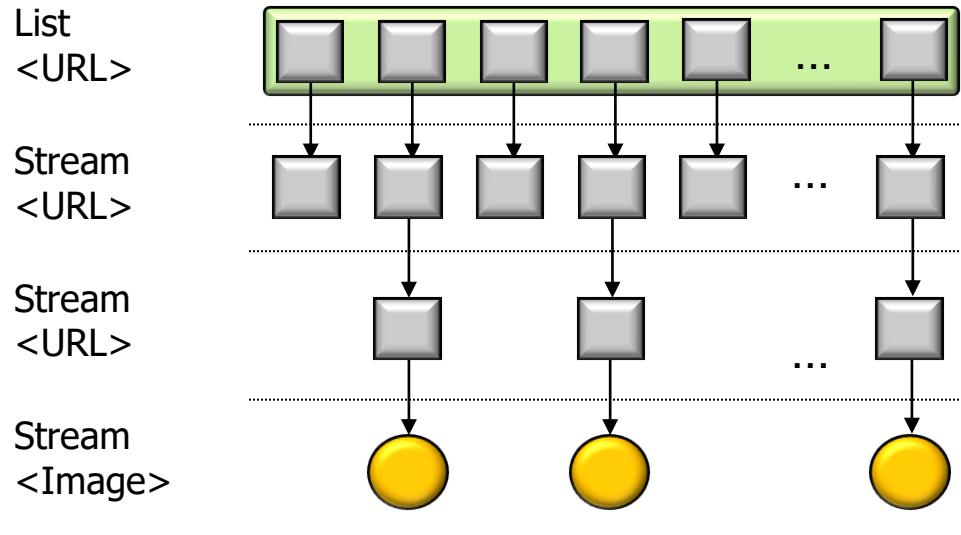


Input a stream of downloaded images



Overview of Parallel Streams in ImageStreamGang

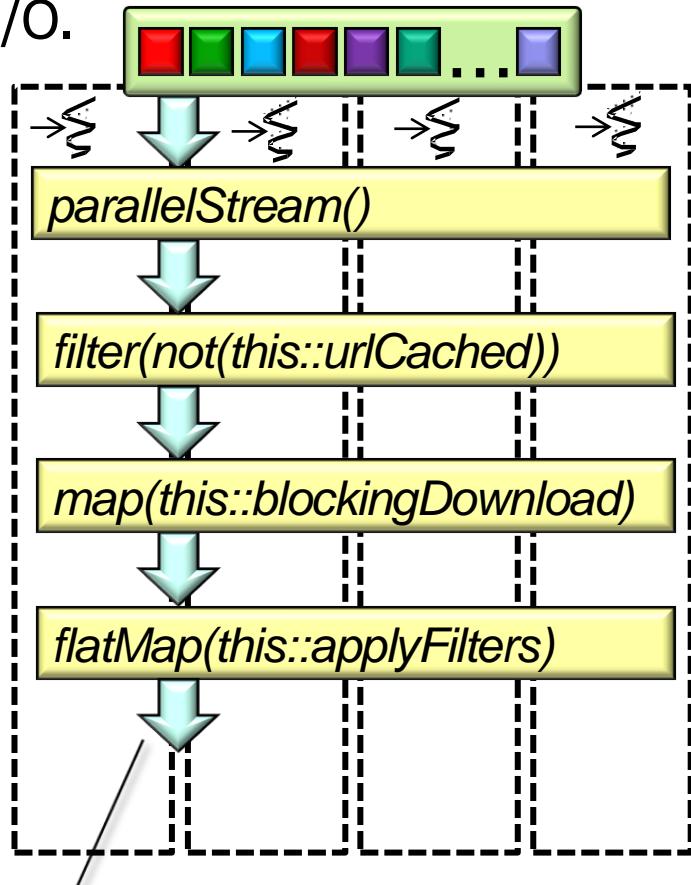
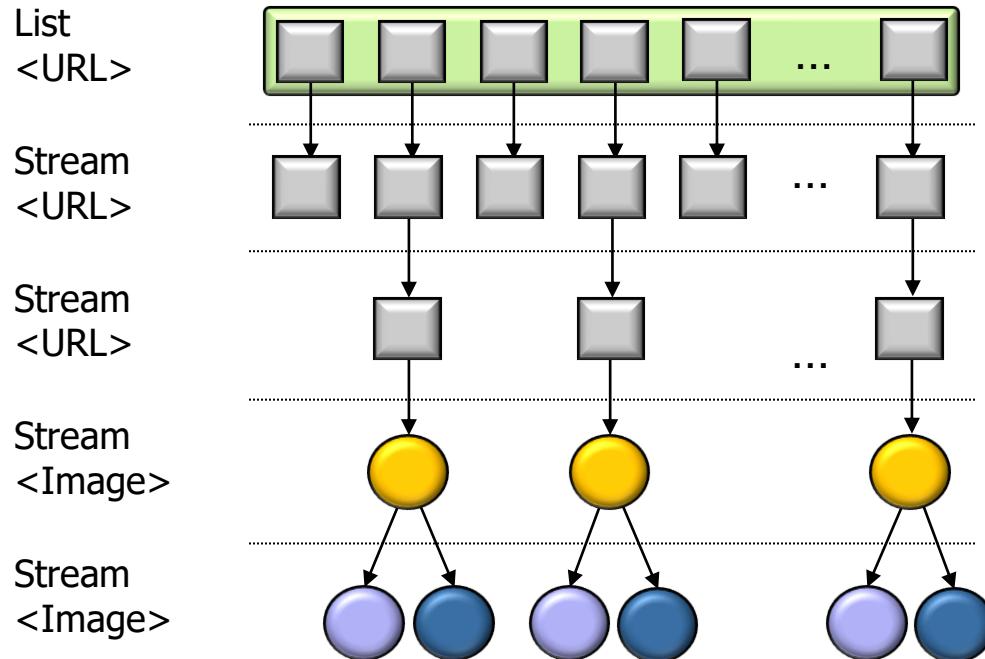
- This app uses a parallel stream with blocking I/O.



Apply filters to each image in parallel and store filtered images in file system

Overview of Parallel Streams in ImageStreamGang

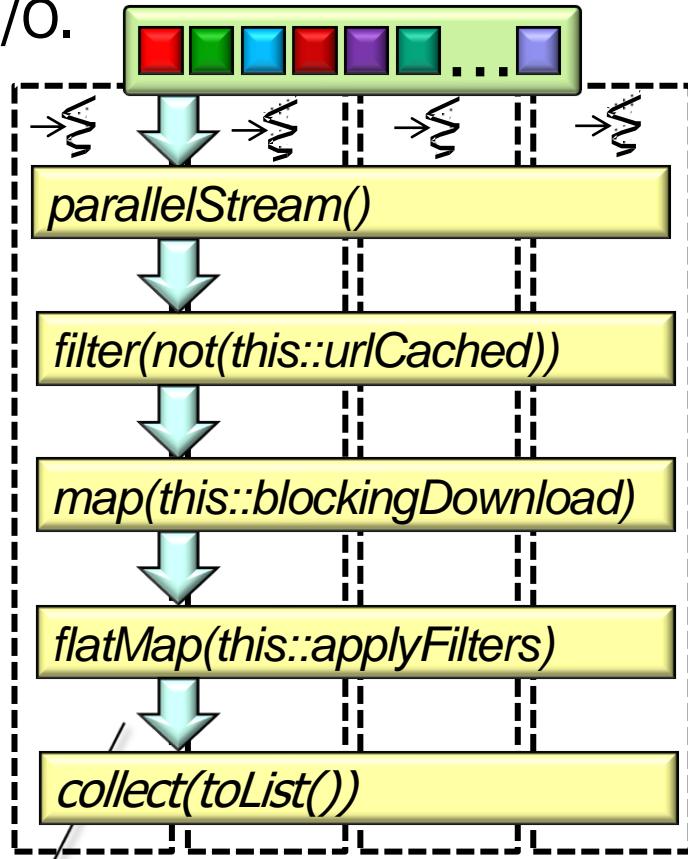
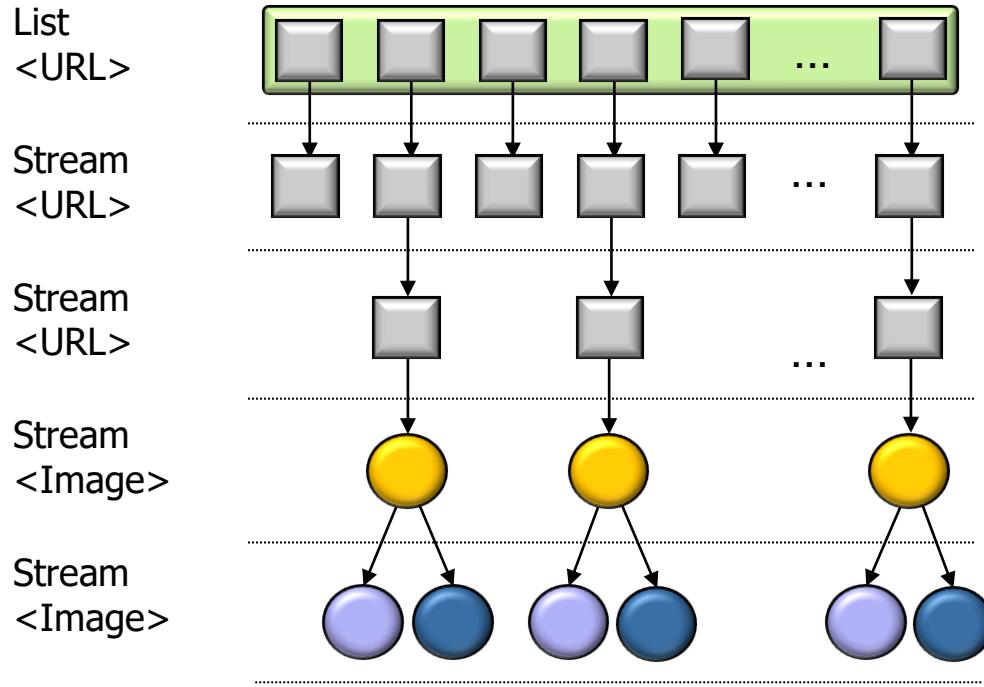
- This app uses a parallel stream with blocking I/O.



Output a stream of filtered and stored images

Overview of Parallel Streams in ImageStreamGang

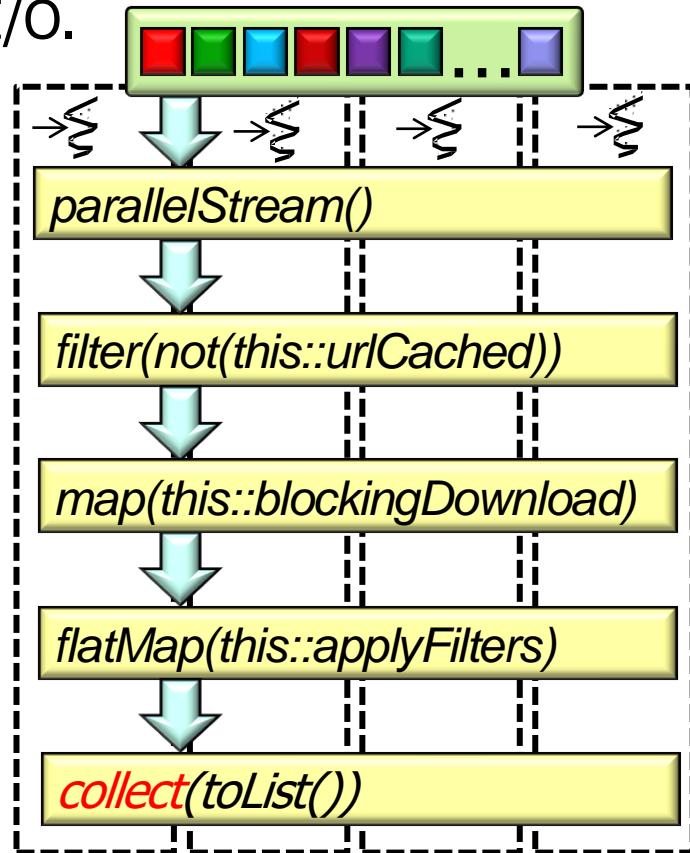
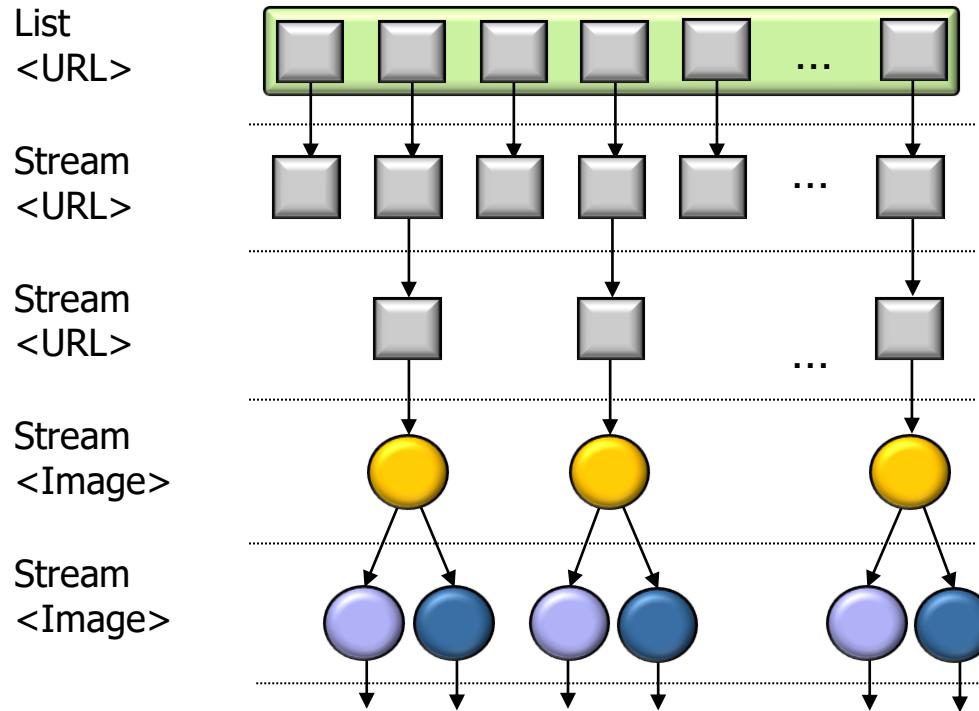
- This app uses a parallel stream with blocking I/O.



Input a stream of filtered and stored images

Overview of Parallel Streams in ImageStreamGang

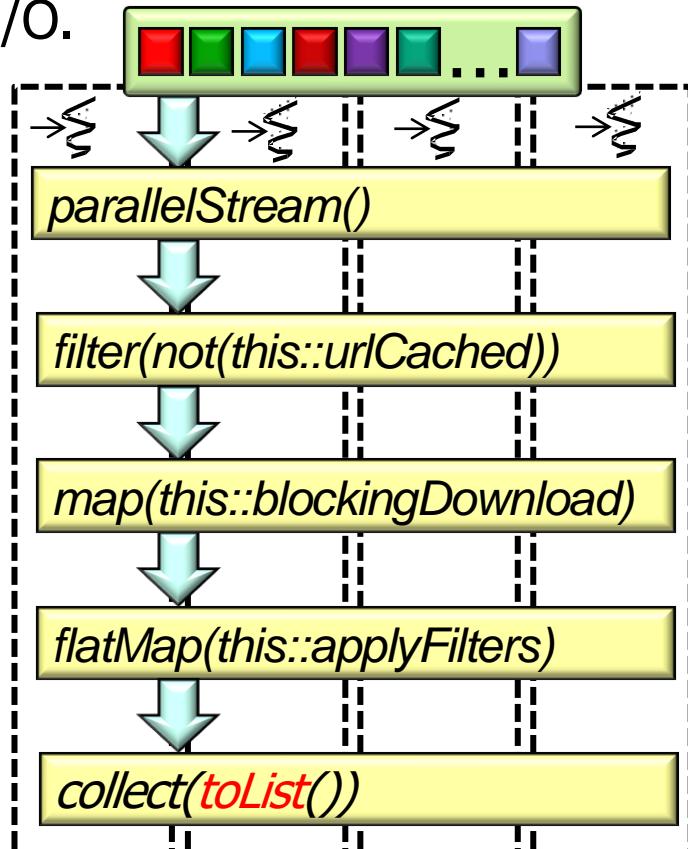
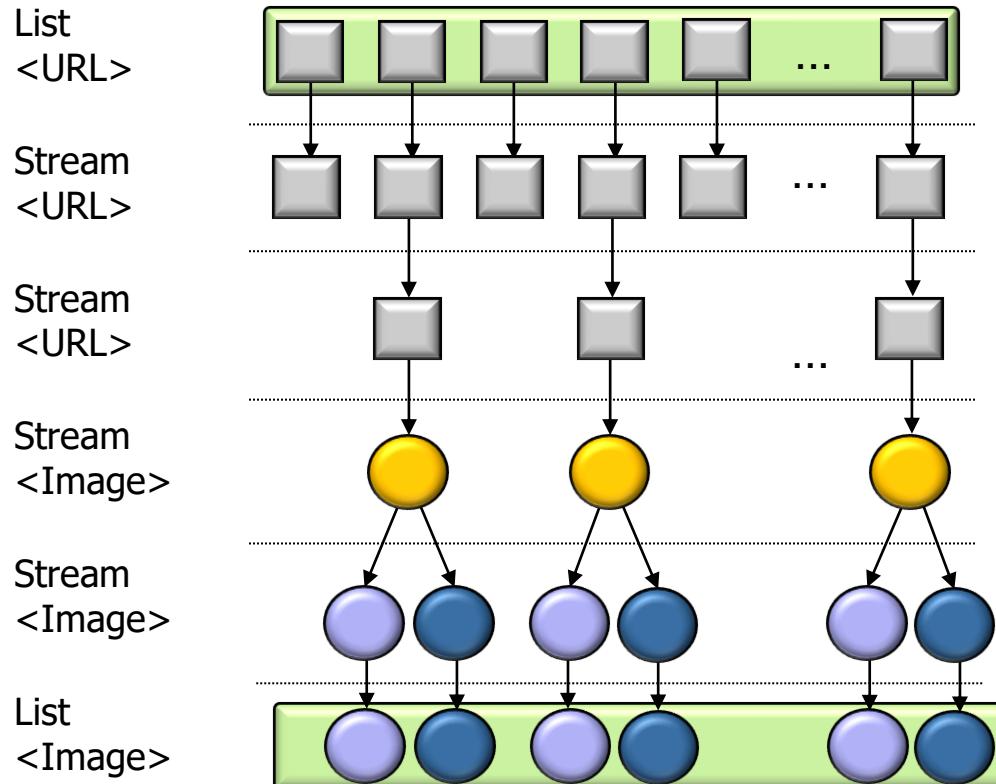
- This app uses a parallel stream with blocking I/O.



Trigger intermediate operation processing

Overview of Parallel Streams in ImageStreamGang

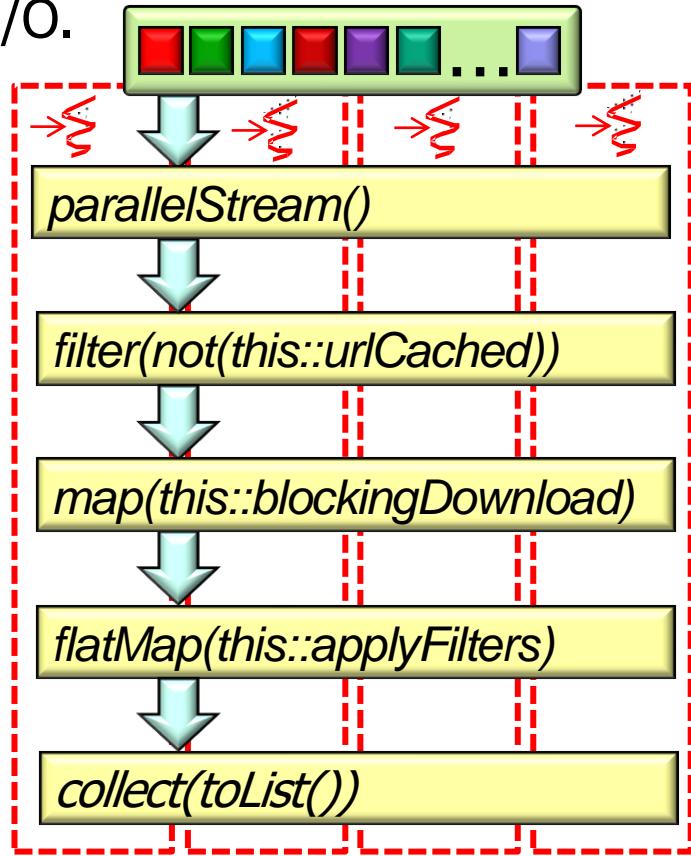
- This app uses a parallel stream with blocking I/O.



Return a list of filtered and stored images

Overview of Parallel Streams in ImageStreamGang

- This app uses a parallel stream with blocking I/O.
 - Ignore cached images
 - Download non-cached images
 - Apply list of filters to each image
 - Store filtered images in the file system
 - Display images to the user (after triggering stream processing)



The Java streams framework orchestrates all these steps in parallel.

Java Parallel ImageStreamGang Example: Visualizing Behaviors

The End

Java Parallel ImageStreamGang Example

Implementing Behaviors

Douglas C. Schmidt

Learning Objectives in This Part of the Lesson

- Recognize the structure/functionality of the ImageStreamGang app.
- Know how Java parallel streams are applied to the ImageStreamGang app.
- Understand the parallel streams implementation of ImageStreamGang.

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages =  
    urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

See [imagestreamgang/streams/ImageStreamParallel.java](#)

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Get a list of URLs

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

GetInput() is defined by the underlying StreamGang framework.

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

*Convert a collection
into a parallel stream*

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream consisting of the URLs in the input stream that are not already cached

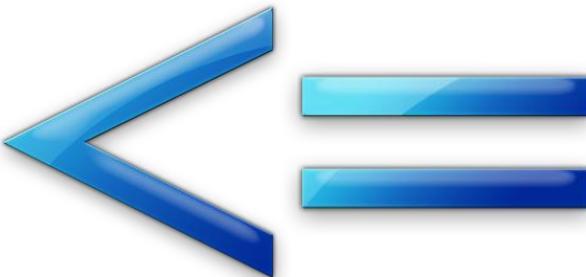
```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream consisting of the URLs in the input stream that are not already cached

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
}  
  
System.out.println(TAG  
    + "Image(s) filtered = "  
    + filteredImages.size());
```



of output stream elements will be \leq # of input stream elements

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
boolean urlCached(URL url) {  
    return mFilters  
        .stream()  
        .filter(filter ->  
            urlCached(url,  
                      filter.getName()))  
        .count() > 0;  
}
```

Determine whether this URL has been downloaded to an image and had filters applied to it yet

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

See [imagestreamgangstreams/ImageStreamGang.java](#)

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
boolean urlCached(URL url,
                   String filterName) {
    File file =
        new File(getPath(),
                  filterName);

    File imageFile =
        new File(file,
                  getNameForUrl(url));

    return imageFile.exists();
}
```

```
void processStream() {
    List<URL> urls = getInput();

    List<Image> filteredImages = urls
        .parallelStream()
        .filter(not(this::urlCached))
        .map(this::blockingDownload)
        .flatMap(this::applyFilters)
        .collect(toList());

    System.out.println(TAG
        + "Image(s) filtered = "
        + filteredImages.size());
}
```

Check if a file with this name already exists

See [imagestreamgang/streams/ImageStreamGang.java](https://github.com/imagestreamgang/streams/tree/main/ImageStreamGang.java)

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.



**ClearlyBetter®
SOLUTIONS**

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

There are clearly better ways of implementing an image cache!

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream consisting of the images that were downloaded from the URLs in the input stream

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream consisting of the images that were downloaded from the URLs in the input stream

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

of output stream elements must match the # of input stream elements

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
Image blockingDownload  
    (URL url) {  
    return BlockingTask  
        .callInManagedBlock  
        (() ->  
            downloadImage(url));  
}
```

*Downloads content from a URL
and converts it into an image*

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

See [imagestreamgangstreamsImageStreamParallel.java](#)

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
Image blockingDownload  
    (URL url) {  
    return BlockingTask  
        .callInManagedBlock  
        (() ->  
            downloadImage(url));  
}
```

Uses a "managed blocker" to ensure sufficient threads are in the common ForkJoinPool

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
Image blockingDownload  
    (URL url) {  
    return BlockingTask  
        .callInManagedBlock  
        (() ->  
            downloadImage(url));  
}
```

I/O-bound tasks on an N-core CPU typically run best with $N(1+WT/ST)$ threads (WT = wait time & ST = service time).*

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream containing the results of applying a list of filters to each image in the input stream and storing the results in the file system

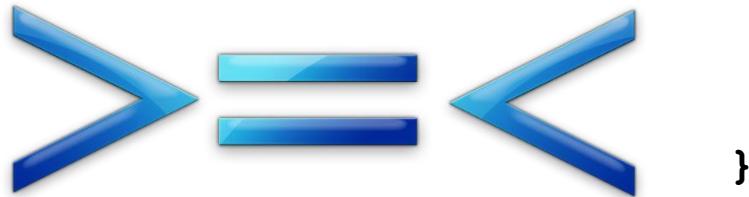


```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Return an output stream containing the results of applying a list of filters to each image in the input stream and storing the results in the file system



```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

of output stream elements may differ from the # of input stream elements

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
Stream<Image> applyFilters  
    (Image image) {  
    return mFilters  
        .parallelStream()  
        .map(filter ->  
            makeFilterWithImage  
            (filter,  
             image) .run())  
}
```

Apply all filters to an image in parallel and store on the device

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

See [imagestreamgangstreamsImageStreamParallel.java](#)

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());
```

Collect() is a "reduction" operation that combines elements into one result.

```
System.out.println(TAG  
    + "Image(s) filtered = "  
    + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

Trigger all intermediate operations

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());
```

```
System.out.println(TAG  
    + "Image(s) filtered = "  
    + filteredImages.size());  
}
```

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Create a list containing all the filtered and stored images

Implementing a Parallel Stream in ImageStreamGang

- We focus on processStream() in ImageStreamParallel.java.

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Logs the # of images that were downloaded, filtered, and stored.

Java Parallel ImageStreamGang Example: Implementing Behaviors

The End

Java Parallel ImageStreamGang Example

Evaluating Pros and Cons

Douglas C. Schmidt

Learning Objectives in This Part of the Lesson

- Recognize the structure/functionality of the ImageStreamGang app.
- Know how Java parallel streams are applied to the ImageStreamGang app.
- Understand the parallel streams implementation of ImageStreamGang.
- Be aware of the pros and cons of the parallel streams solution.



Pros of the Java Parallel Streams Solution

Pros of the Java Parallel Streams Solution

- The parallel stream version is faster than the sequential streams version.

Starting ImageStreamGangTest

Printing 4 results for input file 1 from fastest to slowest

COMPLETABLE_FUTURES_1 executed in 312 msecs

COMPLETABLE_FUTURES_2 executed in 335 msecs

PARALLEL_STREAM executed in 428 msecs

SEQUENTIAL_STREAM executed in 981 msecs

Printing 4 results for input file 2 from fastest to slowest

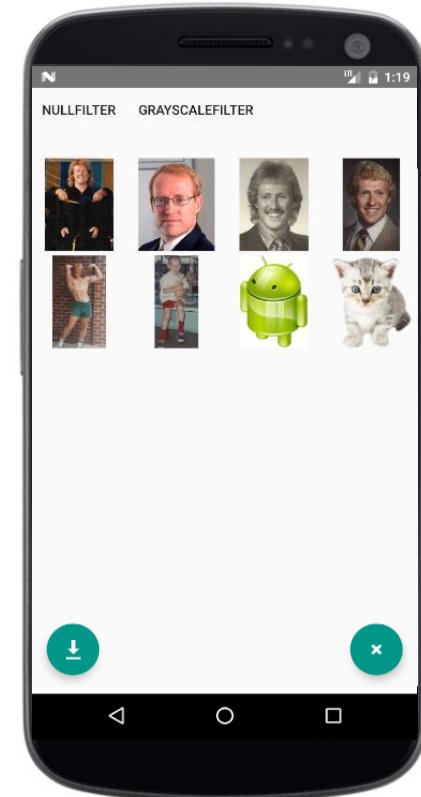
COMPLETABLE_FUTURES_2 executed in 82 msecs

COMPLETABLE_FUTURES_1 executed in 83 msecs

PARALLEL_STREAM executed in 102 msecs

SEQUENTIAL_STREAM executed in 251 msecs

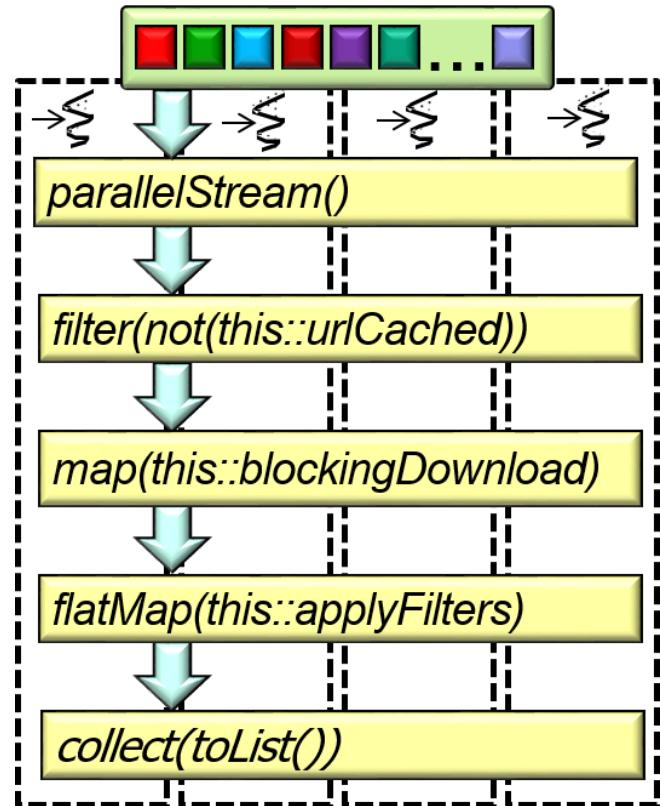
Ending ImageStreamGangTest



Six-core 2.6 GHz Windows Intel computer with 64 GB RAM

Pros of the Java Parallel Streams Solution

- The parallel stream version is faster than the sequential streams version.
 - e.g., images are downloaded and processed in parallel on multiple cores.



Pros of the Java Parallel Streams Solution

- The solution is relatively straightforward to understand.



```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Pros of the Java Parallel Streams Solution

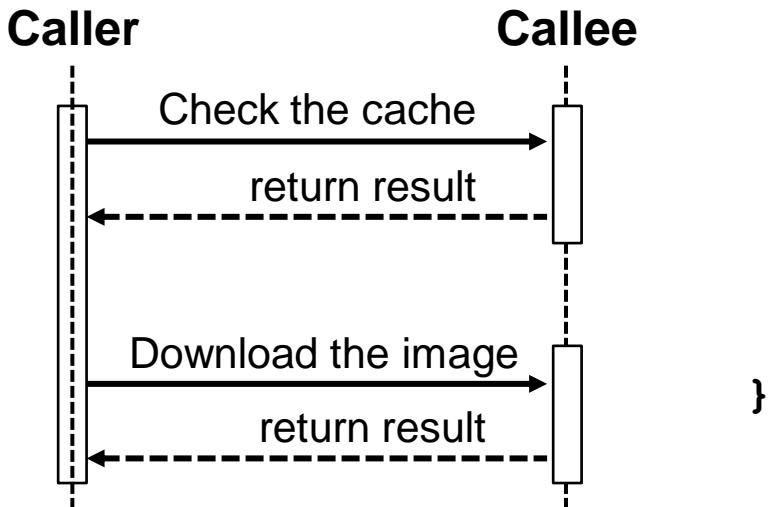
- The solution is relatively straightforward to understand, e.g.
 - The behaviors map cleanly onto the domain intent.



```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
  
    System.out.println(TAG  
        + "Image(s) filtered = "  
        + filteredImages.size());  
}
```

Pros of the Java Parallel Streams Solution

- The solution is relatively straightforward to understand, e.g.
 - The behaviors map cleanly onto the domain intent.
 - Behaviors are all synchronous.

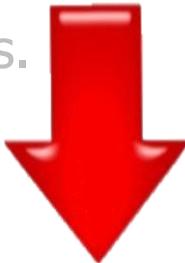


```
void processStream() {
    List<URL> urls = getInput();
    List<Image> filteredImages = urls
        .parallelStream()
        .filter(not(this::urlCached))
        .map(this::blockingDownload)
        .flatMap(this::applyFilters)
        .collect(toList());
    System.out.println(TAG
        + "Image(s) filtered = "
        + filteredImages.size());
}
```

Pros of the Java Parallel Streams Solution

- The solution is relatively straightforward to understand, e.g.
 - The behaviors map cleanly onto the domain intent.
 - Behaviors are all synchronous.
 - The flow of control can be read “linearly.”

```
void processStream() {  
    List<URL> urls = getInput();  
  
    List<Image> filteredImages = urls  
        .parallelStream()  
        .filter(not(this::urlCached))  
        .map(this::blockingDownload)  
        .flatMap(this::applyFilters)  
        .collect(toList());  
}
```



```
System.out.println(TAG  
    + "Image(s) filtered = "  
    + filteredImages.size());  
}
```

Cons of the Java Parallel Streams Solution

Cons of the Java Parallel Streams Solution

- The CompletableFuture versions are faster than the parallel streams version.

Starting ImageStreamGangTest

Printing 4 results for input file 1 from fastest to slowest

COMPLETABLE_FUTURES_1 executed in 312 msecs

COMPLETABLE_FUTURES_2 executed in 335 msecs

PARALLEL_STREAM executed in 428 msecs

SEQUENTIAL_STREAM executed in 981 msecs

Printing 4 results for input file 2 from fastest to slowest

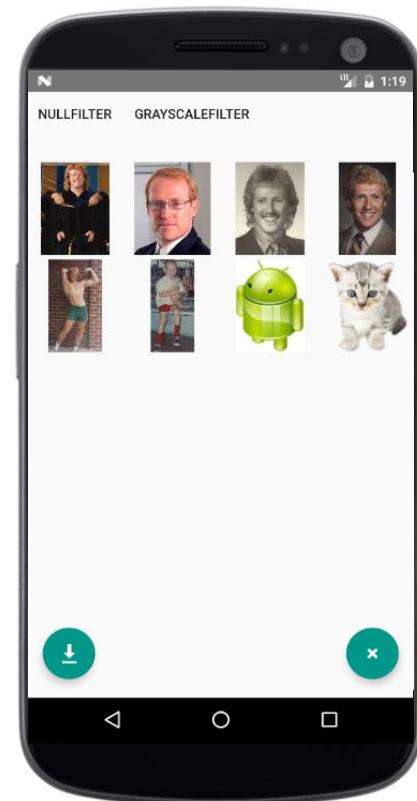
COMPLETABLE_FUTURES_2 executed in 82 msecs

COMPLETABLE_FUTURES_1 executed in 83 msecs

PARALLEL_STREAM executed in 102 msecs

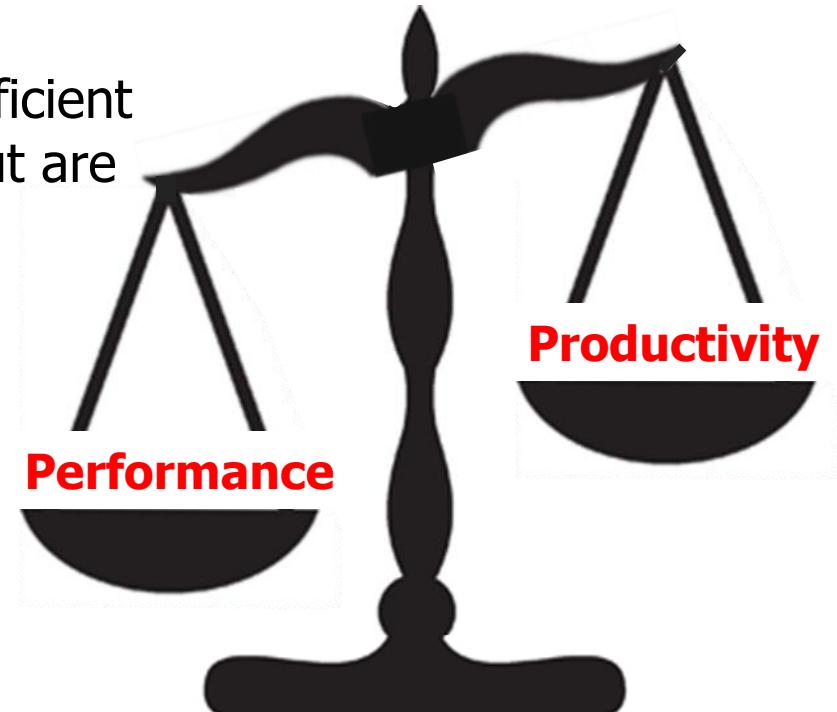
SEQUENTIAL_STREAM executed in 251 msecs

Ending ImageStreamGangTest



Cons of the Java Parallel Streams Solution

- In general, there's a trade-off between computing performance and programmer productivity when choosing amongst Java parallelism frameworks.
 - e.g., `CompletableFuture`s are more efficient and scalable than parallel streams, but are somewhat harder to program.



Java Parallel ImageStreamGang Example:
Evaluating Pros and Cons

The End

Java Parallel Streams

Evaluating the Pros

Douglas C. Schmidt

Learning Objectives in This Lesson

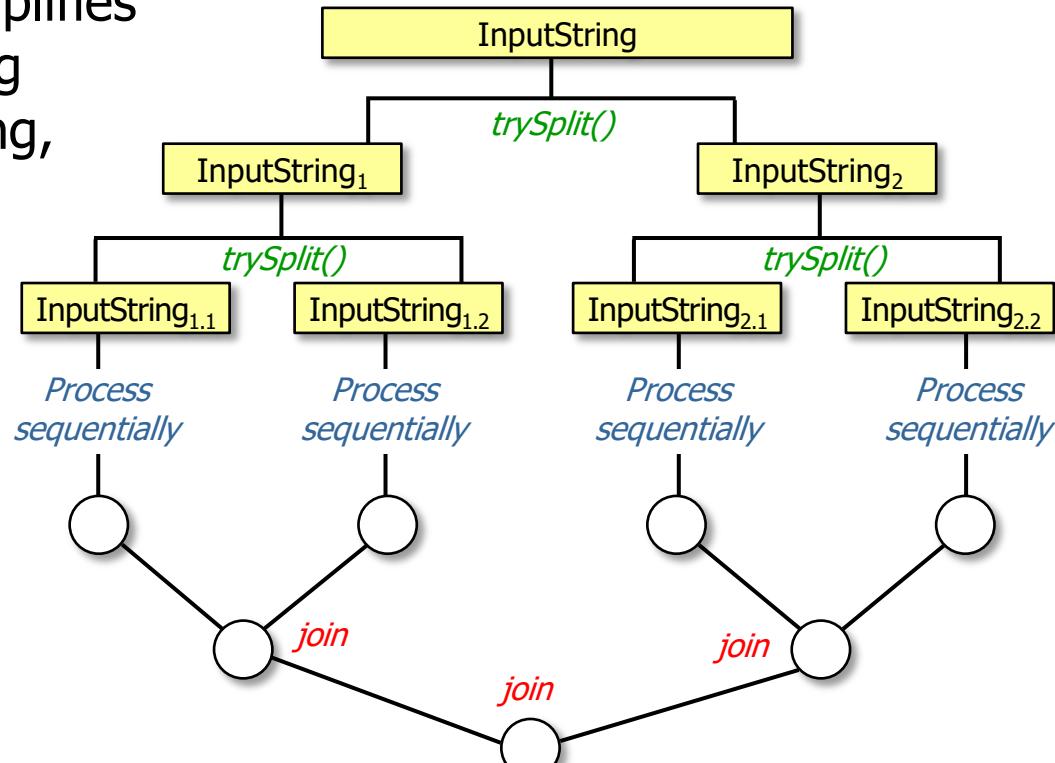
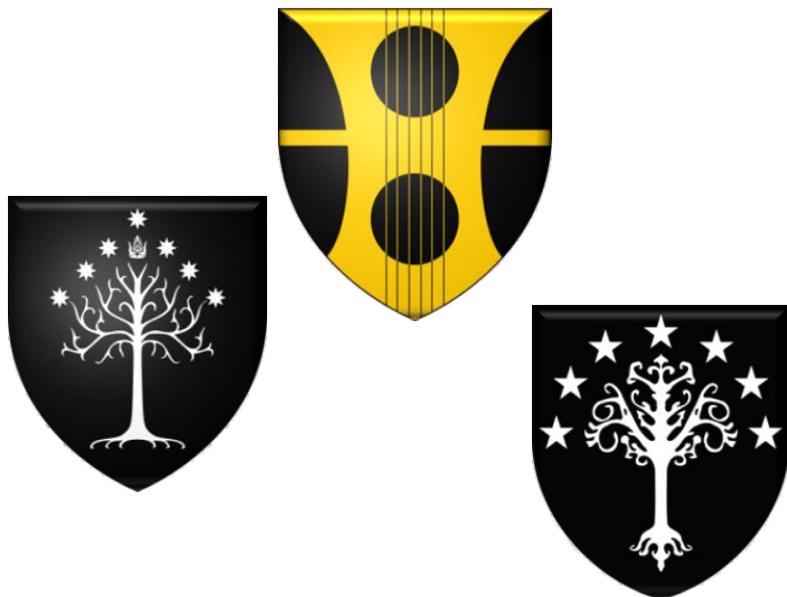
- Evaluate the pros of Java parallel streams.



Pros of Java Parallel Streams

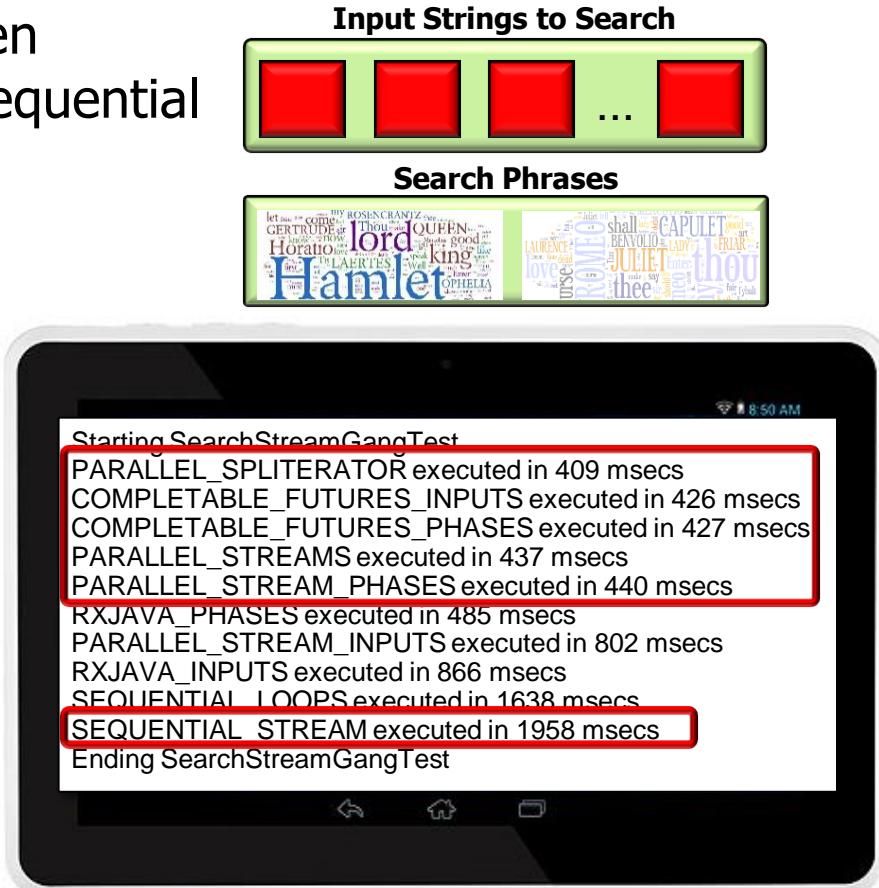
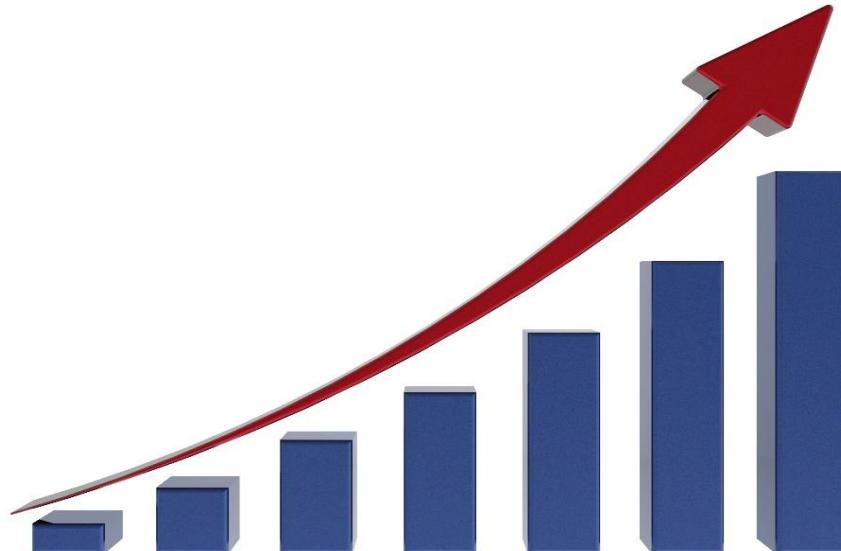
Pros of Java Parallel Streams

- The Java streams framework simplifies parallel programming by shielding developers from details of splitting, applying, and combining results.



Pros of Java Parallel Streams

- Parallel stream implementations are often (much) faster and more scalable than sequential (stream and loops) implementations.

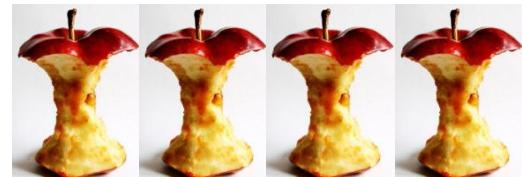
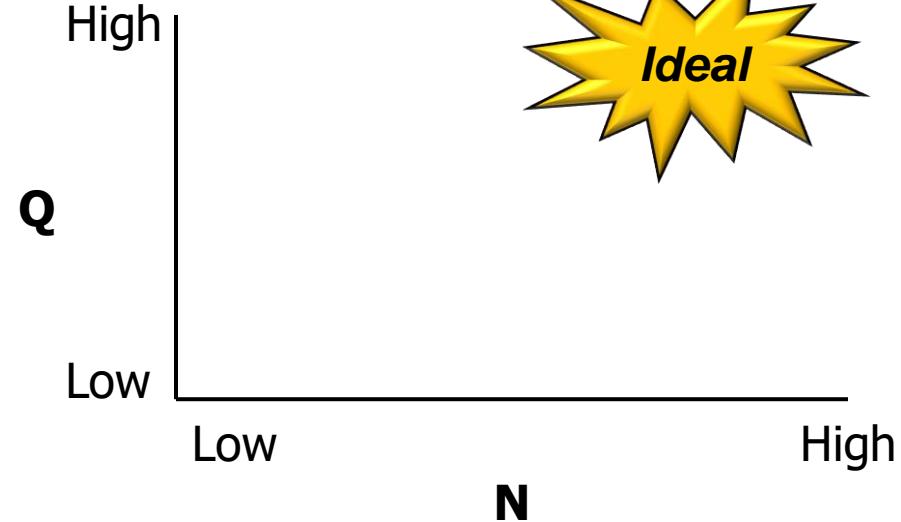


Pros of Java Parallel Streams

- The performance speedup is largely a function of the partitioning strategy for the input (N), the amount of work performed (Q), and the # of cores.

The NQ model

- N is the # of data elements to process per thread.
- Q quantifies how CPU-intensive the processing is.



Pros of Java Parallel Streams

- Apps often don't need explicit synchronization or threading.



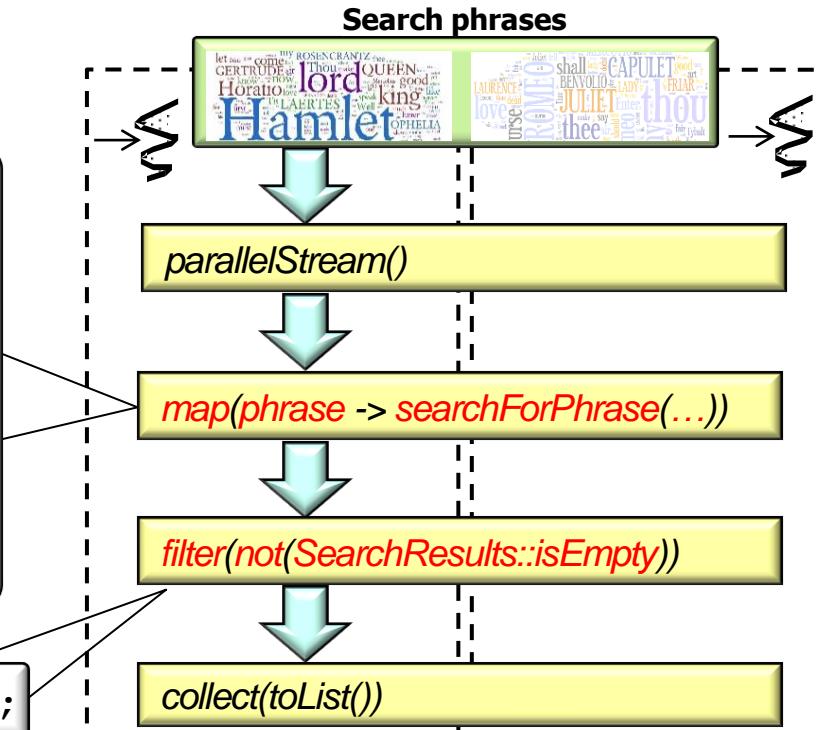
Alleviates many accidental and inherent complexities of concurrency/parallelism

Pros of Java Parallel Streams

- Apps often don't need explicit synchronization or threading.
 - Stateless behaviors alleviate the need to access shared mutable state.

```
return new SearchResults  
(Thread.currentThread().getId(),  
currentCycle(), phrase, title,  
StreamSupport  
.stream(new PhraseMatchSpliterator  
        (input, phrase),  
        parallel)  
.collect(toList()));
```

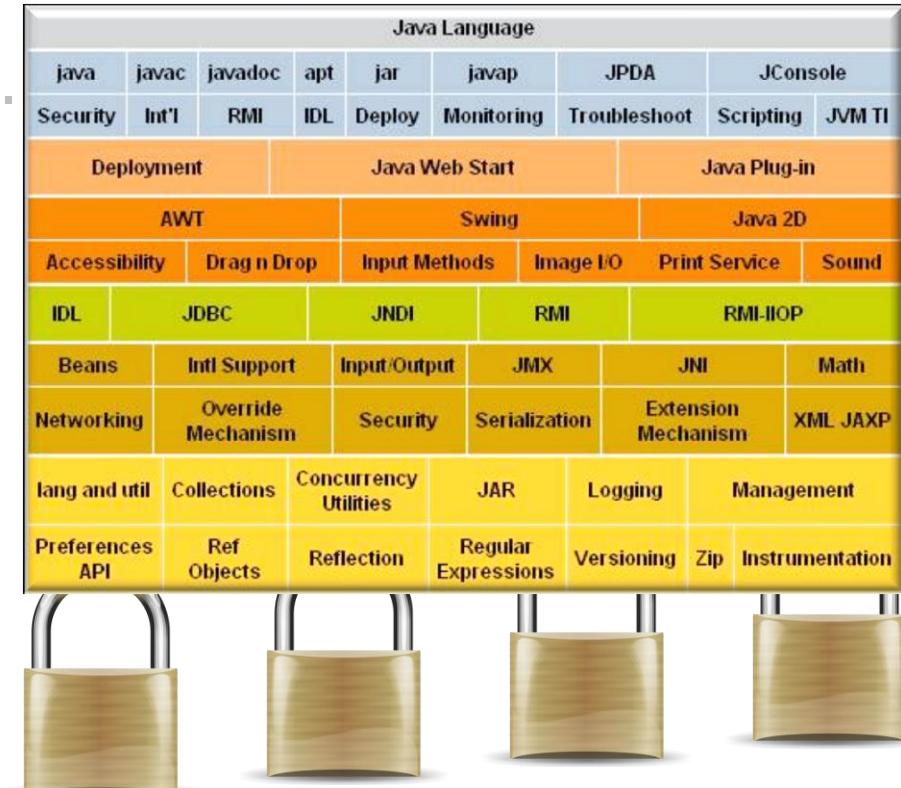
```
return mList.size() == 0;
```



See en.wikipedia.org/wiki/Pure_function

Pros of Java Parallel Streams

- Apps often don't need explicit synchronization or threading.
 - Stateless behaviors alleviate the need to access shared mutable state.
 - Java class library handles locking needed to protect shared mutable state.



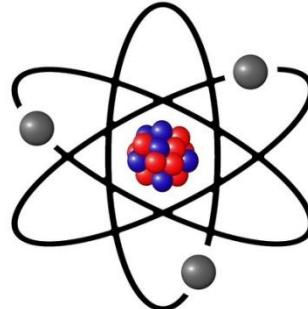
See docs.oracle.com/javase/tutorial/essential/concurrency/collections.html

Pros of Java Parallel Streams

- Streams ensure that the structure of sequential and parallel code is the same.

```
List<List<SearchResults>>
    processStream() {
    return getInput()
        .stream()
        .map(this::processInput)
        .collect(toList());
}
```

```
List<List<SearchResults>>
    processStream() {
    return getInput()
        .parallelStream()
        .map(this::processInput)
        .collect(toList());
}
```



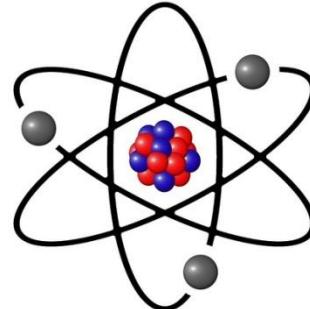
Converting sequential to parallel streams only requires minuscule changes!

Pros of Java Parallel Streams

- Streams ensure that the structure of sequential and parallel code is the same.

```
List<SearchResults> results =  
    mPhrasesToFind  
        .parallelStream()  
        .map(phase ->  
            searchForPhrase(...,  
                false))  
        .filter(not(SearchResults  
            ::isEmpty))  
        .collect(toList());
```

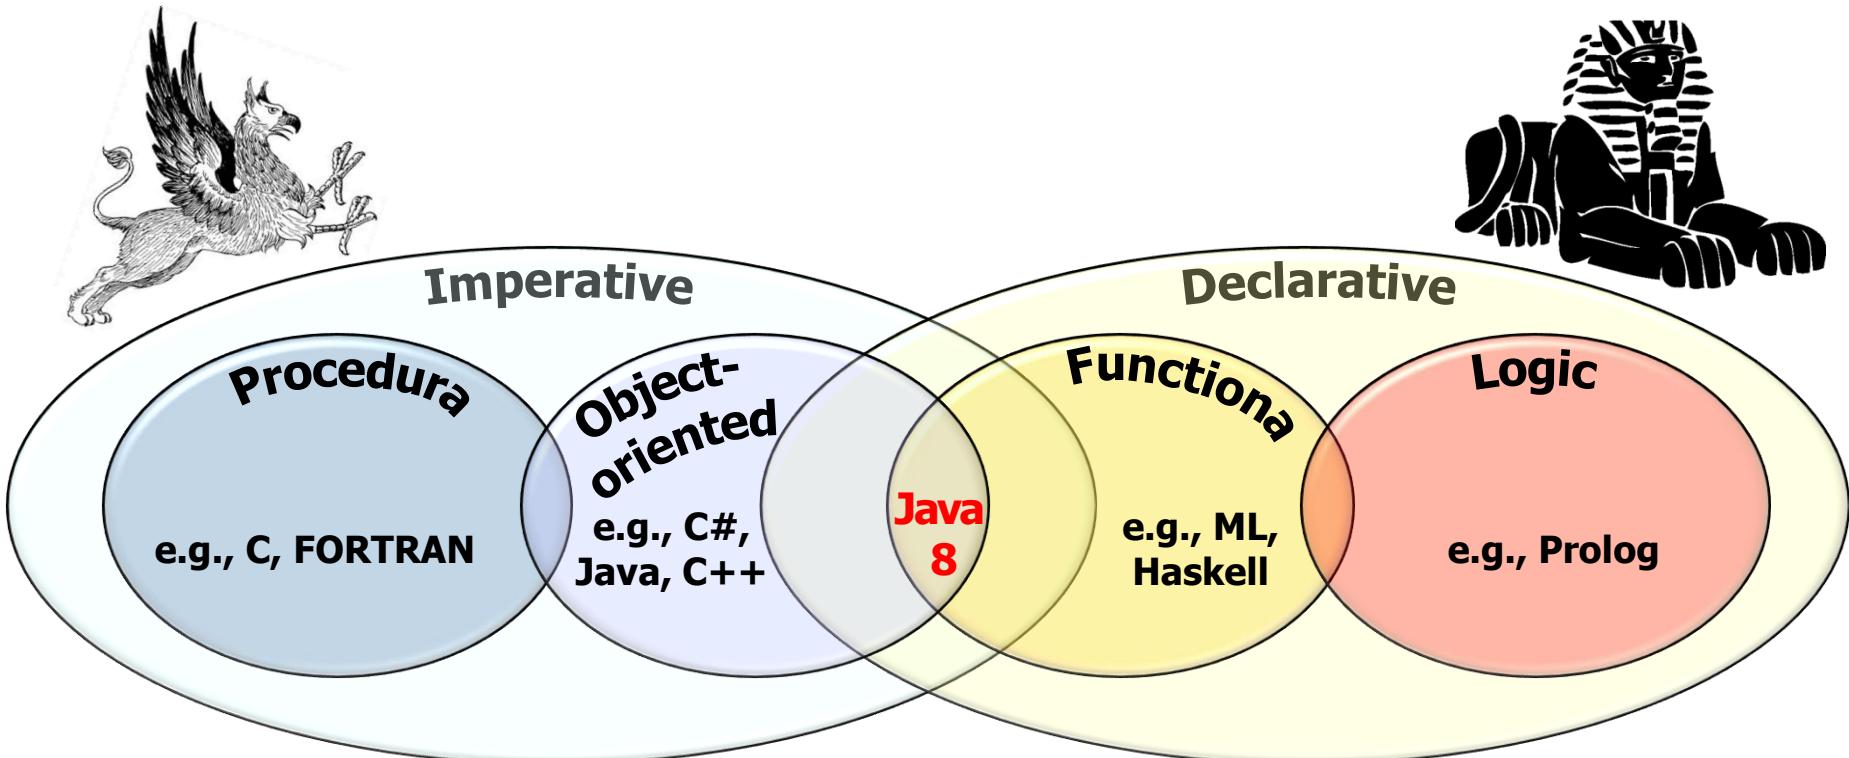
```
List<SearchResults> results =  
    mPhrasesToFind  
        .parallelStream()  
        .map(phase ->  
            searchForPhrase(...,  
                true))  
        .filter(not(SearchResults  
            ::isEmpty))  
        .collect(toList());
```



Converting sequential to parallel streams only requires minuscule changes!

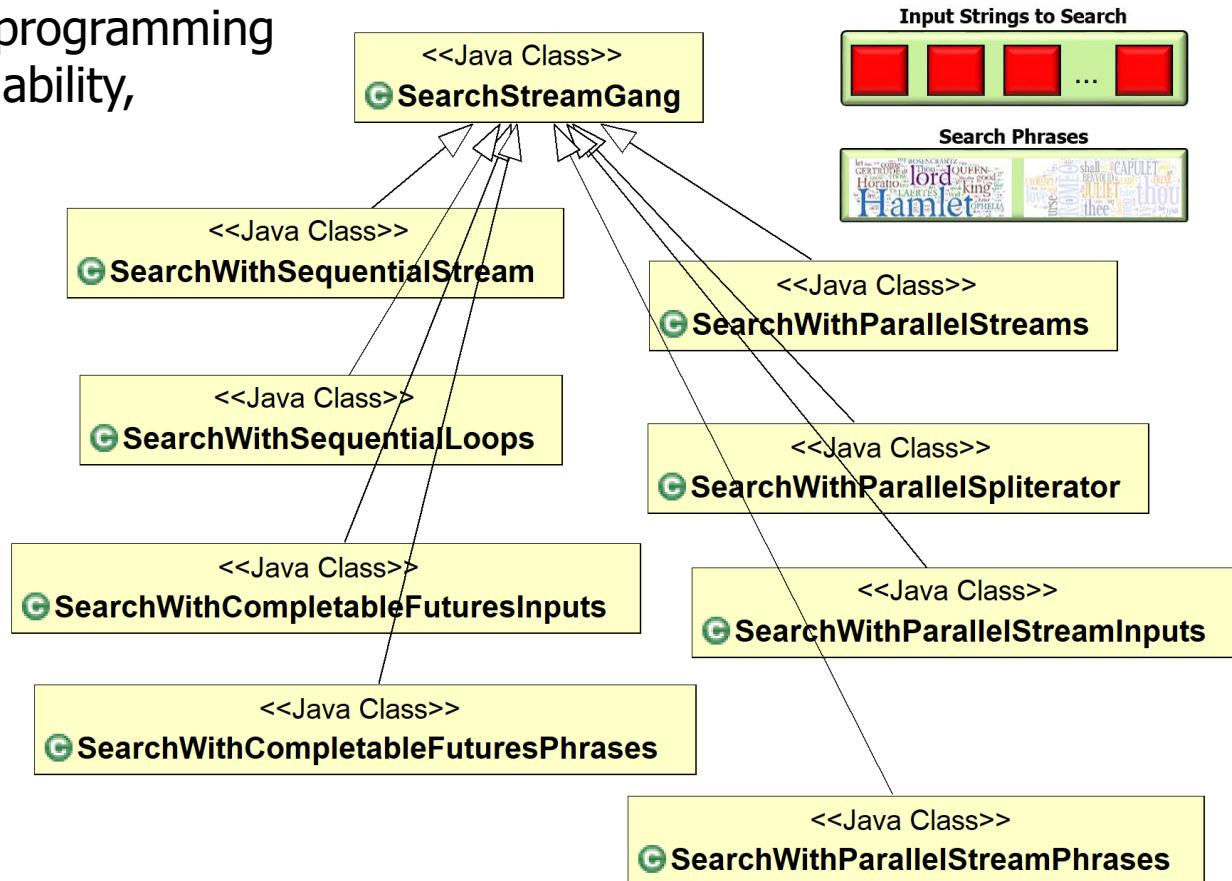
Pros of Java Parallel Streams

- Examples show synergies between functional and object-oriented programming.



Pros of Java Parallel Streams

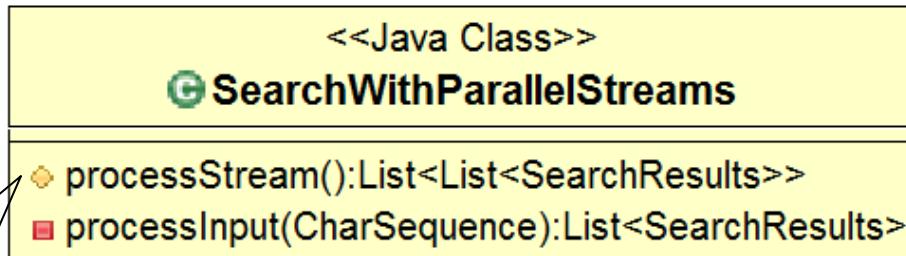
- Object-oriented design and programming features simplify understandability, reusability, and extensibility.



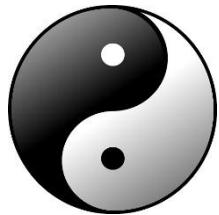
Object-oriented techniques emphasize systematic reuse of *structure*.

Pros of Java Parallel Streams

- Implementing object-oriented hook methods with functional programming features helps to close gap between domain intent and computations.



```
getInput()  
    .parallelStream()  
    .map(this::processInput)  
    .collect(toList());
```



```
return mPhrasesToFind  
    .parallelStream()  
    .map(phrase -> searchForPhrase(phrase, input, title, false))  
    .filter(not(SearchResults::isEmpty))  
    .collect(toList());
```

Java Parallel Streams: Evaluating the Pros

The End

Java Parallel Streams

Evaluating the Cons

Douglas C. Schmidt

Learning Objectives in This Lesson

- Evaluate the cons of Java parallel streams.



Cons of Java Parallel Streams

Cons of Java Parallel Streams

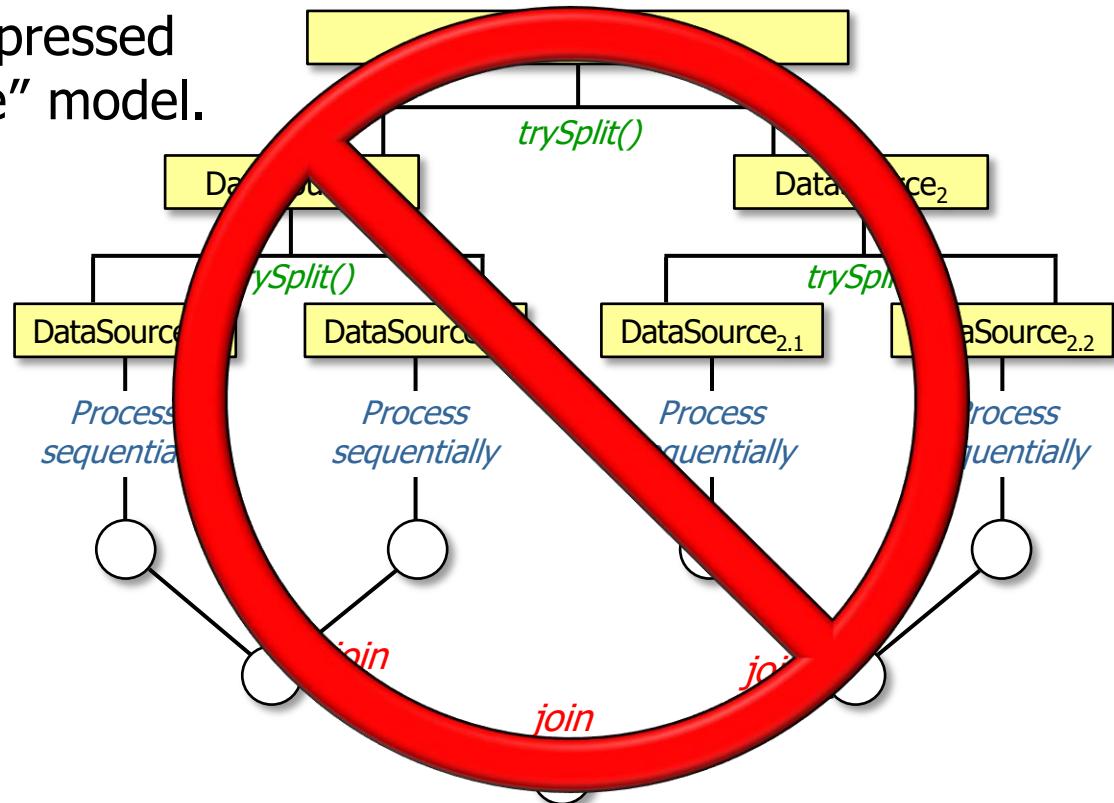
- There are some limitations with Java parallel streams.



The Java parallel streams framework is not all unicorns and rainbows!

Cons of Java Parallel Streams

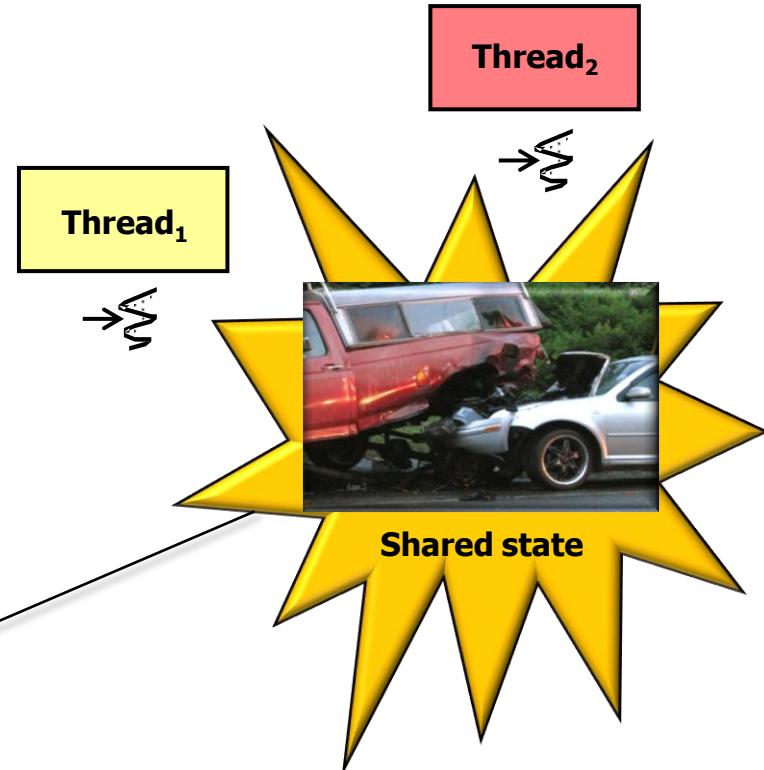
- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.



See dzone.com/articles/whats-wrong-java-8-part-iii

Cons of Java Parallel Streams

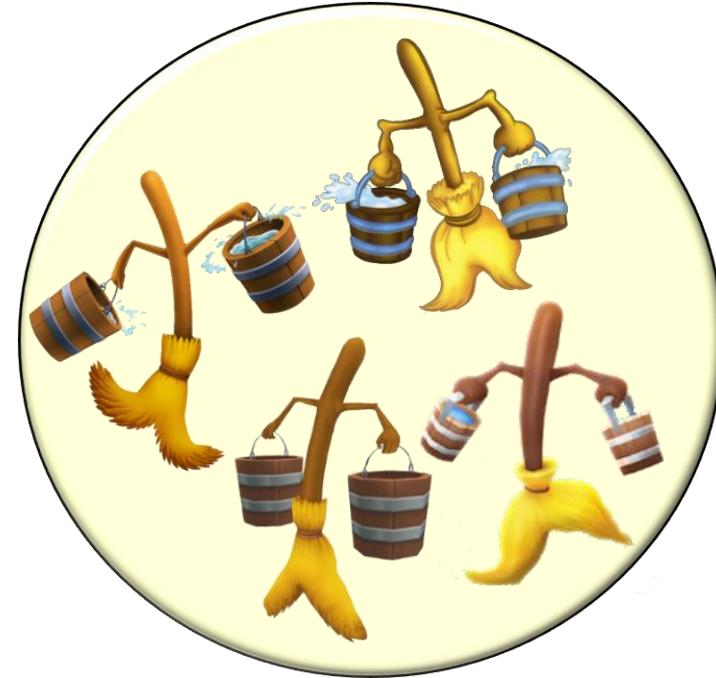
- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.



Race conditions occur when a program depends on the sequence or timing of threads for it to operate properly.

Cons of Java Parallel Streams

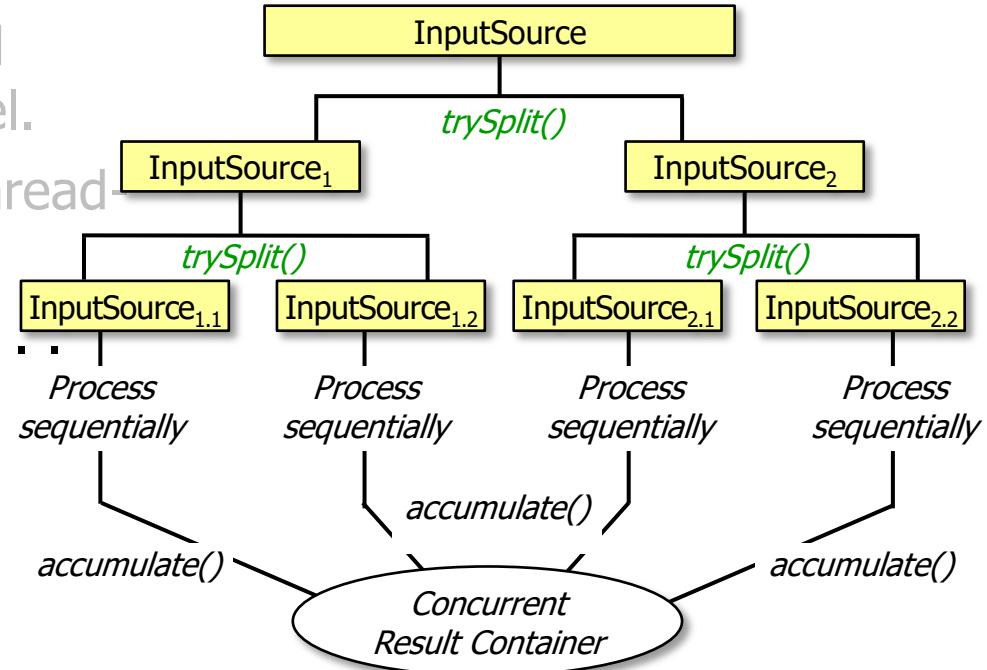
- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .



See lesson on “*Java SearchWithParallelSpliterator Example: trySplit()*.”

Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - Concurrent collectors are easier.



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Java CompletableFuture don't have this limitation.



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Java CompletableFutures don't have this limitation.
 - It's important to know how to apply ManagedBlockers.



See "*The Java ForkJoinPool: Applying the ManagedBlocker Interface.*"

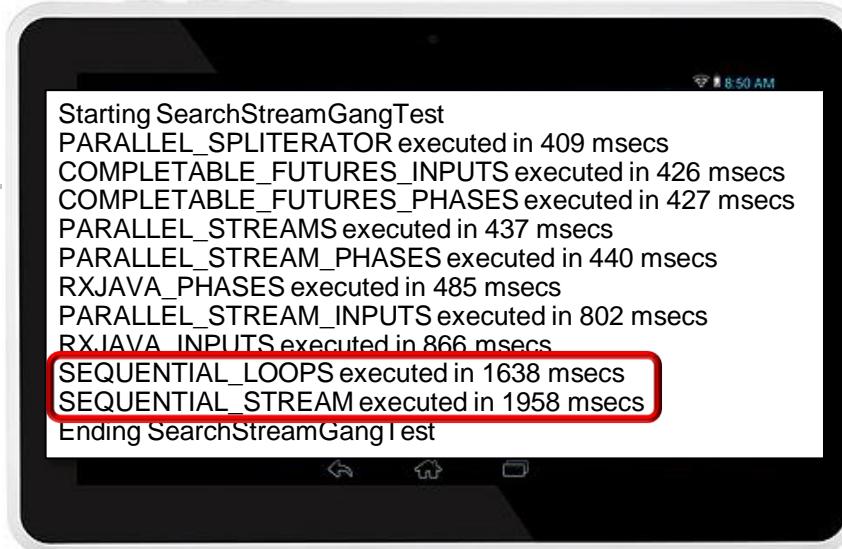
Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead.



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead, e.g.
 - Splitting and combining overhead



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead, e.g.
 - Splitting and combining overhead
 - Fork/join framework

A Java Fork/Join Blunder

Ed Harnd
eh at coopsoft dot com

The F/J framework is a faulty enterprise from the beginning. The basic design is Divide-and Conquer using dyadic recursive decomposition. Simply put, the framework supports tasks that decompose or fork into two tasks, that decompose into two tasks, that decompose... When the decomposing or forking stops, the bottom tasks return a result up the chain. The forking tasks retrieve the results of the forked tasks with an intermediate join¹. Hence, Fork/Join. This is a beautiful design in theory. In the reality of JavaSE it doesn't work well.

It doesn't work well because it is the wrong tool for the job. The F/J framework is the underlying software experiment for the 2000 research paper, "A Java Fork/Join Framework."² That experimental software is not, has never been, and will never be the foundation for a general-purpose application framework. Using such a tool for application development is like using a pocketknife to chisel a granite sculpture. There is just so, so much wrong with the F/J framework as a general-purpose, commercial application development tool that the author wrote two articles³, with seventeen (17) points, to illustrate the calamity. This paper is a consolidation of those articles explaining why the F/J framework is the wrong tool for the job.

There are four major faults with the F/J framework:

1. The use of Deques/Submission queues
2. The use of an intermediate join()
3. The use of academic research standards instead of application development standards
4. The use of the CountedCompleter class

1. The use of Deques/Submission queues

The first design fault with the F/J framework is the use of Deques/Submission queues. Deques/Submission-Queues are a feature primarily for

1. Applications that run on clusters of computers (Cilk for one.)
2. Operating systems that balance the load between CPU's.
3. A number of other environments irrelevant to this discussion.

While dequeues are efficient in limiting contention (there are many academic research papers on work-stealing and dequeues), there is no hint of how new processes (tasks) actually get into the dequeues.

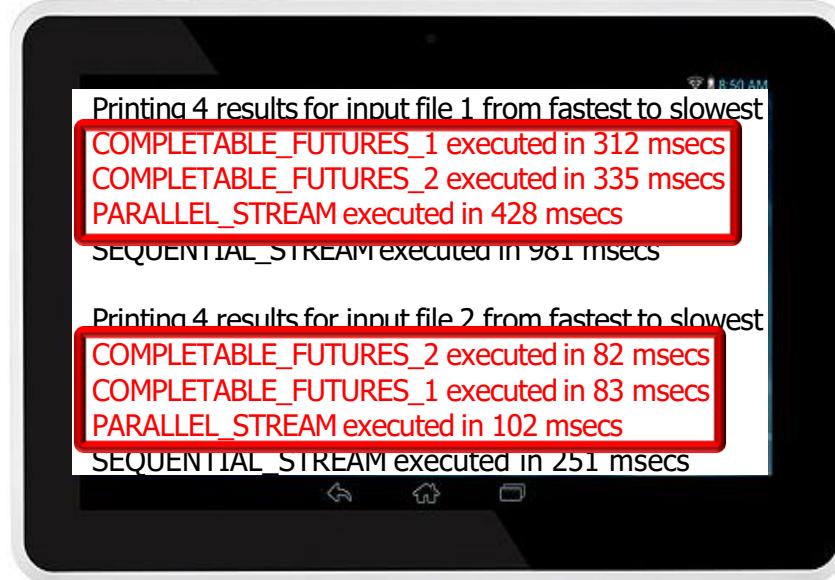
¹ An intermediate join() waits for the fork() to complete and should not be confused with a Thread.join() where the later waits for another Thread to finish.

² <http://gee.cs.oswego.edu/dl/papers/fj.pdf>

³ <http://coopsoft.com/ar/CalamityArticle.html>

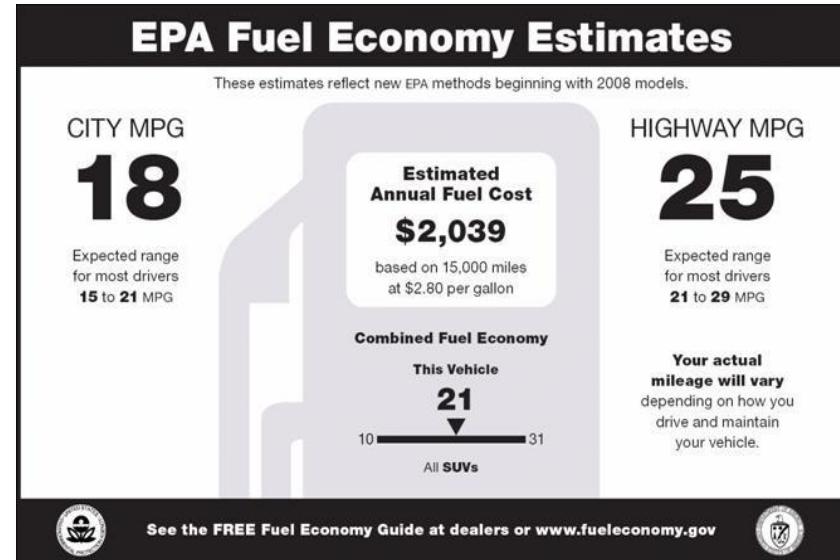
Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead, e.g.
 - Splitting and combining overhead
 - Fork/join framework
 - Java CompletableFuture may be more efficient and scalable



Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead, e.g.
 - Splitting and combining overhead
 - Fork/join framework
 - Java CompletableFuture may be more efficient and scalable



Naturally, your mileage may vary.

Cons of Java Parallel Streams

- There are some limitations with Java parallel streams, e.g.
 - Some problems can't be expressed via the "split-apply-combine" model.
 - If behaviors aren't stateless and thread-safe, race conditions may occur.
 - Parallel spliterators may be tricky . . .
 - All parallel streams share a common ForkJoinPool.
 - Streams incur some overhead.
 - There's no substitute for benchmarking!

algorithms array avoiding worst
practices BigDecimal binary serialization
bitset book review boxing byte buffer

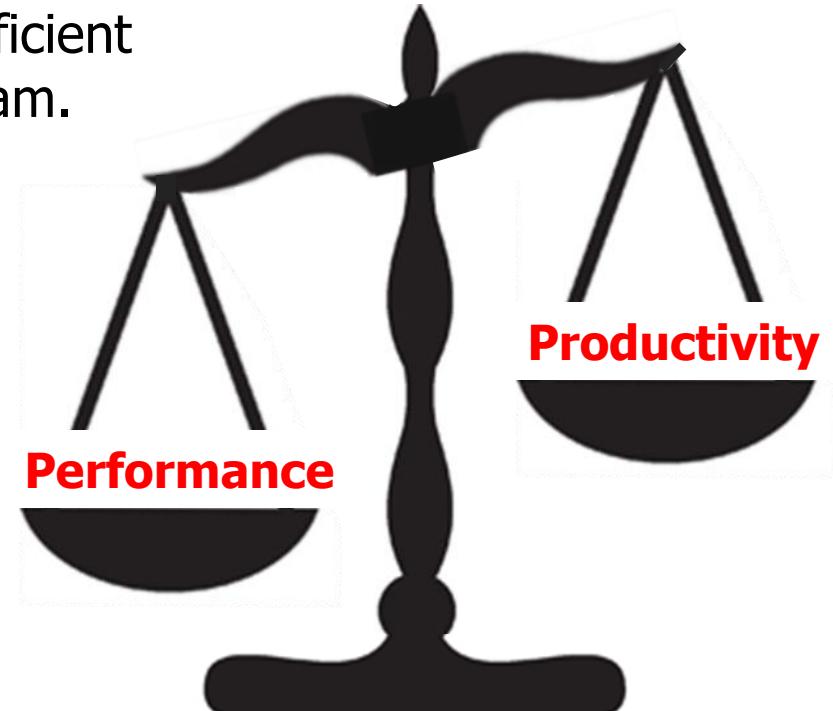
collections cpu
optimization data
compression datatype
optimization date dateformat double
exceptions FastUtil FIX hashCode hashmap
hdd hppc io Java 7 Java 8 java dates jdk
8 JMH JNI Koloboke map memory layout

memory
optimization multithreading
parsing primitive collections profiler ssd
string string concatenation string pool
sun.misc.Unsafe tools trove

Wrapping Up Java Parallel Streams

Wrapping Up Java Parallel Streams

- In general, there's a trade-off between computing performance and programmer productivity when choosing amongst these frameworks.
 - e.g., `CompletableFuture`s are more efficient and scalable, but are harder to program.



Wrapping Up Java Parallel Streams

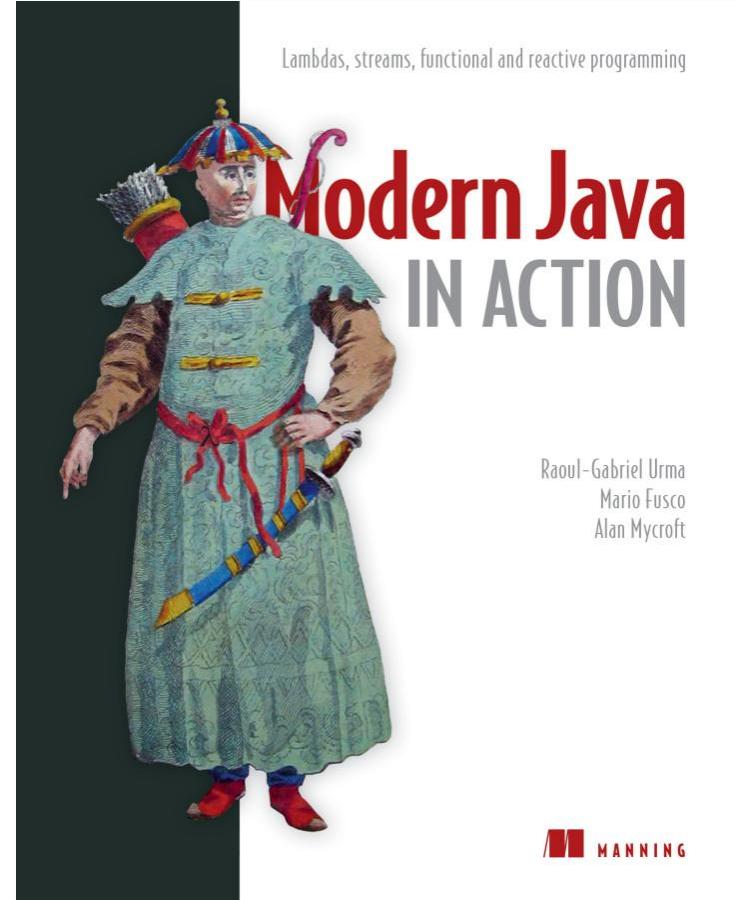
- In general, however, the pros of Java parallel streams far outweigh the cons for many use cases!



See www.ibm.com/developerworks/library/j-jvmc2

Wrapping Up Java Parallel Streams

- Good coverage of parallel streams appears in the book “Modern Java in Action.”



See www.manning.com/books/modern-java-in-action

Java Parallel Streams: Evaluating the Cons

The End